
CS250 Computer Architecture

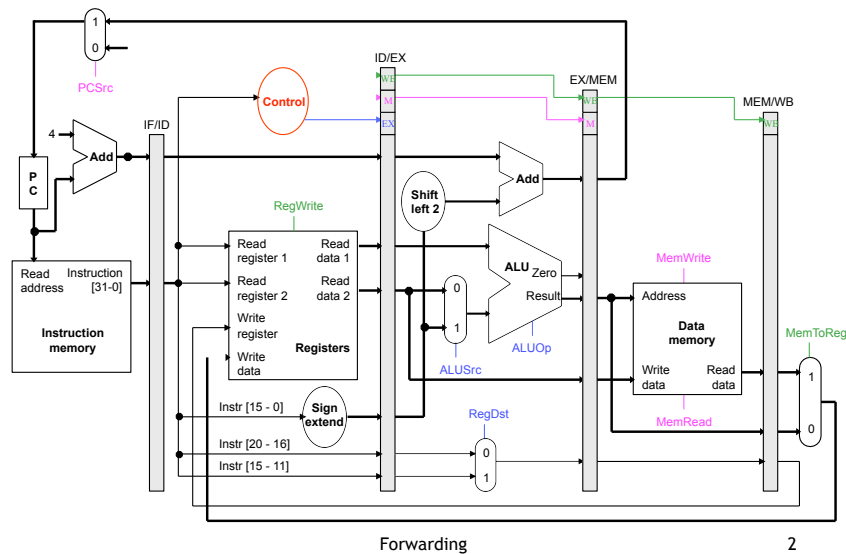
Data Hazards and Forwarding

Forwarding


- Previously, we introduced a **pipelined** MIPS processor which executes several instructions simultaneously.
 - Each instruction requires five stages, and five cycles, to complete.
 - Each stage uses different functional units of the datapath.
 - So we can execute up to five instructions in any clock cycle, with each instruction in a different stage and using different hardware.
- Today we'll introduce some problems that **data hazards** can cause for our pipelined processor, and show how to handle them with **forwarding**.



The pipelined datapath



Pipeline diagram review



- This diagram shows the execution of an ideal code fragment.
 - Each instruction needs a total of five cycles for execution.
 - One instruction begins on every clock cycle for the first five cycles.
 - One instruction completes on each cycle from that time on.

Forwarding

3

Our examples are too simple

- Here is the example instruction sequence used to illustrate pipelining on the previous page.

```
lw    $8, 4($29)
sub   $2, $4, $5
and   $9, $10, $11
or    $16, $17, $18
add   $13, $14, $0
```

- The instructions in this example are **independent**.
 - Each instruction reads and writes completely different registers.
 - Our datapath handles this sequence easily, as we saw last time.
- But most sequences of instructions are *not* independent!

Forwarding

4

An example with dependencies

```
sub   $2, $1, $3
and   $12, $2, $5
or    $13, $6, $2
add   $14, $2, $2
sw    $15, 100($2)
```

Forwarding

5

An example with dependencies

```

sub    $2, $1, $3
and    $12, $2, $5
or     $13, $6, $2
add    $14, $2, $2
sw     $15, 100($2)


```

- There are several **dependencies** in this new code fragment.
 - The first instruction, SUB, stores a value into **\$2**.
 - That register is used as a source in the rest of the instructions.
- This is not a problem for the single-cycle datapath.
 - Each instruction is executed completely before the next one begins.
 - This ensures that instructions 2 through 5 above use the new value of \$2 (the sub result), just as we expect.
- How would this code sequence fare in our pipelined datapath?

Forwarding

6

Data hazards in the pipeline diagram




- The SUB instruction does not write to register \$2 until clock cycle 5. This causes two **data hazards** in our current pipelined datapath.
 - The AND reads register \$2 in cycle 3. Since SUB hasn't modified the register yet, this will be the *old* value of \$2, not the new one.
 - Similarly, the OR instruction uses register \$2 in cycle 4, again before it's actually updated by SUB.

Forwarding

7

Things that are okay




- The ADD instruction is okay, because of the register file design.
 - Registers are written at the beginning of a clock cycle.
 - The new value will be available by the end of that cycle.
- The SW is no problem at all, since it reads \$2 after the SUB finishes.

Forwarding

8

Dependency arrows




- Arrows indicate the flow of data between instructions.
 - The tails of the arrows show when register \$2 is written.
 - The heads of the arrows show when \$2 is read.
- Any arrow that points backwards in time represents a **data hazard** in our basic pipelined datapath. Here, hazards exist between instructions 1 & 2 and 1 & 3.

Forwarding

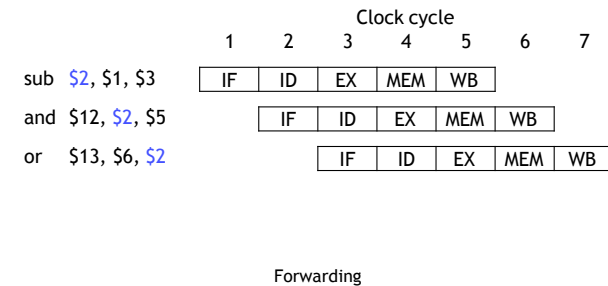
9

A fancier pipeline diagram




A more detailed look at the pipeline

- We have to eliminate the hazards, so the AND and OR instructions in our example will use the correct value for register \$2.
- When is the data actually produced and consumed?
- What can we do?



A more detailed look at the pipeline

- We have to eliminate the hazards, so the AND and OR instructions in our example will use the correct value for register \$2.
- Let's look at when the data is actually produced and consumed.
 - The SUB instruction produces its result in its EX stage, during cycle 3 in the diagram below.
 - The AND and OR need the new value of \$2 in their EX stages, during clock cycles 4-5 here.




Forwarding

12

Bypassing the register file

- The actual result \$1 - \$3 is computed in clock cycle 3, *before* it's needed in cycles 4 and 5.
- If we could somehow bypass the writeback and register read stages when needed, then we can eliminate these data hazards.
 - Today we'll focus on hazards involving arithmetic instructions.
 - Next time, we'll examine the lw instruction.
- Essentially, we need to pass the ALU output from SUB directly to the AND and OR instructions, without going through the register file.




Forwarding

13

Where to find the ALU result

- The ALU result generated in the EX stage is normally passed through the pipeline registers to the MEM and WB stages, before it is finally written to the register file.
- This is an abridged diagram of our pipelined datapath.

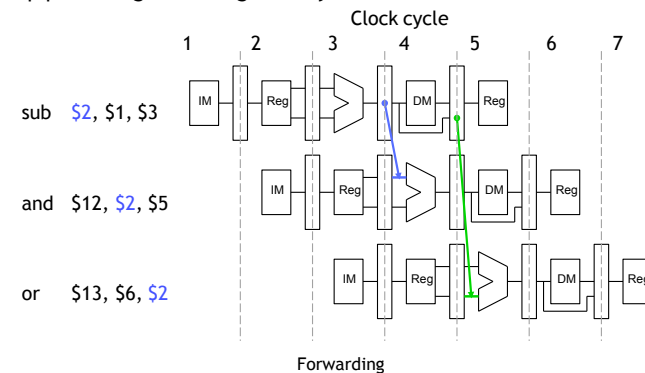


Forwarding

14

Forwarding

- Since the pipeline registers already contain the ALU result, we could just **forward** that value to subsequent instructions, to prevent data hazards.
 - In clock cycle 4, the AND instruction can get the value \$1 - \$3 from the **EX/MEM** pipeline register used by sub.
 - Then in cycle 5, the OR can get that same result from the **MEM/WB** pipeline register being used by SUB.

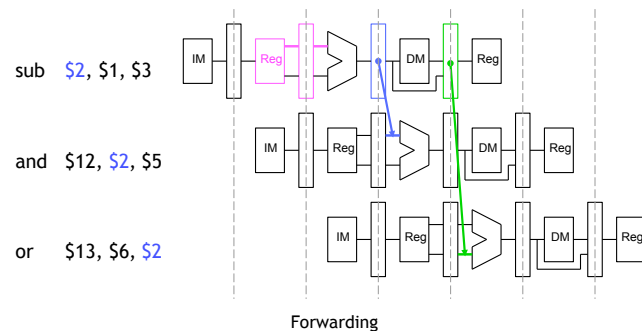


Forwarding

15

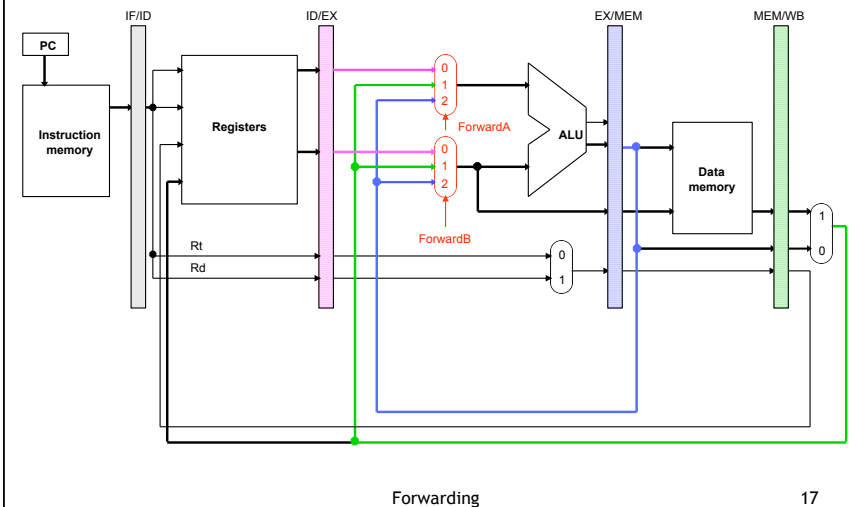
Outline of forwarding hardware

- A **forwarding unit** selects the correct ALU inputs for the EX stage.
 - If there is no hazard, the ALU's operands will come from the **register file**, just like before.
 - If there is a hazard, the operands will come from either the **EX/MEM** or **MEM/WB** pipeline registers instead.
- The ALU sources will be selected by two new multiplexers, with control signals named **ForwardA** and **ForwardB**.



16


Simplified datapath with forwarding muxes



17

Detecting EX/MEM data hazards

- When do we need to know that a hazard exists?
- So how can the hardware determine if a hazard exists?




Forwarding

18

Detecting EX/MEM data hazards

- So how can the hardware determine if a hazard exists?
- An **EX/MEM hazard** occurs between the instruction currently in its EX stage and the previous instruction if:
 1. The previous instruction will write to the register file, *and*
 2. The destination is one of the ALU source registers in the EX stage.
- There is an EX/MEM hazard between the two instructions below.



- Data in a pipeline register can be referenced using a class-like syntax. For example, **ID/EX.RegisterRt** refers to the rt field stored in the ID/EX pipeline.

Forwarding

19

EX/MEM data hazard equations

The first ALU source comes from the pipeline register when necessary.

if ($EX/MEM.RegWrite = 1$ and

$EX/MEM.RegisterRd = ID/EX.RegisterRs$)

then $ForwardA = 2$

- The second ALU source is similar.


if ($EX/MEM.RegWrite = 1$ and

$EX/MEM.RegisterRd = ID/EX.RegisterRt$)

then $ForwardB = 2$

sub \$2, \$1, \$3

and \$12, \$2, \$5



Forwarding

20

Detecting MEM/WB data hazards

- A **MEM/WB hazard** may occur between an instruction in the EX stage and the instruction from *two* cycles ago.
- One new problem is if a register is updated twice in a row.


```
add $1, $2, $3
add $1, $1, $4
sub $5, $5, $1
```

- Register \$1 is written by *both* of the previous instructions; from which instruction should it receive its value?

add \$1, \$2, \$3

add \$1, \$1, \$4

sub \$5, \$5, \$1



Forwarding

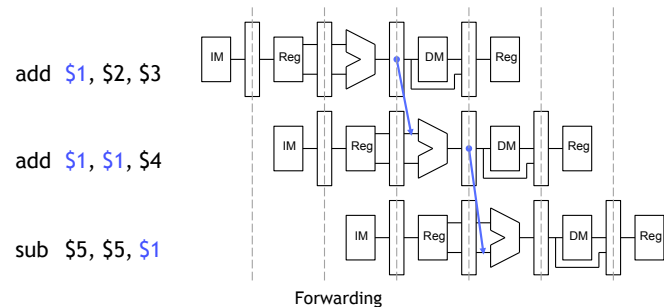
21

Detecting MEM/WB data hazards

- A **MEM/WB hazard** may occur between an instruction in the EX stage and the instruction from *two* cycles ago.
- One new problem is if a register is updated twice in a row.

```
add $1, $2, $3
add $1, $1, $4
sub $5, $5, $1
```

- Register \$1 is written by *both* of the previous instructions, but only the most recent result (from the second ADD) should be forwarded.



22

MEM/WB hazard equations

- Here is an equation for detecting and handling MEM/WB hazards for the first ALU source.

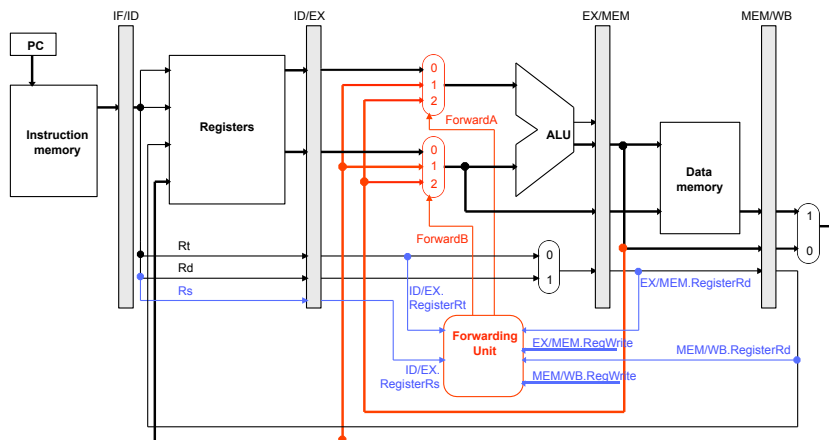

```
if (MEM/WB.RegWrite = 1
    and MEM/WB.RegisterRd = ID/EX.RegisterRs
    and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRs or EX/MEM.RegWrite = 0)
    then ForwardA = 1
```
- The second ALU operand is handled similarly.


```
if (MEM/WB.RegWrite = 1
    and MEM/WB.RegisterRd = ID/EX.RegisterRt
    and (EX/MEM.RegisterRd ≠ ID/EX.RegisterRt or EX/MEM.RegWrite = 0)
    then ForwardB = 1
```

Forwarding

23

Simplified datapath with forwarding



Forwarding

24

The forwarding unit

- The forwarding unit has several control signals as inputs.

ID/EX.RegisterRs EX/MEM.RegisterRd MEM/WB.RegisterRd
ID/EX.RegisterRt EX/MEM.RegWrite MEM/WB.RegWrite

(The two RegWrite signals are not shown in the diagram, but they come from the control unit.)

- The forwarding unit outputs are selectors for the ForwardA and ForwardB multiplexers attached to the ALU. These outputs are generated from the inputs using the equations on the previous pages.
- Some new buses route data from pipeline registers to the new muxes.

Forwarding

25

Example

```

sub  $2, $1, $3
and  $12, $2, $5
or   $13, $6, $2
add  $14, $2, $2
sw   $15, 100($2)

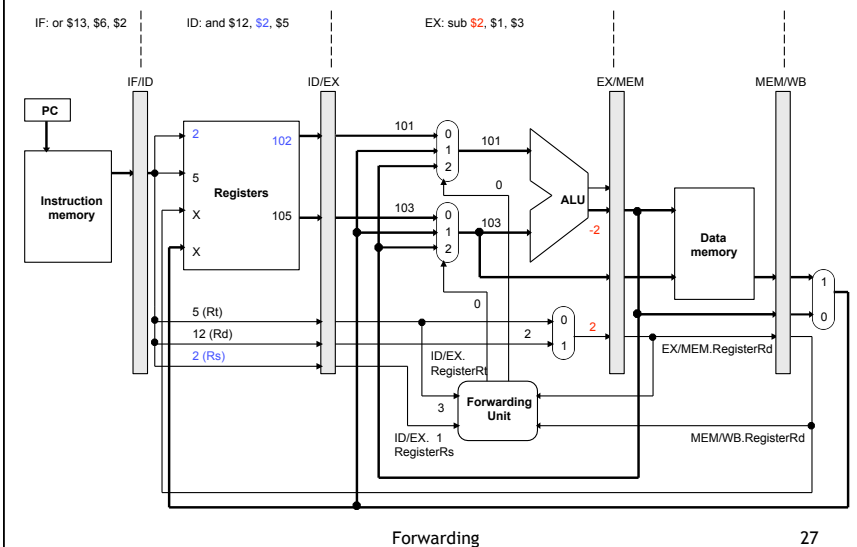
```

- Assume again each register initially contains its number plus 100.
 - After the first instruction, \$2 should contain -2 (101 - 103).
 - The other instructions should all use -2 as one of their operands.
- We'll try to keep the example short.
 - Assume no forwarding is needed except for register \$2.
 - We'll skip the first two cycles, since they're the same as before.

Forwarding

26

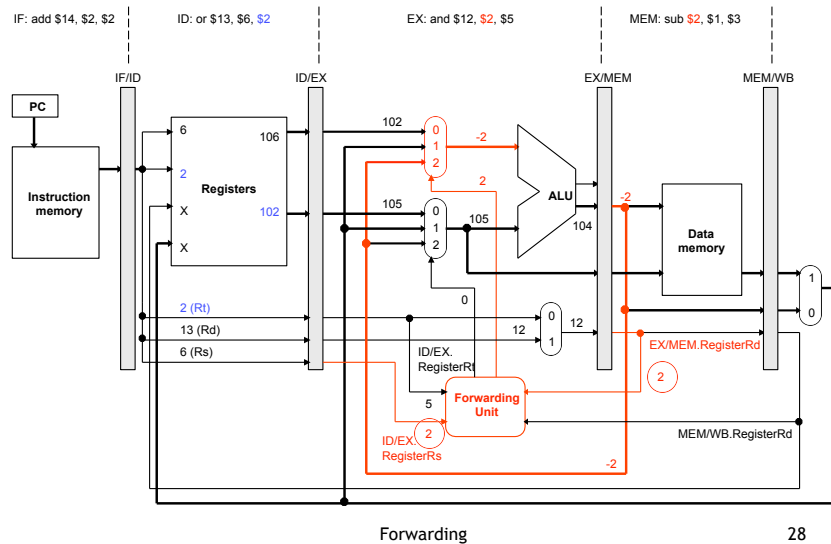
Clock cycle 3



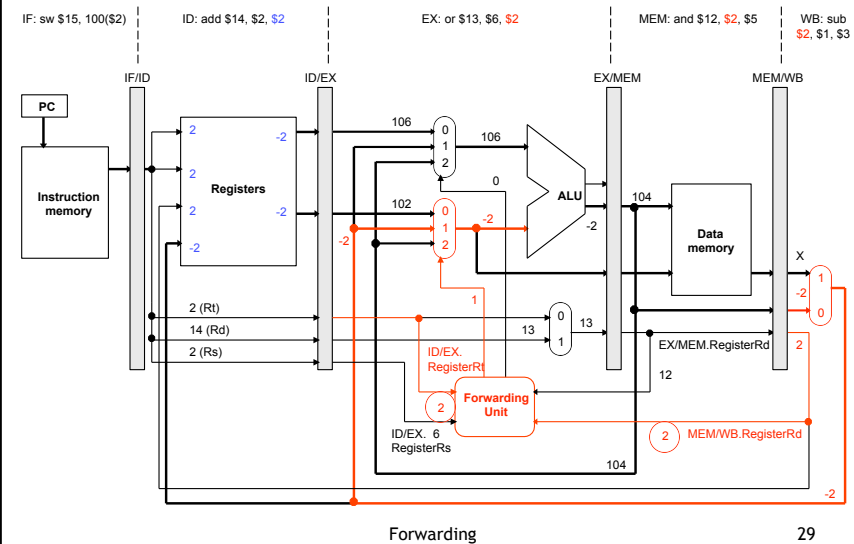
Forwarding

27

Clock cycle 4: forwarding \$2 from EX/MEM



Clock cycle 5: forwarding \$2 from MEM/WB



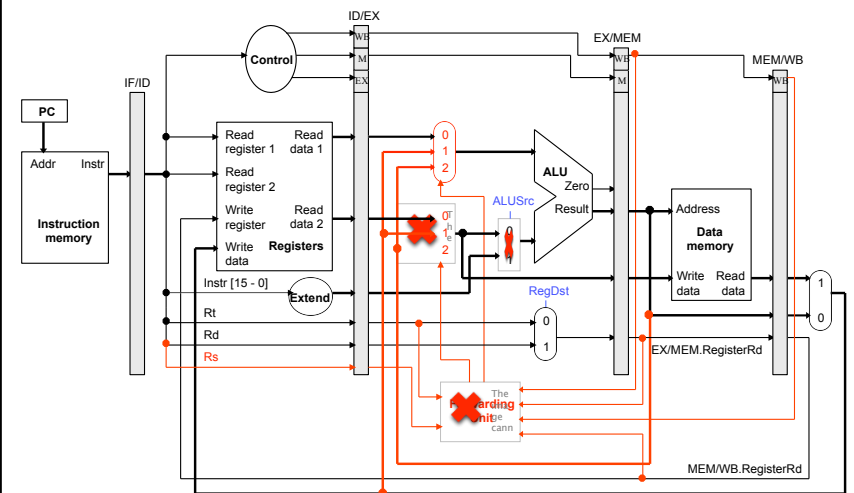
Lots of data hazards

- The first data hazard occurs during cycle 4.
 - The forwarding unit notices that the ALU's first source register for the AND is also the destination of the SUB instruction.
 - The correct value is forwarded from the EX/MEM register, overriding the incorrect old value still in the register file.
- A second hazard occurs during clock cycle 5.
 - The ALU's second source (for OR) is the SUB destination again.
 - This time, the value has to be forwarded from the MEM/WB pipeline register instead.
- There are no other hazards involving the SUB instruction.
 - During cycle 5, SUB writes its result back into register \$2.
 - The ADD instruction can read this new value from the register file in the same cycle.

Forwarding

30

Complete pipelined datapath...so far




Forwarding

31

What about stores?

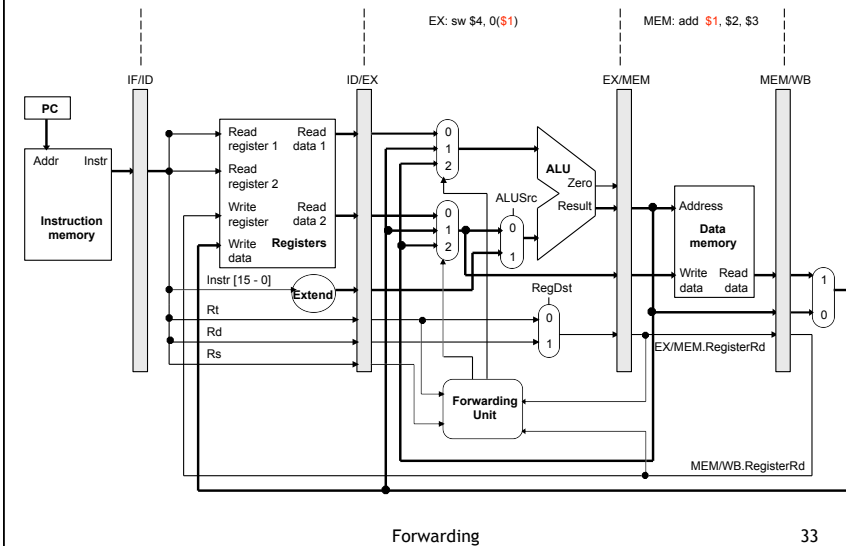
- Two “easy” cases:



Forwarding

32

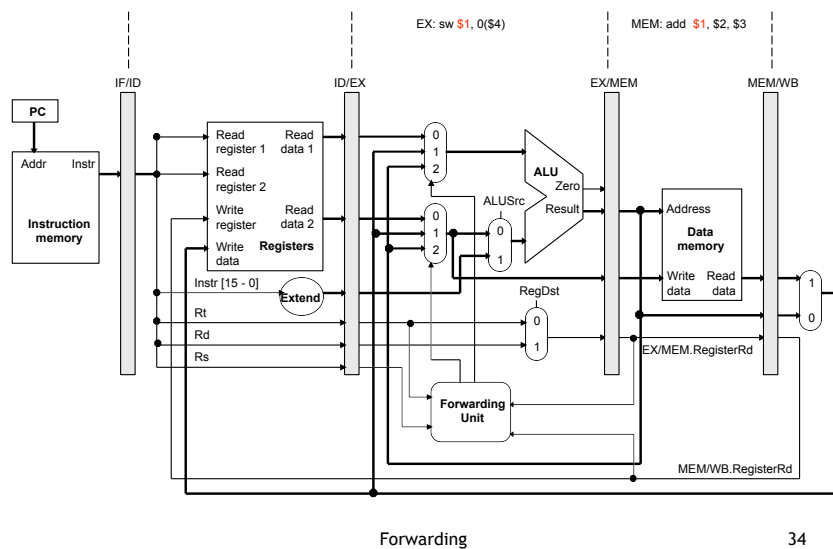
Store Bypassing: Version 1



Forwarding


33

Store Bypassing: Version 2



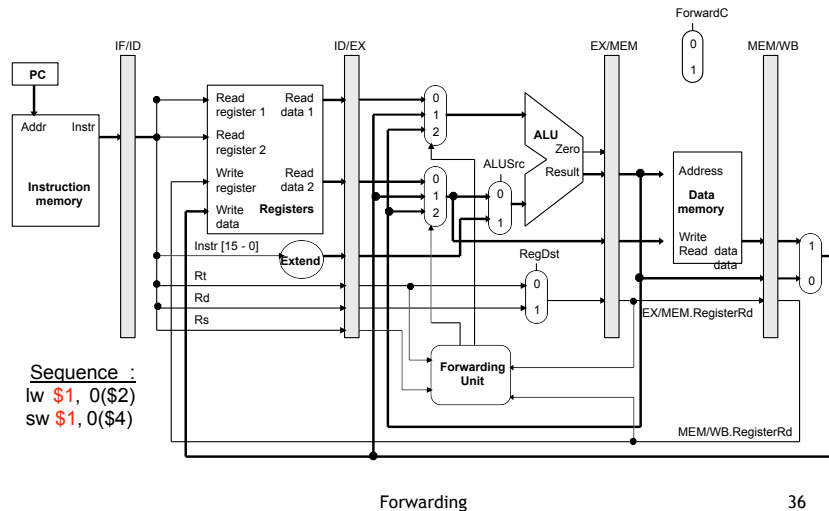
What about stores?

- A harder case:



- In what cycle is:
 - The load value available?
 - The store value needed?
- What do we have to add to the datapath?

Load/Store Bypassing: Extend the Datapath



Miscellaneous comments

- Each MIPS instruction writes to at most one register.
 - This makes the forwarding hardware easier to design, since there is only one destination register that ever needs to be forwarded.
- Forwarding is especially important with deep pipelines like the ones in all current PC processors.
- Section 6.4 of the textbook has some additional material not shown here.
 - Their hazard detection equations also ensure that the source register is not \$0, which can never be modified.
 - There is a more complex example of forwarding, with several cases covered. Take a look at it!

Forwarding

37

Summary

- In real code, most instructions are dependent upon other ones.
 - This can lead to **data hazards** in our original pipelined datapath.
 - Instructions can't write back to the register file soon enough for the next two instructions to read.
- **Forwarding** eliminates data hazards involving arithmetic instructions.
 - The forwarding unit detects hazards by comparing the destination registers of previous instructions to the source registers of the current instruction.
 - Hazards are avoided by grabbing results from the pipeline registers *before* they are written back to the register file.
- Next time we'll finish up pipelining.
 - Forwarding can't save us in some cases involving lw.
 - We still haven't talked about branches for the pipelined datapath.

