

A Tutorial on Data Representation

Integers, Floating-point Numbers, and Characters

1. Number Systems

Human beings use *decimal* (base 10) and *duodecimal* (base 12) number systems for counting and measurements (probably because we have 10 fingers and two big toes). Computers use *binary* (base 2) number system, as they are made from binary digital components (known as transistors) operating in two states - on and off. In computing, we also use *hexadecimal* (base 16) or *octal* (base 8) number systems, as a *compact* form for represent binary numbers.

1.1 Decimal (Base 10) Number System

Decimal number system has ten symbols: 0, 1, 2, 3, 4, 5, 6, 7, 8, and 9, called *digits*. It uses *positional notation*. That is, the least-significant digit (right-most digit) is of the order of 10^0 (units or ones), the second right-most digit is of the order of 10^1 (tens), the third right-most digit is of the order of 10^2 (hundreds), and so on. For example,

$$735 = 7 \times 10^2 + 3 \times 10^1 + 5 \times 10^0$$

We shall denote a decimal number with an optional suffix D if ambiguity arises.

1.2 Binary (Base 2) Number System

Binary number system has two symbols: 0 and 1, called *bits*. It is also a *positional notation*, for example,

$$10110B = 1 \times 2^4 + 0 \times 2^3 + 1 \times 2^2 + 1 \times 2^1 + 0 \times 2^0$$

We shall denote a binary number with a suffix B. Some programming languages denote binary numbers with prefix 0b (e.g., 0b1001000), or prefix b with the bits quoted (e.g., b'10001111').

A binary digit is called a *bit*. Eight bits is called a *byte* (why 8-bit unit? Probably because $8=2^3$).

1.3 Hexadecimal (Base 16) Number System

Hexadecimal number system uses 16 symbols: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, A, B, C, D, E, and F, called *hex digits*. It is a *positional notation*, for example,

$$A3EH = 10 \times 16^2 + 3 \times 16^1 + 14 \times 16^0$$

We shall denote a hexadecimal number (in short, hex) with a suffix H. Some programming languages denote hex numbers with prefix 0x (e.g., 0x1A3C5F), or prefix x with hex digit quoted (e.g., x'C3A4D98B').

Each hexadecimal digit is also called a *hex digit*. Most programming languages accept lowercase 'a' to 'f' as well as uppercase 'A' to 'F'.

Computers uses binary system in their internal operations, as they are built from binary digital electronic components. However, writing or reading a long sequence of binary bits is cumbersome and error-prone. Hexadecimal system is used as a *compact* form or *shorthand* for binary bits. Each hex digit is equivalent to 4 binary bits, i.e., shorthand for 4 bits, as follows:

TABLE OF CONTENTS (HIDE)

1. Number Systems
 - 1.1 Decimal (Base 10) Number
 - 1.2 Binary (Base 2) Number Sys
 - 1.3 Hexadecimal (Base 16) Nun
 - 1.4 Conversion from Hexadecirr
 - 1.5 Conversion from Binary to H
 - 1.6 Conversion from Base r to C
 - 1.7 Conversion from Decimal (B
 - 1.8 General Conversion betwee
 - 1.9 Exercises (Number Systems
2. Computer Memory & Data Repr
3. Integer Representation
 - 3.1 n -bit Unsigned Integers
 - 3.2 Signed Integers
 - 3.3 n -bit Sign Integers in Sign-M
 - 3.4 n -bit Sign Integers in 1's Coi
 - 3.5 n -bit Sign Integers in 2's Coi
 - 3.6 Computers use 2's Complem
 - 3.7 Range of n -bit 2's Complem
 - 3.8 Decoding 2's Complement N
 - 3.9 Big Endian vs. Little Endian
 - 3.10 Exercise (Integer Represer
4. Floating-Point Number Represer
 - 4.1 IEEE-754 32-bit Single-Prec
 - 4.2 Exercises (Floating-point Nu
 - 4.3 IEEE-754 64-bit Double-Pre
 - 4.4 More on Floating-Point Repr
5. Character Encoding
 - 5.1 7-bit ASCII Code (aka US-A
 - 5.2 8-bit Latin-1 (aka ISO/IEC 8i
 - 5.3 Other 8-bit Extension of US-
 - 5.4 Unicode (aka ISO/IEC 1064i
 - 5.5 UTF-8 (Unicode Transforma
 - 5.6 UTF-16 (Unicode Transform
 - 5.7 UTF-32 (Unicode Transform
 - 5.8 Formats of Multi-Byte (e.g.,
 - 5.9 Formats of Text Files
 - 5.10 Windows' CMD Codepage
 - 5.11 Chinese Character Sets
 - 5.12 Collating Sequences (for R
 - 5.13 For Java Programmers - j.
 - 5.14 For Java Programmers - cl
 - 5.15 Displaying Hex Values & H
6. Summary - Why Bother about Di
 - 6.1 Exercises (Data Representa

0H (0000B) (0D)	1H (0001B) (1D)	2H (0010B) (2D)	3H (0011B) (3D)
4H (0100B) (4D)	5H (0101B) (5D)	6H (0110B) (6D)	7H (0111B) (7D)
8H (1000B) (8D)	9H (1001B) (9D)	AH (1010B) (10D)	BH (1011B) (11D)
CH (1100B) (12D)	DH (1101B) (13D)	EH (1110B) (14D)	FH (1111B) (15D)

1.4 Conversion from Hexadecimal to Binary

Replace each hex digit by the 4 equivalent bits, for examples,

```
A3C5H = 1010 0011 1100 0101B
102AH = 0001 0000 0010 1010B
```

1.5 Conversion from Binary to Hexadecimal

Starting from the right-most bit (least-significant bit), replace each group of 4 bits by the equivalent hex digit (pad the left-most bits with zero if necessary), for examples,

```
1001001010B = 0010 0100 1010B = 24AH
10001011001011B = 0010 0010 1100 1011B = 22CBH
```

It is important to note that hexadecimal number provides a *compact form* or *shorthand* for representing binary bits.

1.6 Conversion from Base r to Decimal (Base 10)

Given a n -digit base r number: $d_{n-1} d_{n-2} d_{n-3} \dots d_3 d_2 d_1 d_0$ (base r), the decimal equivalent is given by:

$$d_{n-1} \times r^{(n-1)} + d_{n-2} \times r^{(n-2)} + \dots + d_1 \times r^1 + d_0 \times r^0$$

For examples,

```
A1C2H = 10×16^3 + 1×16^2 + 12×16^1 + 2 = 41410 (base 10)
10110B = 1×2^4 + 1×2^2 + 1×2^1 = 22 (base 10)
```

1.7 Conversion from Decimal (Base 10) to Base r

Use repeated division/remainder. For example,

```
To convert 261D to hexadecimal:
261/16 => quotient=16 remainder=5
16/16 => quotient=1 remainder=0
1/16 => quotient=0 remainder=1 (quotient=0 stop)
Hence, 261D = 105H
```

The above procedure is actually applicable to conversion between any 2 base systems. For example,

```
To convert 1023(base 4) to base 3:
1023(base 4)/3 => quotient=25D remainder=0
25D/3 => quotient=8D remainder=1
8D/3 => quotient=2D remainder=2
2D/3 => quotient=0 remainder=2 (quotient=0 stop)
Hence, 1023(base 4) = 2210(base 3)
```

1.8 General Conversion between 2 Base Systems with Fractional Part

1. Separate the integral and the fractional parts.
2. For the integral part, divide by the target radix repeatedly, and collect the remainder in reverse order.
3. For the fractional part, multiply the fractional part by the target radix repeatedly, and collect the integral part in the same order.

Example 1:

```
Convert 18.6875D to binary
Integral Part = 18D
18/2 => quotient=9 remainder=0
9/2 => quotient=4 remainder=1
4/2 => quotient=2 remainder=0
2/2 => quotient=1 remainder=0
```

```

1/2 => quotient=0 remainder=1 (quotient=0 stop)
Hence, 18D = 10010B
Fractional Part = .6875D
.6875*2=1.375 => whole number is 1
.375*2=0.75  => whole number is 0
.75*2=1.5    => whole number is 1
.5*2=1.0     => whole number is 1
Hence .6875D = .1011B
Therefore, 18.6875D = 10010.1011B

```

Example 2:

```

Convert 18.6875D to hexadecimal
Integral Part = 18D
18/16 => quotient=1 remainder=2
1/16  => quotient=0 remainder=1 (quotient=0 stop)
Hence, 18D = 12H
Fractional Part = .6875D
.6875*16=11.0 => whole number is 11D (BH)
Hence .6875D = .BH
Therefore, 18.6875D = 12.BH

```

1.9 Exercises (Number Systems Conversion)

1. Convert the following *decimal* numbers into *binary* and *hexadecimal* numbers:

- a. 108
- b. 4848
- c. 9000

2. Convert the following binary numbers into hexadecimal and decimal numbers:

- a. 1000011000
- b. 10000000
- c. 1010101010

3. Convert the following hexadecimal numbers into binary and decimal numbers:

- a. ABCDE
- b. 1234
- c. 80F

4. Convert the following decimal numbers into binary equivalent:

- a. 19.25D
- b. 123.456D

Answers: You could use the Windows' Calculator (calc.exe) to carry out number system conversion, by setting it to the scientific mode. (Run "calc" ⇒ Select "View" menu ⇒ Choose "Programmer" or "Scientific" mode.)

- 1. 1101100B, 1001011110000B, 10001100101000B, 6CH, 12F0H, 2328H.
- 2. 218H, 80H, AAAH, 536D, 128D, 2730D.
- 3. 10101011110011011110B, 1001000110100B, 100000001111B, 703710D, 4660D, 2063D.
- 4. ??

2. Computer Memory & Data Representation

Computer uses a *fixed number of bits* to represent a piece of data, which could be a number, a character, or others. A n -bit storage location can represent up to 2^n distinct entities. For example, a 3-bit memory location can hold one of these eight binary patterns: 000, 001, 010, 011, 100, 101, 110, or 111. Hence, it can represent at most 8 distinct entities. You could use them to represent numbers 0 to 7, numbers 8881 to 8888, characters 'A' to 'H', or up to 8 kinds of fruits like apple, orange, banana; or up to 8 kinds of animals like lion, tiger, etc.

Integers, for example, can be represented in 8-bit, 16-bit, 32-bit or 64-bit. You, as the programmer, choose an appropriate bit-length for your integers. Your choice will impose constraint on the range of integers that can be represented. Besides the bit-length, an integer can be represented in various *representation* schemes, e.g., unsigned vs. signed integers. An 8-bit unsigned integer has a range of 0 to 255, while an 8-bit signed integer has a range of -128 to 127 - both representing 256 distinct numbers.

It is important to note that a computer memory location merely *stores a binary pattern*. It is entirely up to you, as the programmer, to decide on how these patterns are to be *interpreted*. For example, the 8-bit binary pattern "0100 0001B" can be interpreted as an unsigned integer 65, or an ASCII character 'A', or some secret information known only to you. In other words, you have to first decide how to represent a piece of data in a binary pattern before the binary patterns make sense. The interpretation of binary pattern is

called *data representation* or *encoding*. Furthermore, it is important that the data representation schemes are agreed-upon by all the parties, i.e., industrial standards need to be formulated and straightly followed.

Once you decided on the data representation scheme, certain constraints, in particular, the precision and range will be imposed. Hence, it is important to understand *data representation* to write *correct* and *high-performance* programs.

Rosette Stone and the Decipherment of Egyptian Hieroglyphs

Egyptian hieroglyphs (next-to-left) were used by the ancient Egyptians since 4000BC. Unfortunately, since 500AD, no one could longer read the ancient Egyptian hieroglyphs, until the re-discovery of the Rosette Stone in 1799 by Napoleon's troop (during Napoleon's Egyptian invasion) near the town of Rashid (Rosetta) in the Nile Delta.

The Rosetta Stone (left) is inscribed with a decree in 196BC on behalf of King Ptolemy V. The decree appears in *three* scripts: the upper text is *Ancient Egyptian hieroglyphs*, the middle portion Demotic script, and the lowest *Ancient Greek*. Because it presents essentially the same text in all three scripts, and Ancient Greek could still be understood, it provided the key to the decipherment of the Egyptian hieroglyphs.



The moral of the story is unless you know the encoding scheme, there is no way that you can decode the data.

Reference and images: Wikipedia.

3. Integer Representation

Integers are *whole numbers* or *fixed-point numbers* with the radix point *fixed* after the least-significant bit. They are contrast to *real numbers* or *floating-point numbers*, where the position of the radix point varies. It is important to take note that integers and floating-point numbers are treated differently in computers. They have different representation and are processed differently (e.g., floating-point numbers are processed in a so-called floating-point processor). Floating-point numbers will be discussed later.

Computers use a *fixed number of bits* to represent an integer. The commonly-used bit-lengths for integers are 8-bit, 16-bit, 32-bit or 64-bit. Besides bit-lengths, there are two representation schemes for integers:

1. *Unsigned Integers*: can represent zero and positive integers.
2. *Signed Integers*: can represent zero, positive and negative integers. Three representation schemes had been proposed for signed integers:
 - a. Sign-Magnitude representation
 - b. 1's Complement representation
 - c. 2's Complement representation

You, as the programmer, need to decide on the bit-length and representation scheme for your integers, depending on your application's requirements. Suppose that you need a counter for counting a small quantity from 0 up to 200, you might choose the 8-bit unsigned integer scheme as there is no negative numbers involved.

3.1 *n*-bit Unsigned Integers

Unsigned integers can represent zero and positive integers, but not negative integers. The value of an unsigned integer is interpreted as "*the magnitude of its underlying binary pattern*".

Example 1: Suppose that $n=8$ and the binary pattern is 0100 0001B, the value of this unsigned integer is $1 \times 2^0 + 1 \times 2^6 = 65D$.

Example 2: Suppose that $n=16$ and the binary pattern is 0001 0000 0000 1000B, the value of this unsigned integer is $1 \times 2^3 + 1 \times 2^{12} = 4104D$.

Example 3: Suppose that $n=16$ and the binary pattern is 0000 0000 0000 0000B, the value of this unsigned integer is 0.

An n -bit pattern can represent 2^n distinct integers. An n -bit unsigned integer can represent integers from 0 to $(2^n) - 1$, as tabulated below:

n	Minimum	Maximum
8	0	$(2^8) - 1$ (=255)
16	0	$(2^{16}) - 1$ (=65,535)
32	0	$(2^{32}) - 1$ (=4,294,967,295) (9+ digits)

```
64      0      (2^64)-1 (=18,446,744,073,709,551,615) (19+
           digits)
```

3.2 Signed Integers

Signed integers can represent zero, positive integers, as well as negative integers. Three representation schemes are available for signed integers:

1. Sign-Magnitude representation
2. 1's Complement representation
3. 2's Complement representation

In all the above three schemes, the *most-significant bit* (msb) is called the *sign bit*. The sign bit is used to represent the *sign* of the integer - with 0 for positive integers and 1 for negative integers. The *magnitude* of the integer, however, is interpreted differently in different schemes.

3.3 n -bit Sign Integers in Sign-Magnitude Representation

In sign-magnitude representation:

- The most-significant bit (msb) is the *sign bit*, with value of 0 representing positive integer and 1 representing negative integer.
- The remaining $n-1$ bits represents the magnitude (absolute value) of the integer. The absolute value of the integer is interpreted as "the magnitude of the $(n-1)$ -bit binary pattern".

Example 1: Suppose that $n=8$ and the binary representation is 0 100 0001B.

Sign bit is 0 \Rightarrow positive

Absolute value is $100\ 0001B = 65D$

Hence, the integer is +65D

Example 2: Suppose that $n=8$ and the binary representation is 1 000 0001B.

Sign bit is 1 \Rightarrow negative

Absolute value is 000 0001B = 1D

Hence, the integer is -1D

Example 3: Suppose that $n=8$ and the binary representation is 0 000 0000B.

Sign bit is 0 \Rightarrow positive

Absolute value is 000 0000B = 0D

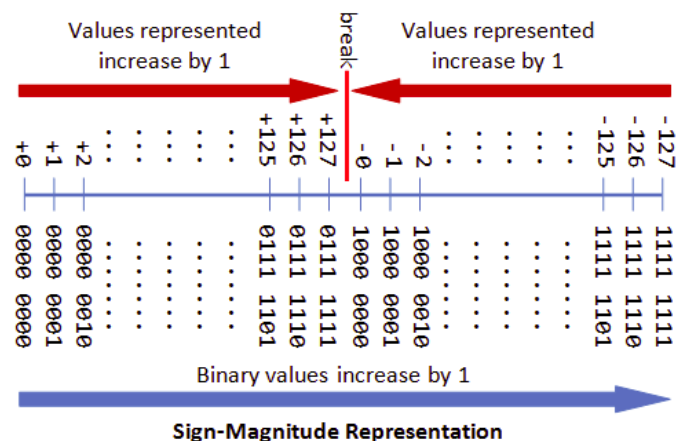
Hence, the integer is +0D

Example 4: Suppose that $n=8$ and the binary representation is 1 000 0000B.

Sign bit is 1 \Rightarrow negative

Absolute value is 000 0000B = 0D

Hence, the integer is $-0D$



The drawbacks of sign-magnitude representation are:

1. There are two representations (0000 0000B and 1000 0000B) for the number zero, which could lead to inefficiency and confusion.
2. Positive and negative integers need to be processed separately.

3.4 n -bit Sign Integers in 1's Complement Representation

In 1's complement representation:

- Again, the most significant bit (msb) is the *sign bit*, with value of 0 representing positive integers and 1 representing negative integers.
- The remaining $n-1$ bits represents the magnitude of the integer, as follows:
 - for positive integers, the absolute value of the integer is equal to "the magnitude of the $(n-1)$ -bit binary pattern".
 - for negative integers, the absolute value of the integer is equal to "the magnitude of the *complement (inverse)* of the $(n-1)$ -bit binary pattern" (hence called 1's complement).

Example 1: Suppose that $n=8$ and the binary representation 0 100 0001B.

Sign bit is 0 = positive

Absolute value is 100 0001B = 65D

Hence, the integer is +65D

Example 2: Suppose that $n=8$ and the binary representation 1 000 0001B.

Sign bit is 1 = negative

Absolute value is the complement of 000 0001B, i.e., 111 1110B = 126D

Hence, the integer is -126D

Example 3: Suppose that $n=8$ and the binary representation 0 000 0000B.

Sign bit is 0 = positive

Absolute value is 000 0000B = 0D

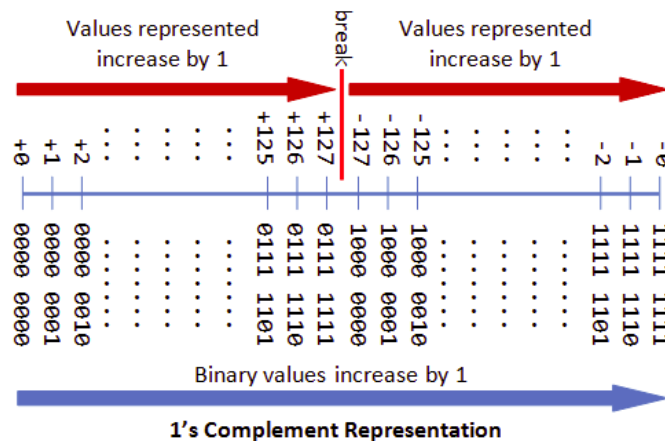
Hence, the integer is +0D

Example 4: Suppose that $n=8$ and the binary representation 1 111 1111B.

Sign bit is 1 = negative

Absolute value is the complement of 111 1111B, i.e., 000 0000B = 0D

Hence, the integer is -0D



Again, the drawbacks are:

1. There are two representations (0000 0000B and 1111 1111B) for zero.
2. The positive integers and negative integers need to be processed separately.

3.5 n -bit Sign Integers in 2's Complement Representation

In 2's complement representation:

- Again, the most significant bit (msb) is the *sign bit*, with value of 0 representing positive integers and 1 representing negative integers.
- The remaining $n-1$ bits represents the magnitude of the integer, as follows:
 - for positive integers, the absolute value of the integer is equal to "the magnitude of the $(n-1)$ -bit binary pattern".
 - for negative integers, the absolute value of the integer is equal to "the magnitude of the *complement* of the $(n-1)$ -bit binary pattern *plus one*" (hence called 2's complement).

Example 1: Suppose that $n=8$ and the binary representation 0 100 0001B.

Sign bit is 0 = positive

Absolute value is 100 0001B = 65D

Hence, the integer is +65D

Example 2: Suppose that $n=8$ and the binary representation 1 000 0001B.

Sign bit is 1 = negative

Absolute value is the complement of 000 0001B plus 1, i.e., 111 1110B + 1B = 127D

Hence, the integer is -127D

Example 3: Suppose that $n=8$ and the binary representation 0 000 0000B.

Sign bit is 0 = positive

Absolute value is 000 0000B = 0D

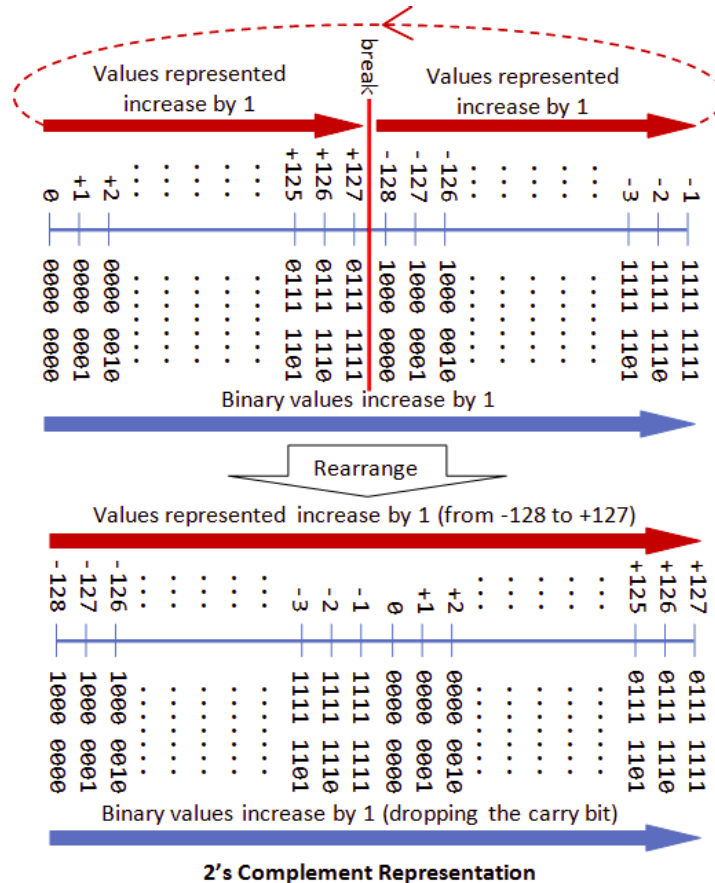
Hence, the integer is +0D

Example 4: Suppose that $n=8$ and the binary representation 1 111 1111B.

Sign bit is 1 = negative

Absolute value is the complement of 111 1111B plus 1, i.e., 000 0000B + 1B = 1D

Hence, the integer is -1D



3.6 Computers use 2's Complement Representation for Signed Integers

We have discussed three representations for signed integers: signed-magnitude, 1's complement and 2's complement. Computers use 2's complement in representing signed integers. This is because:

1. There is only one representation for the number zero in 2's complement, instead of two representations in sign-magnitude and 1's complement.
2. Positive and negative integers can be treated together in addition and subtraction. Subtraction can be carried out using the "addition logic".

Example 1: Addition of Two Positive Integers: Suppose that $n=8$, $65D + 5D = 70D$

```
65D → 0100 0001B
5D  → 0000 0101B(+)
      0100 0110B → 70D (OK)
```

Example 2: Subtraction is treated as Addition of a Positive and a Negative Integers: Suppose that $n=8$, $5D - 5D = 65D + (-5D) = 60D$

```
65D → 0100 0001B
-5D → 1111 1011B(+)
      0011 1100B → 60D (discard carry - OK)
```

Example 3: Addition of Two Negative Integers: Suppose that $n=8$, $-65D - 5D = (-65D) + (-5D) = -70D$

```
-65D → 1011 1111B
-5D  → 1111 1011B(+)
      1011 1010B → -70D (discard carry - OK)
```


Because of the *fixed precision* (i.e., *fixed number of bits*), an n -bit 2's complement signed integer has a certain range. For example, for $n=8$, the range of 2's complement signed integers is -128 to $+127$. During addition (and subtraction), it is important to check whether the result exceeds this range, in other words, whether *overflow* or *underflow* has occurred.

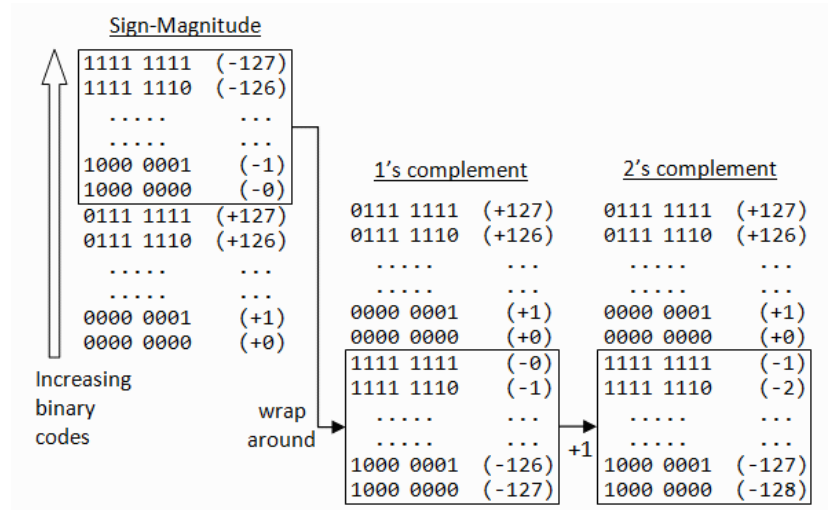
Example 4: Overflow: Suppose that $n=8$, $127D + 2D = 129D$ (overflow - beyond the range)

```
127D → 0111 1111B
 2D → 0000 0010B(+)
      1000 0001B → -127D (wrong)
```

Example 5: Underflow: Suppose that $n=8$, $-125D - 5D = -130D$ (underflow - below the range)

```
-125D → 1000 0011B
 -5D → 1111 1011B(+)
       0111 1110B → +126D (wrong)
```

The following diagram explains how the 2's complement works. By re-arranging the number line, values from -128 to $+127$ are represented contiguously by ignoring the carry bit.



3.7 Range of n -bit 2's Complement Signed Integers

An n -bit 2's complement signed integer can represent integers from $-2^{(n-1)}$ to $+2^{(n-1)}-1$, as tabulated. Take note that the scheme can represent all the integers within the range, without any gap. In other words, there is no missing integers within the supported range.

n	minimum	maximum
8	$-(2^7) \quad (= -128)$	$+(2^7)-1 \quad (= +127)$
16	$-(2^{15}) \quad (= -32,768)$	$+(2^{15})-1 \quad (= +32,767)$
32	$-(2^{31}) \quad (= -2,147,483,648)$	$+(2^{31})-1 \quad (= +2,147,483,647) \text{ (9+ digits)}$
64	$-(2^{63}) \quad (= -9,223,372,036,854,775,808)$	$+(2^{63})-1 \quad (= +9,223,372,036,854,775,807) \text{ (18+ digits)}$

3.8 Decoding 2's Complement Numbers

1. Check the *sign bit* (denoted as S).
2. If $S=0$, the number is positive and its absolute value is the binary value of the remaining $n-1$ bits.
3. If $S=1$, the number is negative. you could "invert the $n-1$ bits and plus 1" to get the absolute value of negative number. Alternatively, you could scan the remaining $n-1$ bits from the right (least-significant bit). Look for the first occurrence of 1. Flip all the bits to the left of that first occurrence of 1. The flipped pattern gives the absolute value. For example,

```
n = 8, bit pattern = 1 100 0100B
S = 1 → negative
Scanning from the right and flip all the bits to the left of the first occurrence of 1 → 011 1100B = 60D
Hence, the value is -60D
```

3.9 Big Endian vs. Little Endian

Modern computers store one byte of data in each memory address or location, i.e., byte addressable memory. An 32-bit integer is,

therefore, stored in 4 memory addresses.

The term "Endian" refers to the *order* of storing bytes in computer memory. In "Big Endian" scheme, the most significant byte is stored first in the lowest memory address (or big in first), while "Little Endian" stores the least significant bytes in the lowest memory address.

For example, the 32-bit integer 12345678H (2215053170₁₀) is stored as 12H 34H 56H 78H in big endian; and 78H 56H 34H 12H in little endian. An 16-bit integer 00H 01H is interpreted as 0001H in big endian, and 0100H as little endian.

3.10 Exercise (Integer Representation)

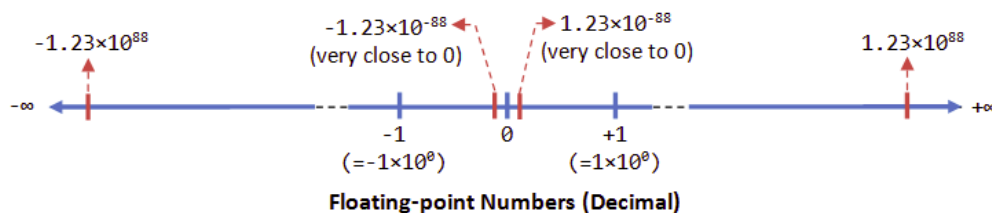
1. What are the ranges of 8-bit, 16-bit, 32-bit and 64-bit integer, in "unsigned" and "signed" representation?
2. Give the value of 88, 0, 1, 127, and 255 in 8-bit unsigned representation.
3. Give the value of +88, -88, -1, 0, +1, -128, and +127 in 8-bit 2's complement signed representation.
4. Give the value of +88, -88, -1, 0, +1, -127, and +127 in 8-bit sign-magnitude representation.
5. Give the value of +88, -88, -1, 0, +1, -127 and +127 in 8-bit 1's complement representation.
6. [TODO] more.

Answers

1. The range of unsigned n -bit integers is $[0, 2^n - 1]$. The range of n -bit 2's complement signed integer is $[-2^{(n-1)}, +2^{(n-1)} - 1]$;
2. 88 (0101 1000), 0 (0000 0000), 1 (0000 0001), 127 (0111 1111), 255 (1111 1111).
3. +88 (0101 1000), -88 (1010 1000), -1 (1111 1111), 0 (0000 0000), +1 (0000 0001), -128 (1000 0000), +127 (0111 1111).
4. +88 (0101 1000), -88 (1101 1000), -1 (1000 0001), 0 (0000 0000 or 1000 0000), +1 (0000 0001), -127 (1111 1111), +127 (0111 1111).
5. +88 (0101 1000), -88 (1010 0111), -1 (1111 1110), 0 (0000 0000 or 1111 1111), +1 (0000 0001), -127 (1000 0000), +127 (0111 1111).

4. Floating-Point Number Representation

A floating-point number (or real number) can represent a very large (1.23×10^{88}) or a very small (1.23×10^{-88}) value. It could also represent very large negative number (-1.23×10^{88}) and very small negative number (-1.23×10^{-88}), as well as zero, as illustrated:



A floating-point number is typically expressed in the scientific notation, with a *fraction* (F), and an *exponent* (E) of a certain *radix* (r), in the form of $F \times r^E$. Decimal numbers use radix of 10 ($F \times 10^E$); while binary numbers use radix of 2 ($F \times 2^E$).

Representation of floating point number is not unique. For example, the number 55.66 can be represented as 5.566×10^1 , 0.5566×10^2 , 0.05566×10^3 , and so on. The fractional part can be *normalized*. In the normalized form, there is only a single non-zero digit before the radix point. For example, decimal number 123.4567 can be normalized as 1.234567×10^2 ; binary number 1010.1011B can be normalized as $1.0101011B \times 2^3$.

It is important to note that floating-point numbers suffer from *loss of precision* when represented with a fixed number of bits (e.g., 32-bit or 64-bit). This is because there are *infinite* number of real numbers (even within a small range of says 0.0 to 0.1). On the other hand, a n -bit binary pattern can represent a *finite* 2^n distinct numbers. Hence, not all the real numbers can be represented. The nearest approximation will be used instead, resulted in loss of accuracy.

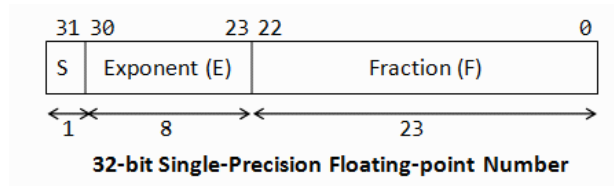
It is also important to note that floating number arithmetic is very much less efficient than integer arithmetic. It could be speed up with a so-called dedicated *floating-point co-processor*. Hence, use integers if your application does not require floating-point numbers.

In computers, floating-point numbers are represented in scientific notation of *fraction* (F) and *exponent* (E) with a *radix* of 2, in the form of $F \times 2^E$. Both E and F can be positive as well as negative. Modern computers adopt IEEE 754 standard for representing floating-point numbers. There are two representation schemes: 32-bit single-precision and 64-bit double-precision.

4.1 IEEE-754 32-bit Single-Precision Floating-Point Numbers

In 32-bit single-precision floating-point representation:

- The most significant bit is the *sign bit* (S), with 0 for positive numbers and 1 for negative numbers.
- The following 8 bits represent *exponent* (E).
- The remaining 23 bits represents *fraction* (F).



Normalized Form

Let's illustrate with an example, suppose that the 32-bit pattern is 1 1000 0001 011 0000 0000 0000 0000, with:

- $S = 1$
- $E = 1000\ 0001$
- $F = 011\ 0000\ 0000\ 0000\ 0000\ 0000$

In the *normalized form*, the actual fraction is normalized with an implicit leading 1 in the form of 1.F. In this example, the actual fraction is $1.011\ 0000\ 0000\ 0000\ 0000\ 0000 = 1 + 1 \times 2^{-2} + 1 \times 2^{-3} = 1.375D$.

The sign bit represents the sign of the number, with $S=0$ for positive and $S=1$ for negative number. In this example with $S=1$, this is a negative number, i.e., $-1.375D$.

In normalized form, the actual exponent is $E - 127$ (so-called excess-127 or bias-127). This is because we need to represent both positive and negative exponent. With an 8-bit E, ranging from 0 to 255, the excess-127 scheme could provide actual exponent of -127 to 128. In this example, $E - 127 = 129 - 127 = 2D$.

Hence, the number represented is $-1.375 \times 2^2 = -5.5D$.

De-Normalized Form

Normalized form has a serious problem, with an implicit leading 1 for the fraction, it cannot represent the number zero! Convince yourself on this!

De-normalized form was devised to represent zero and other numbers.

For $E=0$, the numbers are in the de-normalized form. An implicit leading 0 (instead of 1) is used for the fraction; and the actual exponent is always -126. Hence, the number zero can be represented with $E=0$ and $F=0$ (because $0.0 \times 2^{-126} = 0$).

We can also represent very small positive and negative numbers in de-normalized form with $E=0$. For example, if $S=1$, $E=0$, and $F=011\ 0000\ 0000\ 0000\ 0000\ 0000$. The actual fraction is $0.011 = 1 \times 2^{-2} + 1 \times 2^{-3} = 0.375D$. Since $S=1$, it is a negative number. With $E=0$, the actual exponent is -126. Hence the number is $-0.375 \times 2^{-126} = -4.4 \times 10^{-39}$, which is an extremely small negative number (close to zero).

Summary

In summary, the value (N) is calculated as follows:

- For $1 \leq E \leq 254$, $N = (-1)^S \times 1.F \times 2^{(E-127)}$. These numbers are in the so-called *normalized* form. The sign-bit represents the sign of the number. Fractional part (1.F) are normalized with an implicit leading 1. The exponent is bias (or in excess) of 127, so as to represent both positive and negative exponent. The range of exponent is -126 to +127.
- For $E = 0$, $N = (-1)^S \times 0.F \times 2^{(-126)}$. These numbers are in the so-called *denormalized* form. The exponent of 2^{-126} evaluates to a very small number. Denormalized form is needed to represent zero (with $F=0$ and $E=0$). It can also represents very small positive and negative number close to zero.
- For $E = 255$, it represents special values, such as $\pm INF$ (positive and negative infinity) and NaN (not a number). This is beyond the scope of this article.

Example 1: Suppose that IEEE-754 32-bit floating-point representation pattern is 0 10000000 110 0000 0000 0000 0000 0000.

Sign bit $S = 0 \Rightarrow$ positive number
 $E = 1000\ 0000B = 128D$ (in normalized form)
 Fraction is $1.11B$ (with an implicit leading 1) $= 1 + 1 \times 2^{-1} + 1 \times 2^{-2} = 1.75D$
 The number is $+1.75 \times 2^{(128-127)} = +3.5D$

Example 2: Suppose that IEEE-754 32-bit floating-point representation pattern is 1 01111110 100 0000 0000 0000 0000 0000.

Sign bit $S = 1 \Rightarrow$ negative number

$E = 0111\ 1110_B = 126_D$ (in normalized form)
 Fraction is 1.1_B (with an implicit leading 1) $= 1 + 2^{-1} = 1.5_D$
 The number is $-1.5 \times 2^{(126-127)} = -0.75_D$

Example 3: Suppose that IEEE-754 32-bit floating-point representation pattern is 1 01111110 000 0000 0000 0000 0000 0001.

Sign bit $S = 1 \Rightarrow$ negative number
 $E = 0111\ 1110_B = 126_D$ (in normalized form)
 Fraction is $1.000\ 0000\ 0000\ 0000\ 0000\ 0001_B$ (with an implicit leading 1) $= 1 + 2^{-23}$
 The number is $-(1 + 2^{-23}) \times 2^{(126-127)} = -0.500000059604644775390625$ (may not be exact in decimal!)

Example 4 (De-Normalized Form): Suppose that IEEE-754 32-bit floating-point representation pattern is 1 00000000 000 0000 0000 0000 0001.

Sign bit $S = 1 \Rightarrow$ negative number
 $E = 0$ (in de-normalized form)
 Fraction is $0.000\ 0000\ 0000\ 0000\ 0000\ 0001_B$ (with an implicit leading 0) $= 1 \times 2^{-23}$
 The number is $-2^{-23} \times 2^{(-126)} = -2 \times (-149) \approx -1.4 \times 10^{-45}$

4.2 Exercises (Floating-point Numbers)

1. Compute the largest and smallest positive numbers that can be represented in the 32-bit normalized form.
2. Compute the largest and smallest negative numbers can be represented in the 32-bit normalized form.
3. Repeat (1) for the 32-bit denormalized form.
4. Repeat (2) for the 32-bit denormalized form.

Hints:

1. Largest positive number: $S=0, E=1111\ 1110\ (254), F=111\ 1111\ 1111\ 1111\ 1111\ 1111$.
 Smallest positive number: $S=0, E=0000\ 0001\ (1), F=000\ 0000\ 0000\ 0000\ 0000\ 0000$.
2. Same as above, but $S=1$.
3. Largest positive number: $S=0, E=0, F=111\ 1111\ 1111\ 1111\ 1111\ 1111$.
 Smallest positive number: $S=0, E=0, F=000\ 0000\ 0000\ 0000\ 0000\ 0001$.
4. Same as above, but $S=1$.

Notes For Java Users

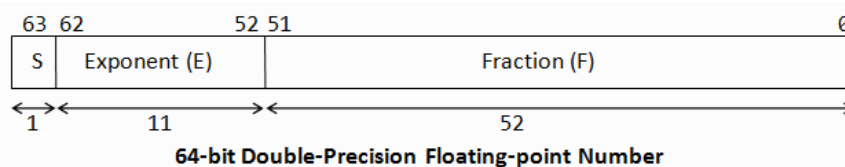
You can use JDK methods `Float.intBitsToFloat(int bits)` or `Double.longBitsToDouble(long bits)` to create a single-precision 32-bit float or double-precision 64-bit double with the specific bit patterns, and print their values. For examples,

```
System.out.println(Float.intBitsToFloat(0x7fffff));  
System.out.println(Double.longBitsToDouble(0x1fffffffffffffL));
```

4.3 IEEE-754 64-bit Double-Precision Floating-Point Numbers

The representation scheme for 64-bit double-precision is similar to the 32-bit single-precision:

- The most significant bit is the *sign bit* (S), with 0 for positive numbers and 1 for negative numbers.
- The following 11 bits represent *exponent* (E).
- The remaining 52 bits represents *fraction* (F).



The value (N) is calculated as follows:

- Normalized form: For $1 \leq E \leq 2046$, $N = (-1)^S \times 1.F \times 2^{(E-1023)}$.
- Denormalized form: For $E = 0$, $N = (-1)^S \times 0.F \times 2^{(-1022)}$. These are in the denormalized form.
- For $E = 2047$, N represents special values, such as $\pm INF$ (infinity), NaN (not a number).

4.4 More on Floating-Point Representation

There are three parts in the floating-point representation:

- The *sign bit* (S) is self-explanatory (0 for positive numbers and 1 for negative numbers).
- For the *exponent* (E), a so-called *bias* (or *excess*) is applied so as to represent both positive and negative exponent. The bias is set at half of the range. For single precision with an 8-bit exponent, the bias is 127 (or excess-127). For double precision with a 11-bit exponent, the bias is 1023 (or excess-1023).
- The *fraction* (F) (also called the *mantissa* or *significand*) is composed of an implicit leading bit (before the radix point) and the fractional bits (after the radix point). The leading bit for normalized numbers is 1; while the leading bit for denormalized numbers is 0.

Normalized Floating-Point Numbers

In normalized form, the radix point is placed after the first non-zero digit, e.g., $9.8765D \times 10^{-23D}$, $1.001011B \times 2^{11B}$. For binary number, the leading bit is always 1, and need not be represented explicitly - this saves 1 bit of storage.

In IEEE 754's normalized form:

- For single-precision, $1 \leq E \leq 254$ with excess of 127. Hence, the actual exponent is from -126 to +127. Negative exponents are used to represent small numbers (< 1.0); while positive exponents are used to represent large numbers (> 1.0).

$$N = (-1)^S \times 1.F \times 2^{(E-127)}$$

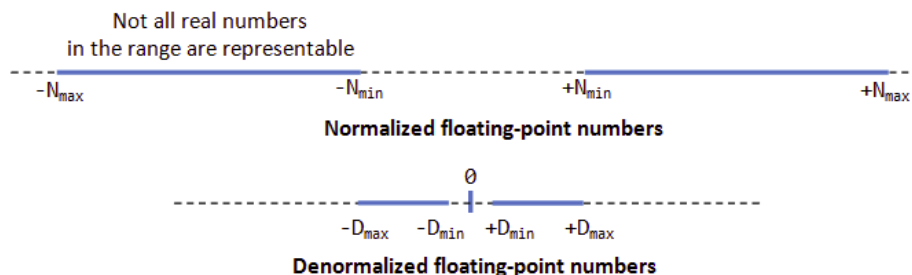
- For double-precision, $1 \leq E \leq 2046$ with excess of 1023. The actual exponent is from -1022 to +1023, and

$$N = (-1)^S \times 1.F \times 2^{(E-1023)}$$

Take note that n-bit pattern has a *finite* number of combinations ($=2^n$), which could represent *finite* distinct numbers. It is not possible to represent the *infinite* numbers in the real axis (even a small range says 0.0 to 1.0 has infinite numbers). That is, not all floating-point numbers can be accurately represented. Instead, the closest approximation is used, which leads to *loss of accuracy*.

The *minimum* and *maximum* normalized floating-point numbers are:

Precision	Normalized N(min)	Normalized N(max)
Single	0080 0000H 0 00000001 000000000000000000000000B E = 1, F = 0 N(min) = $1.0B \times 2^{-126}$ ($\approx 1.17549435 \times 10^{-38}$)	7F7F FFFFH 0 11111110 000000000000000000000000B E = 254, F = 0 N(max) = $1.1...1B \times 2^{127} = (2 - 2^{-23}) \times 2^{127}$ ($\approx 3.4028235 \times 10^{38}$)
Double	0010 0000 0000 0000H N(min) = $1.0B \times 2^{-1022}$ ($\approx 2.2250738585072014 \times 10^{-308}$)	7FEF FFFF FFFF FFFFH N(max) = $1.1...1B \times 2^{1023} = (2 - 2^{-52}) \times 2^{1023}$ ($\approx 1.7976931348623157 \times 10^{308}$)



Denormalized Floating-Point Numbers

If $E = 0$, but the fraction is non-zero, then the value is in denormalized form, and a leading bit of 0 is assumed, as follows:

- For single-precision, $E = 0$,
 $N = (-1)^S \times 0.F \times 2^{(-126)}$
- For double-precision, $E = 0$,
 $N = (-1)^S \times 0.F \times 2^{(-1022)}$

Denormalized form can represent very small numbers closed to zero, and zero, which cannot be represented in normalized form, as shown in the above figure.

The minimum and maximum of *denormalized floating-point numbers* are:

Precision	Denormalized D(min)	Denormalized D(max)
Single	0000 0001H 0 00000000 000000000000000000000001B E = 0, F = 0000000000000000000000001B	007F FFFFH 0 00000000 11111111111111111111111111111111B

	D(min) = 0.0...1 × 2 ⁻¹²⁶ = 1 × 2 ⁻²³ × 2 ⁻¹²⁶ = 2 ⁻¹⁴⁹ (≈ 1.4 × 10 ⁻⁴⁵)	E = 0, F = 11111111111111111111111111111B D(max) = 0.1...1 × 2 ⁻¹²⁶ = (1-2 ⁻²³)×2 ⁻¹²⁶ (≈ 1.1754942 × 10 ⁻³⁸)
Double	0000 0000 0000 0001H D(min) = 0.0...1 × 2 ⁻¹⁰²² = 1 × 2 ⁻⁵² × 2 ⁻¹⁰²² = 2 ⁻¹⁰⁷⁴ (≈ 4.9 × 10 ⁻³²⁴)	001F FFFF FFFF FFFFFH D(max) = 0.1...1 × 2 ⁻¹⁰²² = (1-2 ⁻⁵²)×2 ⁻¹⁰²² (≈ 4.4501477170144023 × 10 ⁻³⁰⁸)

Special Values

Zero: Zero cannot be represented in the normalized form, and must be represented in denormalized form with $E=0$ and $F=0$. There are two representations for zero: $+0$ with $S=0$ and -0 with $S=1$.

Infinity: The value of +infinity (e.g., $1/0$) and -infinity (e.g., $-1/0$) are represented with an exponent of all 1's ($E = 255$ for single-precision and $E = 2047$ for double-precision), $F=0$, and $S=0$ (for +INF) and $S=1$ (for -INF).

Not a Number (NaN): NaN denotes a value that cannot be represented as real number (e.g. $0/0$). NaN is represented with Exponent of all 1's ($E = 255$ for single-precision and $E = 2047$ for double-precision) and any non-zero fraction.

5. Character Encoding

In computer memory, character are "encoded" (or "represented") using a chosen "character encoding schemes" (aka "character set", "charset", "character map", or "code page").

For example, in ASCII (as well as Latin1, Unicode, and many other character sets):

- code numbers 65D (41H) to 90D (5AH) represents 'A' to 'Z', respectively.
- code numbers 97D (61H) to 122D (7AH) represents 'a' to 'z', respectively.
- code numbers 48D (30H) to 57D (39H) represents '0' to '9', respectively.

It is important to note that the representation scheme must be known before a binary pattern can be interpreted. E.g., the 8-bit pattern "0100 0010B" could represent anything under the sun known only to the person encoded it.

The most commonly-used character encoding schemes are: 7-bit ASCII (ISO/IEC 646) and 8-bit Latin-x (ISO/IEC 8859-x) for western european characters, and Unicode (ISO/IEC 10646) for internationalization (i18n).

A 7-bit encoding scheme (such as ASCII) can represent 128 characters and symbols. An 8-bit character encoding scheme (such as Latin-x) can represent 256 characters and symbols; whereas a 16-bit encoding scheme (such as Unicode UCS-2) can represents 65,536 characters and symbols.

5.1 7-bit ASCII Code (aka US-ASCII, ISO/IEC 646, ITU-T T.50)

- ASCII (American Standard Code for Information Interchange) is one of the earlier character coding schemes.
- ASCII is originally a 7-bit code. It has been extended to 8-bit to better utilize the 8-bit computer memory organization. (The 8th-bit was originally used for *parity check* in the early computers.)
- Code numbers 32D (20H) to 126D (7EH) are printable (displayable) characters as tabulated:

Hex	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
2	SP	!	"	#	\$	%	&	'	()	*	+	,	-	.	/
3	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5	P	Q	R	S	T	U	V	W	X	Y	Z	[\]	^	_
6	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7	p	q	r	s	t	u	v	w	x	y	z	{		}	~	

- Code number 32D (20H) is the *blank* or *space* character.
- '0' to '9': 30H-39H (0011 0001B to 0011 1001B) or (0011 xxxxB where xxxx is the equivalent integer value)
- 'A' to 'Z': 41H-5AH (0101 0001B to 0101 1010B) or (010x xxxxB). 'A' to 'Z' are continuous without gap.
- 'a' to 'z': 61H-7AH (0110 0001B to 0111 1010B) or (011x xxxxB). 'A' to 'Z' are also continuous without gap. However, there is a gap between uppercase and lowercase letters. To convert between upper and lowercase, flip the value of bit-5.
- Code numbers 0D (00H) to 1D (1FH), and 127D (7FH) are special control characters, which are non-printable (non-displayable), as tabulated below. Many of these characters were used in the early days for transmission control (e.g., STX,

ETX) and printer control (e.g., Form-Feed), which are now obsolete. The remaining meaningful codes today are:

- 09H for Tab ('`\t`').
- 0AH for Line-Feed or newline (LF or '`\n`') and 0DH for Carriage-Return (CR or '`\r`'), which are used as *line delimiter* (aka *line separator*, *end-of-line*) for text files. There is unfortunately no standard for line delimiter: Unixes and Mac use 0AH (LF or "`\n`"), Windows use 0D0AH (CR+LF or "`\r\n`"). Programming languages such as C/C++/Java (which was created on Unix) use 0AH (LF or "`\n`").
- In programming languages such as C/C++/Java, line-feed (0AH) is denoted as '`\n`', carriage-return (0DH) as '`\r`', tab (09H) as '`\t`'.

DEC	HEX		Meaning	DEC	HEX		Meaning
0	00	NUL	Null	17	11	DC1	Device Control 1
1	01	SOH	Start of Heading	18	12	DC2	Device Control 2
2	02	STX	Start of Text	19	13	DC3	Device Control 3
3	03	ETX	End of Text	20	14	DC4	Device Control 4
4	04	EOT	End of Transmission	21	15	NAK	Negative Ack.
5	05	ENQ	Enquiry	22	16	SYN	Sync. Idle
6	06	ACK	Acknowledgment	23	17	ETB	End of Transmission
7	07	BEL	Bell	24	18	CAN	Cancel
8	08	BS	Back Space ' <code>\b</code> '	25	19	EM	End of Medium
9	09	HT	Horizontal Tab ' <code>\t</code> '	26	1A	SUB	Substitute
10	0A	LF	Line Feed ' <code>\n</code> '	27	1B	ESC	Escape
11	0B	VT	Vertical Feed	28	1C	IS4	File Separator
12	0C	FF	Form Feed ' <code>\f</code> '	29	1D	IS3	Group Separator
13	0D	CR	Carriage Return ' <code>\r</code> '	30	1E	IS2	Record Separator
14	0E	SO	Shift Out	31	1F	IS1	Unit Separator
15	0F	SI	Shift In				
16	10	DLE	Datalink Escape	127	7F	DEL	Delete

5.2 8-bit Latin-1 (aka ISO/IEC 8859-1)

ISO/IEC-8859 is a *collection* of 8-bit character encoding standards for the western languages.

ISO/IEC 8859-1, aka Latin alphabet No. 1, or Latin-1 in short, is the most commonly-used encoding scheme for western european languages. It has 191 printable characters from the latin script, which covers languages like English, German, Italian, Portuguese and Spanish. Latin-1 is backward compatible with the 7-bit US-ASCII code. That is, the first 128 characters in Latin-1 (code numbers 0 to 127 (7FH)), is the same as US-ASCII. Code numbers 128 (80H) to 159 (9FH) are not assigned. Code numbers 160 (A0H) to 255 (FFH) are given as follows:

Hex	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
A	NBSP	¡	¢	£	¤	¥	¦	§	¨	©	ª	«	¬	SHY	®	¯
B	°	±	²	³	´	µ	¶	·	¸	¹	º	»	¼	½	¾	¿
C	À	Á	Â	Ã	Ä	Å	Æ	Ç	È	É	Ê	Ë	Ì	Í	Î	Ï
D	Ð	Ñ	Ò	Ó	Ô	Õ	Ö	×	Ø	Ù	Ú	Û	Ü	Ý	Þ	ß
E	à	á	â	ã	ä	å	æ	ç	è	é	ê	ë	ì	í	î	ï
F	ð	ñ	ò	ó	ô	õ	ö	÷	ø	ù	ú	û	ü	ý	þ	ÿ

ISO/IEC-8859 has 16 parts. Besides the most commonly-used Part 1, Part 2 is meant for Central European (Polish, Czech, Hungarian, etc), Part 3 for South European (Turkish, etc), Part 4 for North European (Estonian, Latvian, etc), Part 5 for Cyrillic, Part 6 for Arabic, Part 7 for Greek, Part 8 for Hebrew, Part 9 for Turkish, Part 10 for Nordic, Part 11 for Thai, Part 12 was abandon, Part 13 for Baltic Rim, Part 14 for Celtic, Part 15 for French, Finnish, etc. Part 16 for South-Eastern European.

5.3 Other 8-bit Extension of US-ASCII (ASCII Extensions)

Beside the standardized ISO-8859-x, there are many 8-bit ASCII extensions, which are not compatible with each others.

ANSI (American National Standards Institute) (aka **Windows-1252**, or Windows Codepage 1252): for Latin alphabets used in the legacy DOS/Windows systems. It is a superset of ISO-8859-1 with code numbers 128 (80H) to 159 (9FH) assigned to displayable characters, such as "smart" single-quotes and double-quotes. A common problem in web browsers is that all the quotes and apostrophes (produced by "smart quotes" in some Microsoft software) were replaced with question marks or some strange symbols. It it because the document is labeled as ISO-8859-1 (instead of Windows-1252), where these code numbers are undefined. Most modern browsers and e-mail clients treat charset ISO-8859-1 as Windows-1252 in order to accommodate such mis-labeling.

Hex	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
8	€		,	f	„	…	†	‡	^	%	Š	<	Ɛ		Ž	
9		'	'	"	"	•	—	—		™	š	>	œ		ž	Ÿ

EBCDIC (Extended Binary Coded Decimal Interchange Code): Used in the early IBM computers.

5.4 Unicode (aka ISO/IEC 10646 Universal Character Set)

Before Unicode, no single character encoding scheme could represent characters in all languages. For example, western european uses several encoding schemes (in the ISO-8859-x family). Even a single language like Chinese has a few encoding schemes (GB2312/GBK, BIG5). Many encoding schemes are in conflict of each other, i.e., the same code number is assigned to different characters.

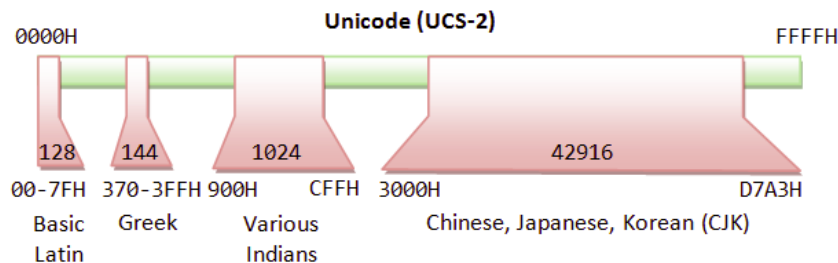
Unicode aims to provide a standard character encoding scheme, which is universal, efficient, uniform and unambiguous. Unicode standard is maintained by a non-profit organization called the Unicode Consortium (@ www.unicode.org). Unicode is an ISO/IEC standard 10646.

Unicode is backward compatible with the 7-bit US-ASCII and 8-bit Latin-1 (ISO-8859-1). That is, the first 128 characters are the same as US-ASCII; and the first 256 characters are the same as Latin-1.

Unicode originally uses 16 bits (called UCS-2 or Unicode Character Set - 2 byte), which can represent up to 65,536 characters. It has since been expanded to more than 16 bits, currently stands at 21 bits. The range of the legal codes in ISO/IEC 10646 is now from U+0000H to U+10FFFFH (21 bits or about 2 million characters), covering all current and ancient historical scripts. The original 16-bit range of U+0000H to U+FFFFH (65536 characters) is known as *Basic Multilingual Plane* (BMP), covering all the major languages in use currently. The characters outside BMP are called *Supplementary Characters*, which are not frequently-used.

Unicode has two encoding schemes:

- **UCS-2** (Universal Character Set - 2 Byte): Uses 2 bytes (16 bits), covering 65,536 characters in the BMP. BMP is sufficient for most of the applications. UCS-2 is now obsolete.
- **UCS-4** (Universal Character Set - 4 Byte): Uses 4 bytes (32 bits), covering BMP and the supplementary characters.



5.5 UTF-8 (Unicode Transformation Format - 8-bit)

The 16/32-bit Unicode (UCS-2/4) is grossly inefficient if the document contains mainly ASCII characters, because each character occupies two bytes of storage. Variable-length encoding schemes, such as UTF-8, which uses 1-4 bytes to represent a character, was devised to improve the efficiency. In UTF-8, the 128 commonly-used US-ASCII characters use only 1 byte, but some less-commonly characters may require up to 4 bytes. Overall, the efficiency improved for document containing mainly US-ASCII texts.

The transformation between Unicode and UTF-8 is as follows:

Bits	Unicode	UTF-8 Code	Bytes
------	---------	------------	-------

7	00000000 0xxxxxxx	0xxxxxxx	1
			(ASCII)
11	00000yyy yyxxxxxx	110yyyyy 10xxxxxx	2
16	zzzzyyyy yyxxxxxx	1110zzzz 10yyyyyy	3
		10xxxxxx	
21	000uuuuu zzzzyyyy	11110uuu 10uuzzzz	4
	yyxxxxxx	10yyyyyy 10xxxxxx	

In UTF-8, Unicode numbers corresponding to the 7-bit ASCII characters are padded with a leading zero; thus has the same value as ASCII. Hence, UTF-8 can be used with all software using ASCII. Unicode numbers of 128 and above, which are less frequently used, are encoded using more bytes (2-4 bytes). UTF-8 generally requires less storage and is compatible with ASCII. The drawback of UTF-8 is more processing power needed to unpack the code due to its variable length. UTF-8 is the most popular format for Unicode.

Notes:

- UTF-8 uses 1-3 bytes for the characters in BMP (16-bit), and 4 bytes for supplementary characters outside BMP (21-bit).
- The 128 ASCII characters (basic Latin letters, digits, and punctuation signs) use one byte. Most European and Middle East characters use a 2-byte sequence, which includes extended Latin letters (with tilde, macron, acute, grave and other accents), Greek, Armenian, Hebrew, Arabic, and others. Chinese, Japanese and Korean (CJK) use three-byte sequences.
- All the bytes, except the 128 ASCII characters, have a leading '1' bit. In other words, the ASCII bytes, with a leading '0' bit, can be identified and decoded easily.

Example: 您好 (Unicode: 60A8H 597DH)

```
Unicode (UCS-2) is 60A8H = 0110 0000 10 101000B
→ UTF-8 is 11100110 10000010 10101000B = E6 82 A8H
Unicode (UCS-2) is 597DH = 0101 1001 01 111101B
→ UTF-8 is 11100101 10100101 10111101B = E5 A5 BDH
```

5.6 UTF-16 (Unicode Transformation Format - 16-bit)

UTF-16 is a variable-length Unicode character encoding scheme, which uses 2 to 4 bytes. UTF-16 is not commonly used. The transformation table is as follows:

Unicode	UTF-16 Code	Bytes
xxxxxxx xxxxxxxx	Same as UCS-2 - no encoding	2
000uuuuu zzzzyyyy yyxxxxxx (uuuuu≠0)	110110ww wwzzzyy yyxxxxxx (www = uuuuu - 1)	4

Take note that for the 65536 characters in BMP, the UTF-16 is the same as UCS-2 (2 bytes). However, 4 bytes are used for the supplementary characters outside the BMP.

For BMP characters, UTF-16 is the same as UCS-2. For supplementary characters, each character requires a pair 16-bit values, the first from the high-surrogates range, (\uD800-\uDBFF), the second from the low-surrogates range (\uDC00-\uDFFF).

5.7 UTF-32 (Unicode Transformation Format - 32-bit)

Same as UCS-4, which uses 4 bytes for each character - unencoded.

5.8 Formats of Multi-Byte (e.g., Unicode) Text Files

Endianess (or byte-order): For a multi-byte character, you need to take care of the order of the bytes in storage. In *big endian*, the most significant byte is stored at the memory location with the lowest address (big byte first). In *little endian*, the most significant byte is stored at the memory location with the highest address (little byte first). For example, 您 (with Unicode number of 60A8H) is stored as 60 A8 in big endian; and stored as A8 60 in little endian. Big endian, which produces a more readable hex dump, is more commonly-used, and is often the default.

BOM (Byte Order Mark): BOM is a special Unicode character having code number of FEFFH, which is used to differentiate big-endian and little-endian. For big-endian, BOM appears as FE FFH in the storage. For little-endian, BOM appears as FF FEH. Unicode reserves these two code numbers to prevent it from crashing with another character.

Unicode text files could take on these formats:

- Big Endian: UCS-2BE, UTF-16BE, UTF-32BE.
- Little Endian: UCS-2LE, UTF-16LE, UTF-32LE.
- UTF-16 with BOM. The first character of the file is a BOM character, which specifies the endianness. For big-endian, BOM appears as FE FFH in the storage. For little-endian, BOM appears as FF FEH.

UTF-8 file is always stored as big endian. BOM plays no part. However, in some systems (in particular Windows), a BOM is added as the first character in the UTF-8 file as the signature to identify the file as UTF-8 encoded. The BOM character (FEFFH) is encoded in UTF-8 as EF BB BF. Adding a BOM as the first character of the file is not recommended, as it may be incorrectly interpreted in other system. You can have a UTF-8 file without BOM.

5.9 Formats of Text Files

Line Delimiter or End-Of-Line (EOL): Sometimes, when you use the Windows NotePad to open a text file (created in Unix or Mac), all the lines are joined together. This is because different operating platforms use different character as the so-called *line delimiter* (or *end-of-line* or EOL). Two non-printable control characters are involved: 0AH (Line-Feed or LF) and 0DH (Carriage-Return or CR).

- Windows/DOS uses 0D0AH (CR+LF or "\r\n") as EOL.
- Unix and Mac use 0AH (LF or "\n") only.

End-of-File (EOF): [TODO]

5.10 Windows' CMD Codepage

Character encoding scheme (charset) in Windows is called *codepage*. In CMD shell, you can issue command "chcp" to display the current codepage, or "chcp codepage-number" to change the codepage.

Take note that:

- The default codepage 437 (used in the original DOS) is an 8-bit character set called *Extended ASCII*, which is different from Latin-1 for code numbers above 127.
- Codepage 1252 (Windows-1252), is not exactly the same as Latin-1. It assigns code number 80H to 9FH to letters and punctuation, such as smart single-quotes and double-quotes. A common problem in browser that display quotes and apostrophe in question marks or boxes is because the page is supposed to be Windows-1252, but mislabelled as ISO-8859-1.
- For internationalization and chinese character set: codepage 65001 for UTF8, codepage 1201 for UCS-2BE, codepage 1200 for UCS-2LE, codepage 936 for chinese characters in GB2312, codepage 950 for chinese characters in Big5.

5.11 Chinese Character Sets

Unicode supports all languages, including asian languages like Chinese (both simplified and traditional characters), Japanese and Korean (collectively called CJK). There are more than 20,000 CJK characters in Unicode. Unicode characters are often encoded in the UTF-8 scheme, which unfortunately, requires 3 bytes for each CJK character, instead of 2 bytes in the unencoded UCS-2 (UTF-16).

Worse still, there are also various chinese character sets, which is not compatible with Unicode:

- GB2312/GBK: for *simplified* chinese characters. GB2312 uses 2 bytes for each chinese character. The most significant bit (MSB) of both bytes are set to 1 to co-exist with 7-bit ASCII with the MSB of 0. There are about 6700 characters. GBK is an extension of GB2312, which include more characters as well as traditional chinese characters.
- BIG5: for *traditional* chinese characters BIG5 also uses 2 bytes for each chinese character. The most significant bit of both bytes are also set to 1. BIG5 is not compatible with GBK, i.e., the same code number is assigned to different character.

For example, the world is made more interesting with these many standards:

	Standard	Characters	Codes
Simplified	GB2312	和谐	BACD D0B3
	UCS-2	和谐	548C 8C10
	UTF-8	和谐	E5928C E8B090
Traditional	BIG5	和諧	A94D BFD3
	UCS-2	和諧	548C 8AE7
	UTF-8	和諧	E5928C E8ABA7

Notes for Windows' CMD Users: To display the chinese character correctly in CMD shell, you need to choose the correct codepage, e.g., 65001 for UTF8, 936 for GB2312/GBK, 950 for Big5, 1201 for UCS-2BE, 1200 for UCS-2LE, 437 for the original DOS. You can use command "chcp" to display the current code page and command "chcp codepage-number" to change the codepage. You also have to choose a font that can display the characters (e.g., Courier New, Consolas or Lucida Console, NOT Raster font).

5.12 Collating Sequences (for Ranking Characters)

A string consists of a sequence of characters in upper or lower cases, e.g., "apple", "BOY", "Cat". In sorting or comparing strings, if we order the characters according to the underlying code numbers (e.g., US-ASCII) character-by-character, the order for the example would be "BOY", "apple", "Cat" because uppercase letters have a smaller code number than lowercase letters. This does not agree with the so-called *dictionary order*, where the same uppercase and lowercase letters have the same rank. Another common

problem in ordering strings is "10" (ten) at times is ordered in front of "1" to "9".

Hence, in sorting or comparison of strings, a so-called *collating sequence* (or *collation*) is often defined, which specifies the ranks for letters (uppercase, lowercase), numbers, and special symbols. There are many collating sequences available. It is entirely up to you to choose a collating sequence to meet your application's specific requirements. Some *case-insensitive dictionary-order collating sequences* have the same rank for same uppercase and lowercase letters, i.e., 'A', 'a' ⇒ 'B', 'b' ⇒ ... ⇒ 'Z', 'z'. Some *case-sensitive dictionary-order collating sequences* put the uppercase letter before its lowercase counterpart, i.e., 'A' ⇒ 'B' ⇒ 'C'... ⇒ 'a' ⇒ 'b' ⇒ 'c'.... Typically, space is ranked before digits '0' to '9', followed by the alphabets.

Collating sequence is often language dependent, as different languages use different sets of characters (e.g., á, é, a, α) with their own orders.

5.13 For Java Programmers - java.nio.Charset

JDK 1.4 introduced a new java.nio.charset package to support encoding/decoding of characters from UCS-2 used internally in Java program to any supported charset used by external devices.

Example: The following program encodes some Unicode texts in various encoding scheme, and display the Hex codes of the encoded byte sequences.

```
import java.nio.ByteBuffer;
import java.nio.CharBuffer;
import java.nio.charset.Charset;

public class TestCharsetEncodeDecode {
    public static void main(String[] args) {
        // Try these charsets for encoding
        String[] charsetNames = {"US-ASCII", "ISO-8859-1", "UTF-8", "UTF-16",
                                "UTF-16BE", "UTF-16LE", "GBK", "BIG5"};

        String message = "Hi,您好!"; // message with non-ASCII characters
        // Print UCS-2 in hex codes
        System.out.printf("%10s: ", "UCS-2");
        for (int i = 0; i < message.length(); i++) {
            System.out.printf("%04X ", (int)message.charAt(i));
        }
        System.out.println();

        for (String charsetName: charsetNames) {
            // Get a Charset instance given the charset name string
            Charset charset = Charset.forName(charsetName);
            System.out.printf("%10s: ", charset.name());

            // Encode the Unicode UCS-2 characters into a byte sequence in this charset.
            ByteBuffer bb = charset.encode(message);
            while (bb.hasRemaining()) {
                System.out.printf("%02X ", bb.get()); // Print hex code
            }
            System.out.println();
            bb.rewind();
        }
    }
}
```

```
UCS-2: 0048 0069 002C 60A8 597D 0021 [16-bit fixed-length]
      H i , 您 好 !
US-ASCII: 48 69 2C 3F 3F 21 [8-bit fixed-length]
      H i , ? ? !
ISO-8859-1: 48 69 2C 3F 3F 21 [8-bit fixed-length]
      H i , ? ? !
UTF-8: 48 69 2C E6 82 A8 E5 A5 BD 21 [1-4 bytes variable-length]
      H i , 您 好 !
UTF-16: FE FF 00 48 00 69 00 2C 60 A8 59 7D 00 21 [2-4 bytes variable-length]
      BOM H i , 您 好 ! [Byte-Order-Mark indicates Big-Endian]
UTF-16BE: 00 48 00 69 00 2C 60 A8 59 7D 00 21 [2-4 bytes variable-length]
      H i , 您 好 !
UTF-16LE: 48 00 69 00 2C 00 A8 60 7D 59 21 00 [2-4 bytes variable-length]
      H i , 您 好 !
GBK: 48 69 2C C4 FA BA C3 21 [1-2 bytes variable-length]
      H i , 您 好 !
Big5: 48 69 2C B1 7A A6 6E 21 [1-2 bytes variable-length]
      H i , 您 好 !
```

5.14 For Java Programmers - char and String

The char data type are based on the *original* 16-bit Unicode standard called UCS-2. The Unicode has since evolved to 21 bits, with code range of U+0000 to U+10FFFF. The set of characters from U+0000 to U+FFFF is known as the *Basic Multilingual Plane (BMP)*. Characters above U+FFFF are called *supplementary* characters. A 16-bit Java char cannot hold a supplementary character.

Recall that in the UTF-16 encoding scheme, a BMP characters uses 2 bytes. It is the same as UCS-2. A supplementary character uses 4 bytes. and requires a pair of 16-bit values, the first from the high-surrogates range, (\uD800-\uDBFF), the second from the low-surrogates range (\uDC00-\uDFFF).

In Java, a String is a sequences of Unicode characters. Java, in fact, uses UTF-16 for String and StringBuffer. For BMP characters, they are the same as UCS-2. For supplementary characters, each characters requires a pair of char values.

Java methods that accept a 16-bit char value does not support supplementary characters. Methods that accept a 32-bit int value support all Unicode characters (in the lower 21 bits), including supplementary characters.

This is meant to be an academic discussion. I have yet to encounter the use of supplementary characters!

5.15 Displaying Hex Values & Hex Editors

At times, you may need to display the hex values of a file, especially in dealing with Unicode characters. A *Hex Editor* is a handy tool that a good programmer should possess in his/her toolbox. There are many freeware/shareware Hex Editor available. Try google "Hex Editor".

I used the followings:

- NotePad++ with Hex Editor Plug-in: Open-source and free. You can toggle between Hex view and Normal view by pushing the "H" button.
- PSPad: Freeware. You can toggle to Hex view by choosing "View" menu and select "Hex Edit Mode".
- TextPad: Shareware without expiration period. To view the Hex value, you need to "open" the file by choosing the file format of "binary" (??).
- UltraEdit: Shareware, not free, 30-day trial only.

Let me know if you have a better choice, which is fast to launch, easy to use, can toggle between Hex and normal view, free,

The following Java program can be used to display hex code for Java Primitives (integer, character and floating-point):

```

1 public class PrintHexCode {
2
3     public static void main(String[] args) {
4         int i = 12345;
5         System.out.println("Decimal is " + i);                // 12345
6         System.out.println("Hex is " + Integer.toHexString(i)); // 3039
7         System.out.println("Binary is " + Integer.toBinaryString(i)); // 11000000111001
8         System.out.println("Octal is " + Integer.toOctalString(i)); // 30071
9         System.out.printf("Hex is %x\n", i);                // 3039
10        System.out.printf("Octal is %o\n", i);                // 30071
11
12        char c = 'a';
13        System.out.println("Character is " + c);              // a
14        System.out.printf("Character is %c\n", c);              // a
15        System.out.printf("Hex is %x\n", (short)c);            // 61
16        System.out.printf("Decimal is %d\n", (short)c);         // 97
17
18        float f = 3.5f;
19        System.out.println("Decimal is " + f);                // 3.5
20        System.out.println(Float.toHexString(f));              // 0x1.cp1 (Fraction=1.c, Exponent=1)
21
22        f = -0.75f;
23        System.out.println("Decimal is " + f);                // -0.75
24        System.out.println(Float.toHexString(f));              // -0x1.8p-1 (F=-1.8, E=-1)
25
26        double d = 11.22;
27        System.out.println("Decimal is " + d);                // 11.22
28        System.out.println(Double.toHexString(d));              // 0x1.670a3d70a3d71p3 (F=1.670a3d70a3d71 E=3)
29    }
30 }
```

In Eclipse, you can view the hex code for *integer* primitive Java variables in debug mode as follows: In debug perspective, "Variable" panel = Select the "menu" (inverted triangle) = Java = Java Preferences... = Primitive Display Options = Check "Display hexadecimal values (byte, short, char, int, long)".

6. Summary - Why Bother about Data Representation?

Integer number 1, floating-point number 1.0 character symbol '1', and string "1" are totally different inside the computer memory. You need to know the difference to write good and high-performance programs.

- In 8-bit *signed integer*, integer number 1 is represented as 00000001B.
- In 8-bit *unsigned integer*, integer number 1 is represented as 00000001B.
- In 16-bit *signed integer*, integer number 1 is represented as 00000000 00000001B.
- In 32-bit *signed integer*, integer number 1 is represented as 00000000 00000000 00000000 00000001B.
- In 32-bit *floating-point representation*, number 1.0 is represented as 0 01111111 00000000 00000000 00000000B, i.e., S=0, E=127, F=0.
- In 64-bit *floating-point representation*, number 1.0 is represented as 0 0111111111 0000 00000000 00000000 00000000 00000000 00000000B, i.e., S=0, E=1023, F=0.
- In 8-bit Latin-1, the character symbol '1' is represented as 00110001B (or 31H).
- In 16-bit UCS-2, the character symbol '1' is represented as 00000000 00110001B.
- In UTF-8, the character symbol '1' is represented as 00110001B.

If you "add" a 16-bit signed integer 1 and Latin-1 character '1' or a string "1", you could get a surprise.

6.1 Exercises (Data Representation)

For the following 16-bit codes:

```
0000 0000 0010 1010;  
1000 0000 0010 1010;
```

Give their values, if they are representing:

1. a 16-bit unsigned integer;
2. a 16-bit signed integer;
3. two 8-bit unsigned integers;
4. two 8-bit signed integers;
5. a 16-bit Unicode characters;
6. two 8-bit ISO-8859-1 characters.

Ans: (1) 42, 32810; (2) 42, -32726; (3) 0, 42; 128, 42; (4) 0, 42; -128, 42; (5) '*', '磅'; (6) NUL, '*', PAD, '*'.

REFERENCES & RESOURCES

1. (Floating-Point Number Specification) IEEE 754 (1985), "IEEE Standard for Binary Floating-Point Arithmetic".
2. (ASCII Specification) ISO/IEC 646 (1991) (or ITU-T T.50-1992), "Information technology - 7-bit coded character set for information interchange".
3. (Latin-I Specification) ISO/IEC 8859-1, "Information technology - 8-bit single-byte coded graphic character sets - Part 1: Latin alphabet No. 1".
4. (Unicode Specification) ISO/IEC 10646, "Information technology - Universal Multiple-Octet Coded Character Set (UCS)".
5. Unicode Consortium @ <http://www.unicode.org>.

Last modified: January, 2014

Feedback, comments, corrections, and errata can be sent to Chua Hock-Chuan (ehchua@ntu.edu.sg) | [HOME](#)