

## 3. Process Management in xv6

We begin understanding xv6 process management by looking at the `proc` data structure (line 2353), that corresponds to the PCB. Especially note the fields corresponding to the kernel stack pointer, the trap frame (which stores context before handling a trap), and the context structure (which stores context during a context switch). We will now walk through some of the paths in the code: interrupt handling (sheets 32–34), new process creation using `fork` (sheets 24–25), and the workings of the first `init` and `shell` processes (sheets 24–25, sheets 83–88).

### 3.0: Review of x86 architecture

- The architecture of the underlying CPU defines the instructions and registers available during process execution. Details of the CPU architecture are useful in understanding the architecture-dependent parts of an OS. We will review the basics here.
- Several CPU registers are referenced in the discussion of operating system concepts. While the names and number of registers differ based on the CPU architecture, here are some names of x86 registers that we will find useful. (Note: this list of registers is by no means exhaustive.)
  1. EAX, EBX, ECX, EDX, ESI, and EDI are general purpose registers used to store variables during computations.
  2. The general purpose registers **EBP and ESP store the base and top of the current stack frame.**
  3. The **program counter (PC) is also referred to as EIP (instruction pointer).**
  4. The segment registers CS, DS, ES, FS, GS, and SS store pointers to various segments (e.g., code segment, data segment, stack segment) of the process memory.
  5. The control registers like CR0 hold control information. For example, the CR3 register holds the address of the page table of the current running process, which is used to translate virtual addresses to physical addresses.
- Of these registers, EAX, ECX, and EDX are called caller-save registers, and the rest (EBX, ESI, EDI, ESP, EBP, EIP) are callee-save registers. The callee-save registers must be preserved across function calls and context switches. Suppose an executing process pushes a new stack frame onto a stack during a function call, or the CPU moves away to another process. When the function call returns, or the CPU is switched back to the process, the callee-save registers must have the same values as before the event. The ESP register should be pointing to the old stack as before, and EIP should contain the return address from where the process resumes execution. The EAX register is used to pass arguments and return values. The caller save registers may have changed, and no assumptions would be made of them. This convention is followed by compilers and operating systems.

## 3.1 Handling Interrupts

- Interrupts are assigned unique numbers (called IRQ) on every machine (all system calls get the same IRQ). The interrupt descriptor table (IDT) stores information on what to do for every interrupt number, and is setup during initialization (line 3317). The IDT entries for all interrupts point to a generic function to handle all traps (line 3254) that we will examine soon.
- When an interrupt / trap / system call occurs in xv6, the interrupting hardware causes the processor to execute the `int` instruction, with a specific interrupt number (IRQ) as argument. This special instruction moves the CPU to kernel mode (if it was in user mode previously). The CPU has a *task state segment* for the current running task, that lets the CPU find the kernel stack for the current process and switch to it. As part of the execution of the instruction, the CPU saves some registers on the kernel stack (pointers to the old code segment, stack segment, program counter etc.), which will eventually form a part of the trap frame. Next, the CPU fetches the IDT entry corresponding to the interrupt number, which has a pointer to the interrupt handler. Now, the control is transferred to the kernel, and the CPU starts executing the kernel code that is responsible for handling interrupts. Note that this switch to the kernel code segment and kernel stack need not happen if the process was already in kernel mode at the time the interrupt occurred, and this switch is done by the CPU before transferring control to the kernel interrupt handler.
- The kernel code pointed at by the IDT entry (sheet 32) pushes the interrupt number onto the stack (e.g., line 3233), and completes saving various CPU registers (lines 3256-3260), with the result that the kernel stack now contains a trap frame (line 0602). The trap frame serves two purposes: it saves the execution context that existed just before the trap (so that the process can resume where it left off), and it provides arguments to handle the trap (e.g., interrupt number). In addition to creating a trap frame, the kernel sets up various segment registers to point to the correct code (lines 3263-3268). Then the stack pointer (which is also a pointer to the trap frame, since the top of the stack is the trap frame) is pushed onto the stack, and the generic function to handle all traps (line 3350) is executed. The trap frame is an argument to this function. This function looks at the number of the interrupt in the trap frame, and executes the corresponding action. For example, if it is a system call, the corresponding system call function is executed.
- After trap returns at `trapret`, (line 3277), the registers that were pushed onto the stack are popped (by the OS in sheet 32, and by the CPU using the `iret` instruction). Thus the context of the process prior to the interrupt is restored, and the process can resume execution.
- If the interrupt was a timer interrupt, the trap function tries to yield the CPU to another process (line 3424). If the scheduler does decide to context switch, the return from trap does not happen right away. Instead, the kernel context of the process is pushed onto the kernel stack by the code that does a context switch (we will read it later), and another process starts executing. The return from trap happens after this process gets the CPU again.
- Note that if the trap handled by a process was an interrupt, it could have lead to some blocked processes being unblocked (e.g., some process waiting for disk data). However, note that the

process that serviced the interrupt to unblock a process does not immediately switch to the unblocked process. Instead, it continues execution and gives up the CPU upon a timer interrupt or when it blocks itself. It is important to note that unblocked processes do not run right away, and will wait in the ready state to be scheduled by the scheduler.

- **The contents of the kernel stack are key to understanding the workings of xv6.** When servicing an interrupt, the kernel stack has the trap frame. When the process is switched out by the scheduler after servicing an interrupt, the kernel stack has the trap frame, followed by the process context that is written during the context switch.
- Interrupt handling in most Unix systems is conceptually similar, though more complicated. For example, in Linux, interrupt handling is split into two parts: every interrupt handler has a *top half* and a *bottom half*. The top half consists of the basic necessities that must be performed on receiving an interrupt (e.g., copy packets from the network card's memory to kernel memory), while the bottom half that is executed at a later time does the not-so-time-critical tasks (e.g., TCP/IP processing). Linux also has a separate interrupt context and associated stack that the kernel uses while servicing interrupts, instead of running on the kernel stack of the interrupted process itself.

## 3.2 Process creation

- When a process calls fork, the generic trap handling function calls the specific system call fork function (line 2554). Fork calls the allocproc function (line 2455) to allocate an unused proc data structure.
- The allocproc function creates a new process data structure, and allocates a new kernel stack. On the kernel stack of a new process, it pushes a trap frame and a context structure. The fork function copies the parent's trap frame onto the child's trap frame (line 2572). As a result, both child and the parent return to the same point in the program, right after the fork system call. The only difference is that the fork function zeroes out the EAX register in the child's trap frame (line 2575), which is the return value from fork, causing the child to return with the value 0. On the other hand, the parent returns with the pid of the child.
- In addition to the trap frame, the child's kernel stack also has the context data structure. Normally, processes have this context data structure pushed onto the kernel stack when they are being switched out, so that this data structure can be used to restore context when the process has to resume again. For a new process, this context structure is pushed onto the stack by allocproc, so that the scheduler can start scheduling this new process like any other process. The instruction pointer in this (somewhat artificially created) context data structure is set (line 2491) to point to a piece of code called forkret, which is where this process starts executing when the scheduler swaps it in. The function forkret mainly calls trapret (and does a few other things we'll see later). So, when the child process resumes, it also starts executing the code corresponding to the return from a trap, much like its parent. This is the reason why the child process returns to exactly same instruction after the fork system call, much like the parent, with only a different return value.
- Memory management functions like copying the user space memory image, setting up separate page tables, and other initializations are also done by fork. We will study these later.

### 3.3 Boot procedure

- The first part of the boot process concerns itself with setting up the kernel code in memory, setting up suitable page tables to translate the kernel's memory addresses, and initializing devices (sheet 12). After the kernel starts running, it starts setting up user processes, starting with the first `init` process in the `userinit` function (line 2502).
- To create the first process, the `allocproc` function is used to allocate a `proc` structure, much like in `fork`. Next, the trap frame of this child process is hand-created (lines 2514-2520) to look like the process encountered a trap right on the first instruction in its memory (i.e., the EIP saved in the trap frame is address 0). As a result, when this process executes, its context is restored from the trap frame, and it starts executing at logical address 0. The `userinit` function would have set up the page table of the new process, and loaded the `"/init"` binary at logical address 0. So the CPU starts executing the main function of `"/init"` (sheet 83).
- The `init` process sets up the first three file descriptors (`stdin`, `stdout`, `stderr`), and all point to the `console`. All subsequent processes that are spawned inherit these descriptors. Next, it forks a child, and `exec`'s the shell executable in the child. While the child runs the shell (and forks various other processes), the parent waits and reaps zombies.
- The shell program (sheets 83–88, main function at line 8501) gets the user's input, parses it into a command (`parsecmd` at line 8718), and runs it (`runcmd` at line 8406). For commands that involve multiple sub-commands (e.g., list of commands or pipes), new children are forked and the function `runcmd` is called recursively for each subcommand. The simplest subcommands will eventually call `exec` when the recursion terminates. Also note the complex manipulations on file descriptors in the pipe command (line 8450).