

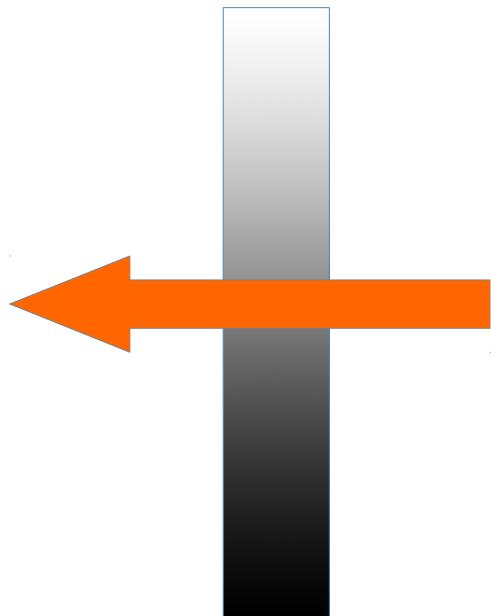
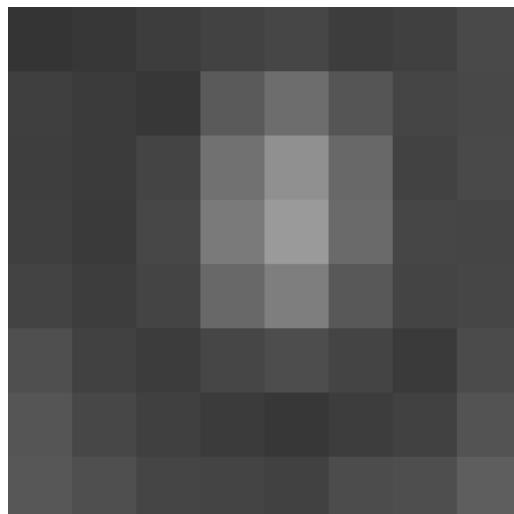
Digital Image Processing

Image Enhancement : Gray-Level Transformations

Suyash P. Awate

Digital Image

- Assumptions
 - 8-bit per pixel
 - 256 intensity levels or gray levels
 - 0, 1, ..., 255 : this is called the grayscale
 - 0 = black
 - 255 = white



52	55	61	66	70	61	64	73
63	59	55	90	109	85	69	72
62	59	68	113	144	104	66	73
63	58	71	122	154	106	70	69
67	61	68	104	126	88	68	70
79	65	60	70	77	68	58	75
85	71	64	59	55	61	65	83
87	79	69	68	65	76	78	94

Enhancement

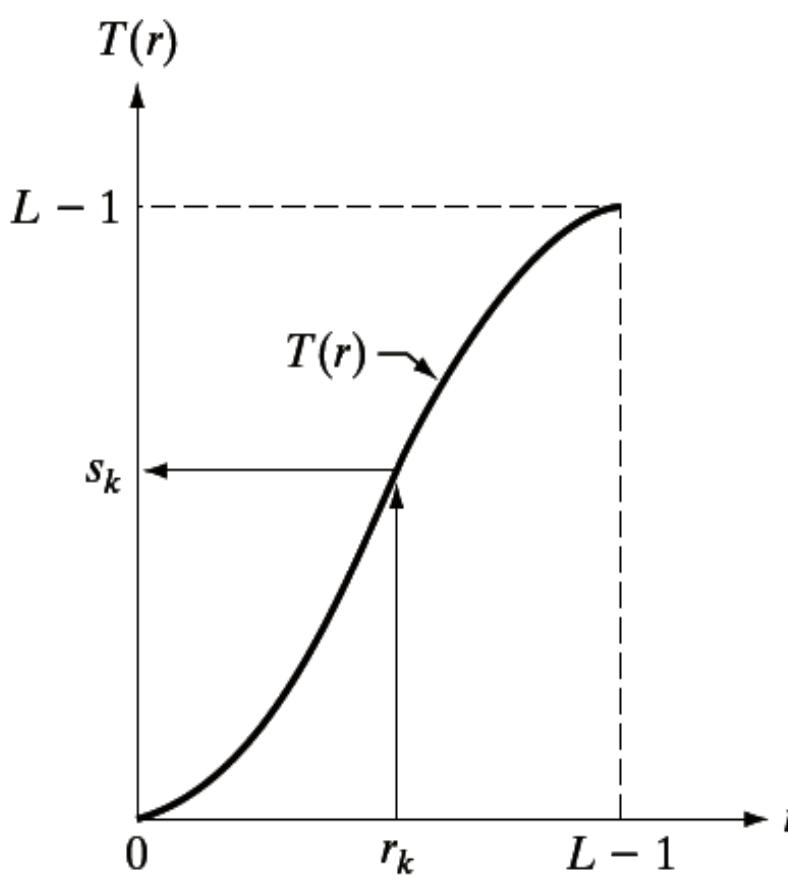
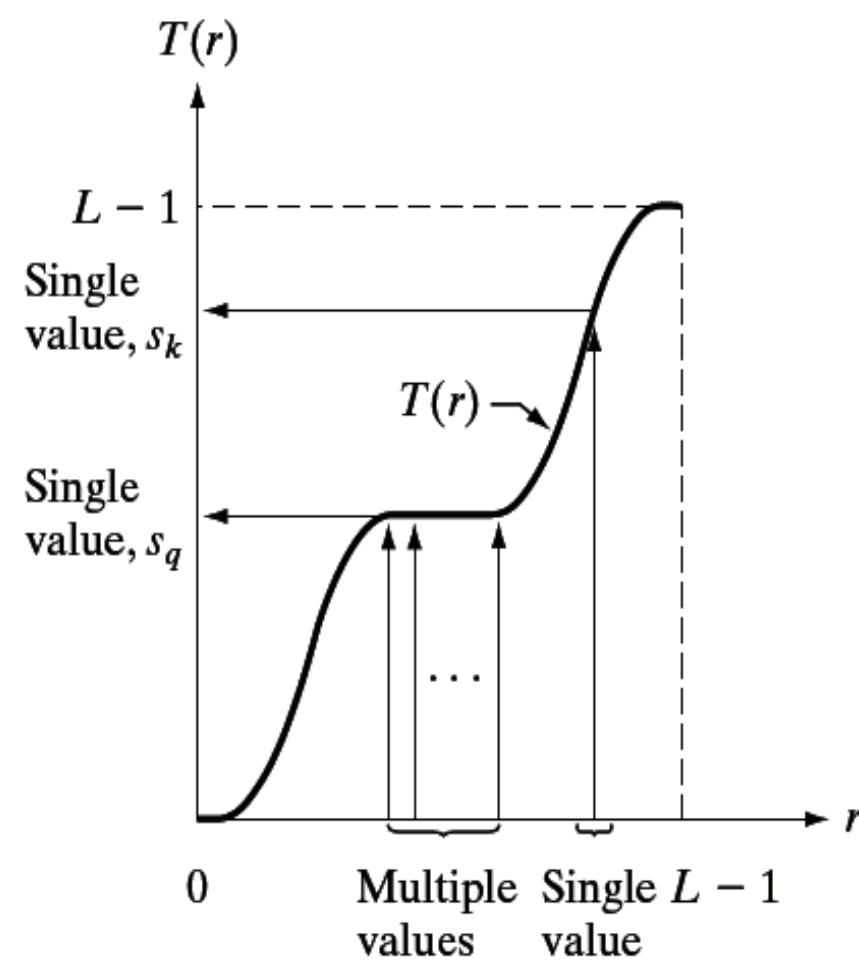
- What does it mean ?
 - Processing an image so that the result is more suitable (compared to the original image) for a **specific** application
- Evaluation
 - Visual, subjective
 - Quantitative

Gray-Level Transformations

- Design **functions** to transform gray levels within an image by directly operating on them
 - Example
 - Input Image : $f(x,y)$
 - Transformation : $T[.]$ operates on individual intensities
 - Output Image : $g(x,y) = T[f(x,y)]$
 - Another example
 - Transformation :
 $U[.]$ operates on
intensities in a neighborhood
 - Output Image :
 $g(x,y) = U[f(x,y), f(x+1,y), f(x-1,y)]$

Gray-Level Transformations

- We want transformations to be **monotonic** functions
 - Monotonically increasing function: $a > b \rightarrow T(a) \geq T(b)$
 - Monotonically decreasing function: $a > b \rightarrow T(a) \leq T(b)$

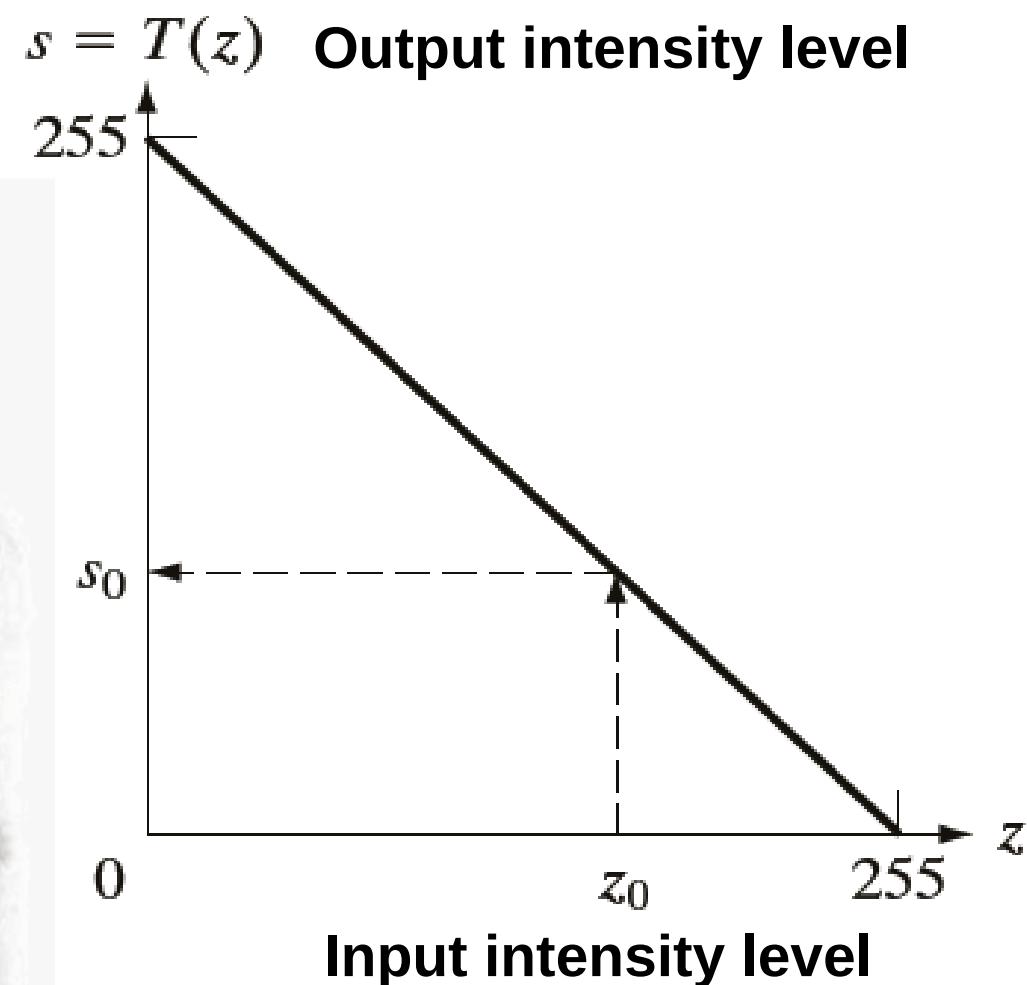
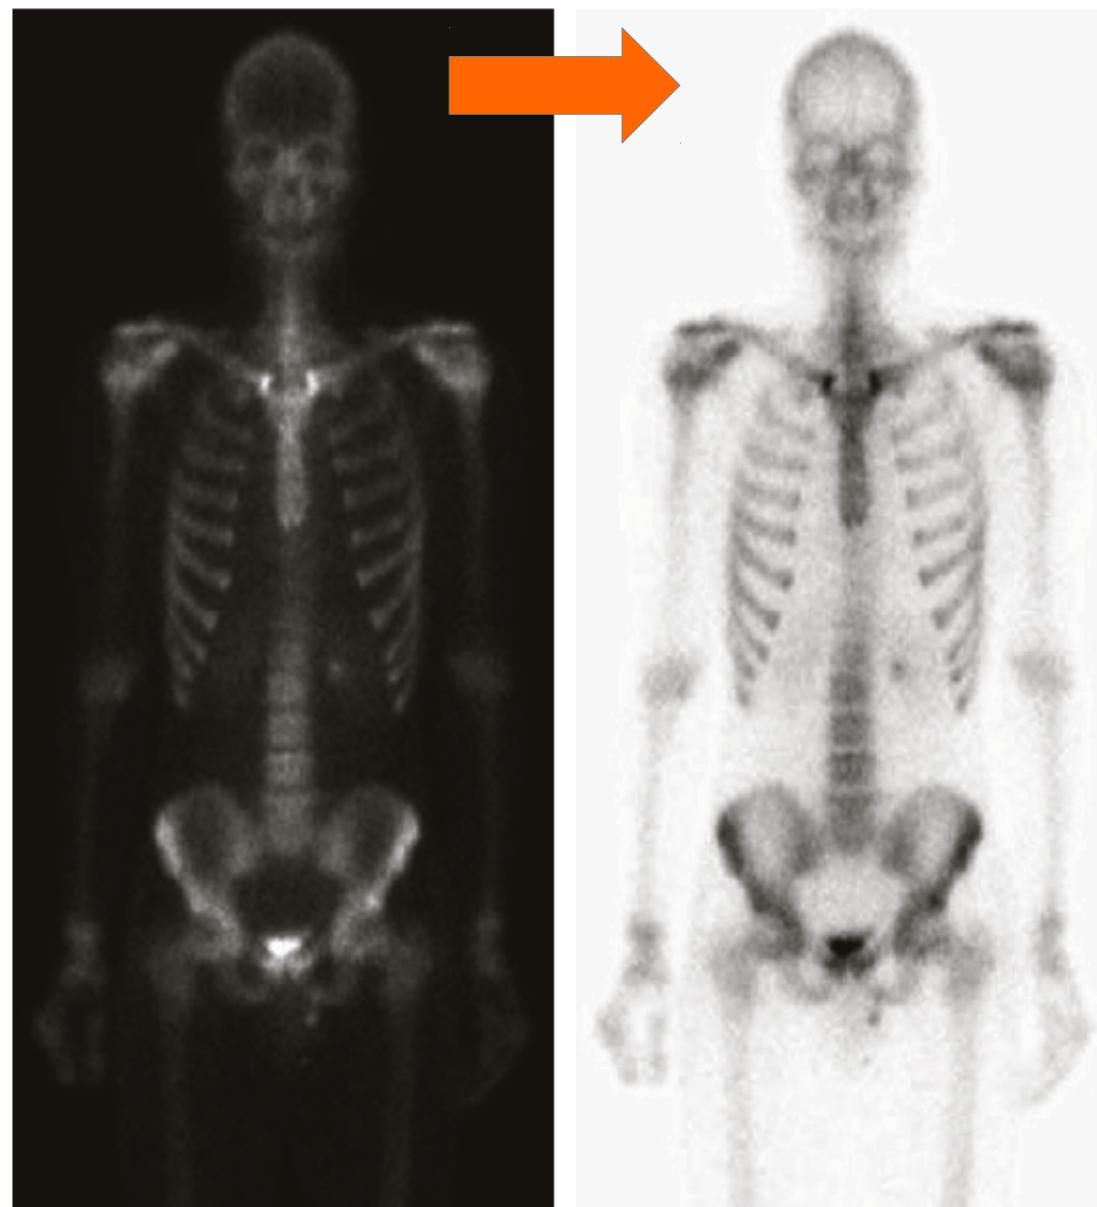


a b

FIGURE 3.17
(a) Monotonically increasing function, showing how multiple values can map to a single value.
(b) Strictly monotonically increasing function. This is a one-to-one mapping, both ways.

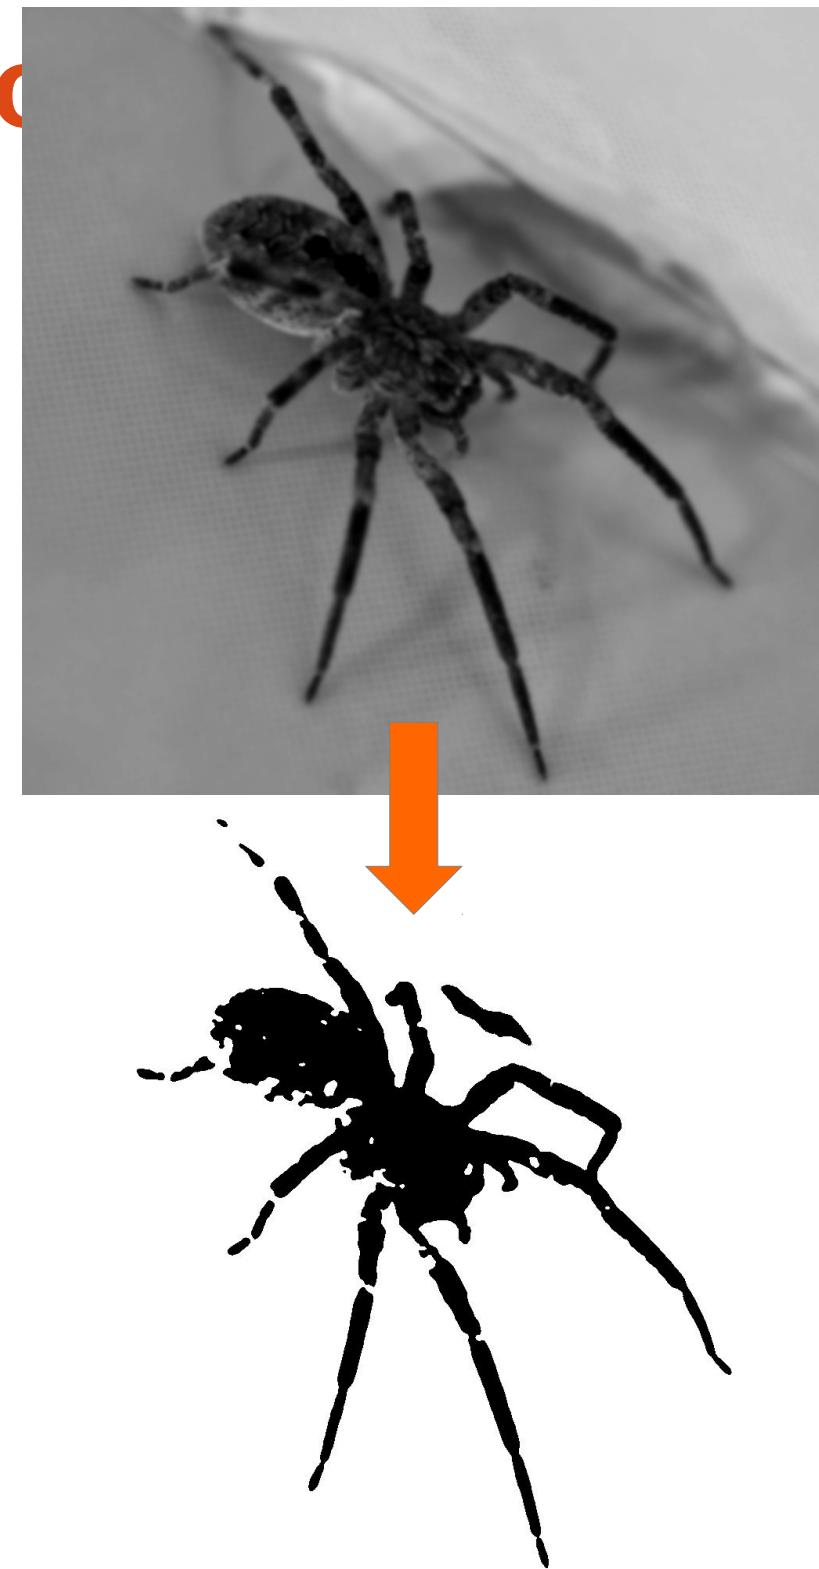
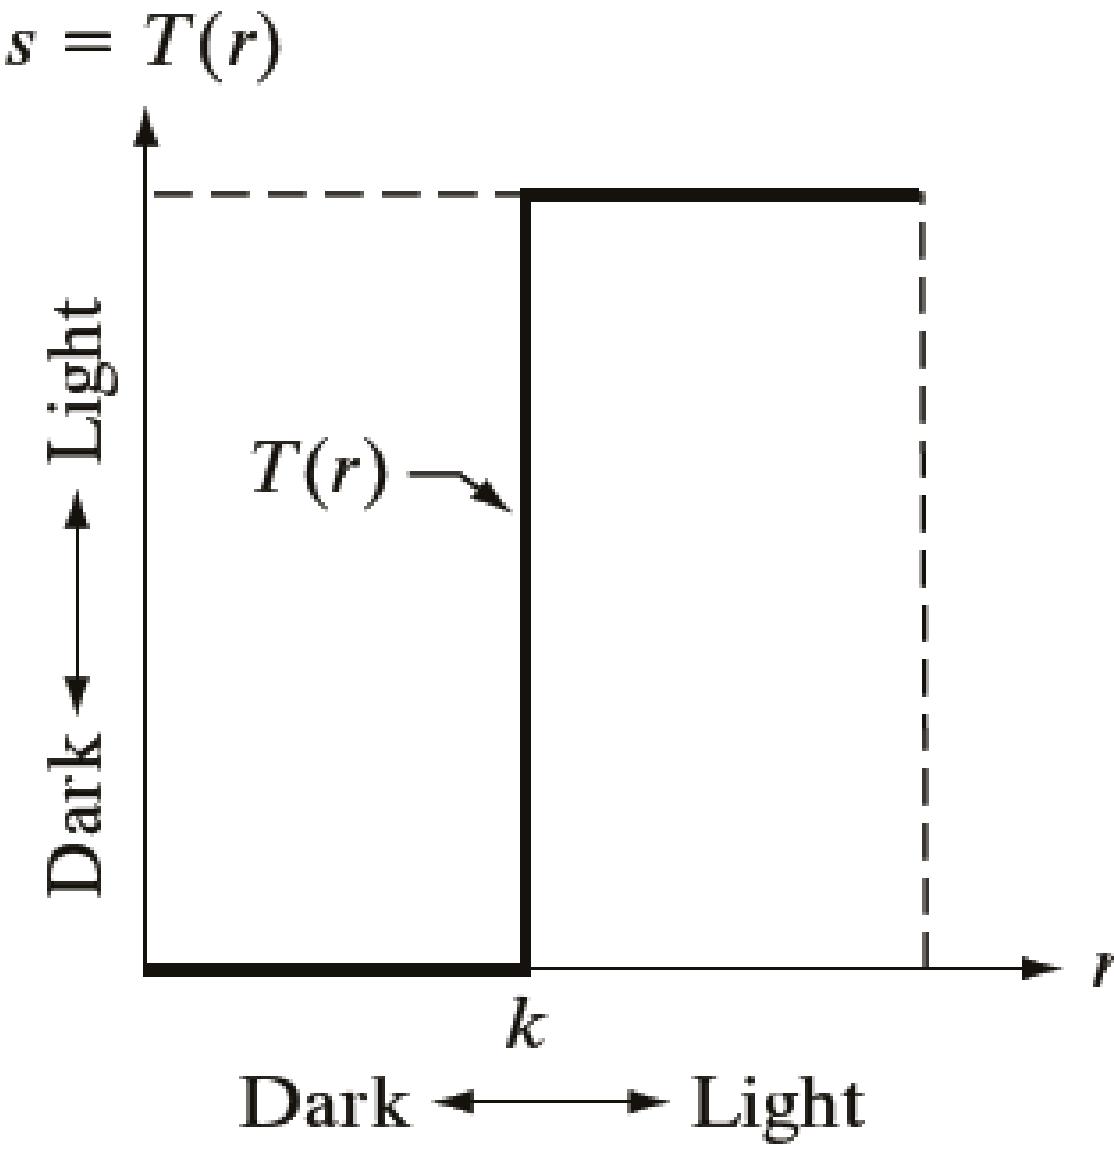
Gray-Level Transformations

- Negative



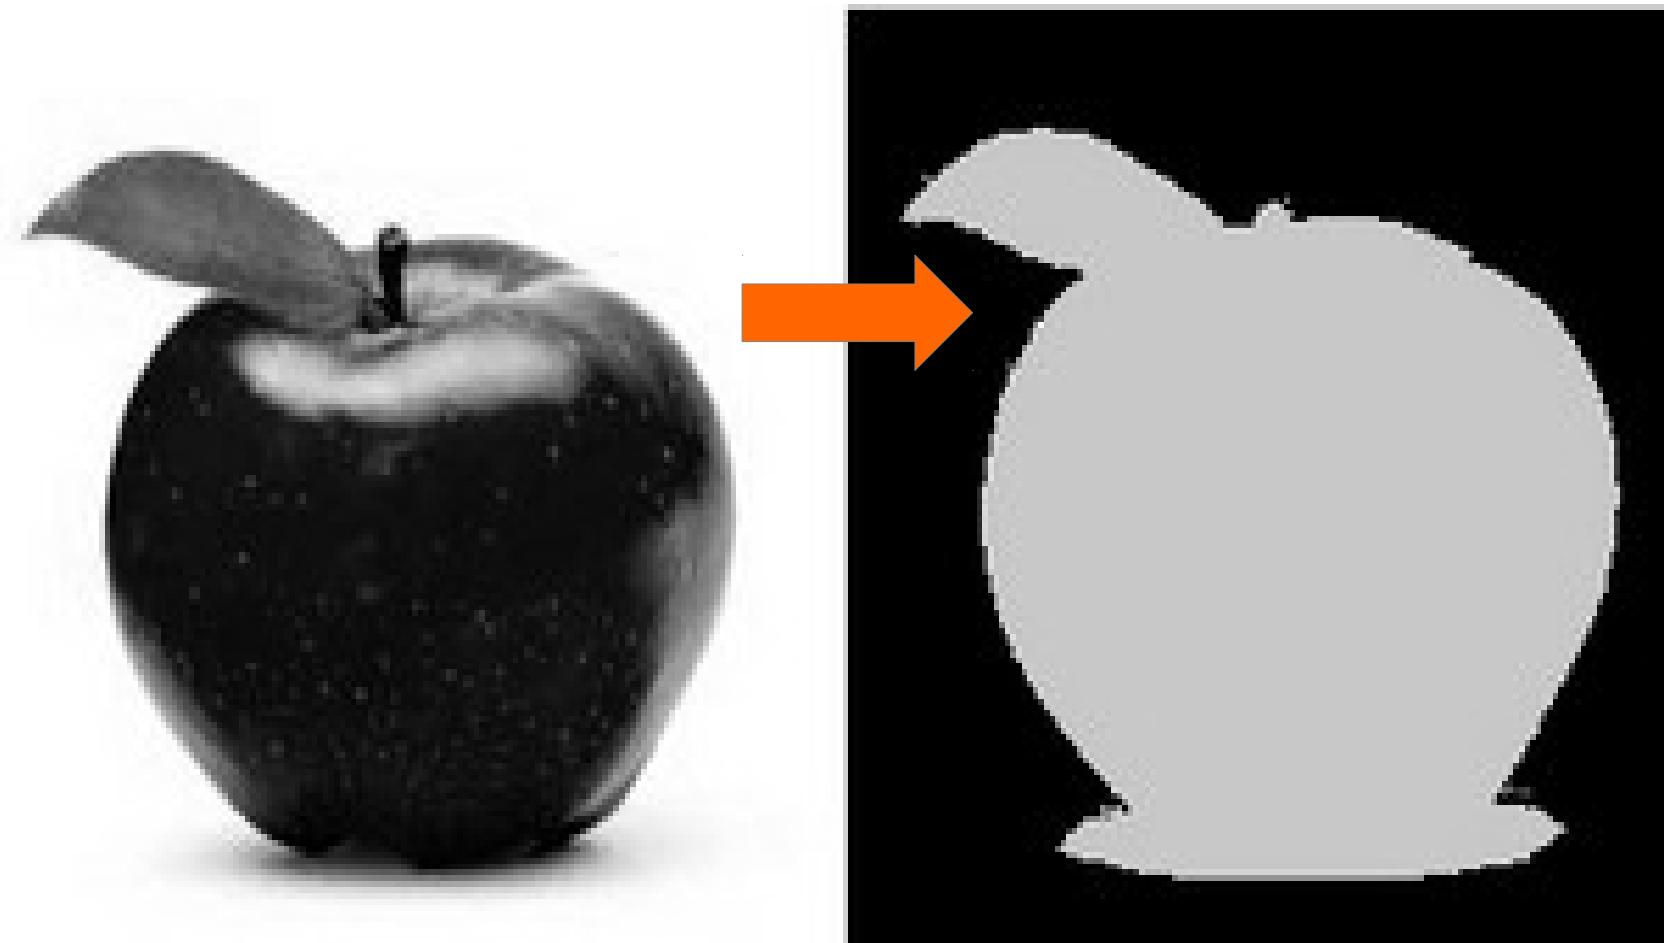
Gray-Level Transformation

- Thresholding



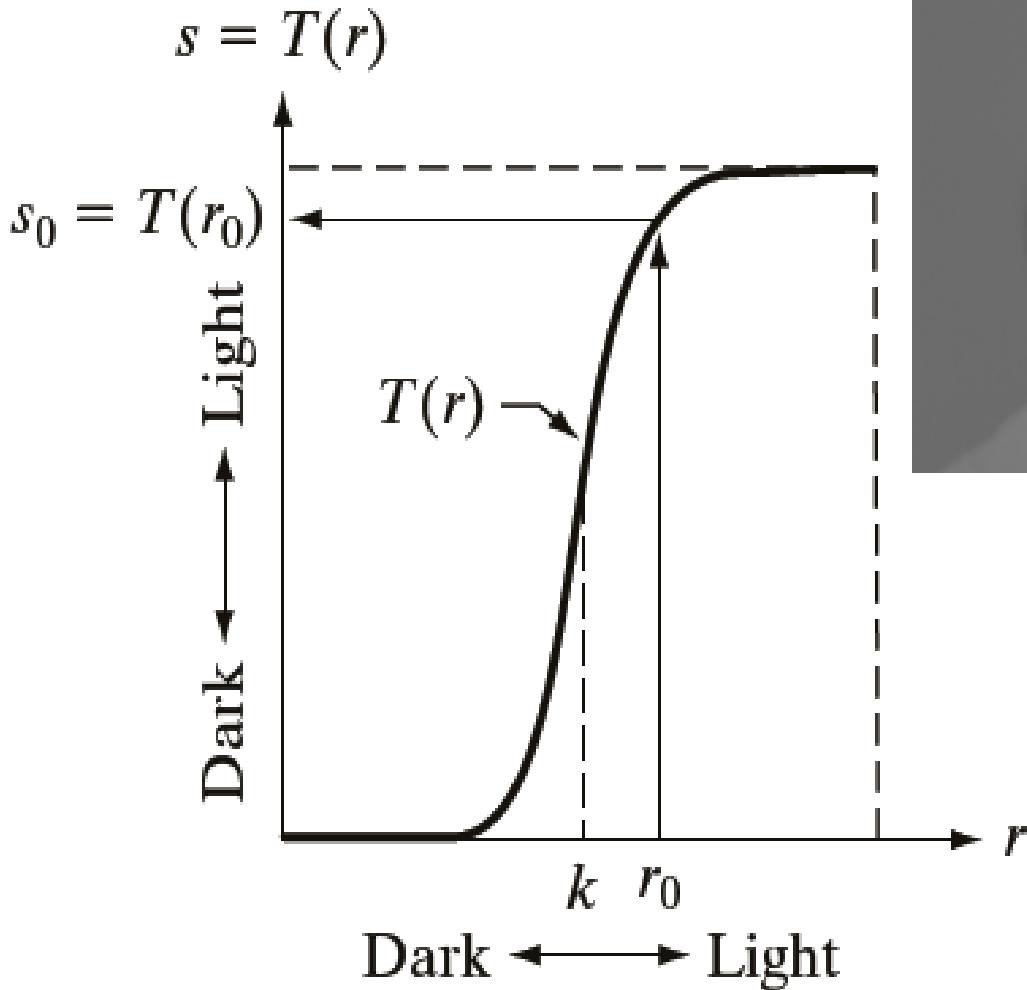
Gray-Level Transformations

- Thresholding
 - What is the transformation function for this ?



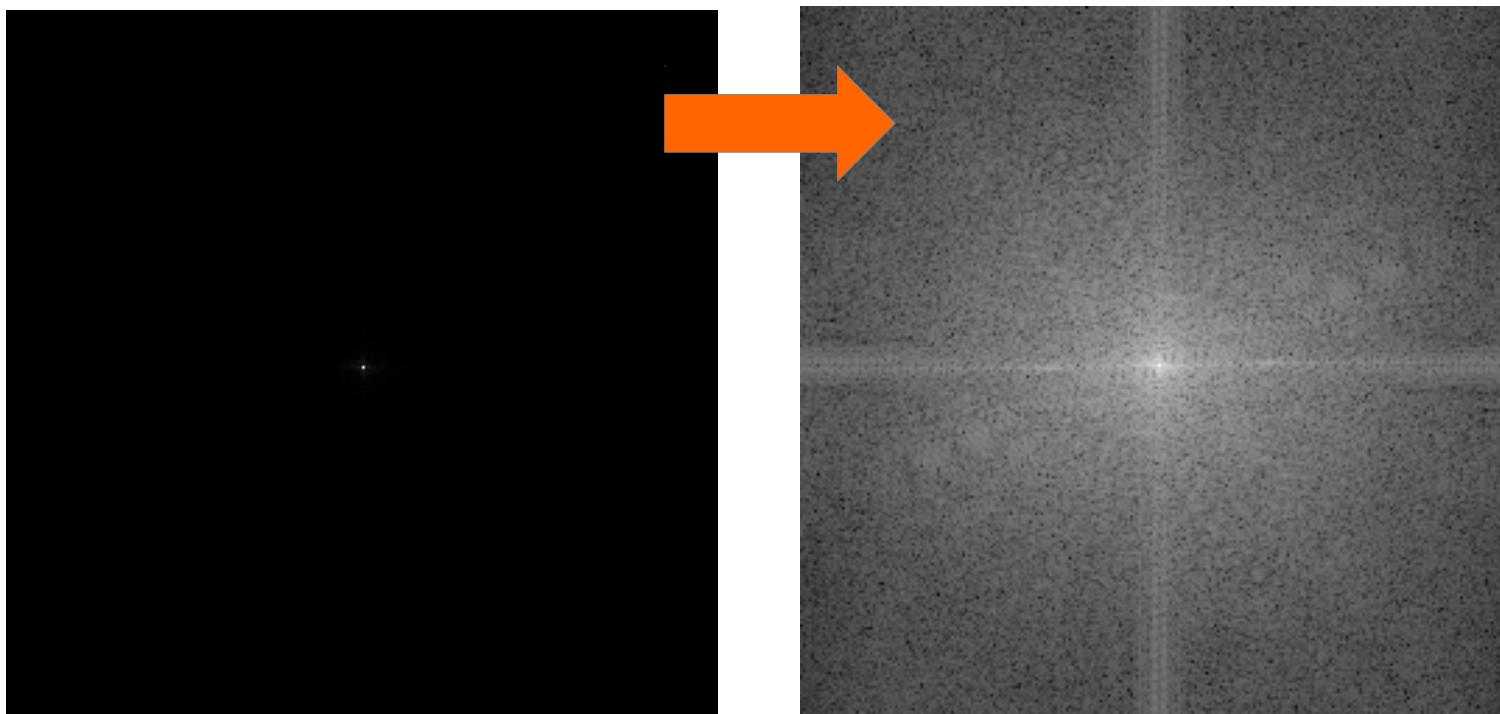
Gray-Level Transformations

- Contrast stretching



Gray-Level Transformations

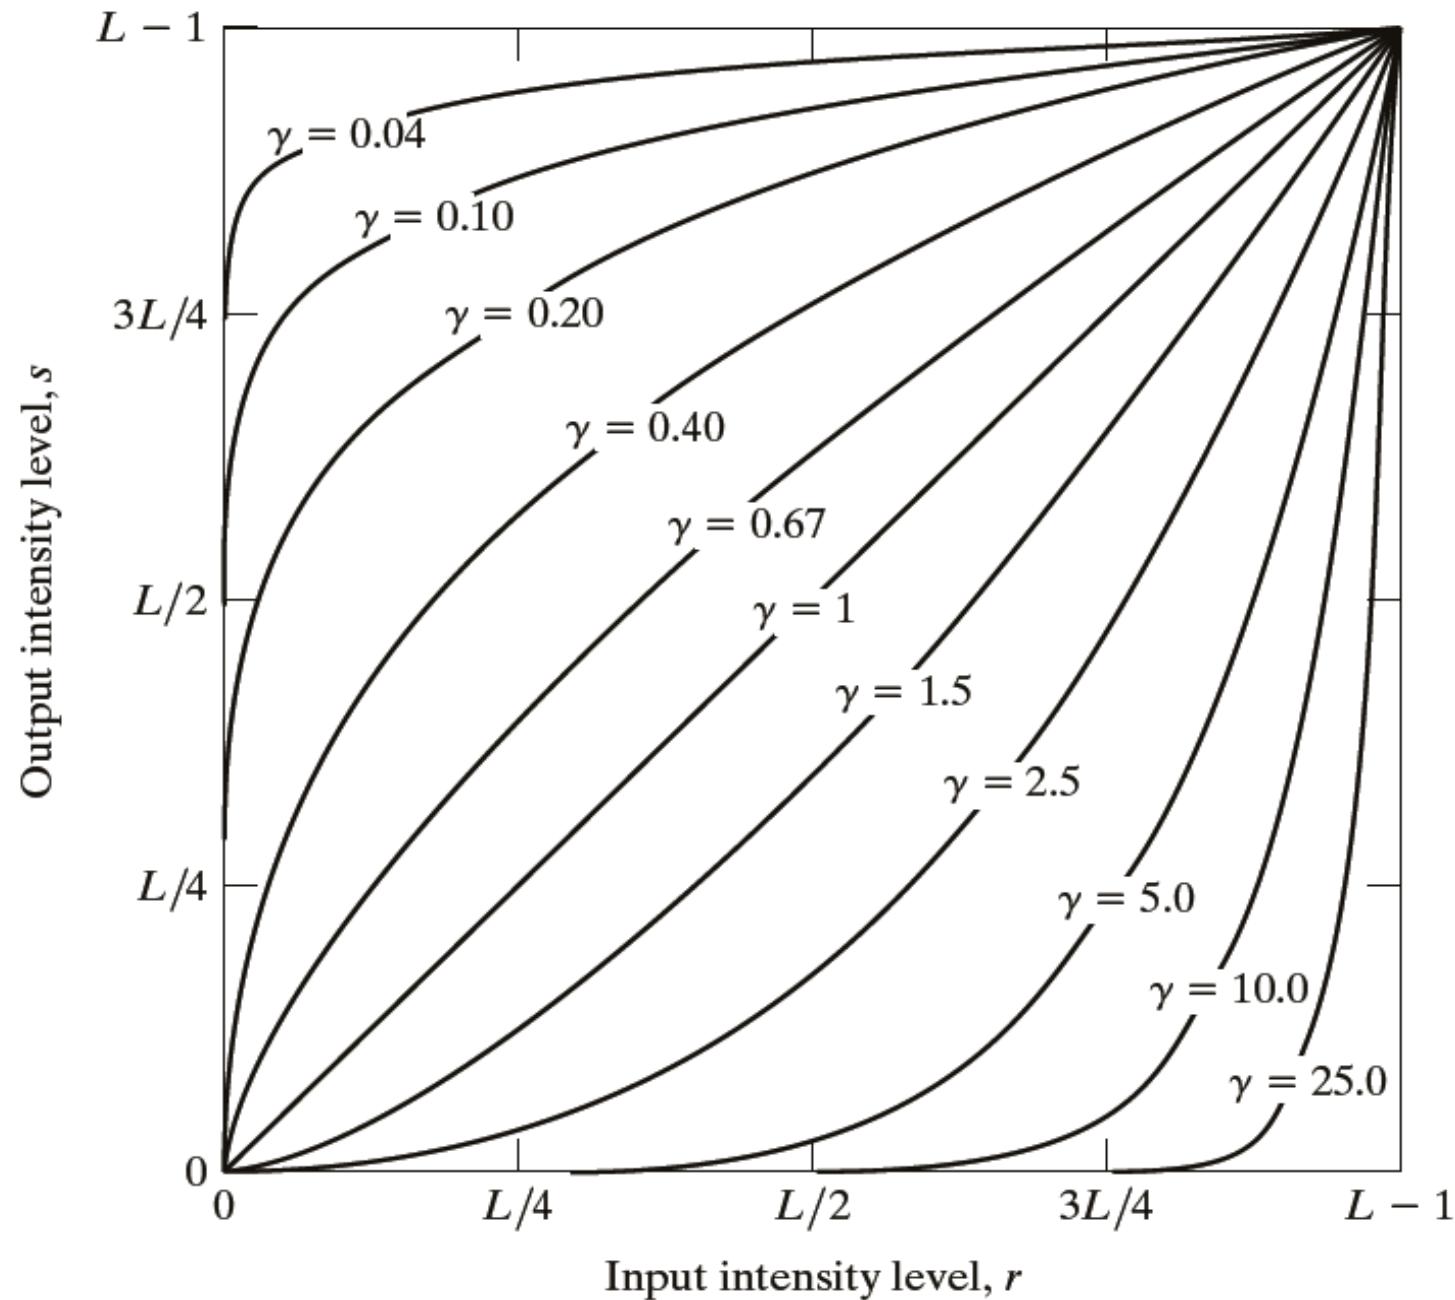
- Logarithm
 - $g(x,y) = c \log [1 + f(x,y)]$, where $c = \text{constant}$
 - Input image has some values very large
 - Remaining values map to few colors (left) \rightarrow details lost
 - Log \rightarrow compress large values, spread small values



Gray-Level Transformations

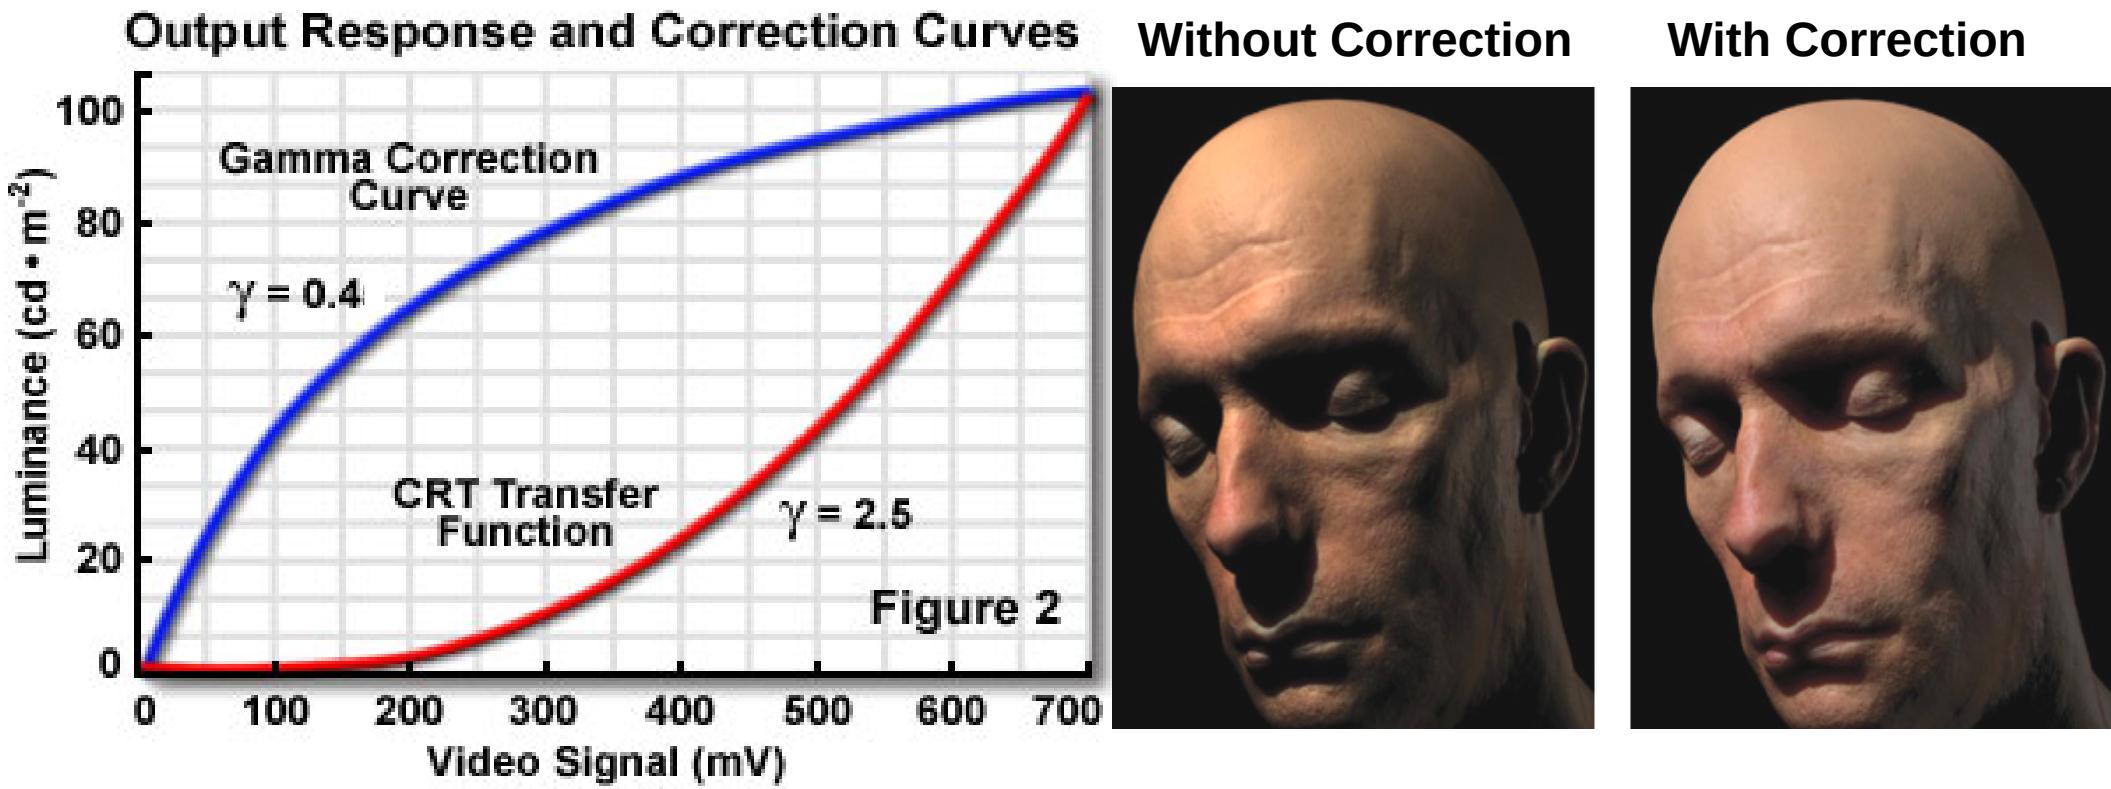
- Power and root transformations

- $s(x,y) = d \left(r(x,y) / c \right)^\gamma$
- For graph →
 - r in $[0, L-1]$
 - $c = L-1$
 - $d = L-1$



Gray-Level Transformations

- Gamma correction for display monitors
 - Mapping function from intensity (data) to voltage (screen appearance) is a power function
 - gamma within [1.8, 2.5]
 - Correction applies power transform: $\text{gamma}' = 1 / \text{gamma}$

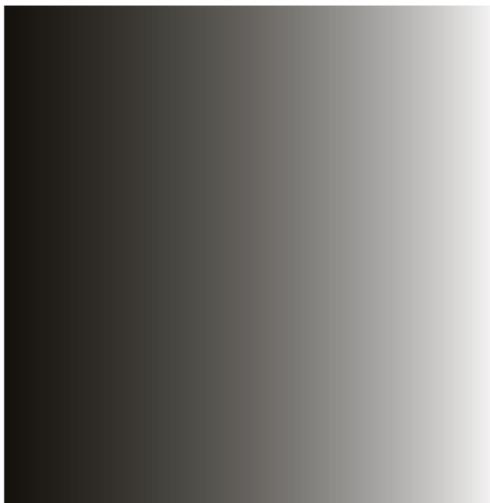


Gray-Level Transformations

Without Correction

With Correction

- Gamma correction



Original image

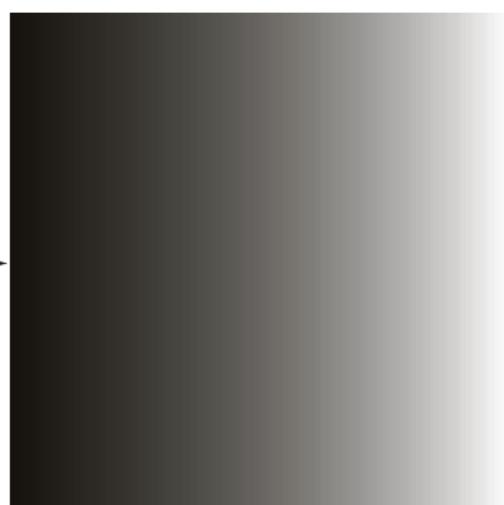


Gamma
correction

Original image as viewed
on monitor



Gamma-corrected image

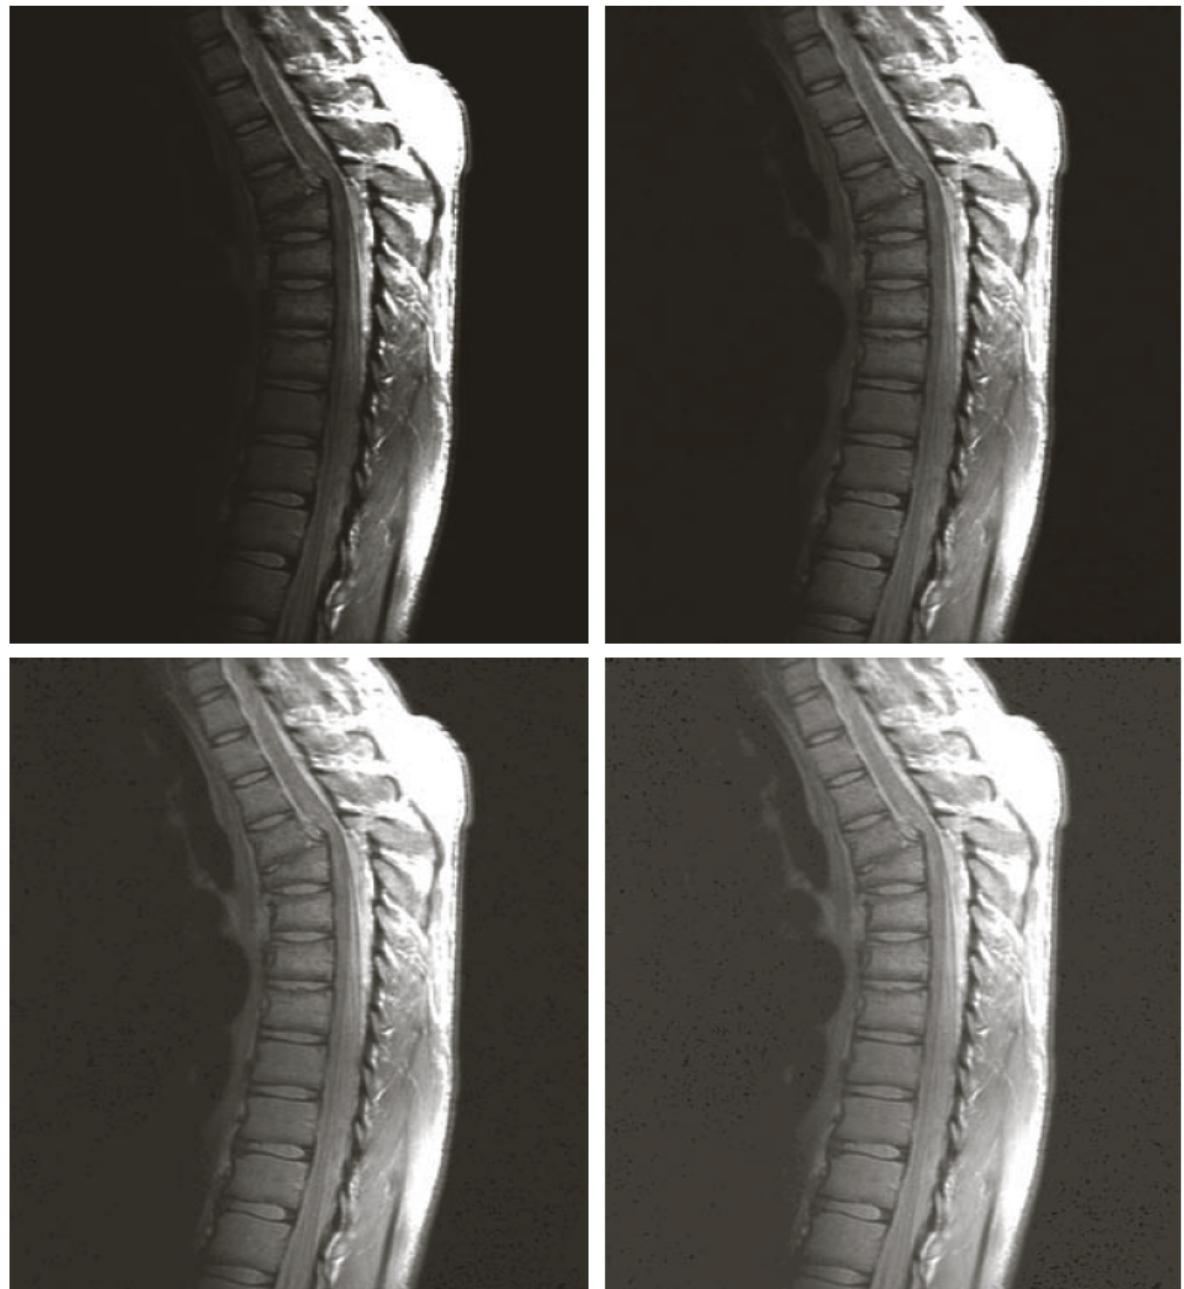


Gamma-corrected image as
viewed on the same monitor

Gray-Level Transformations

- Gamma correction in medical imaging (e.g., MRI)

- gamma =
1.0 , 0.6
0.4 , 0.3



Gray-Level Transformations

- Gamma correction to improve washed-out images

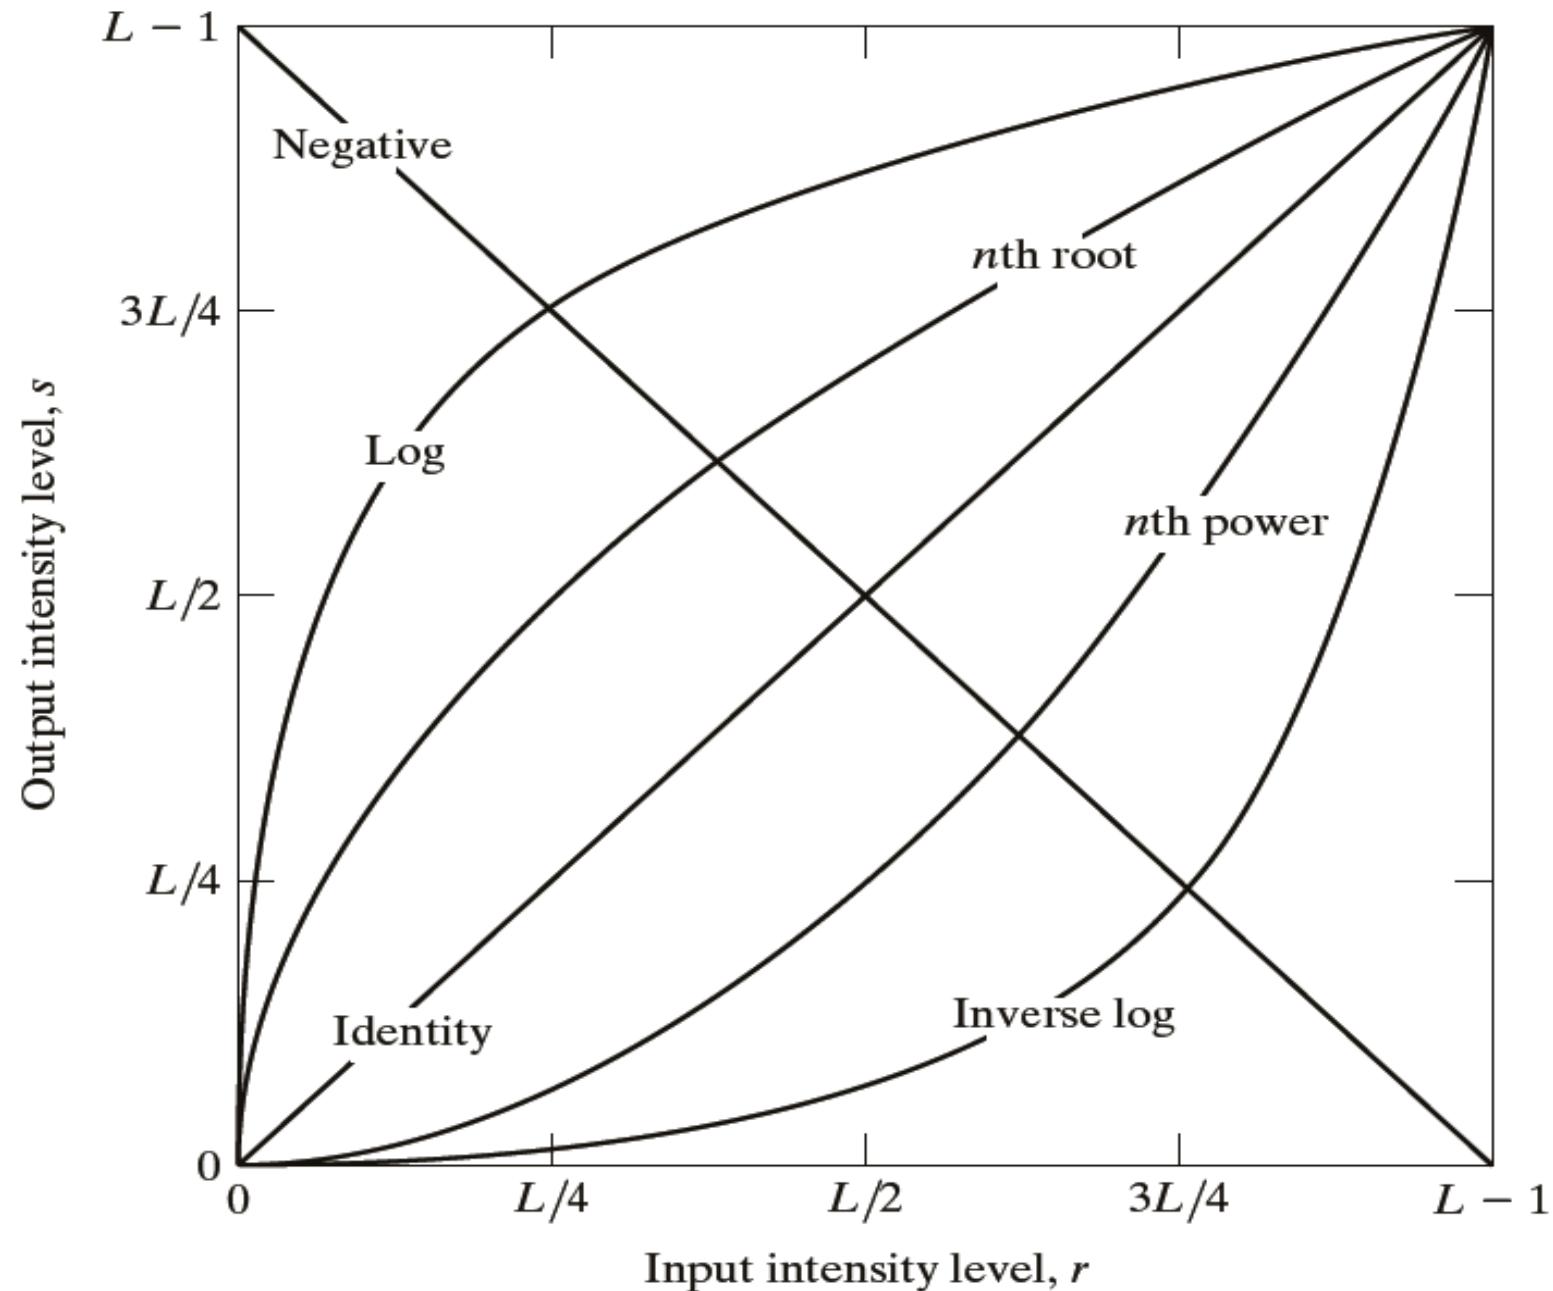
- Input image
= top left
 - For other images,
is gamma
more than 1 or
less than 1 ?



Gray-Level Transformations

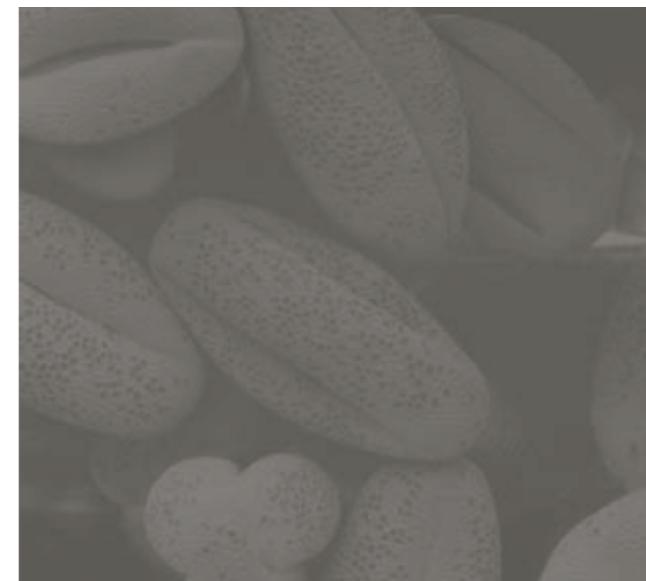
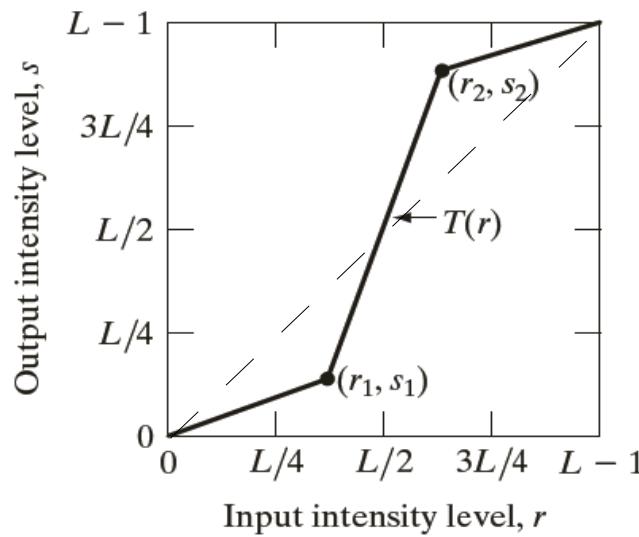
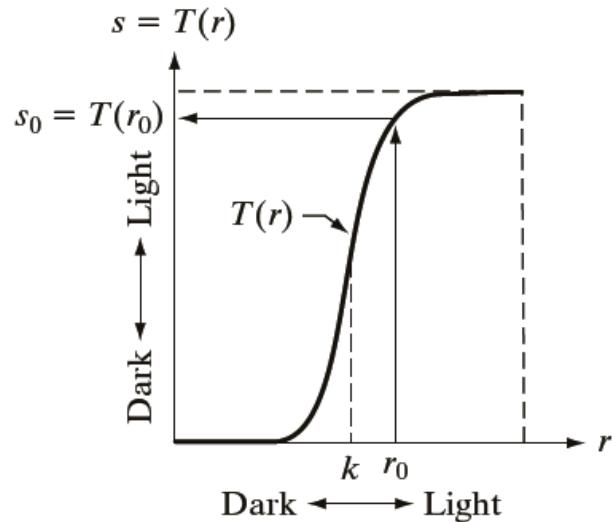
- Typical transformation functions

- Log, exp
- Power, root
 - $n > 1$

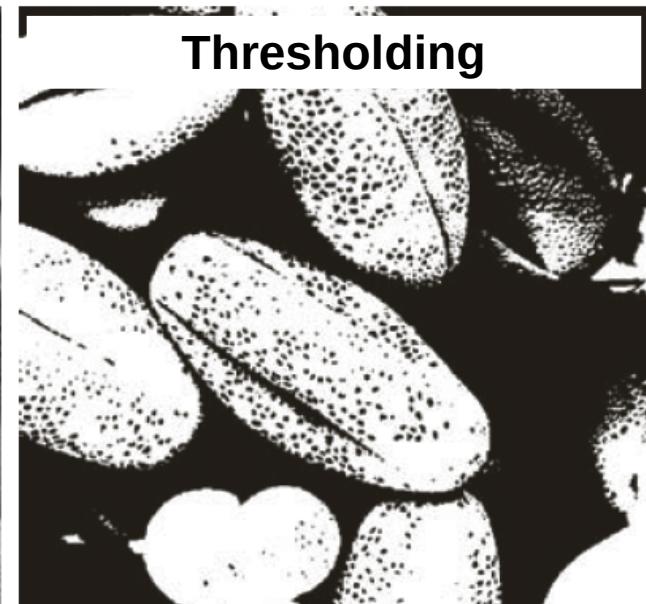
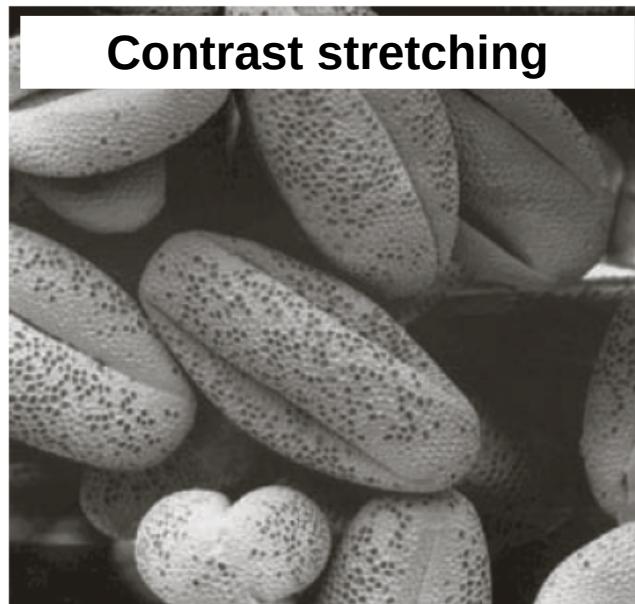


Gray-Level Transformations

- Piecewise-linear transformation functions
 - Easier to design than complex nonlinear functions



- Can model complex transformations
- Contrast stretching



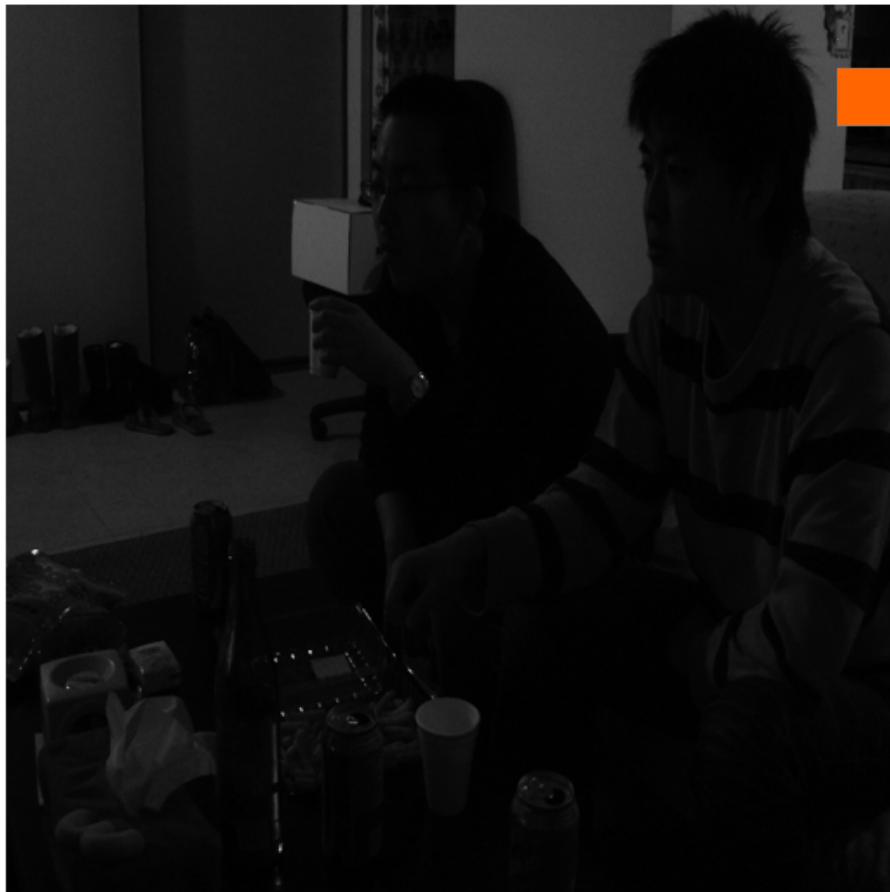
Thresholding

Gray-Level Transformations

- Histogram processing
 - What is a histogram ?
 - Probability theory refresher
 - What can it achieve ?
 - Contrast enhancement
 - Automatic, without tuning parameters

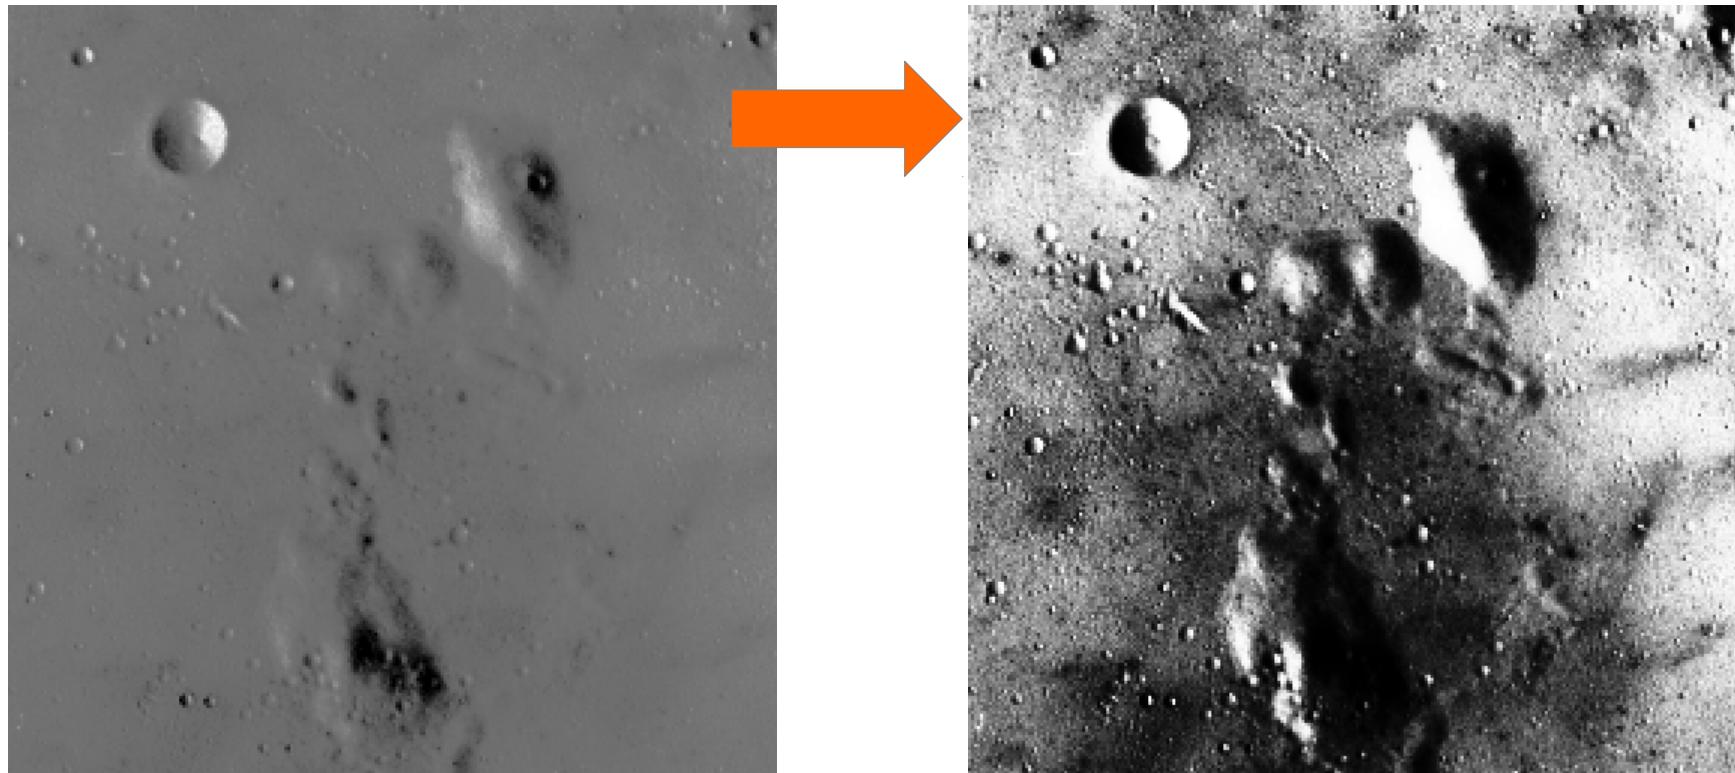
Gray-Level Transformations

- Histogram processing
 - What can it achieve ?



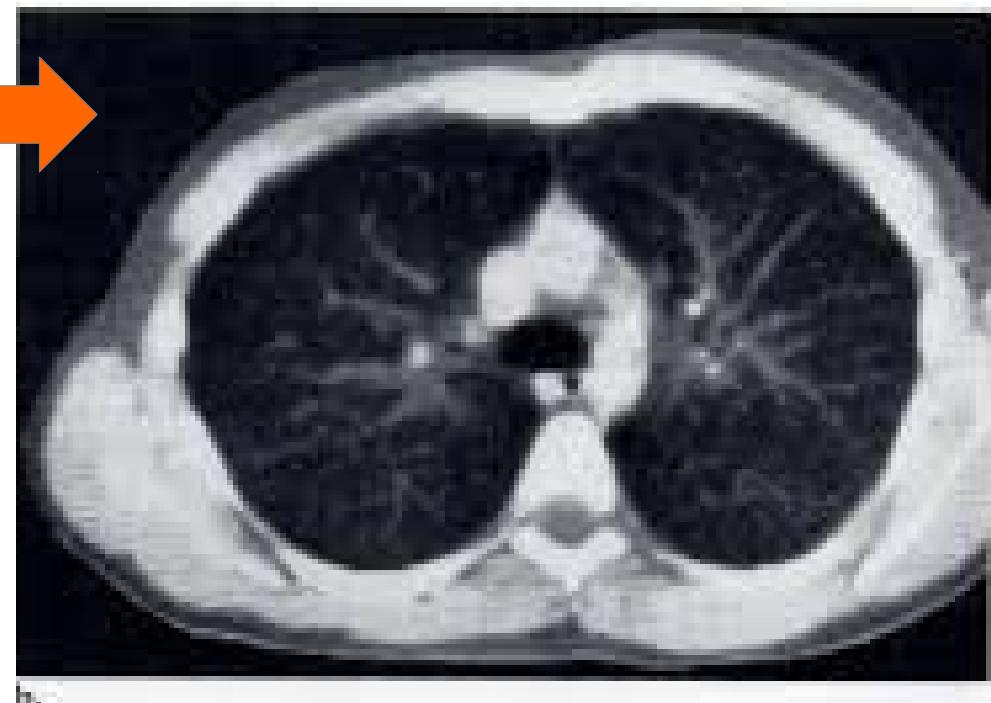
Gray-Level Transformations

- Histogram processing
 - What can it achieve ?



Gray-Level Transformations

- Histogram processing
 - What can it achieve ?



Gray-Level Transformations

- Histogram processing
 - What can it achieve ?







Gray-Level Transformations

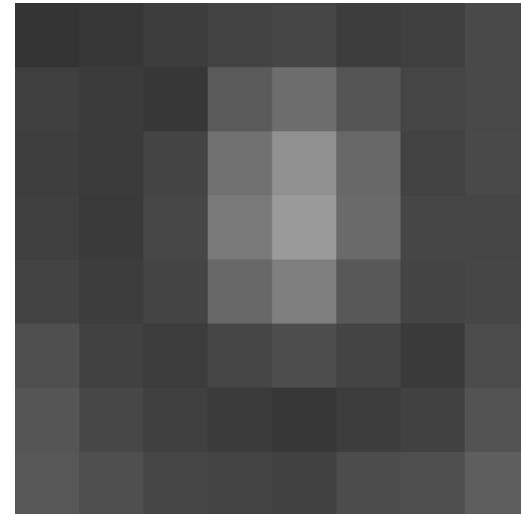
- Histogram processing
 - What is a histogram ?
 - Probability theory refresher (very brief)
 - Probability mass/density function
 - Random variables
 - Random experiments
 - Sample space, events, probability functions
 - Transformation of a random variable

Gray-Level Transformations

- **Random experiment** = an experiment whose outcome isn't certain
 - Flip of a coin
 - Throw of a die
- **Sample space** = set of all possible outcomes of a random experiment
 - Coin flip : { head, tail }
 - Die throw : { 1, 2, 3, 4, 5, 6 }

Gray-Level Transformations

- Random experiment
 - Select a pixel location in an image
- Sample space
 - Set of all pixel locations
$$\{ (1,1), (1,2), \dots, (1,8), (2,1), (2,2), \dots, (2,8), \dots, (8,1), (8,2), \dots, (8,8) \}$$



52	55	61	66	70	61	64	73
63	59	55	90	109	85	69	72
62	59	68	113	144	104	66	73
63	58	71	122	154	106	70	69
67	61	68	104	126	88	68	70
79	65	60	70	77	68	58	75
85	71	64	59	55	61	65	83
87	79	69	68	65	76	78	94



Gray-Level Transformations

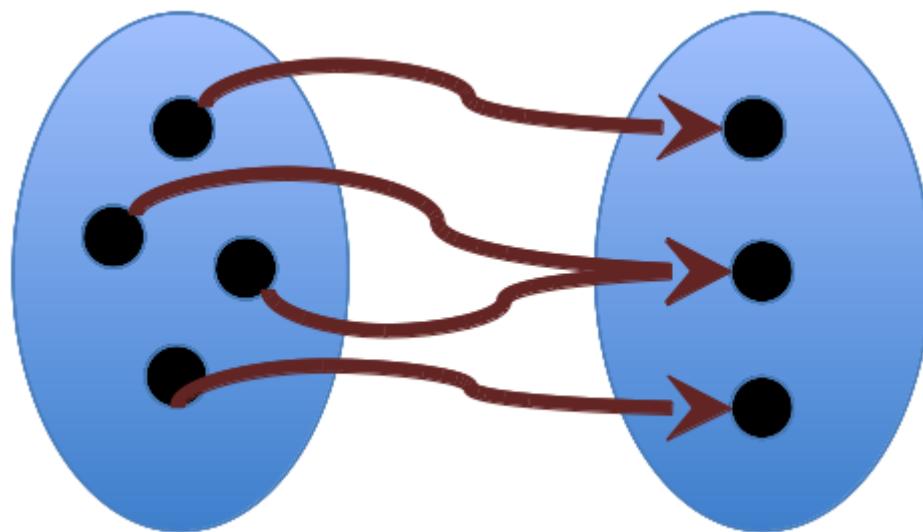
- **Event** is a subset of a sample space
 - Event of getting a head on a coin toss = { head }
 - Event of getting an even number on a die roll = ?
= { 2, 4, 6 }
 - Event of getting a pixel in the left half of the image
- **Event space** = set of all possible events

Gray-Level Transformations

- **Probability function $P(\cdot)$**
 - A probability function on **sample space Ω** assigns every event A in Ω a number in $[0,1]$ s.t.
 - $P(\Omega) = 1$
 - $P(A \cup B) = P(A) + P(B)$ when $A \cap B = \emptyset$
 - $P(A)$ is the **probability** that **event A occurs**
 - Examples
 - Probability function for coin toss = ?
 - Probability function for die roll = ?

Gray-Level Transformations

- **Random variable** = a function
 - $X : \Omega \rightarrow \mathbb{R}$
 - Domain = sample space
 - Range = set of real numbers



Gray-Level Transformations

- Random variable is an **abstraction**
 - We don't care about outcomes (directly)
 - We only care about values mapped to outcomes
 - Die example
 - We don't care about die value, just even or odd
 - Image example
 - We don't care about pixel location, just pixel intensity

Gray-Level Transformations

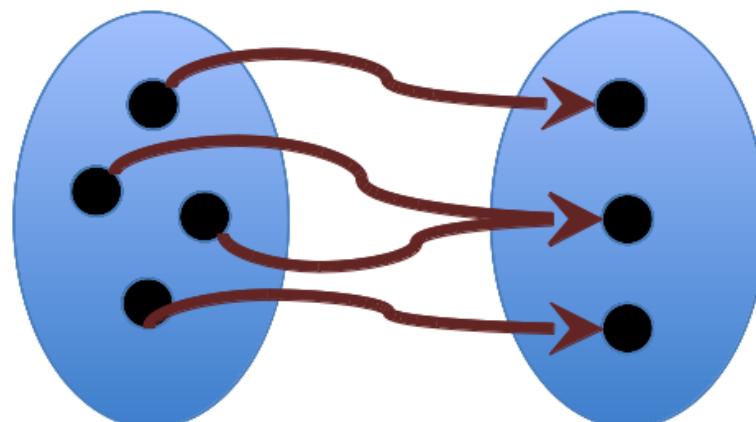
- Defining events using random variables
 - Example: Sum of a pair of dice
 - Sample Space: $\Omega = \{ (i,j) : 1 \leq i,j \leq 6 \}$
 - Probability function: $P((i,j)) = 1/36$
 - Random variable: $X((i,j)) = i + j$
 - Event $\{ X=a \} = \{ w \in \Omega : X(w)=a \}$
 - $P(X=5)$
= $P(\{(1,4), (2,3), (3,2), (4,1)\})$
= $P((1,4)) + P((2,3)) + P((3,2)) + P((4,1))$
= $4 / 36$

Gray-Level Transformations

- Defining events using random variables
 - Example:
Intensity at a randomly-chosen pixel in 8x8 image $f(x,y)$
 - Sample Space: $\Omega = \{ (x,y) : 1 \leq x, y \leq 8 \}$
 - Probability function: $P((x,y)) = 1/64$
 - Random variable: $Z((x,y)) = f(x,y)$
 - Event $\{ Z > 100 \} = \{ w \in \Omega : Z(w) > 100 \}$

Gray-Level Transformations

- **Discrete random variable (RV)**
 - Maps outcomes to values in a countable set
 - **Probability mass function** $p : R \rightarrow [0,1]$
 - Probability that the RV takes a value 'a'
 - $p(a) = p(X=a) = \sum_{w: X(w)=a} P(w)$
 - **Cumulative distribution function** $F : R \rightarrow [0,1]$
 - Probability that the RV takes a value **less than** x
 - $F(a) = p(X \leq a) = \sum_{w: X(w) \leq a} P(w)$
 - $F(a) = p(X \leq a) = \sum_{b: b \leq a} p(X=b)$



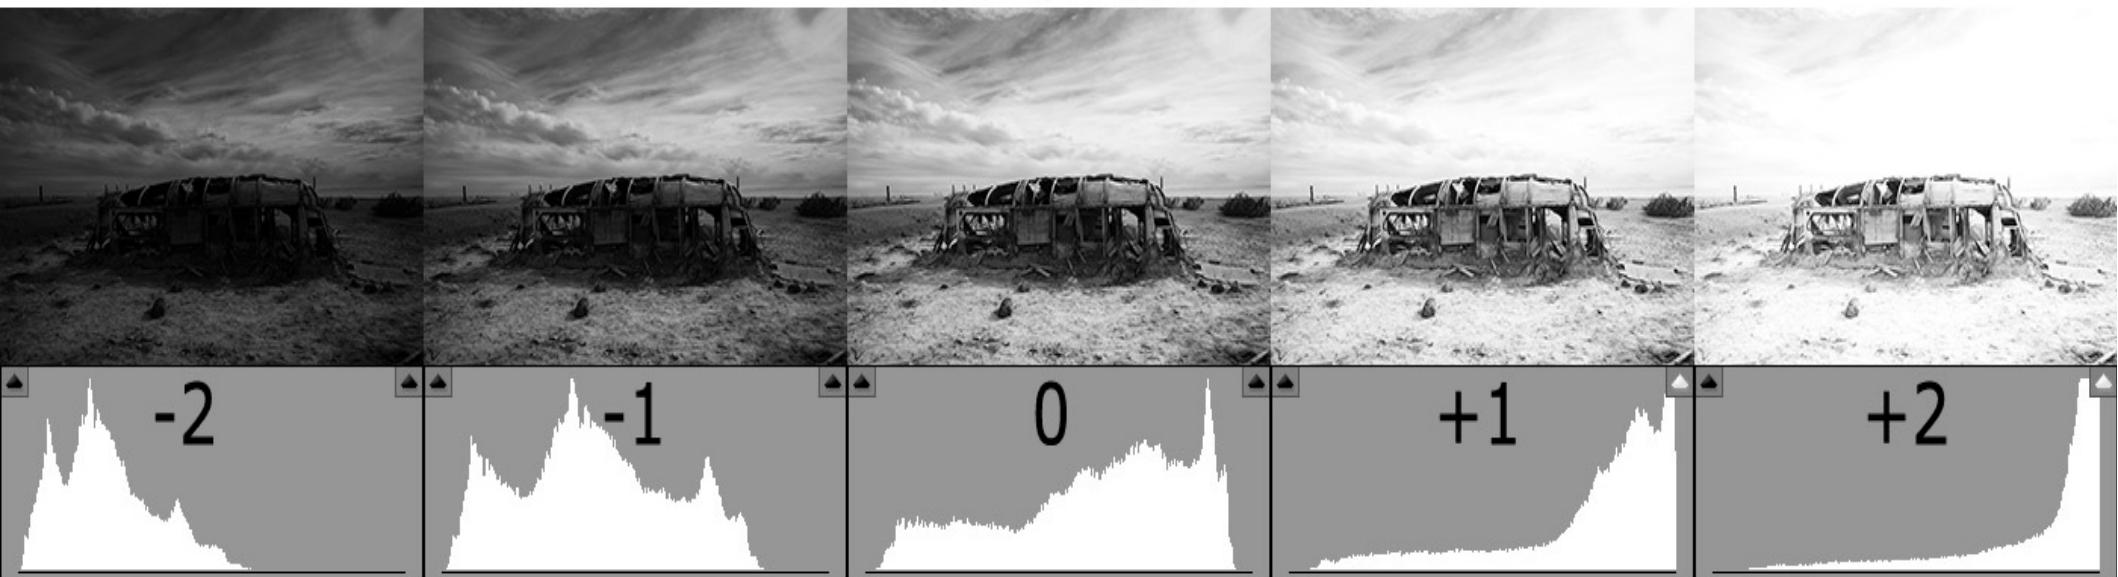
Gray-Level Transformations

- **Continuous random variable (RV)**
 - Maps outcomes to values in an uncountable set
 - Sample space itself is uncountable
 - e.g., continuous image function $X(\cdot)$ defined on $[0,1] \times [0,1]$
 - **Cumulative distribution function** $F : \mathbb{R} \rightarrow [0,1]$
 - Probability that the RV takes a value less than c
 - $F(c) = p(x < c) = \int_{w : X(w) \leq c} P(w) dw$
 - $F(c) = p(x < c) = \int_{x < c} p(x) dx$
 - **Probability density function** $p : \mathbb{R} \rightarrow \mathbb{R}$
 - $p(\cdot)$ is derivative of $F(\cdot)$
 - $p(x)$ is NOT probability ($p(x)$ can be > 1)
 - Events, of interest, are subsets of the form $\{x : a < x < b\}$
 - $p(a < x < b) = \int_{a < x < b} p(x) dx$

Gray-Level Transformations

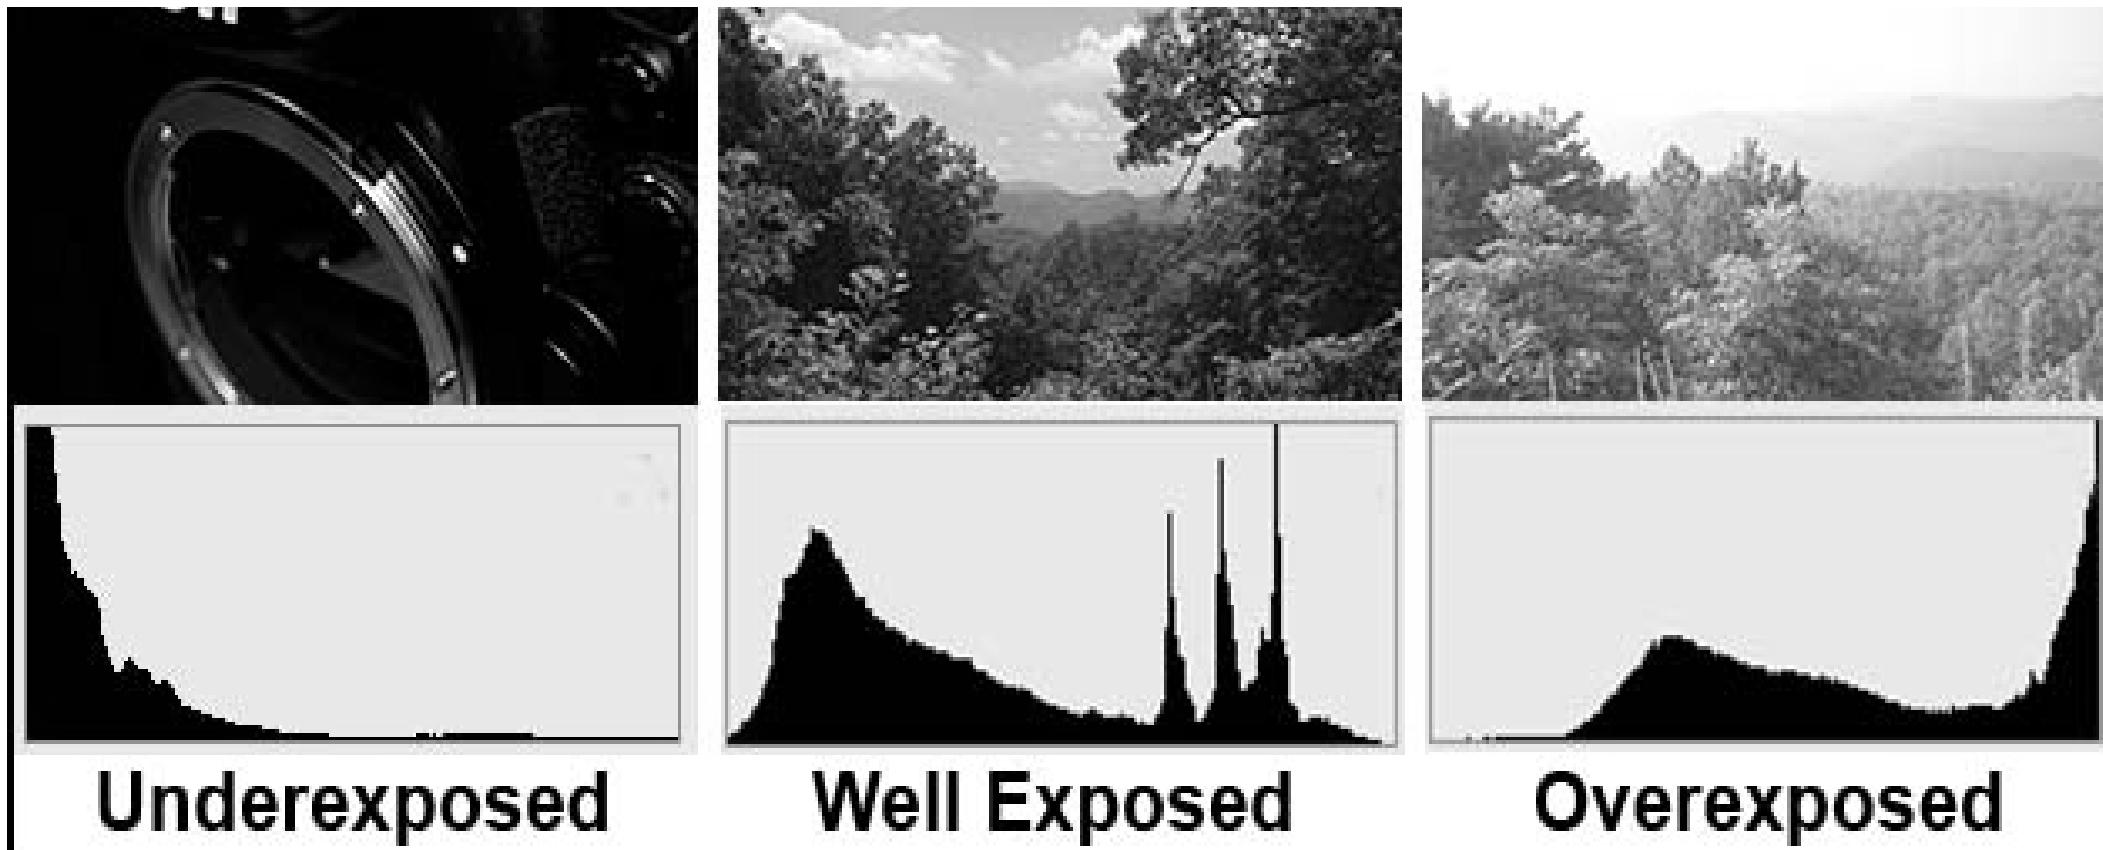
- Histograms are related to photograph exposure
 - Aperture size (amount of light per unit time)
 - Shutter speed (time of exposure)

EXPOSURE



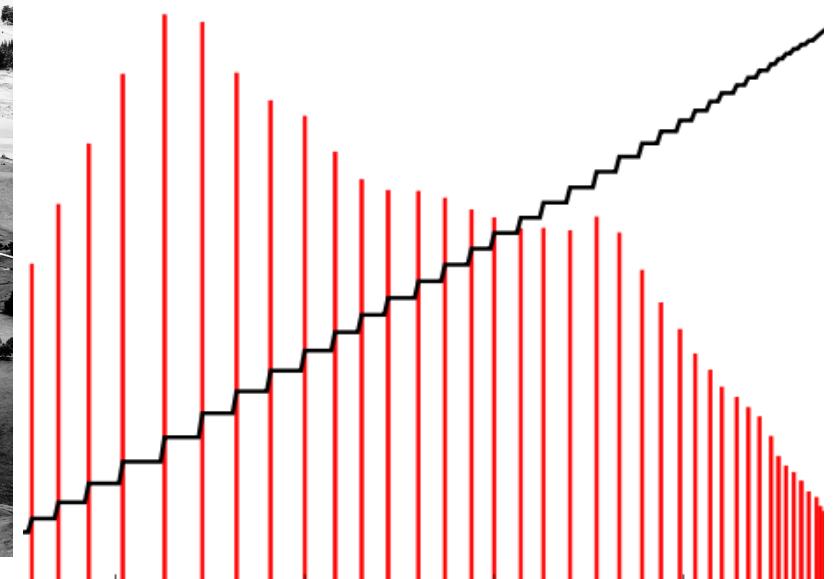
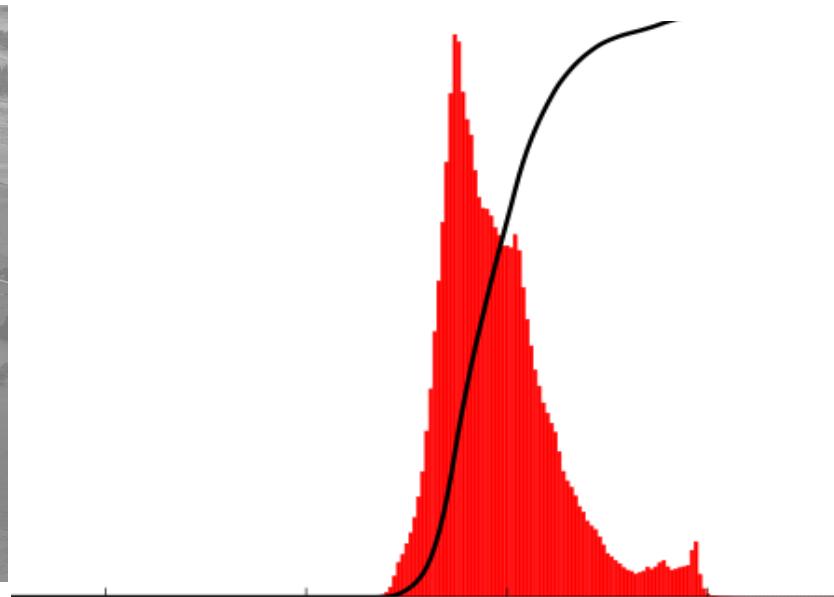
Gray-Level Transformations

- Histograms are related to photograph exposure
 - Aperture size
 - Shutter speed



Gray-Level Transformations

- Histogram equalization : Motivation

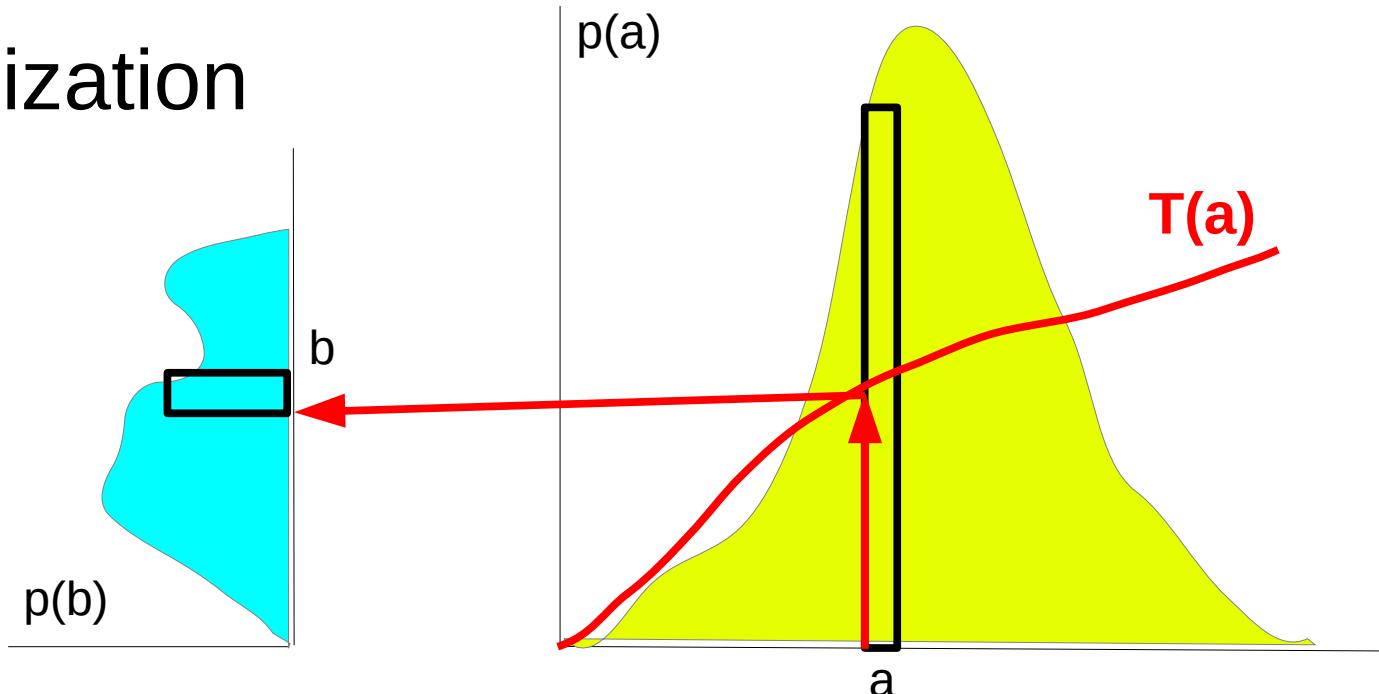


Gray-Level Transformations

- Histogram equalization
 - Assume **continuous** distributions (and image)
 - Given histogram $p(a)$
 - **Design function $T(a) = b$ such that $p(b) = \text{uniform distribution}$**
 - Assume $0 < a < 1, 0 < b < 1$
 - Then, we want $p(b) = 1$
 - Consider a **monotonically increasing** intensity-transformation function $T(\cdot)$

Gray-Level Transformations

- Histogram equalization



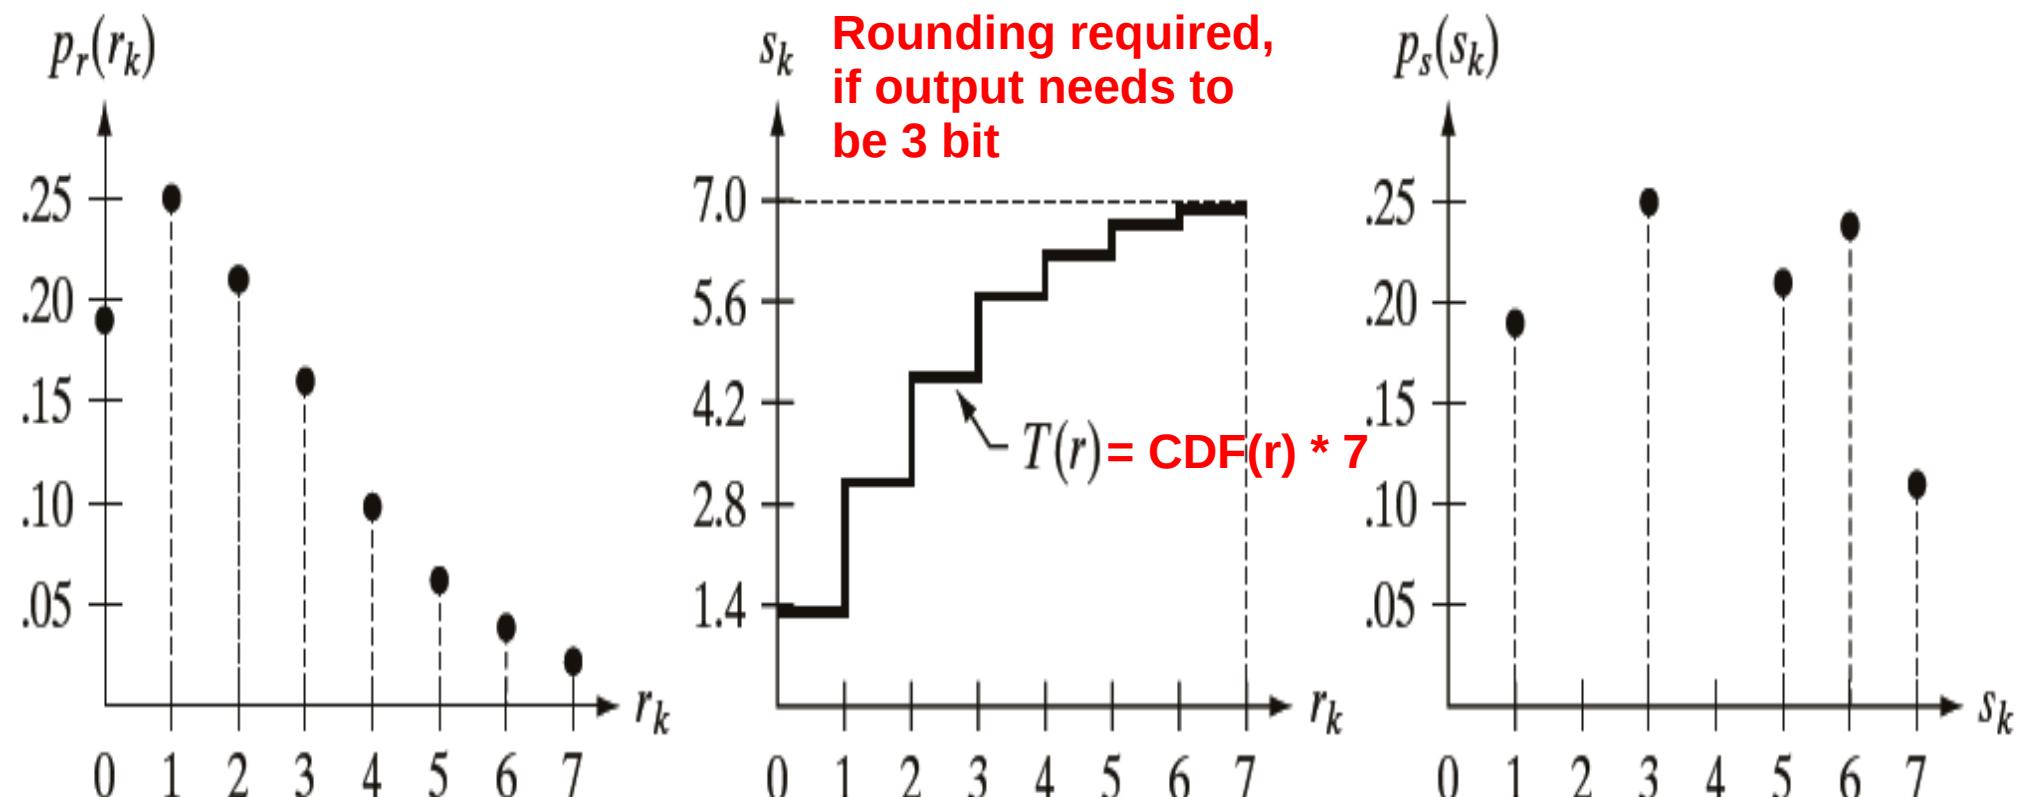
- Let $b = T(a)$
- Mass conservation** implies: $p(a) da = p(b) db$
- Integrate left hand side from 0 to a'
- Integrate right hand side from $0 = T(a') = b'$
- Left hand side = $CDF_a(a')$
- Right hand side = $CDF_b(b')$

Gray-Level Transformations

- Histogram equalization
 - If we want $p(b) = 1$, then $CDF_b(b') = b'$
 - So, $CDF_a(a') = CDF_b(b') = b'$
 - So, transformation function $T(.) =$
CDF of intensities in original image $CDF_a(.)$

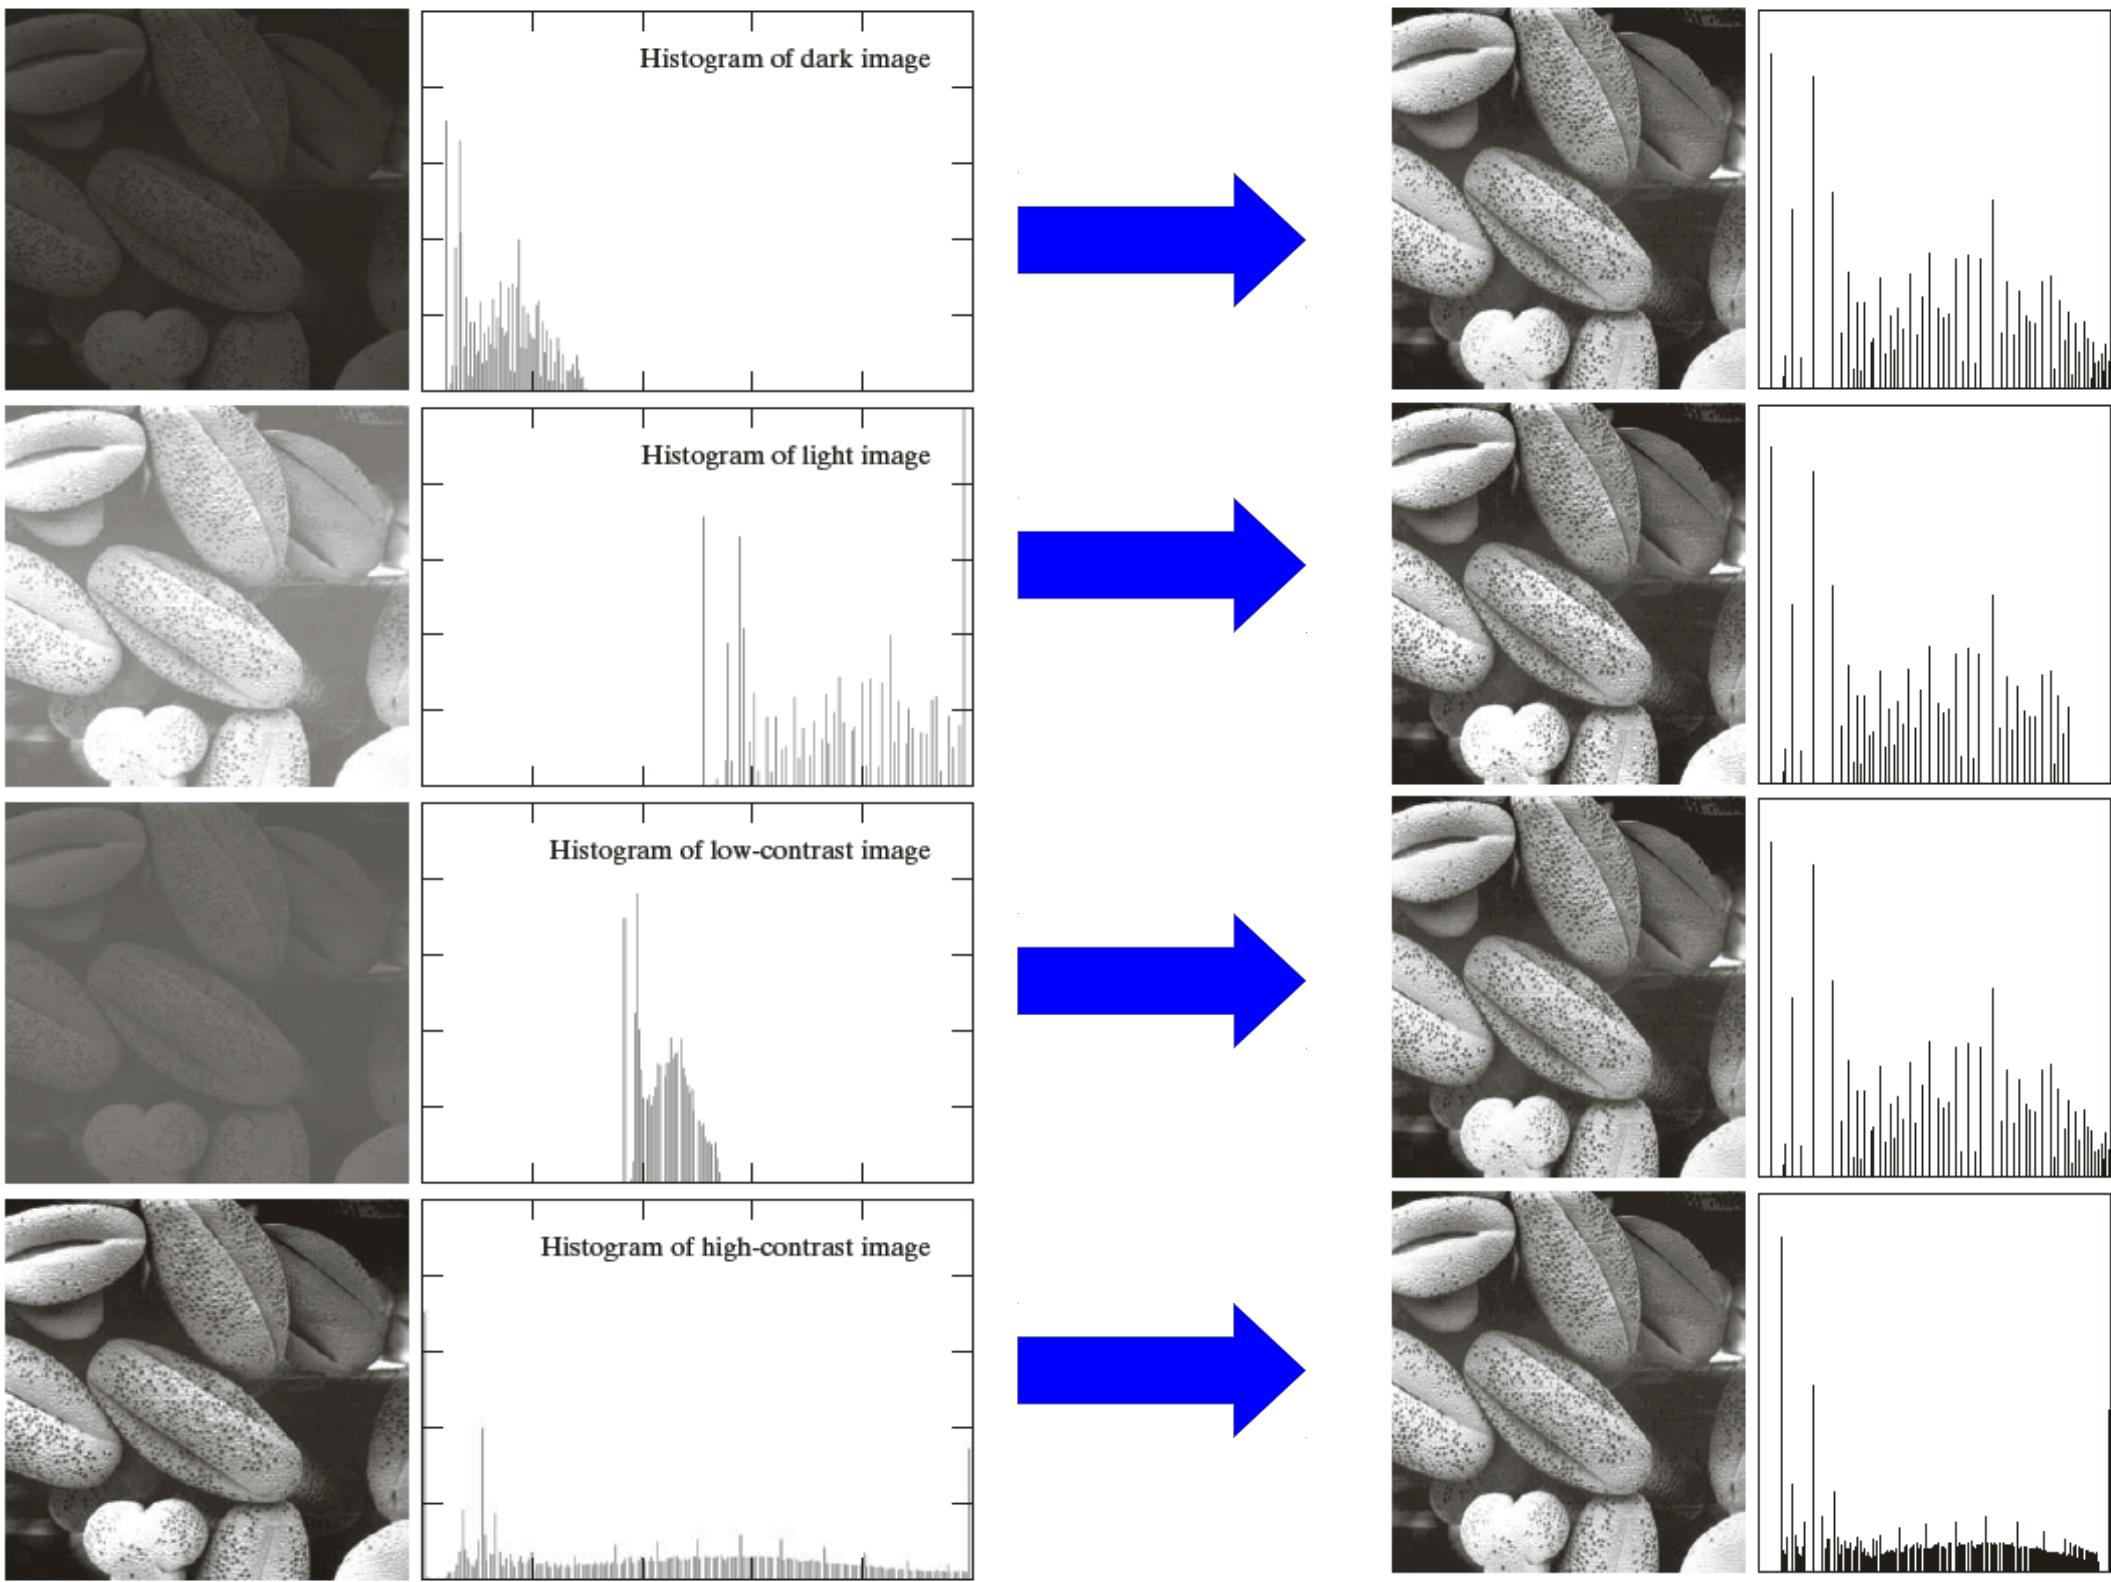
Gray-Level Transformations

- Histogram equalization example (3-bit image)



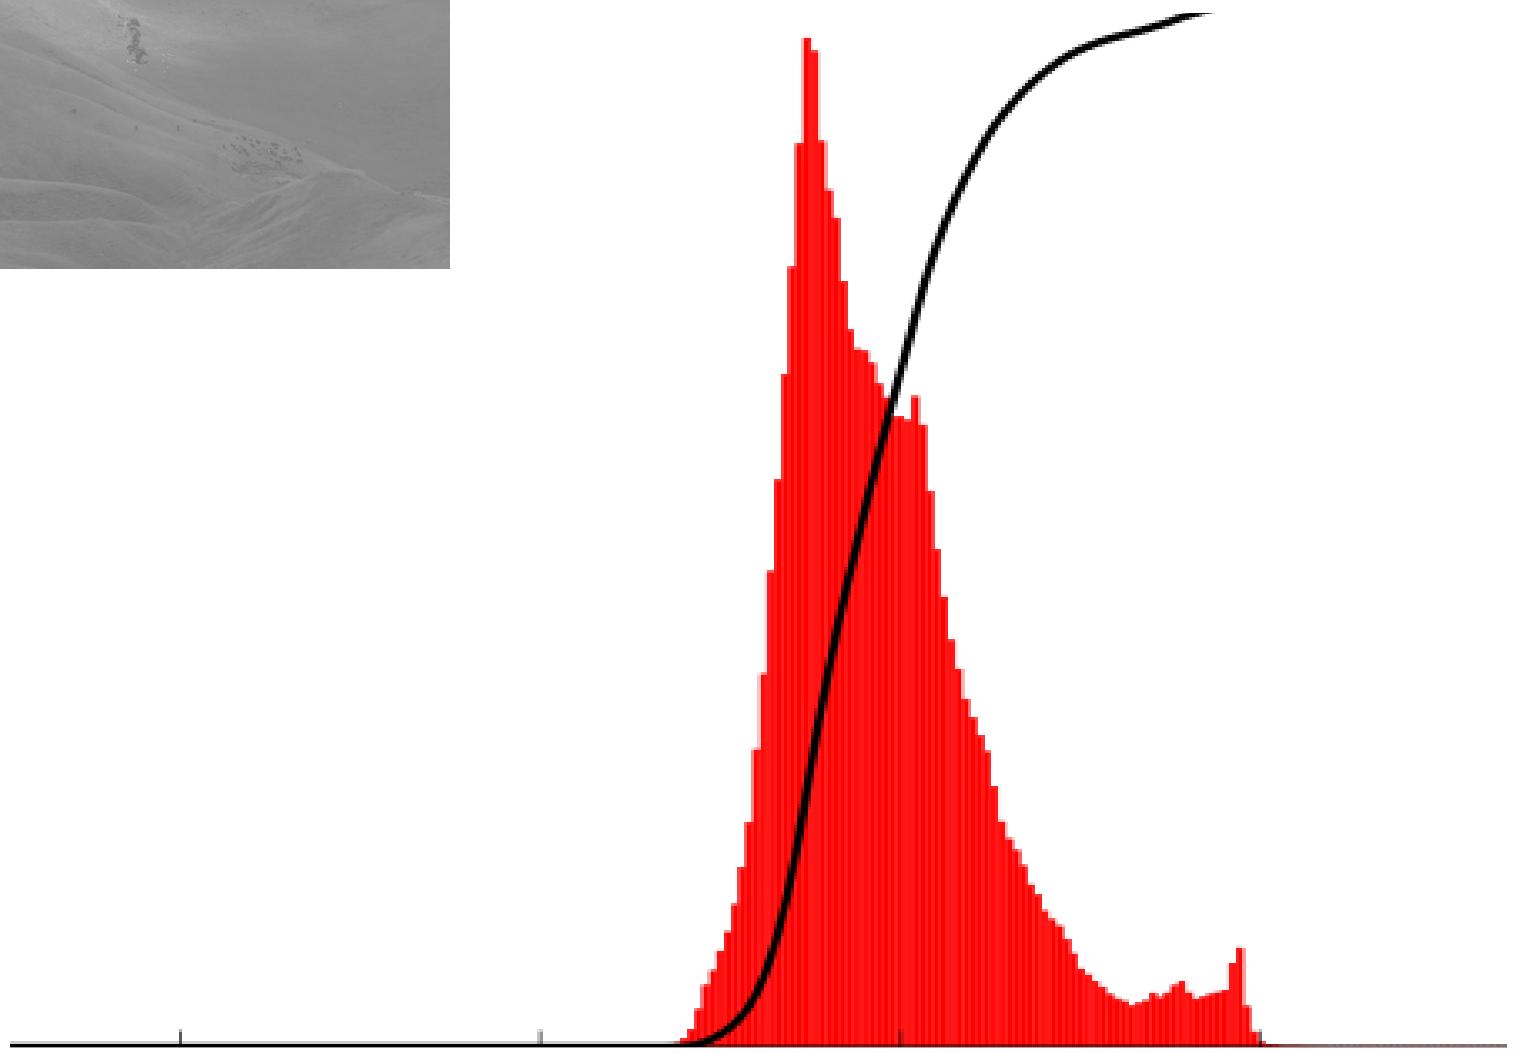
a b c

FIGURE 3.19 Illustration of histogram equalization of a 3-bit (8 intensity levels) image. (a) Original histogram. (b) Transformation function. (c) Equalized histogram.



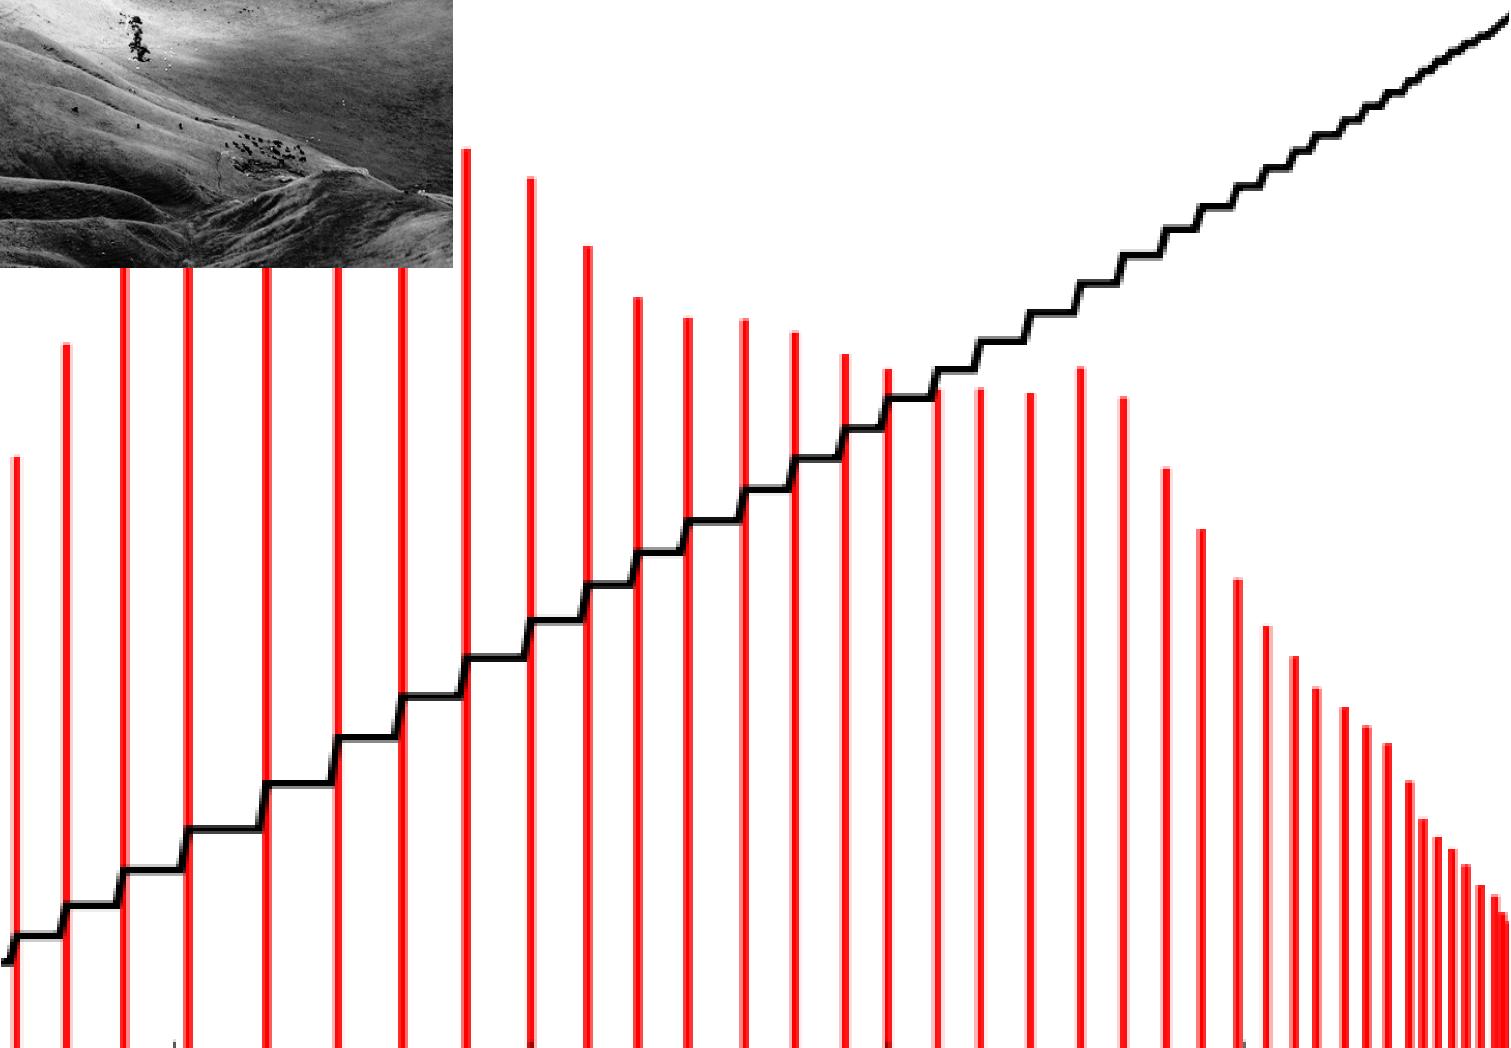


Transformations





Transformations



Gray-Level Transformations

- Why is the output histogram NOT exactly uniform ?
 - Because of discretization of space and intensity
 - Theoretically
 - Contradicts the assumption of continuous distributions (i.e., continuous spatial domain, continuous intensities)
 - We used that to derive transformation for histogram equalization
 - Practical problem (1 of 2)
 - Imagine 8-bit images → 256 intensities
 - Imagine an input image containing only 100 different intensities
 - Uniform histogram must contain **some** pixels for each intensity
 - How can we map 100 values to 256 values ?
 - No systematic and straightforward way to do that

Gray-Level Transformations

- Why is the output histogram NOT exactly uniform ?
 - Because the set of intensities in digital images is discrete
 - Practical problem (2 of 2)
 - Imagine 8-bit images → 256 intensities
 - Imagine an input image containing 256 different intensities and 2560 pixels
 - Uniform histogram must contain **an equal number of (10) pixels** for each intensity
 - If some intensity “a” in input image has 20 pixels, then intensity $b=T(a)$ in output image will have ≥ 20 pixels
 - How can we decide which 10 of those ≥ 20 pixels should be assigned intensity “b” ?
 - No systematic and straightforward way to do that

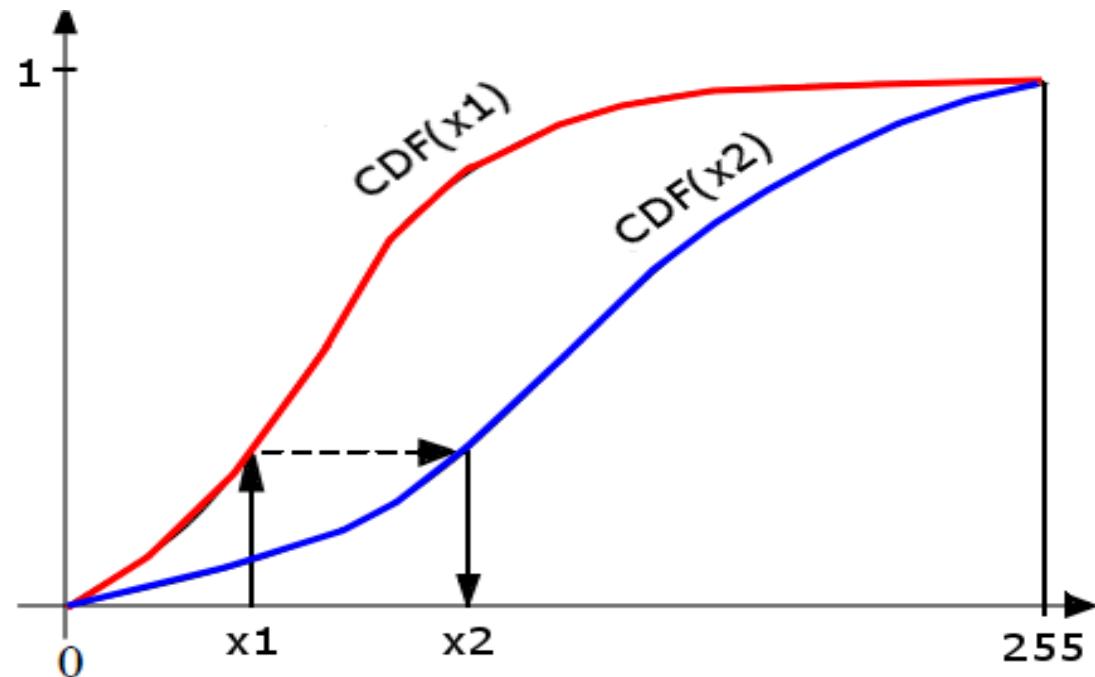
Gray-Level Transformations

- Histogram matching
 - Sometimes, we don't want to equalize
 - But, we want to match the histogram of one image to another chosen histogram (NOT uniform)
 - e.g., for standardizing the intensity
 - In that case,

$$CDF_a(a) = CDF_b(b)$$

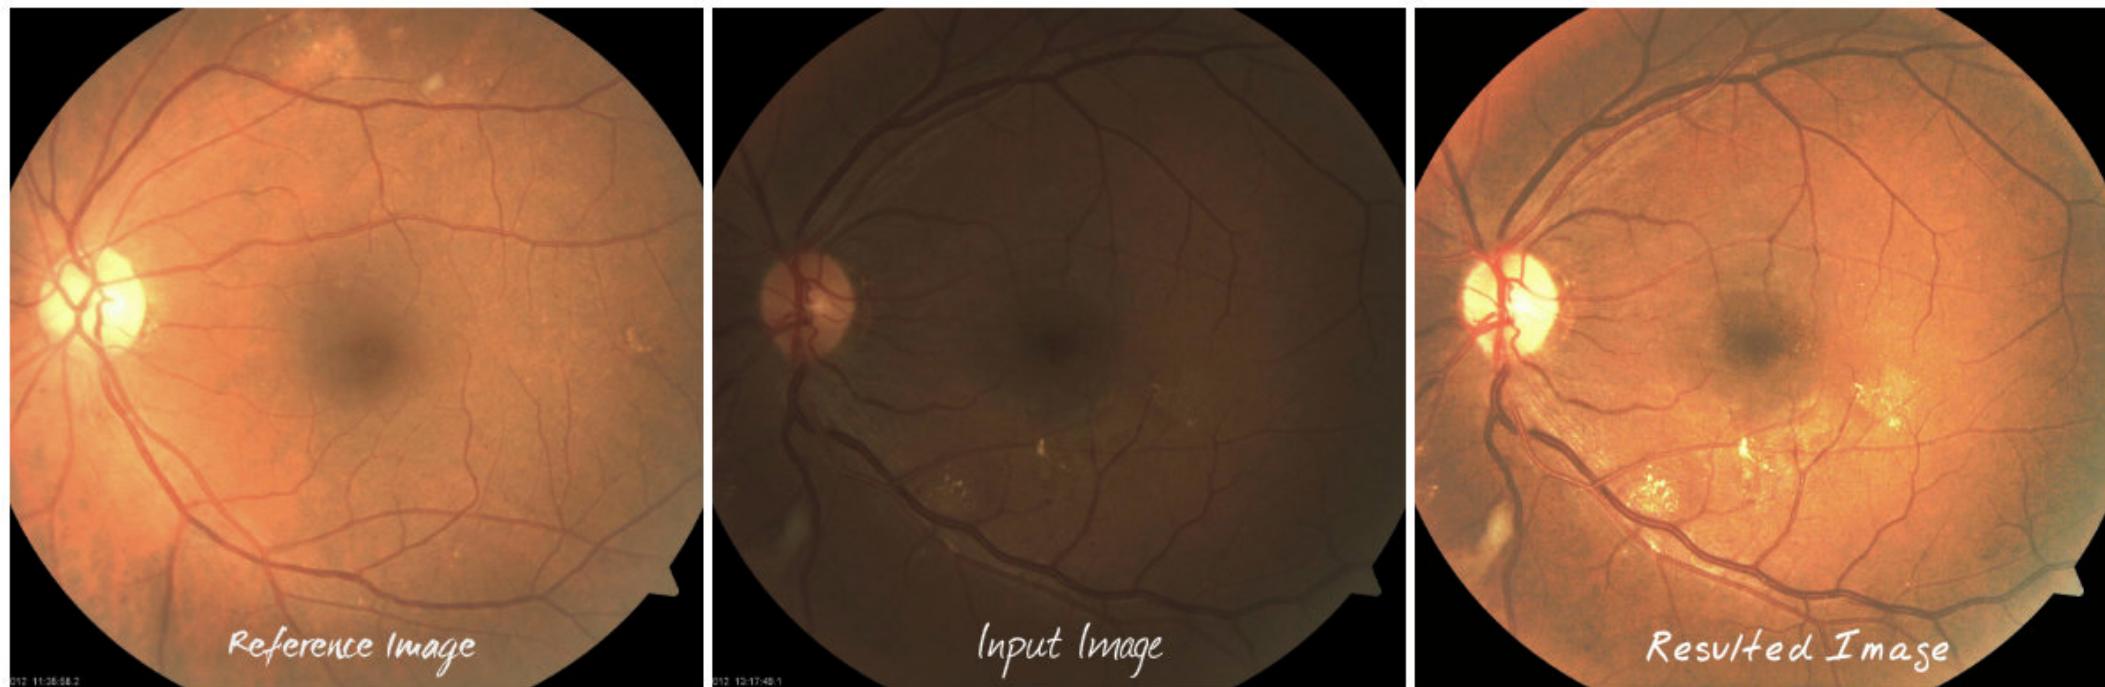
implies

$$b = CDF_b^{-1}(CDF_a(a))$$



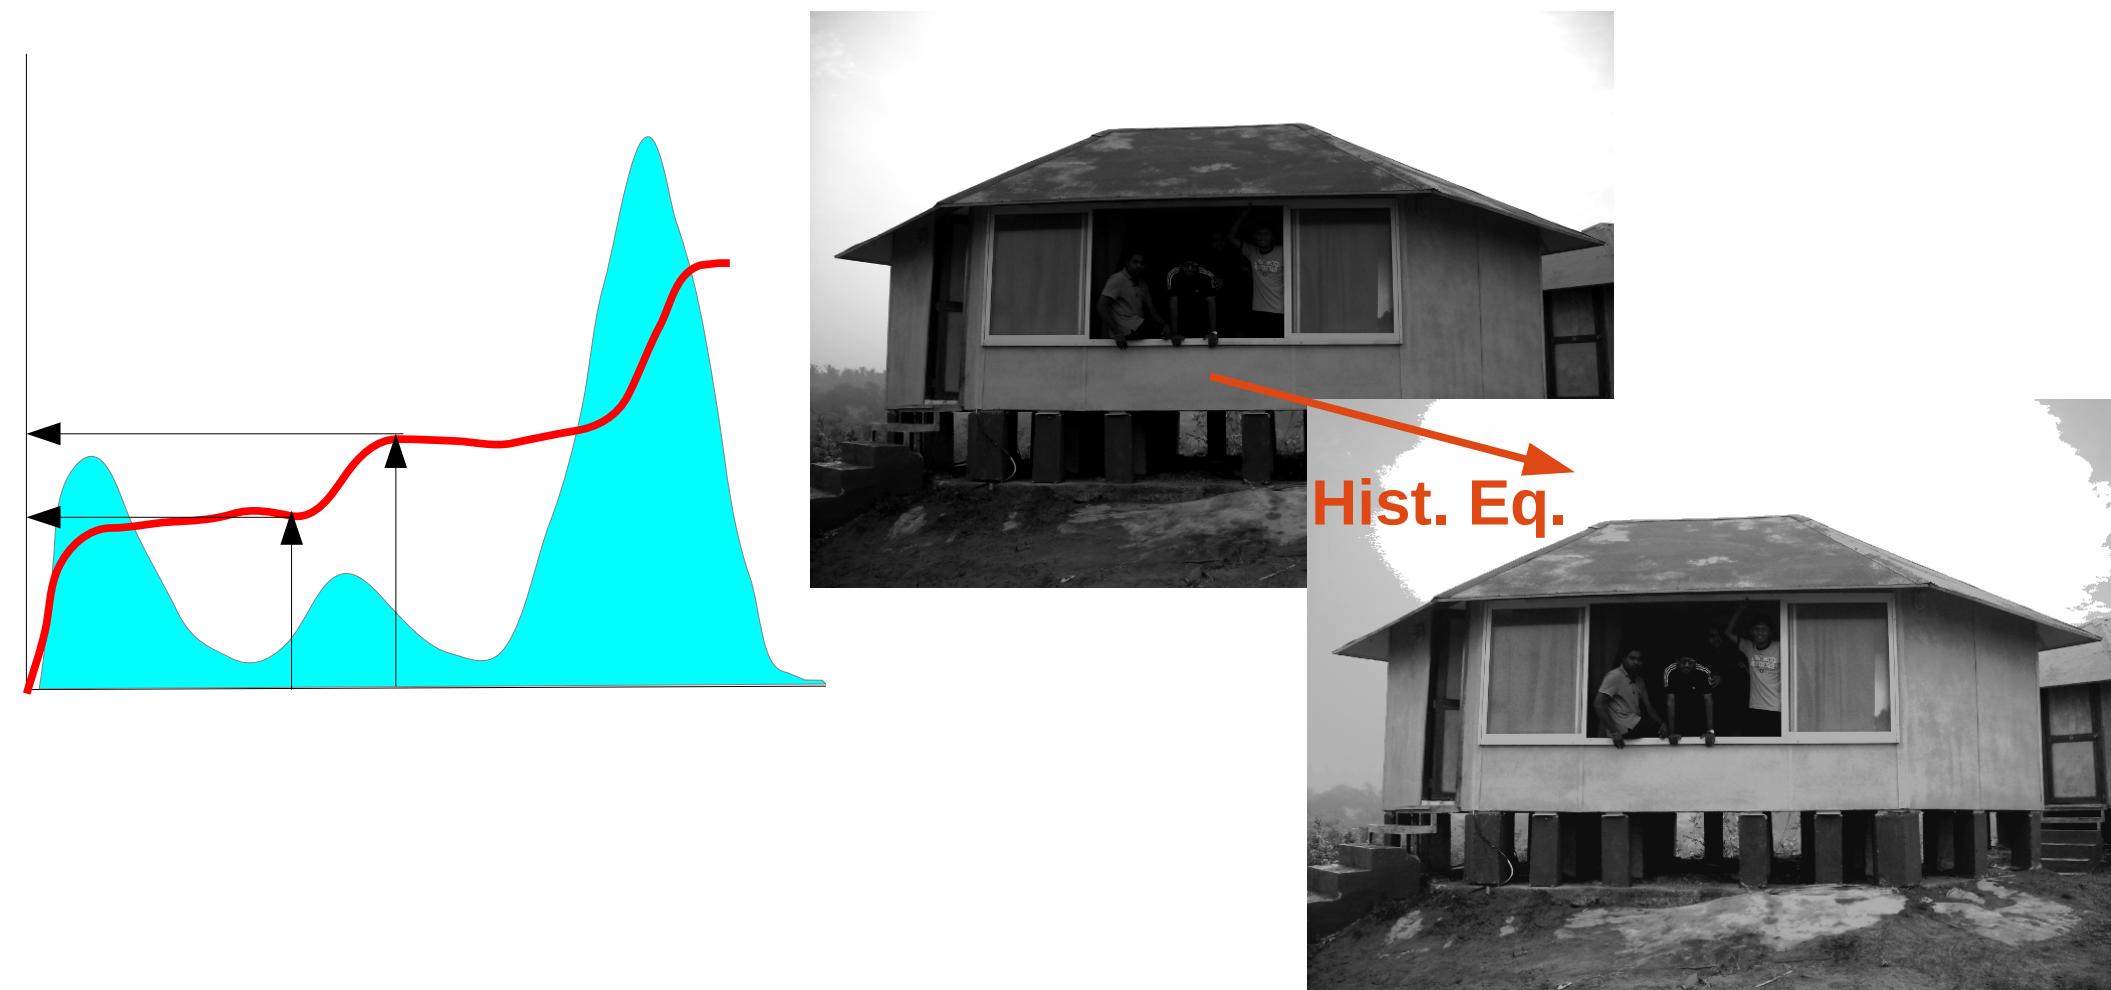
Gray-Level Transformations

- Histogram matching



Gray-Level Transformations

- Limitations of histogram equalization
 - Images has lots of dark pixels or / and lots of bright pixels
 - Then, the “middle” intensities don't get contrast enhanced



Gray-Level Transformations

- Limitations of histogram equalization
 - Images has lots of dark pixels or / and lots of bright pixels
 - Then, the “middle” intensities don't get contrast enhanced
 - How can we get better output ? e.g., top right image

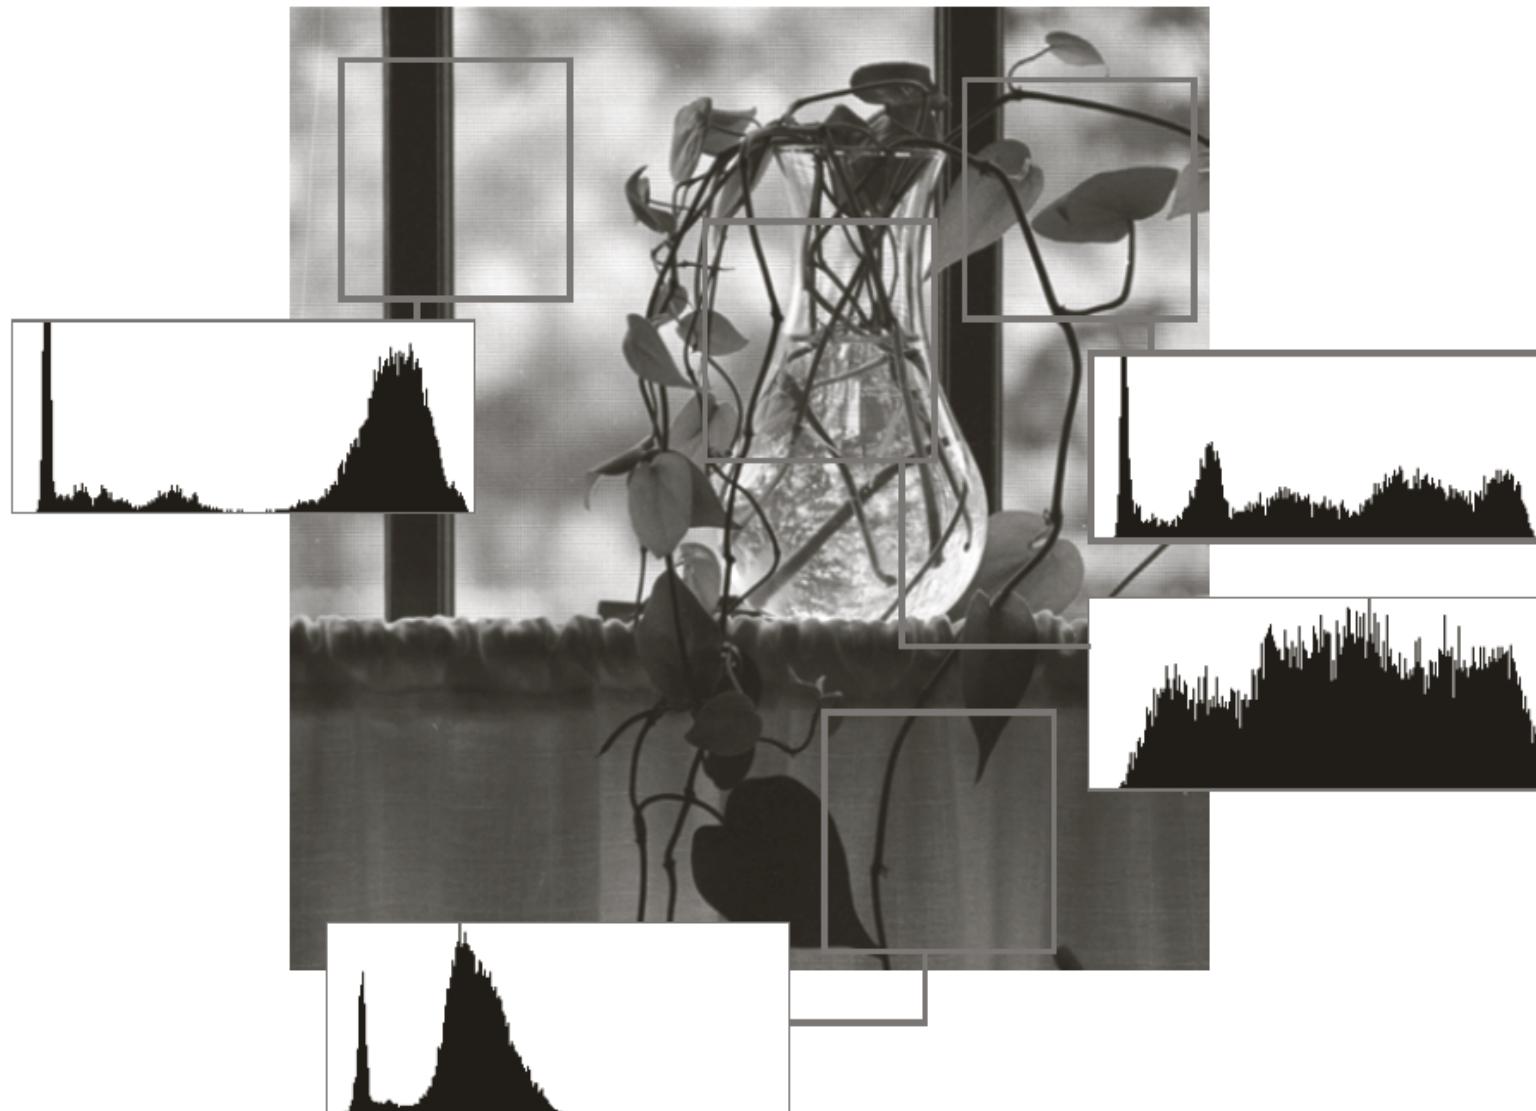


Gray-Level Transformations

- Adaptive histogram equalization (AHE)
 - Localized analysis
 - Algorithm:
 - For each pixel “p”
 - Construct a window of size NxN around the pixel
 - Perform histogram equalization within that window
 - (1) Compute histogram within window
 - (2) Compute CDF within window
 - (3) Map intensity of center pixel “p” based on the CDF
 - Window size N is a user-defined parameter
 - User-defined parameter is also called “free” parameter
 - How to choose N ?

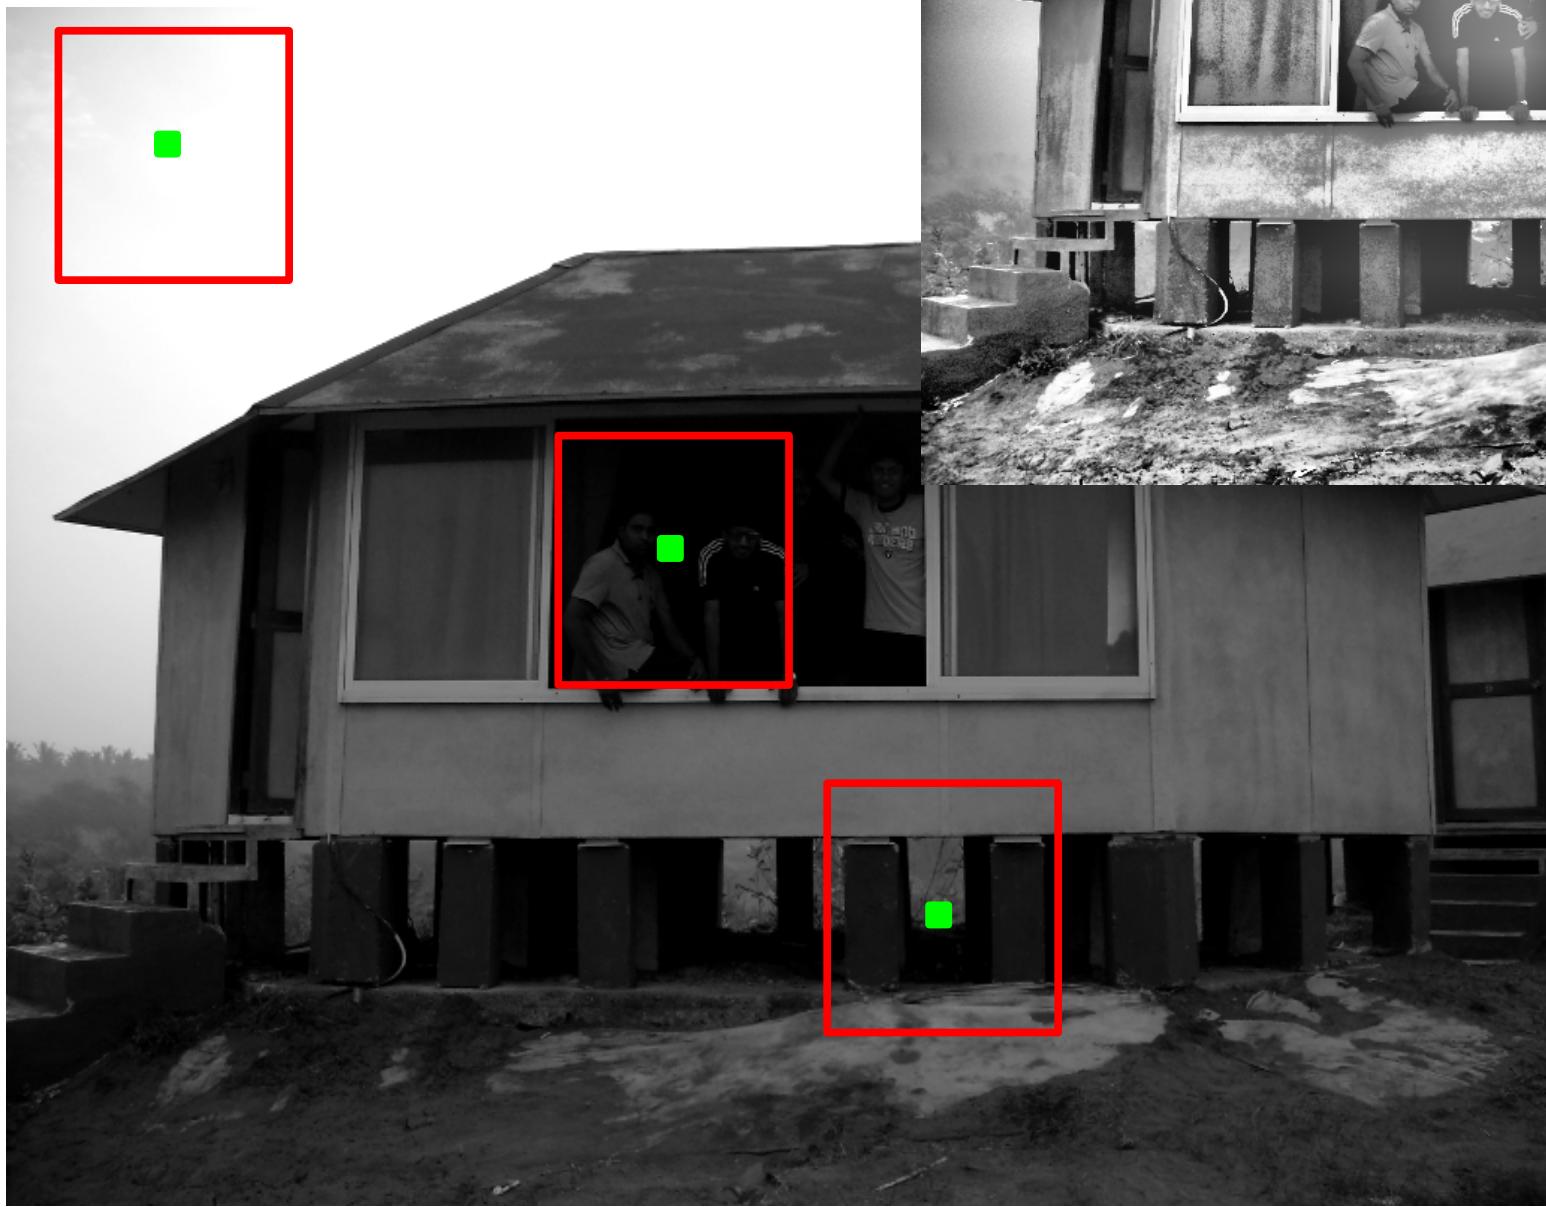
Gray-Level Transformations

- Adaptive histogram equalization (AHE)
 - Variations in local histograms



Gray-Level

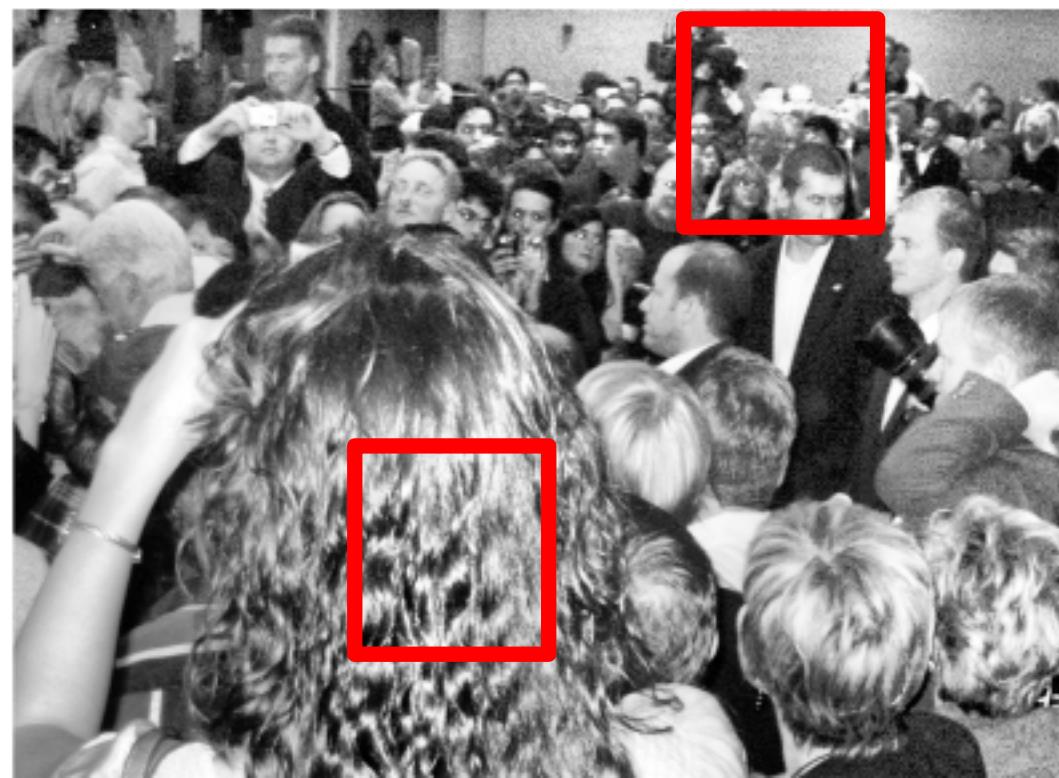
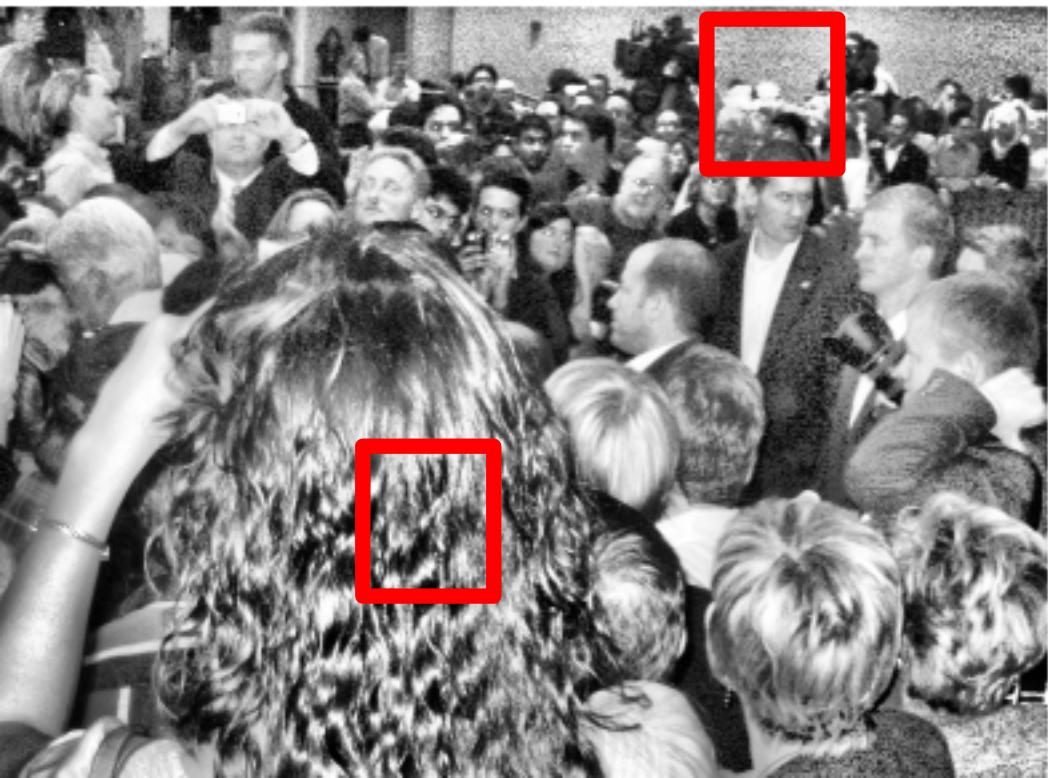
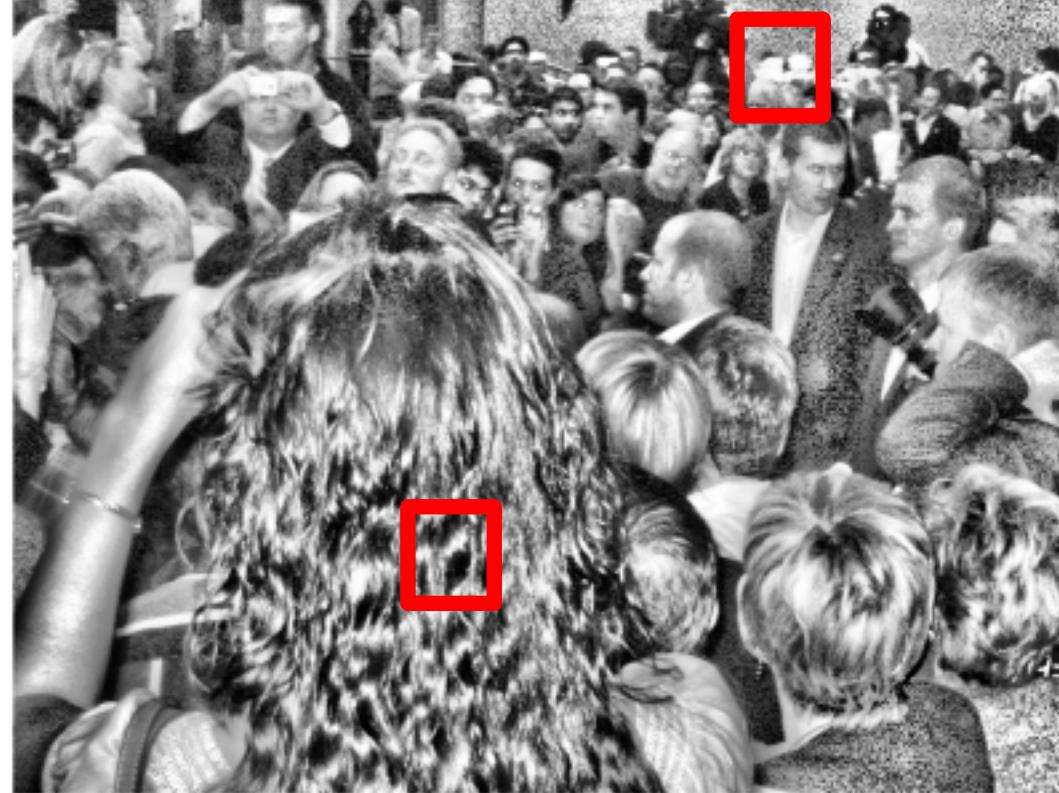
- Adaptive histogram equalization



Gray-Level Transformations

- Input image Global Hist. Eq.
 - On next slide:
AHE with square windows of widths 25, 64, 100, 200
- 
- 

AHE



AHE

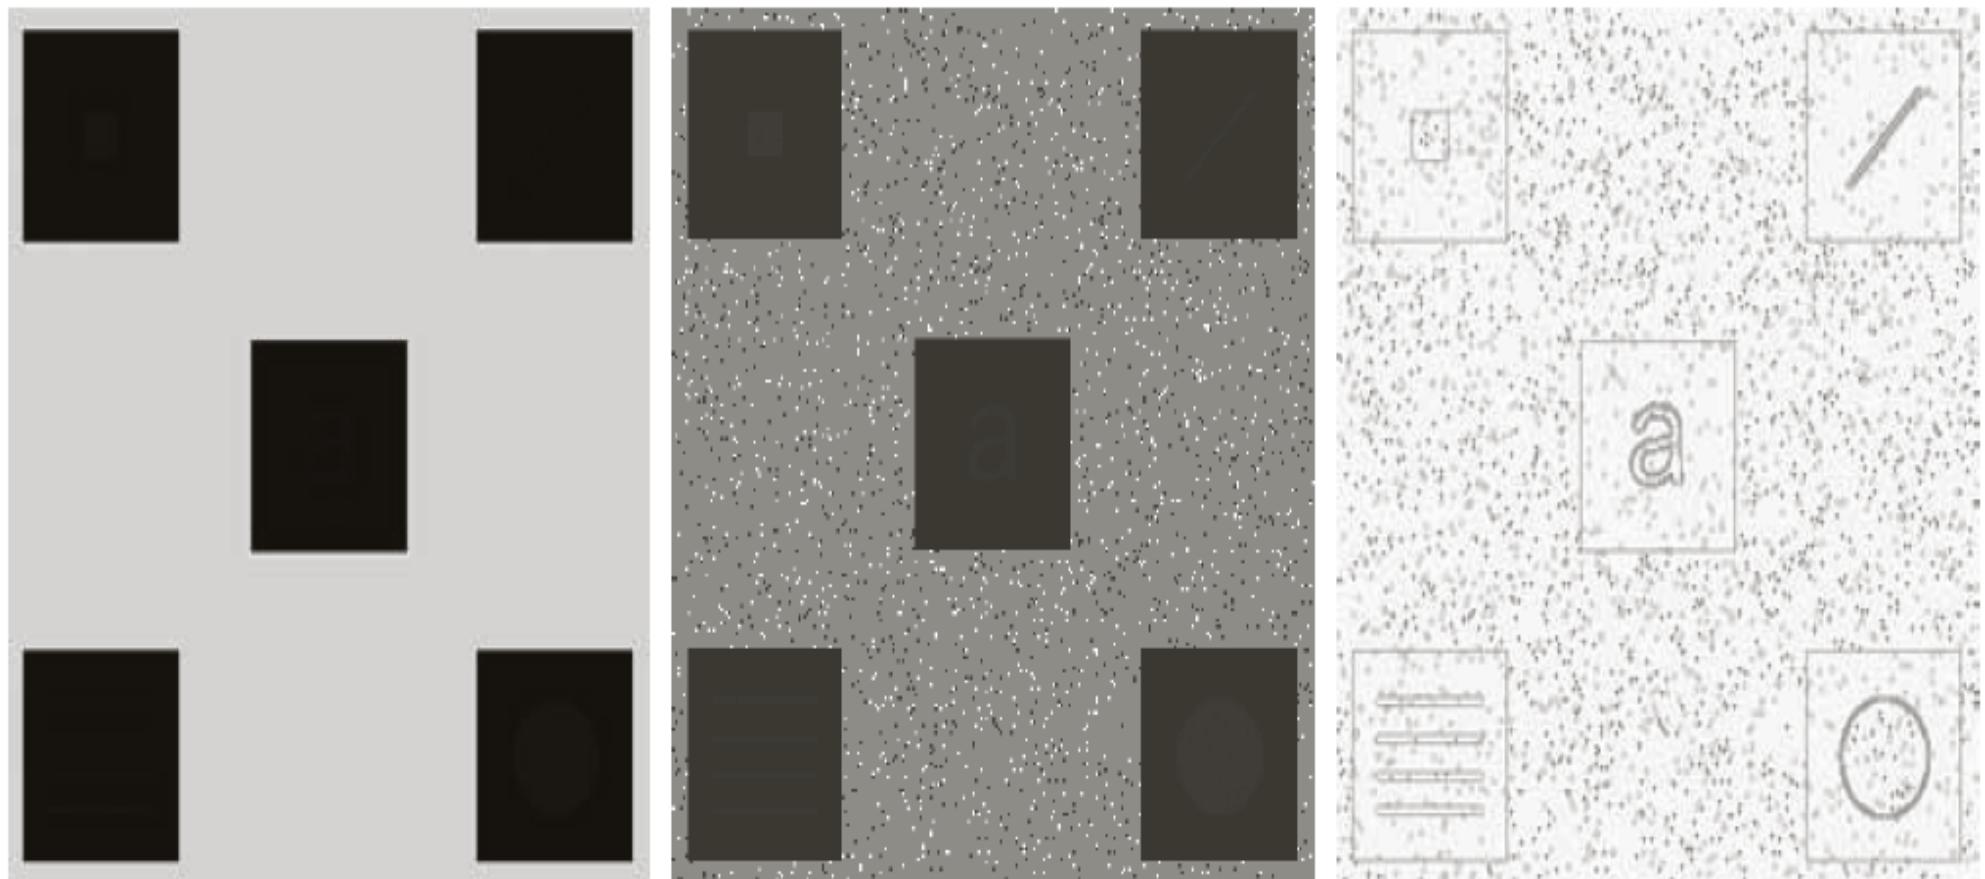
Input Image

Hist. Eq.



Gray-Level Transformations

- Adaptive histogram equalization (AHE)



a | b | c

FIGURE 3.26 (a) Original image. (b) Result of global histogram equalization. (c) Result of local histogram equalization applied to (a), using a neighborhood of size 3×3 .

Gray-Level Transformations

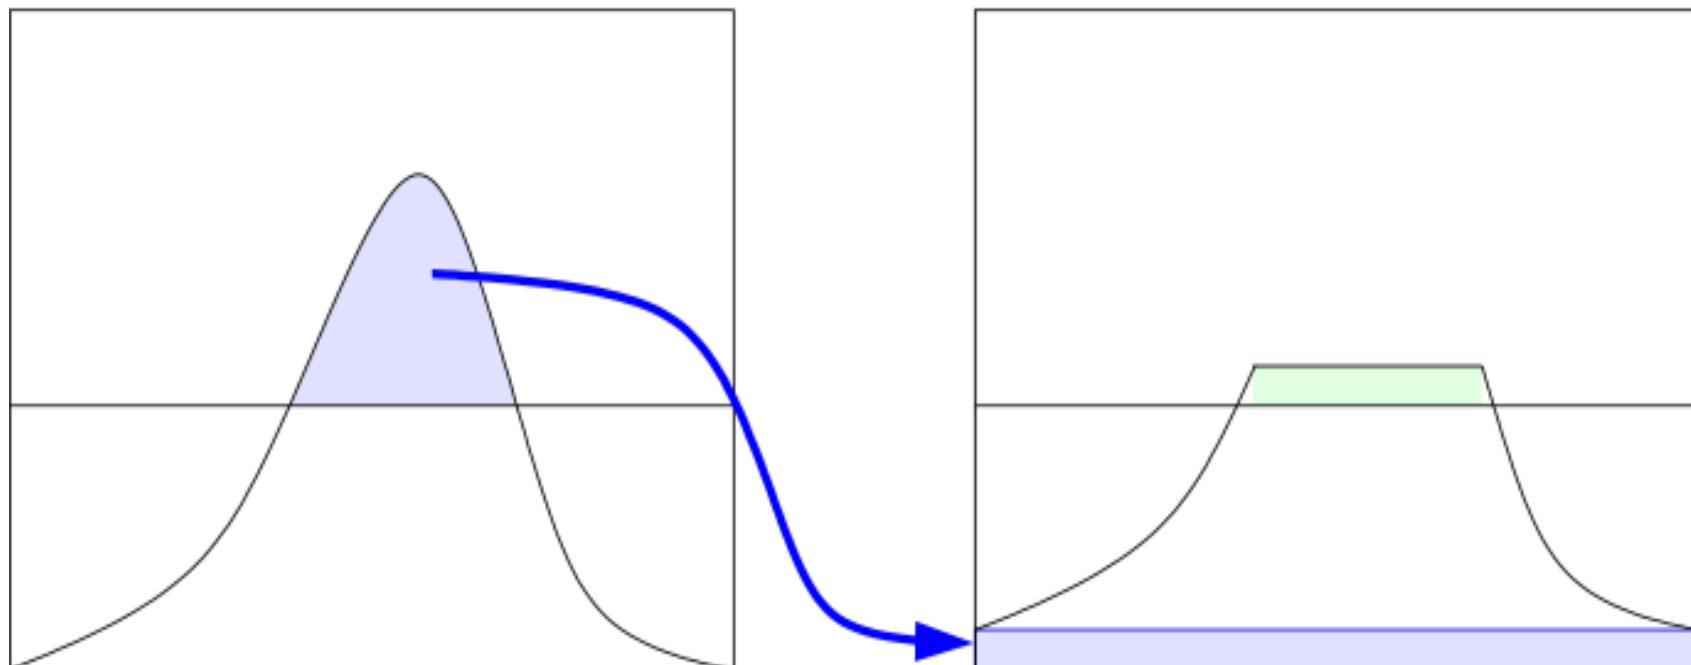
- Limitations of adaptive histogram equalization
 - Small window size → bring out more detail
 - Small window size → amplify noise in constant patches
- How to try to get the best of both worlds ?
 - We want to enhance texture, without amplifying noise

Gray-Level Transformations

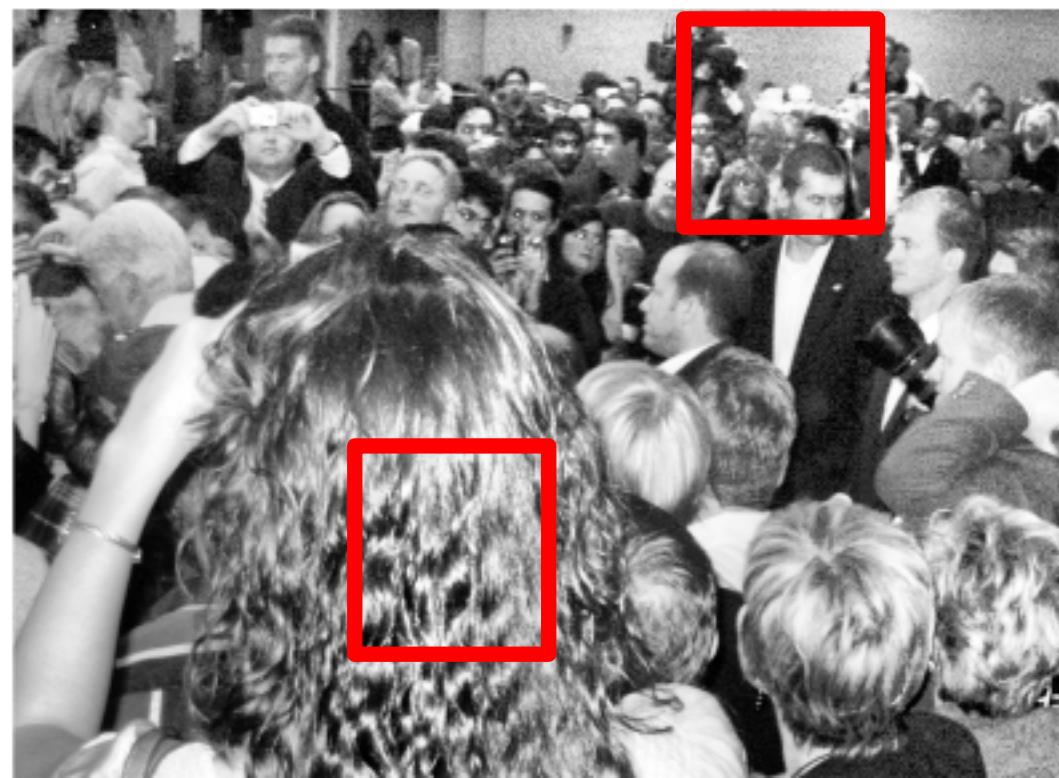
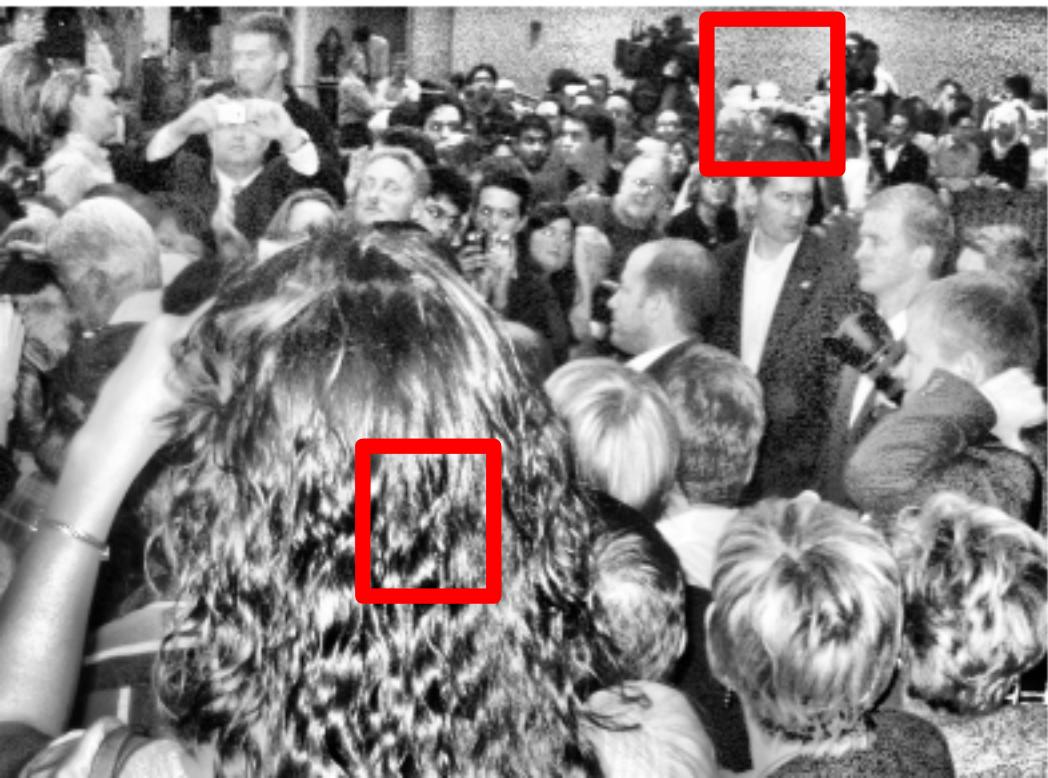
- Contrast-limited adaptive histogram equalization (CLAHE)
 - **Principle:** Contrast amplification in the vicinity of a given pixel value is given by:
 - Slope of the transformation function
= Slope of the CDF
 - But CDF slope depends on histogram (PDF) value

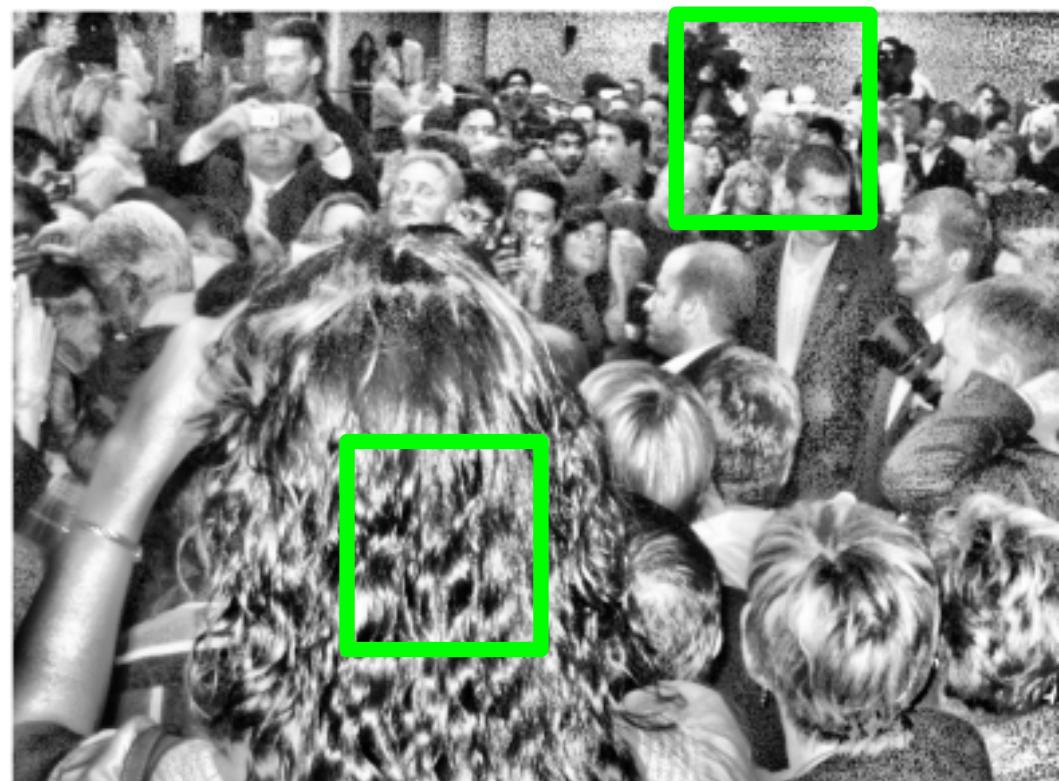
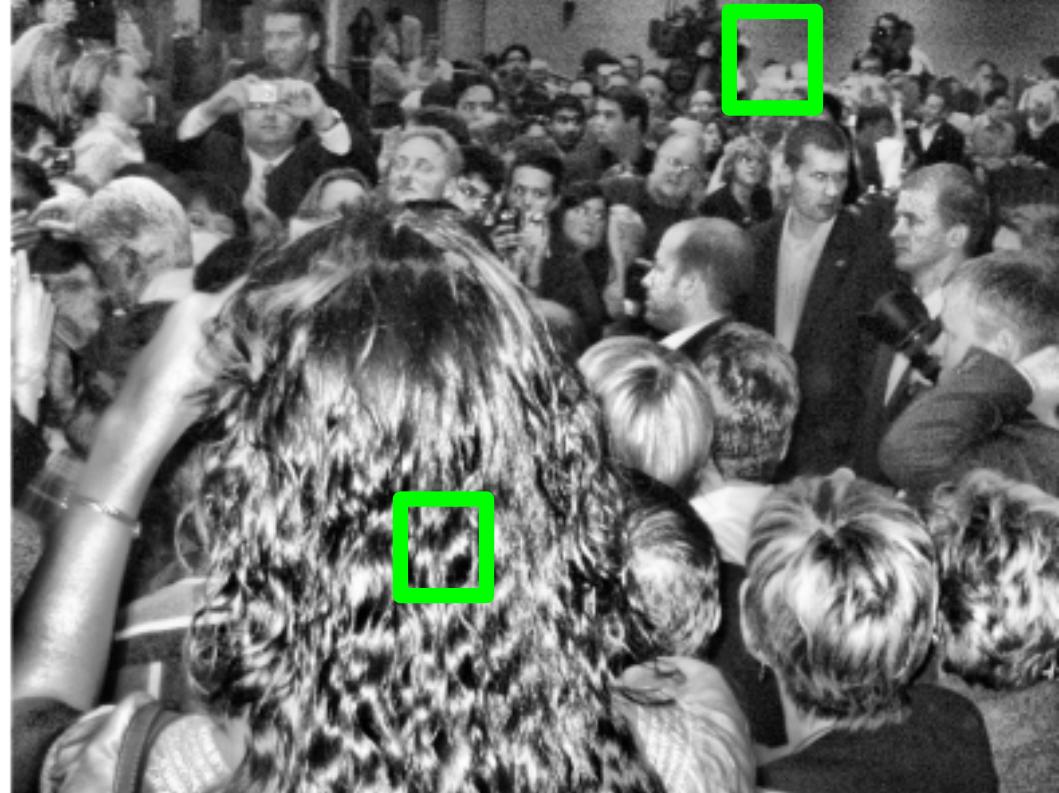
Gray-Level Transformations

- Contrast-limited adaptive histogram equalization (CLAHE)
 - Limits the amplification by:
 - (1) Clip the histogram at a predefined value (free parameter)
 - (2) Redistribute the mass uniformly throughout the range
 - (3) Then, compute the CDF



AHE





Gray-Level Transformations

- (CL)AHE is computationally expensive
 - Computes PDF (+ CDF) at window around at **every** pixel
- Can update histograms in a smart way
 - As neighborhood shifts,
only a small number of pixel change
- Approximations to (CL)AHE for speedup
 - See next slide ...

Gray-Level Transformations

- Approximations to CLAHE for speedup
 - (1) Split image into **MxM non-overlapping tiles**
 - (2) Compute CDF for each tile
 - (3) Bilinearly interpolate transformation for all pixels based on distances from tile centers

