

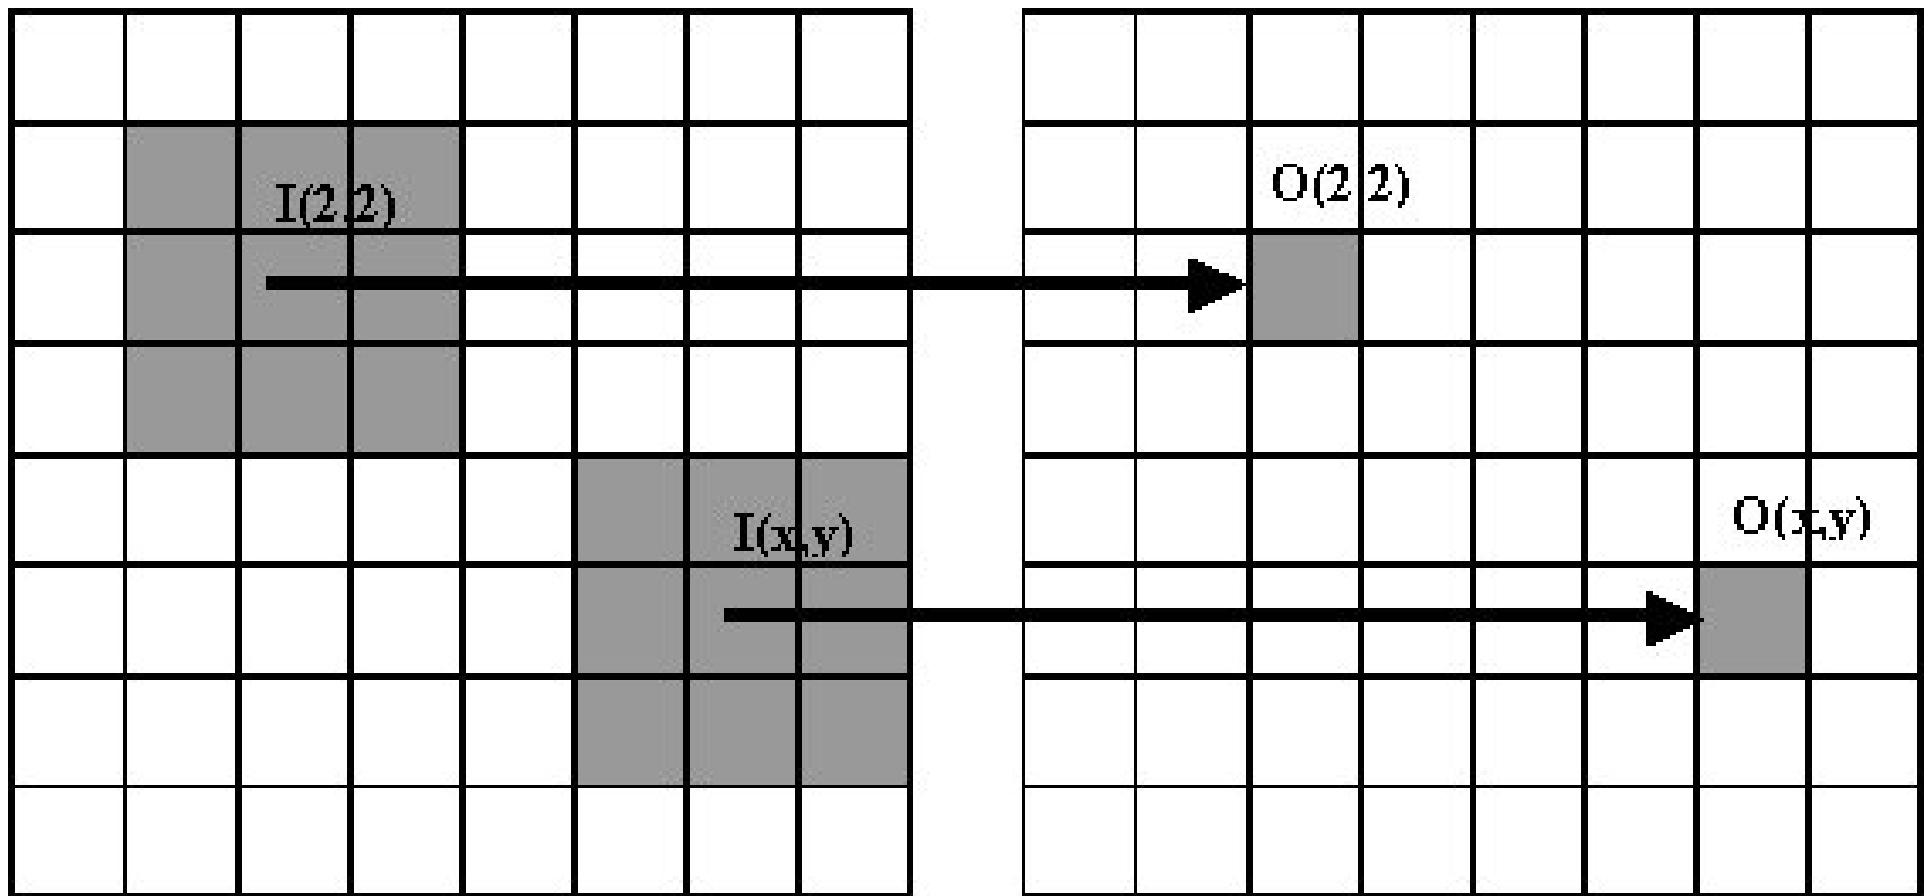
Digital Image Processing

Image Enhancement : Spatial Filtering

Suyash P. Awate

Spatial Filtering

- Key idea
 - Intensity at pixel “p” in output image is a function of the pixel intensities in the neighborhood of “p” in input image

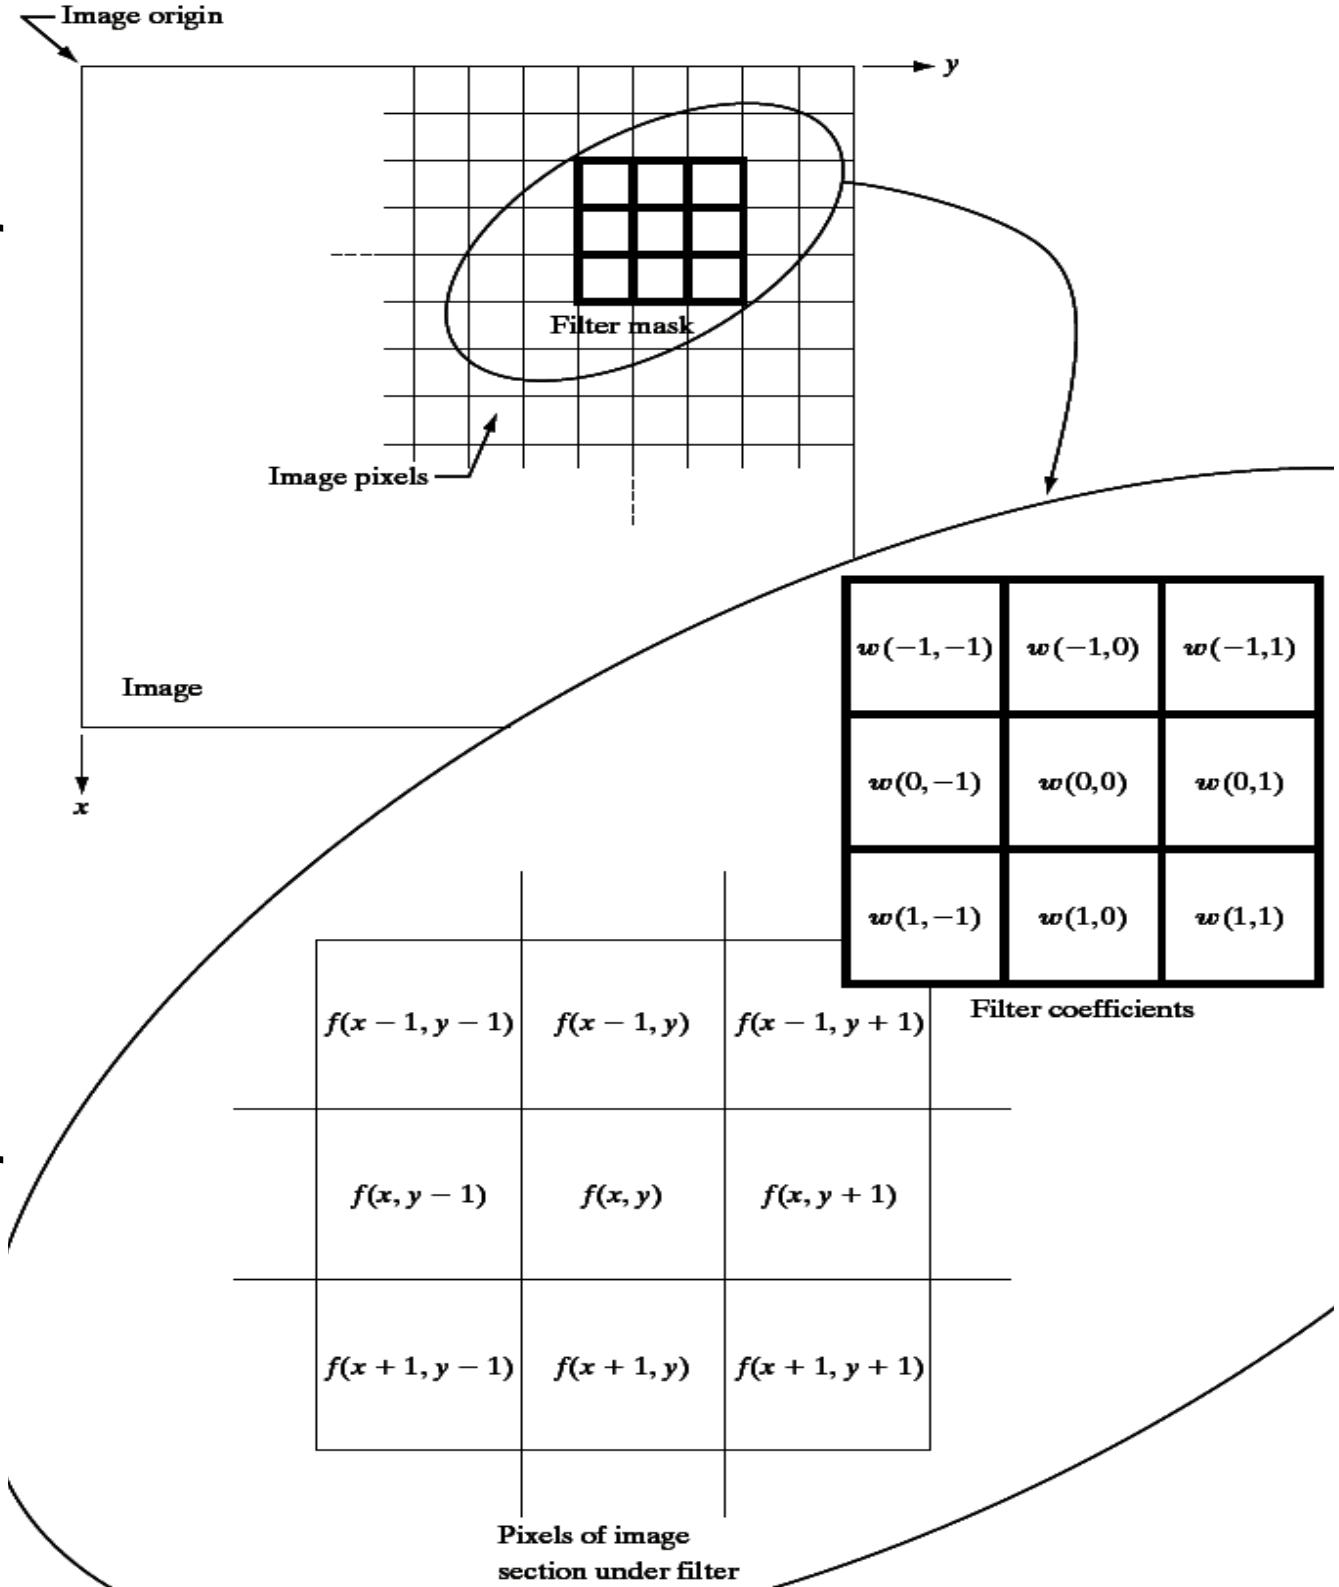


Spatial Filtering

- **Linear** spatial filtering
 - Spatial filtering where the **output** pixel intensity = **linear function** of the **input** pixel intensities

Spatial Filter

- Linear spatial filter
 - $g(x,y)$ typically designed as =
$$\sum_{i=-1,0,1} \sum_{j=-1,0,1} w(i,j) * f(x+i, y+j)$$
- w = **filter mask**
- $w(i,j)$ = **filter coefficients**
- Linear spatial filter performs
“convolution”



Spatial Filtering

- Consider a black-box system

- Takes input $x(t)$ →
Produces output $y(t)$



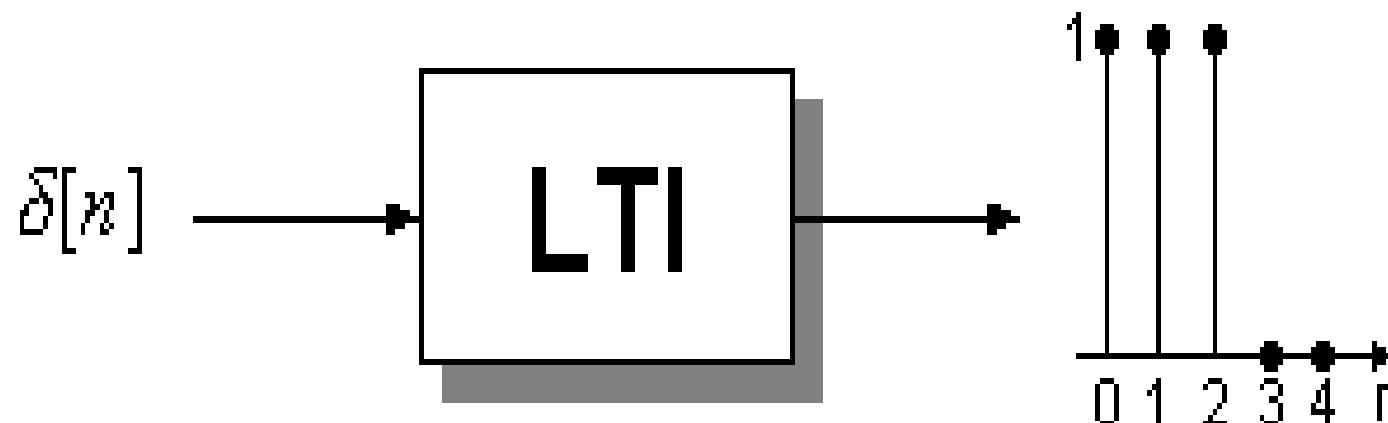
- Let input = **(discrete) unit-impulse function** at time $t=0$

- Kronecker delta:** $\delta(t) = 1$ if $t = 0$; Otherwise $\delta(t) = 0$

- Continuous impulse function = Dirac delta**

- Derivative of the unit-step function
- Rectangular pulse with width → 0, height → ∞, area = 1
- Limit of a Gaussian PDF as variance → 0

- For input = $\delta(t)$, output = $h(t)$ = **impulse response**



Spatial Filtering

- Continuous impulse function (Dirac delta)

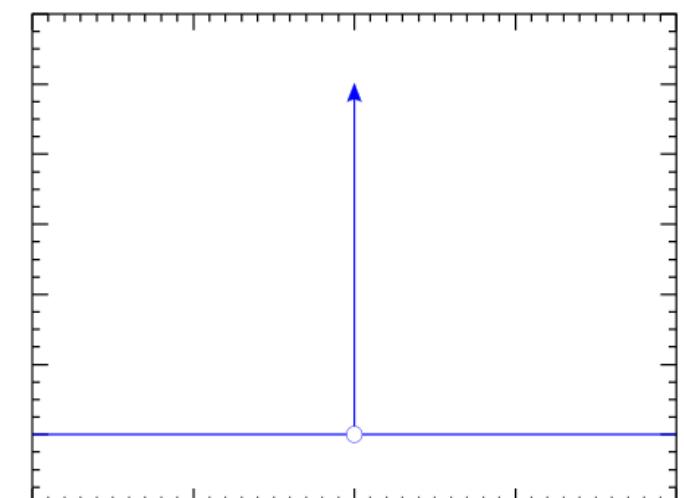
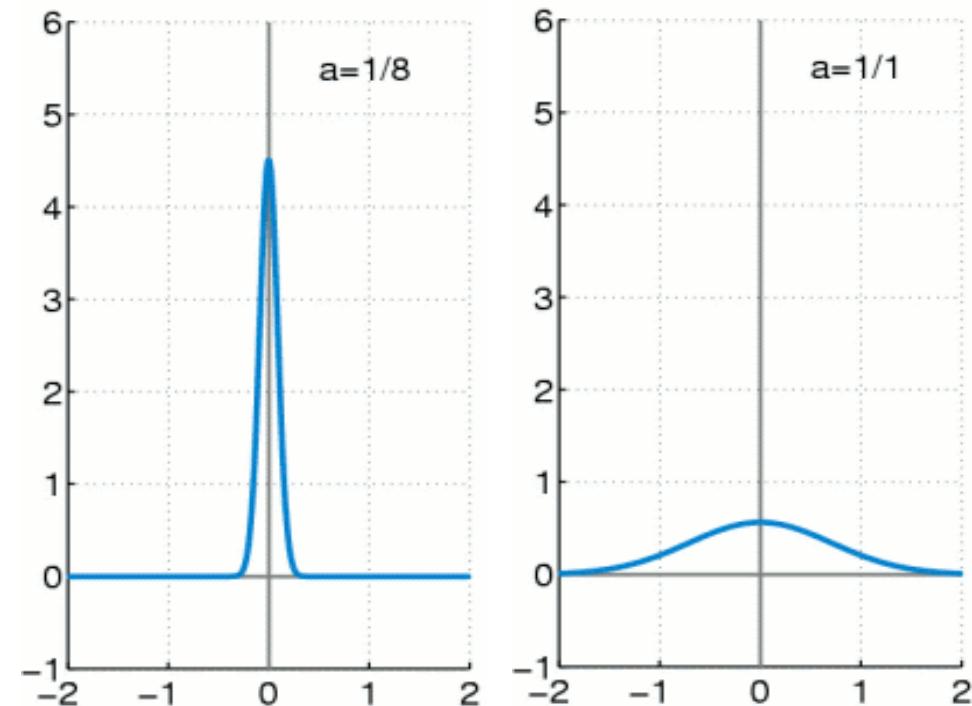
- Definition : $\delta_a(x) = \frac{1}{a\sqrt{\pi}}e^{-x^2/a^2}$
as $a \rightarrow 0$.

- $\delta(x) = 0$ for all $|x| > 0$
 - $\delta(x)$ integrates to 1

$$\int_{-\infty}^{\infty} \delta(x) dx = 1$$

- Sifting / sampling property

$$\int_{-\infty}^{\infty} f(t)\delta(t - T) dt = f(T)$$



Spatial Filtering

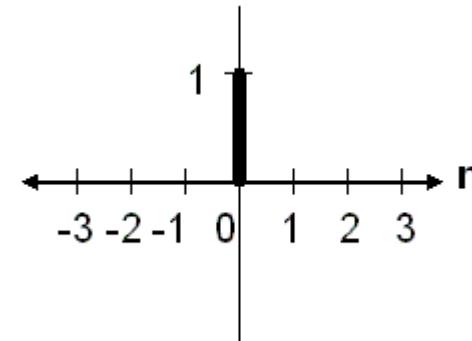
- Discrete impulse function (Kronecker delta)

- Definition :

$$\delta[n] = \begin{cases} 0, & n \neq 0 \\ 1, & n = 0 \end{cases}$$

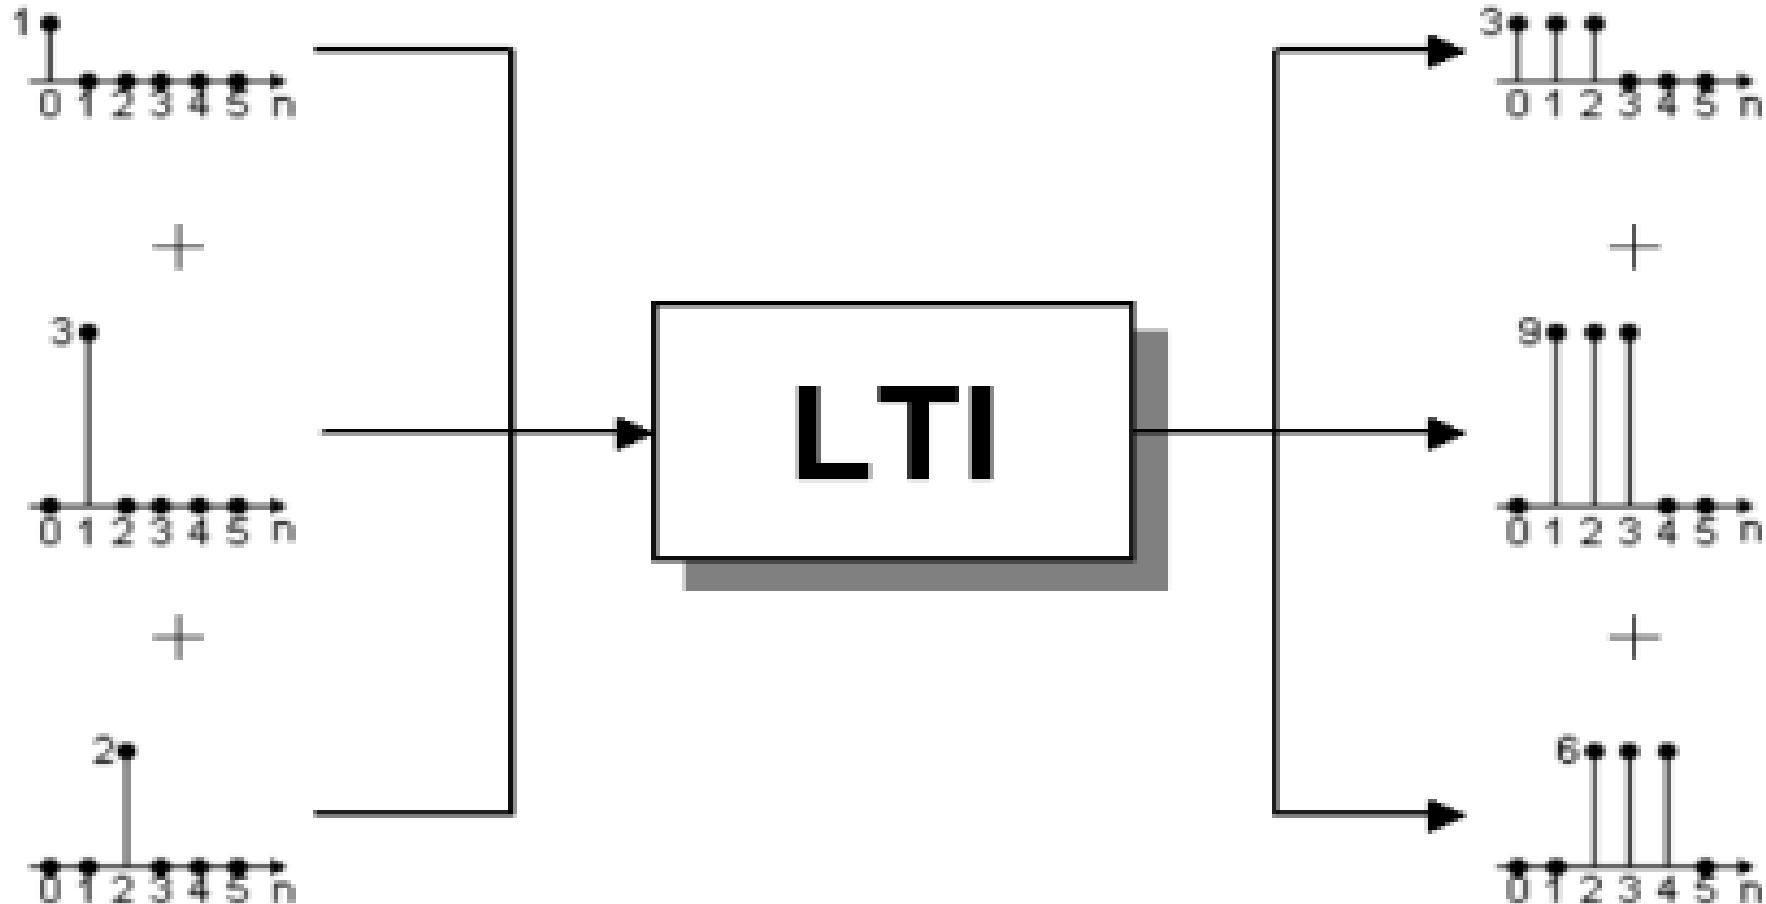
- Sifting property

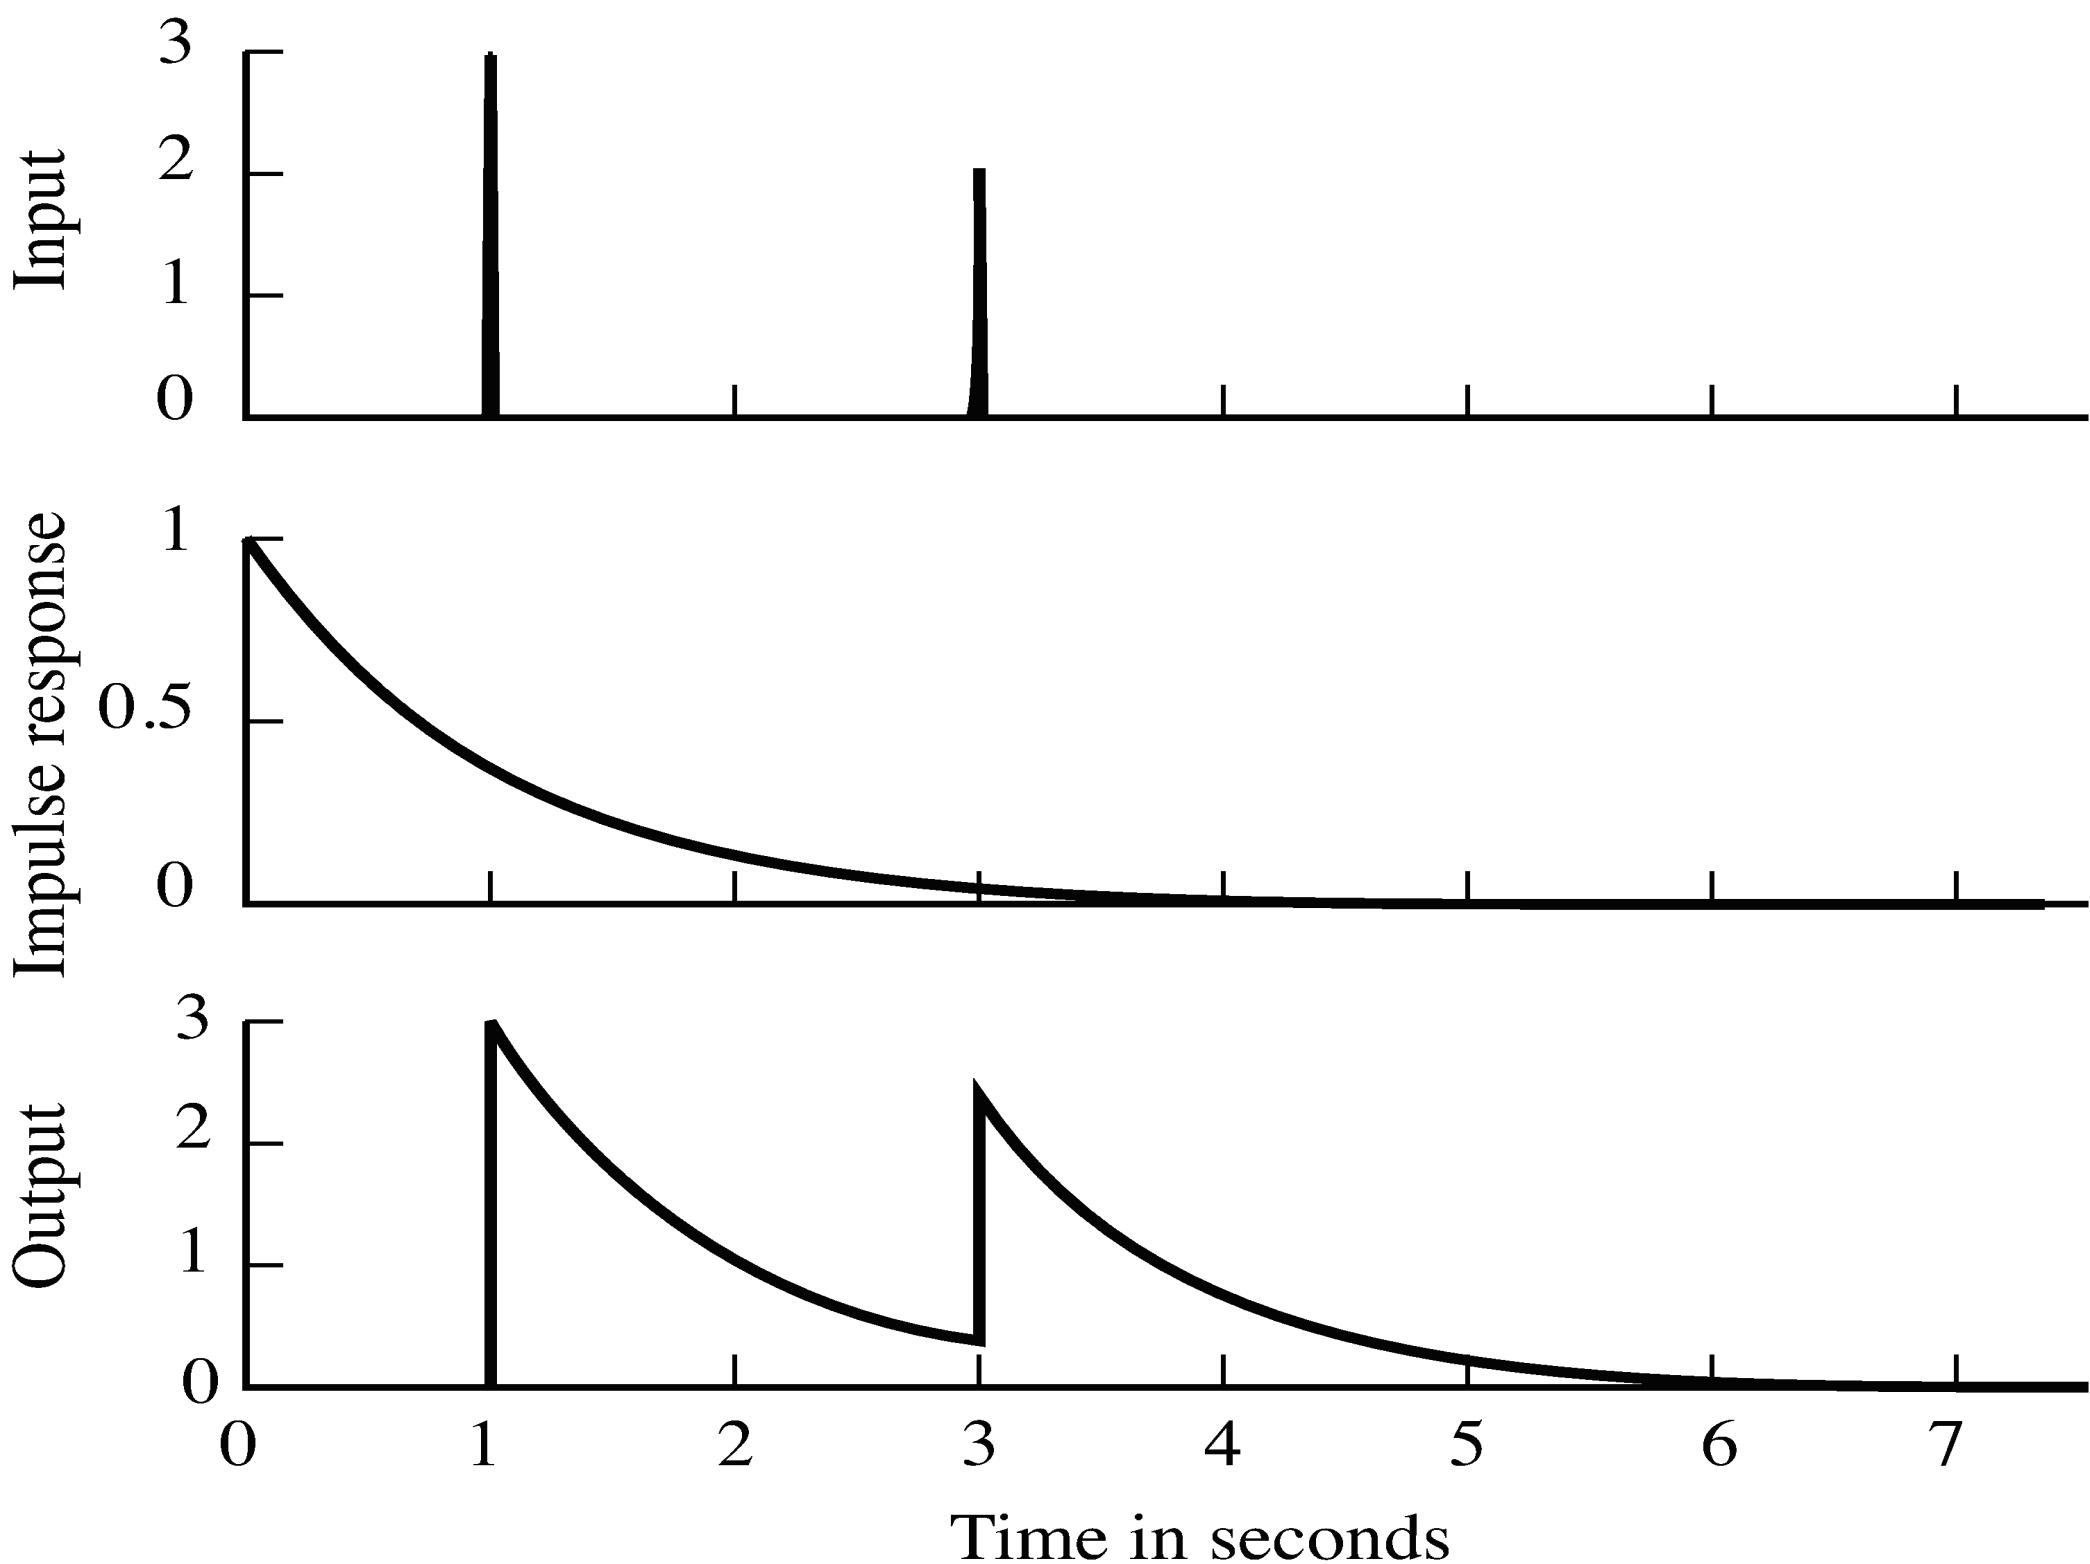
$$\sum_{i=-\infty}^{\infty} a_i \delta_{ij} = a_j$$



Spatial Filtering

- Input-output system
 - For a **time-invariant** system:
 - When input gets **delayed / shifted** by some time interval, output gets shifted by the same time interval
 - $\delta(t-t') \rightarrow h(t-t')$
 - For a **linear** system:
 - When input gets **scaled** by a factor, output gets scaled by same factor
 - $a \delta(t) \rightarrow a h(t)$
 - When the input is a **sum** of 2 (shifted) impulse functions, the output is the sum of 2 (shifted) impulse responses
 - $\delta(t) + \delta(t-t') \rightarrow h(t) + h(t-t')$
 - See picture on next slide ...





Spatial Filtering

- Linear time-invariant system
 - $f(\text{shifted candle}) = \text{shifted}(f(\text{candle}))$
 - $f(\text{candle1} + \text{candle2}) = f(\text{candle1}) + f(\text{candle2})$



Spatial Filtering

- Linear time-invariant system
 - In equations:
 - Impulse response = $H(t)$
 - Input = $F(t)$
 - Output = $G(t)$
 - We can **rewrite any input $F(t)$** as
 - $F(t) = F(0) \delta(t) + F(1) \delta(t-1) + F(2) \delta(t-2) + \dots$
 - By **definition of impulse response**,
if input = $\delta(t)$,
then output = $H(t)$
 - By **definition of LTI system**,
if input is $F(t) = \text{sum of scaled shifted } \delta(t) \text{ functions}$,
then output $G(t) = \text{sum of scaled shifted } H(t) \text{ responses}$

$$\int_{-\infty}^{\infty} f(t) \delta(t - T) dt = f(T)$$

Spatial Filtering

- Linear time-invariant system
 - By the definition of the LTI system,
output $G(t) = \text{sum of scaled shifted impulse responses}$
 - $G(t) = F(0) H(t) + F(1) H(t-1) + F(2) H(t-2) + \dots$
 - $G(t) = \sum_{k=0,1,2} F(k) H(t-k)$
 - Output $G(t)$
= **convolution** of input $F(t)$ with impulse response $H(t)$
 - In general
 - Input function can be arbitrarily long
 - Impulse response function can be arbitrarily long
 - $G(t) = \sum_{k=-\infty, \dots, \infty} F(k) H(t-k)$
 - $\mathbf{G} = \mathbf{F} * \mathbf{H}$
 - $\mathbf{G}(t) = (\mathbf{F} * \mathbf{H})(t)$

Spatial Filtering

- Convolution
 - What if impulse response $H(t) = \delta(t)$?
 - If, impulse response = impulse function,
then $F * H = ?$

Spatial Filtering

- Convolution
 - Convolution is a **symmetric / commutative** operation
 - $F * H = H * F$
 - Observe
$$G(t) = (F * H)(t) = \sum_{k=-\infty, \dots, \infty} F(k) H(t-k)$$
$$= \sum_{m=-\infty, \dots, \infty} F(t-m) H(m) = (H * F)(t)$$

- Just a change of variables:

Define $m = t - k$

Then,

$$k = t - m$$

$$k = \infty \rightarrow m = (-\infty)$$

$$k = (-\infty) \rightarrow m = \infty$$

Spatial Filtering

- Convolution
 - Convolution is a **linear** operation on its arguments
 - What if $H(t) = A(t) + B(t)$?
 - $F * H = F * (A + B) = ?$
 - What if $H(t) = c A(t)$? where $c = \text{constant}$
 - $F * H = F * (c A) = ?$

Spatial Filtering

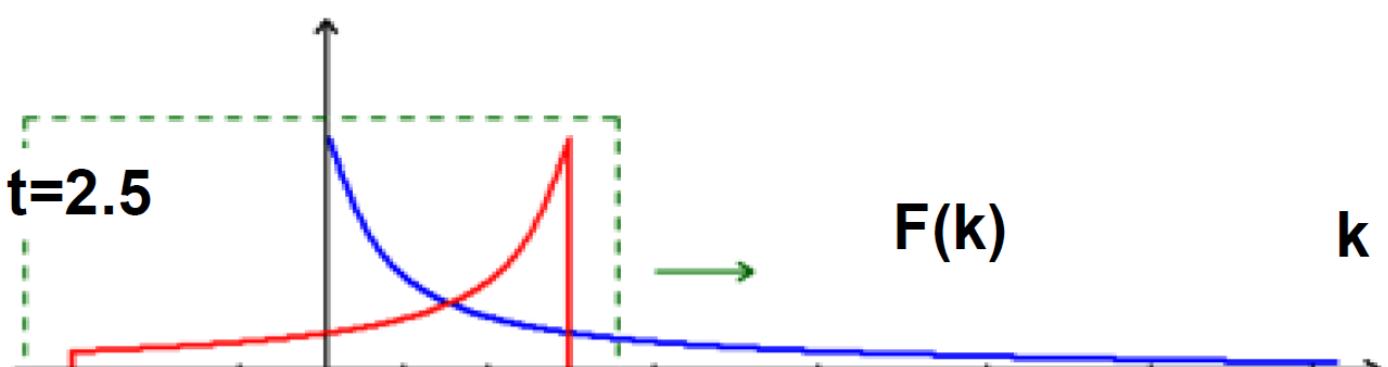
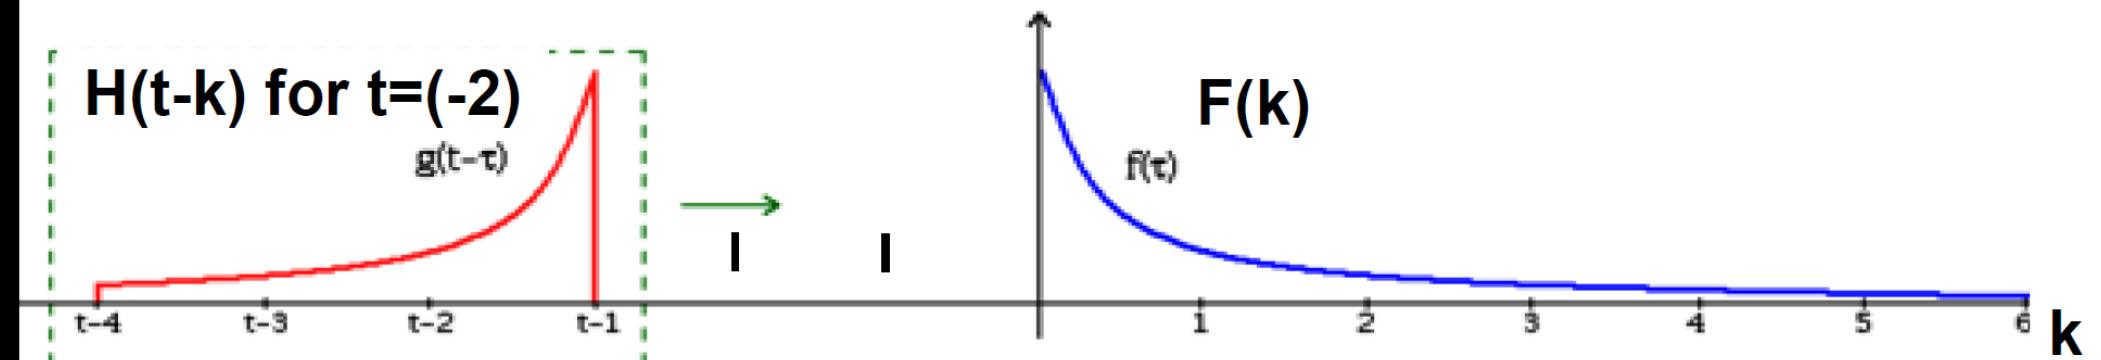
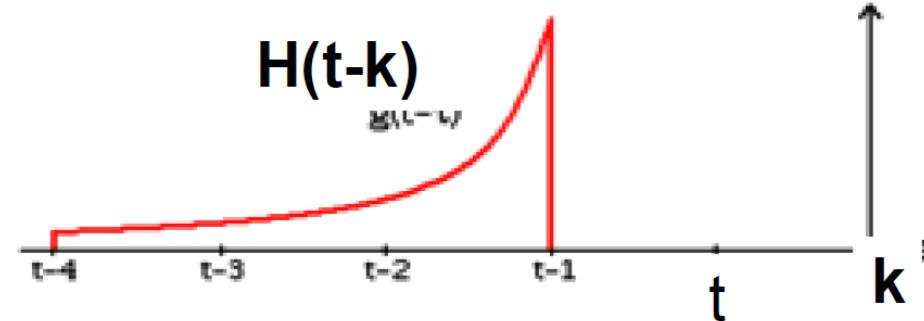
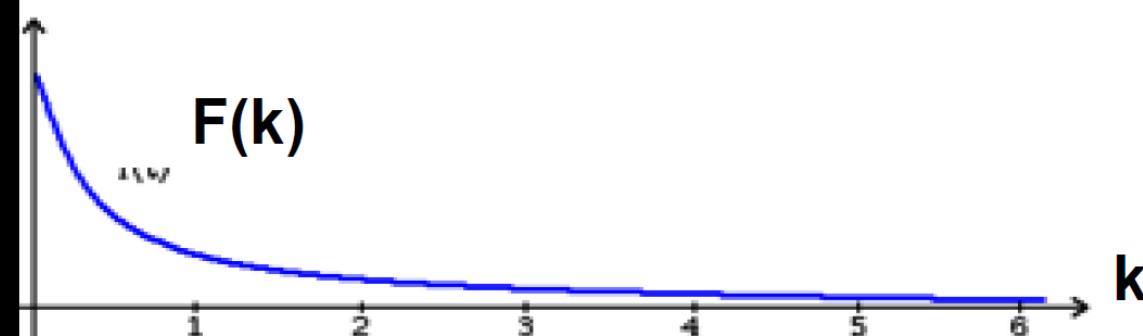
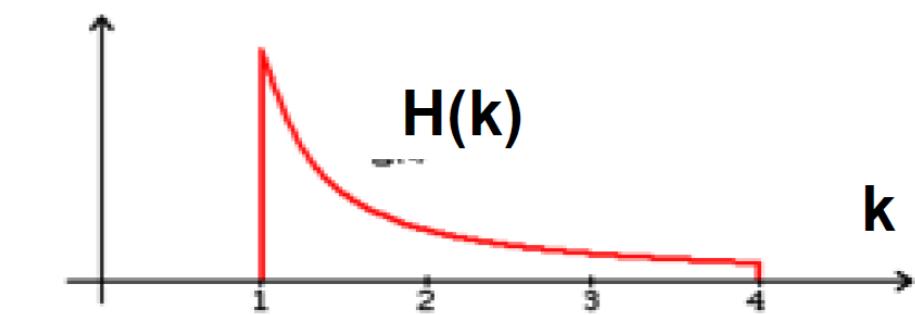
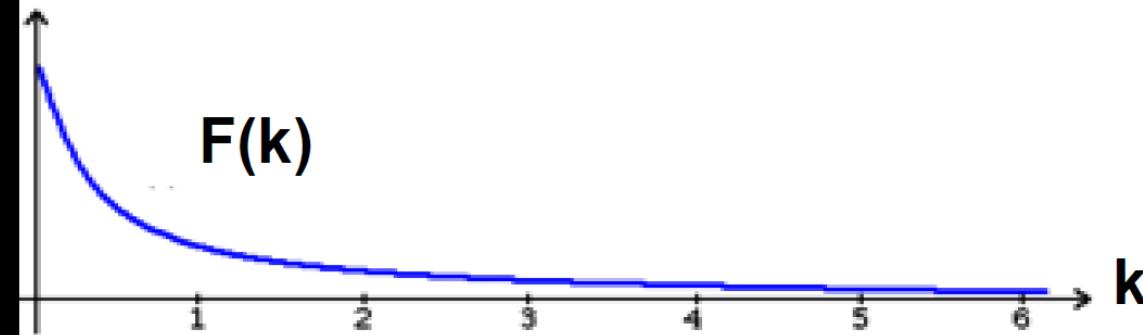
- Convolution
 - Convolution is **associative** (order doesn't matter)
 - $(x * h_1) * h_2 = x * (h_1 * h_2)$
 - Proof: Left hand side $(t) = \int_{-\infty}^{\infty} x(\tau)h_1(t - \tau)d\tau * h_2(t)$
 - (expanding) = $\int_{-\infty}^{\infty} \left[\int_{-\infty}^{\infty} x(\tau)h_1(\gamma - \tau)d\tau \right] h_2(t - \gamma)d\gamma$
 - (reverse order) = $\int_{-\infty}^{\infty} x(\tau) \left[\int_{-\infty}^{\infty} h_1(\gamma - \tau)h_2(t - \gamma)d\gamma \right] d\tau$
 - $\varphi = \gamma - \tau$
 - (substitution) = $\int_{-\infty}^{\infty} x(\tau) \left[\int_{-\infty}^{\infty} h_1(\varphi)h_2(t - \tau - \varphi)d\varphi \right] d\tau$
 - Inner integral = $h_1 * h_2$ evaluated at $(t - \tau)$
 - Outer integral = $x * (h_1 * h_2)$ evaluated at t

Spatial Filtering

- Convolution
 - How to **compute** convolution for discrete images ?
 - Consider time $t = 0$
 - We can redefine any arbitrary time-point as “0” by translating the coordinate frame for input F (and output G)
 - The origin of the time coordinate frame is NOT important in LTI systems
 - $G(t=0) = \sum_{k=-\infty, \dots, \infty} F(k) H(0-k)$
 - What does this mean algorithmically ?
 - To compute $G(0)$:
 - (1) **Flip** the impulse response about the origin $\rightarrow H(-k)$
 - (2) Perform **pointwise multiplication** of F and flipped- H
 - (3) **Sum** up
 - How to compute $G(t)$ at some arbitrary timepoint “ t ” ?

Spatial Filtering

- Convolution
 - What does this mean algorithmically ?
 - We want to compute $G(t) = \sum_{k=-\infty, \dots, \infty} F(k) H(t-k)$ for an arbitrary timepoint t
 - (1) **Flip** the impulse response about the origin $\rightarrow H(-k)$
 - (2) **Shift** the flipped impulse response by $t \rightarrow H(t-k)$
 - (3) Perform **pointwise multiplication** of F and shifted-flipped- H
 - (4) **Sum** up
 - See next slide for pictures



Spatial Filteri

convolution

/kən'velju:ʃ(ə)n/ ⓘ

noun

noun: convolution; plural noun: convolutions; noun: convolution integral; plural noun: convolution integrals; plural noun: convolutions integral

1. a coil or twist.

"crosses adorned with elaborate convolutions"

synonyms: twist, turn, coil, spiral, twirl, curl, helix, whorl, loop, curlicue, kink, sinuosity; [More](#)

- the state of being or process of becoming coiled or twisted.
"the flexibility of the polymer chain allows extensive convolution"

2. a thing that is complex and difficult to follow.

"the convolutions of farm policy"

synonyms: complexity, intricacy, complication, twist, turn, entanglement, contortion; [More](#)

3. a sinuous fold in the surface of the brain.

4. MATHEMATICS

a function derived from two given functions by integration which expresses how the shape of one is modified by the other.

- a method of determination of the sum of two random variables by integration or summation.

Origin

MEDIEVAL LATIN

convolvere
roll together

MEDIEVAL LATIN

convolutio

ENGLISH

convolve

convolution
mid 16th century

mid 16th century: from medieval Latin *convolutio*(*n-*), from *convolvere* 'roll together' (see [convolve](#)).

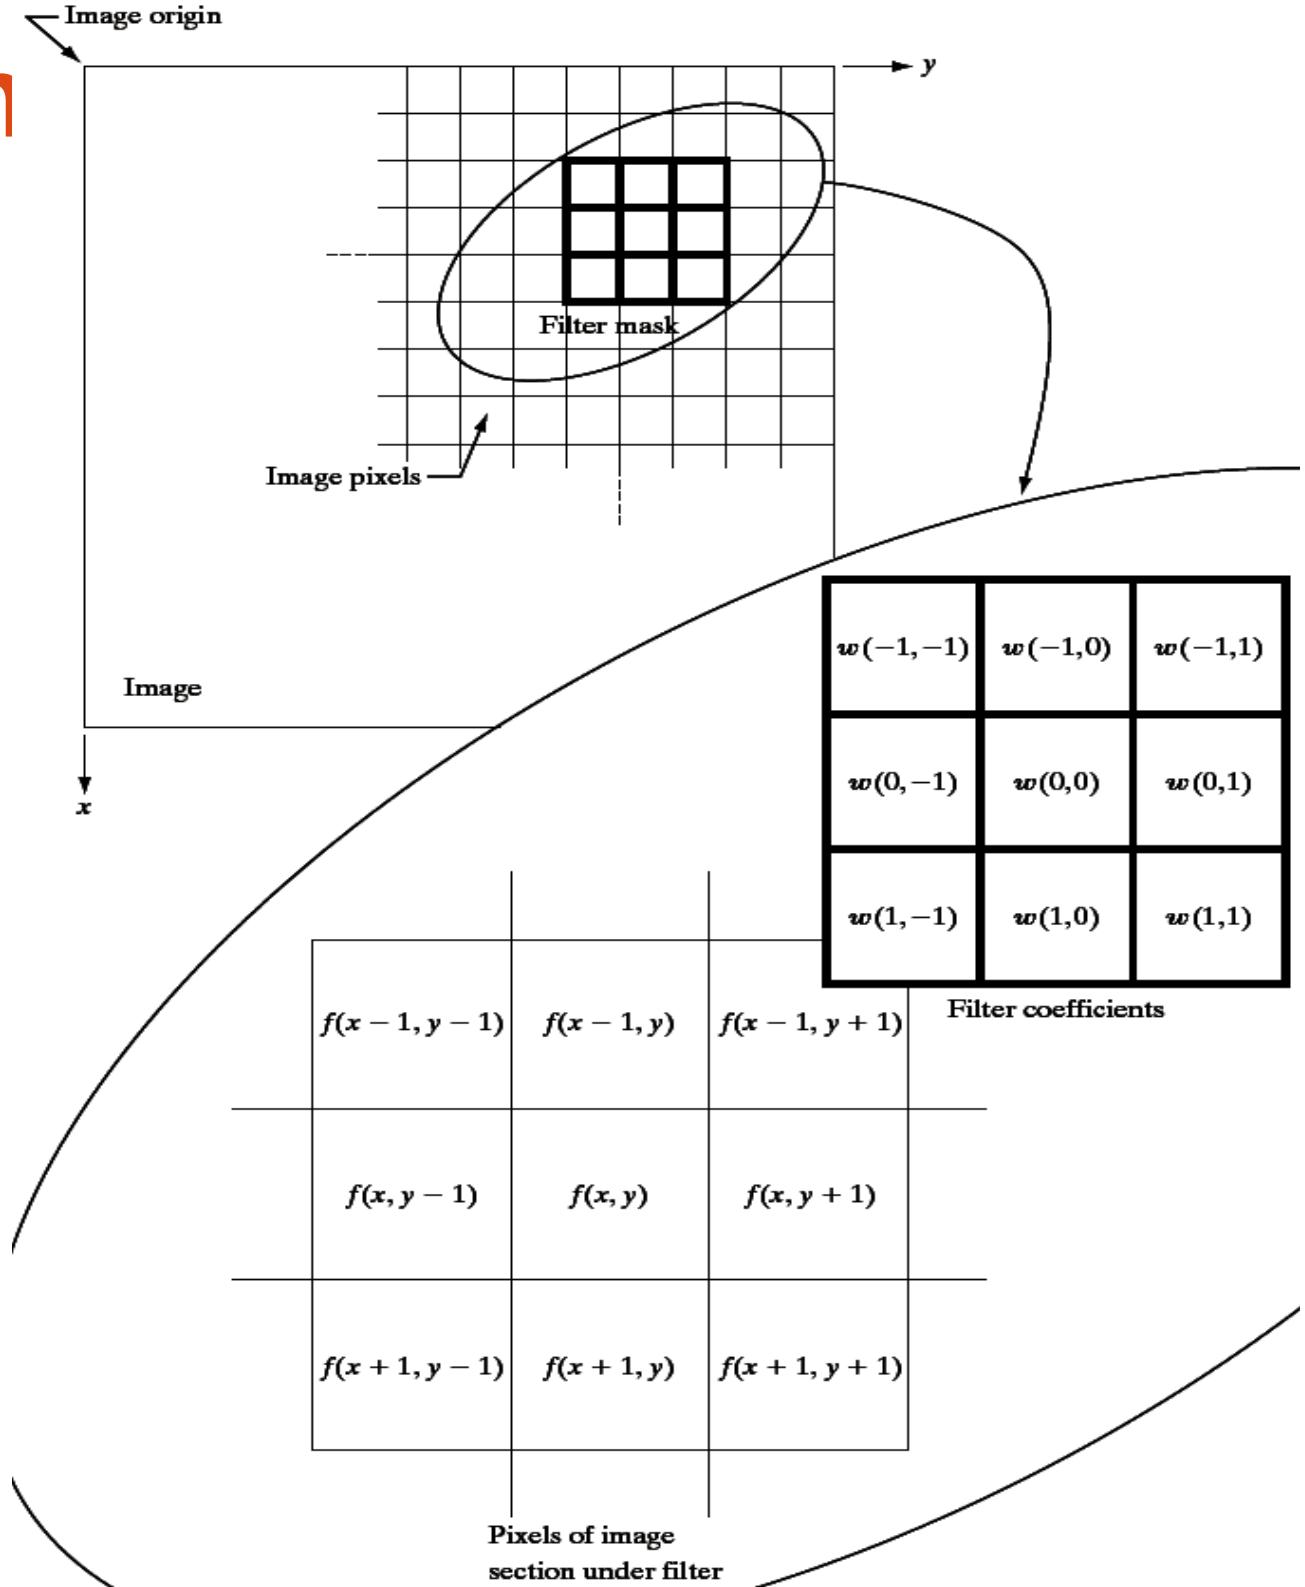




The sign was just as convoluted as the road.

Spatial Filtering

- Convolving images
 - Domain = spatial
 - System is **linear space-invariant / shift-invariant**
 - Impulse response = **point-spread function**
 - Typically, $w(x,y) = w(-x,-y)$
 - $g(x,y) = \sum_{i=-1,0,1} \sum_{j=-1,0,1} w(i,j) * f(x+i,y+j)$



Spatial Filtering

- **Cross-Correlation**

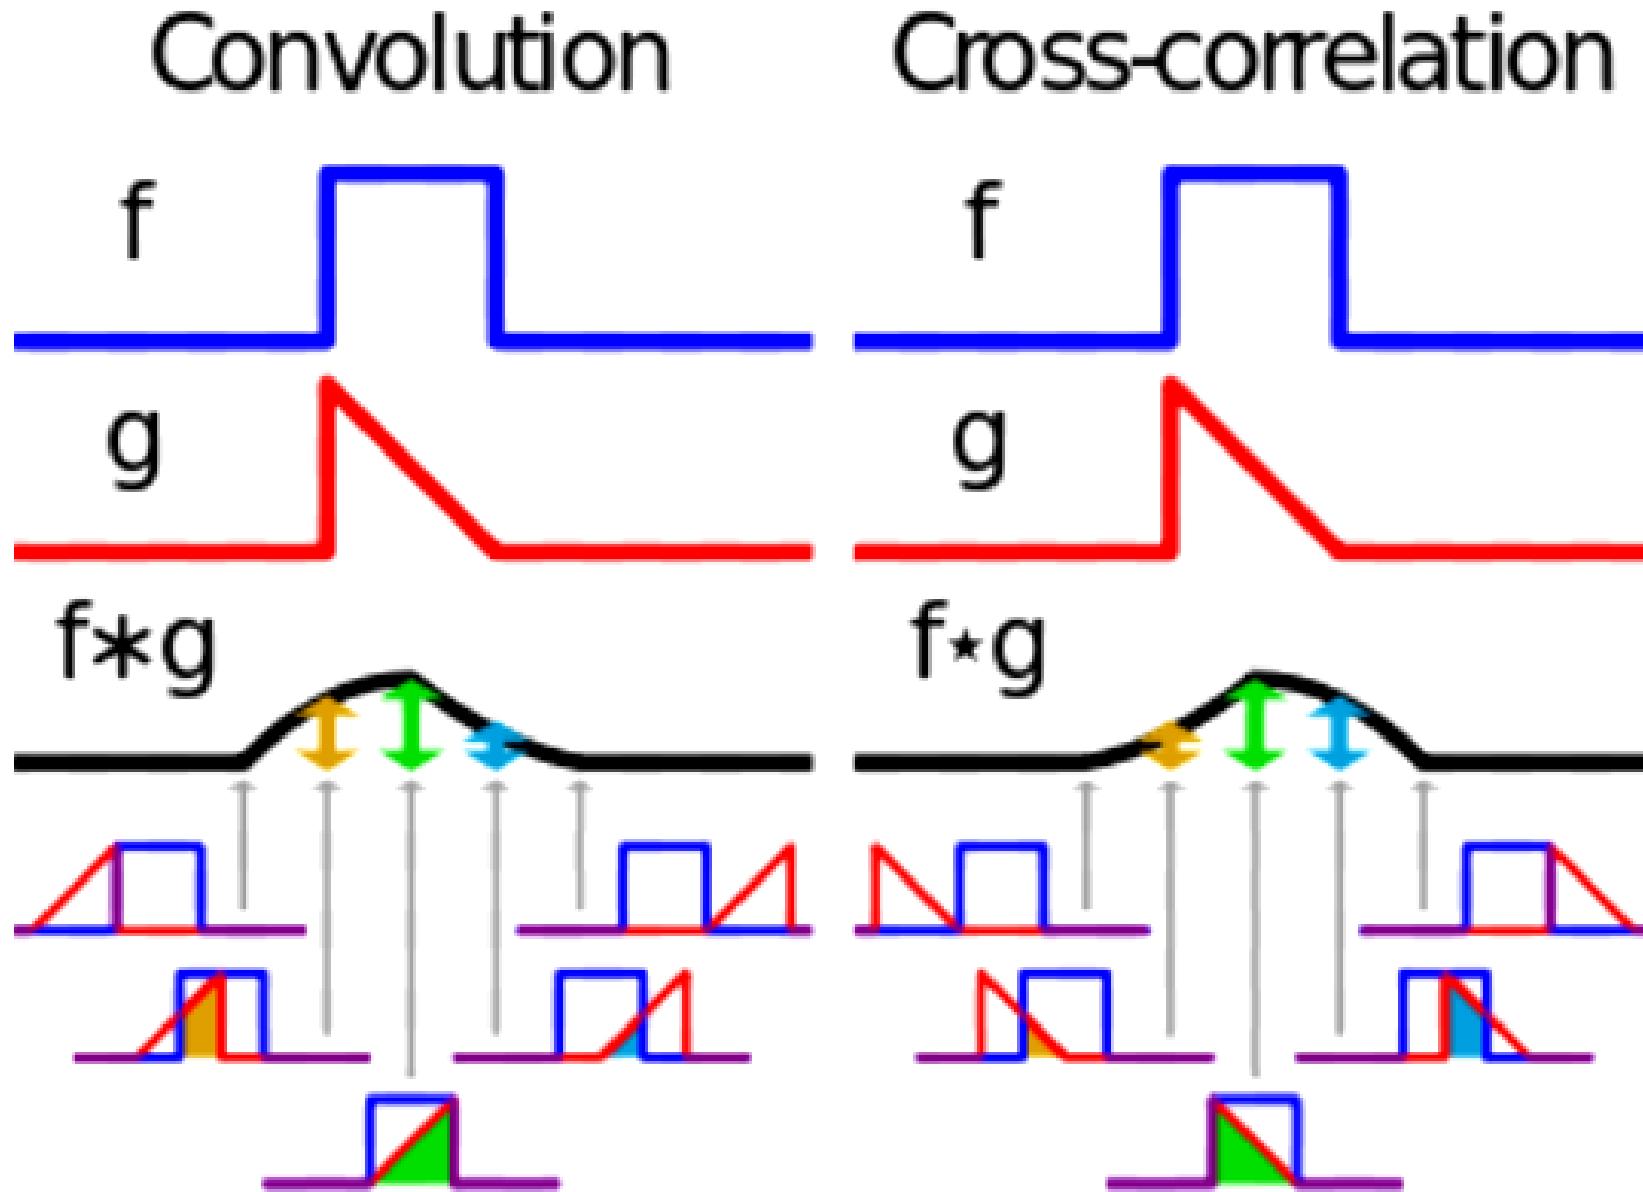
- Signals: $F(k)$ and $H(k)$
- **Measure of similarity** of two signals / images
 - Cross-correlation of F and H is $G(0) = \sum_{k=-\infty, \dots, \infty} F(k) H(k)$
- Similarity between F and a shifted version of H
 - Cross-correlation of F and H is $G(t) = \sum_{k=-\infty, \dots, \infty} F(k) H(k+t)$
- What does this mean algorithmically ?
 - (1) **Shift** one of the signals by $t \rightarrow H(k+t)$
 - (2) Perform **pointwise multiplication** of F and shifted-flipped- H
 - (3) **Sum** up
- Cross-correlation **doesn't involve the flip** !

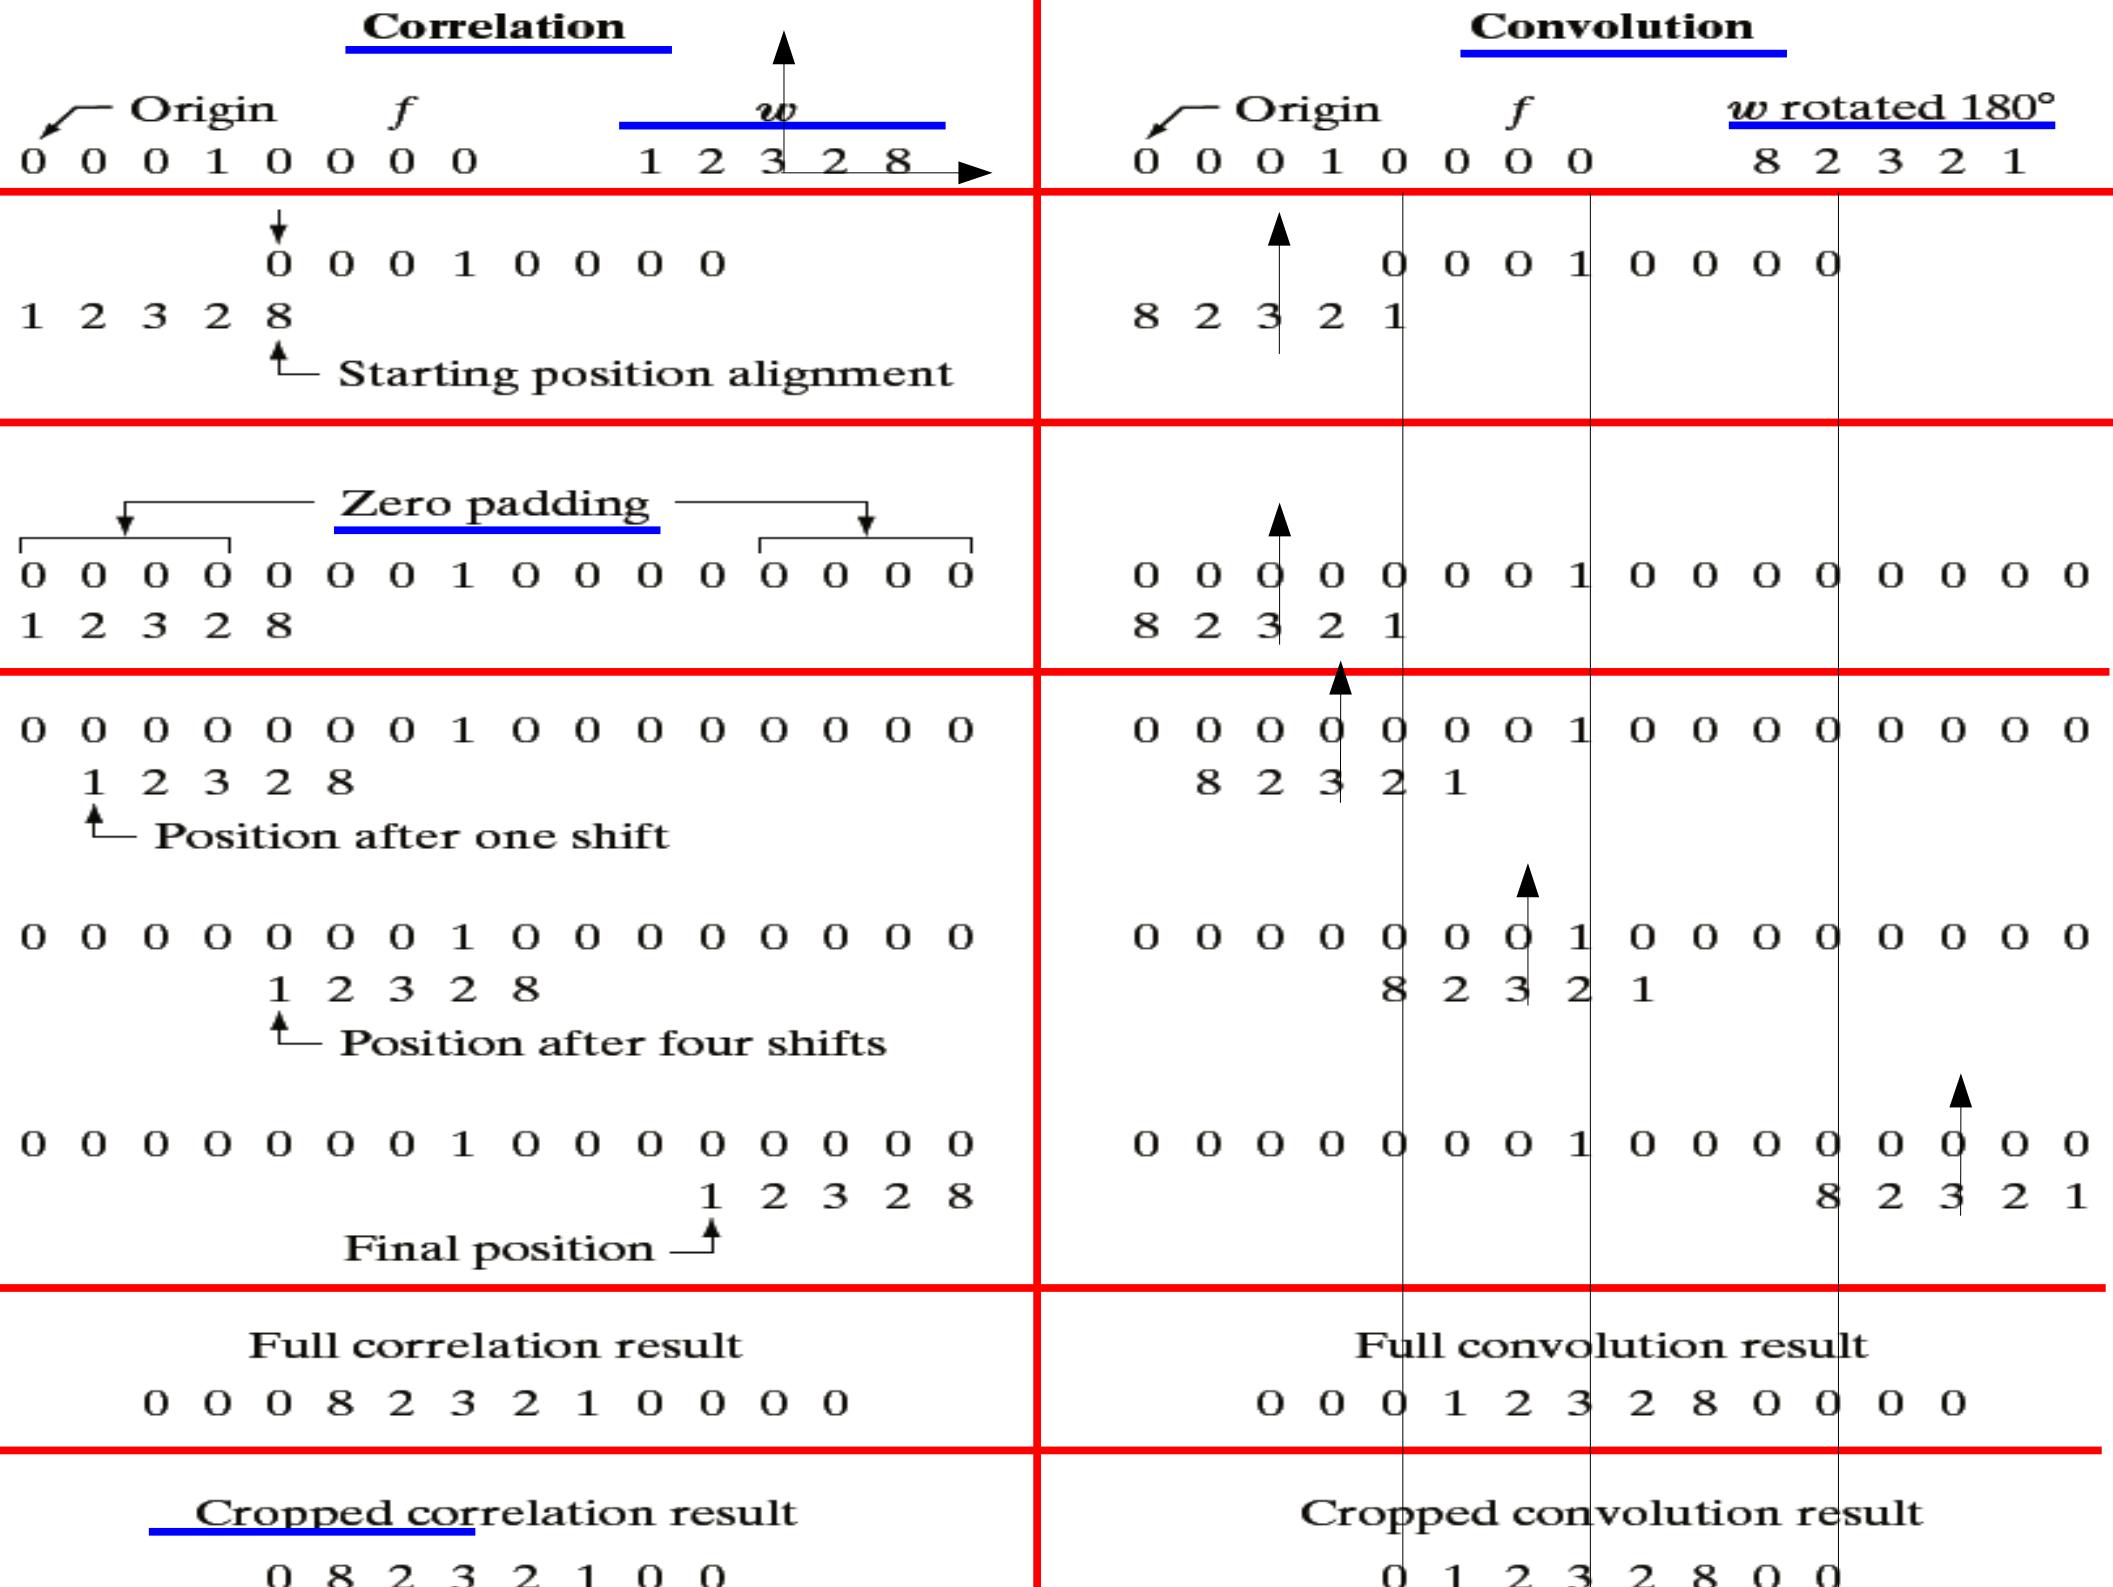
Spatial Filtering

- Cross-Correlation
 - Why does this measure similarity between signals ?
 - Cross-correlation is large when F and shifted-H signals have similar signs
 - When F is positive, shifted-H is also positive
 - When F is negative, shifted-H is also negative
 - i.e., F and shifted-H increase and decrease in sync

Spatial Filtering

- Cross-Correlation





Padded f

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Origin $f(x, y)$

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

$w(x, y)$

1	2	3
4	5	6
7	8	9

(a)

(b)

Initial position for w

1	2	3	0	0	0	0	0	0
4	5	6	0	0	0	0	0	0
7	8	9	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

(c)

Full correlation result

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Cropped correlation result

0	0	0	0	0
0	9	8	7	0
0	6	5	4	0
0	3	2	1	0
0	0	0	0	0

(e)

Full convolution result

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	1	2	3	0	0
0	0	0	0	4	5	6	0	0
0	0	0	0	0	7	8	9	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Cropped convolution result

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

(f)

(g)

(h)

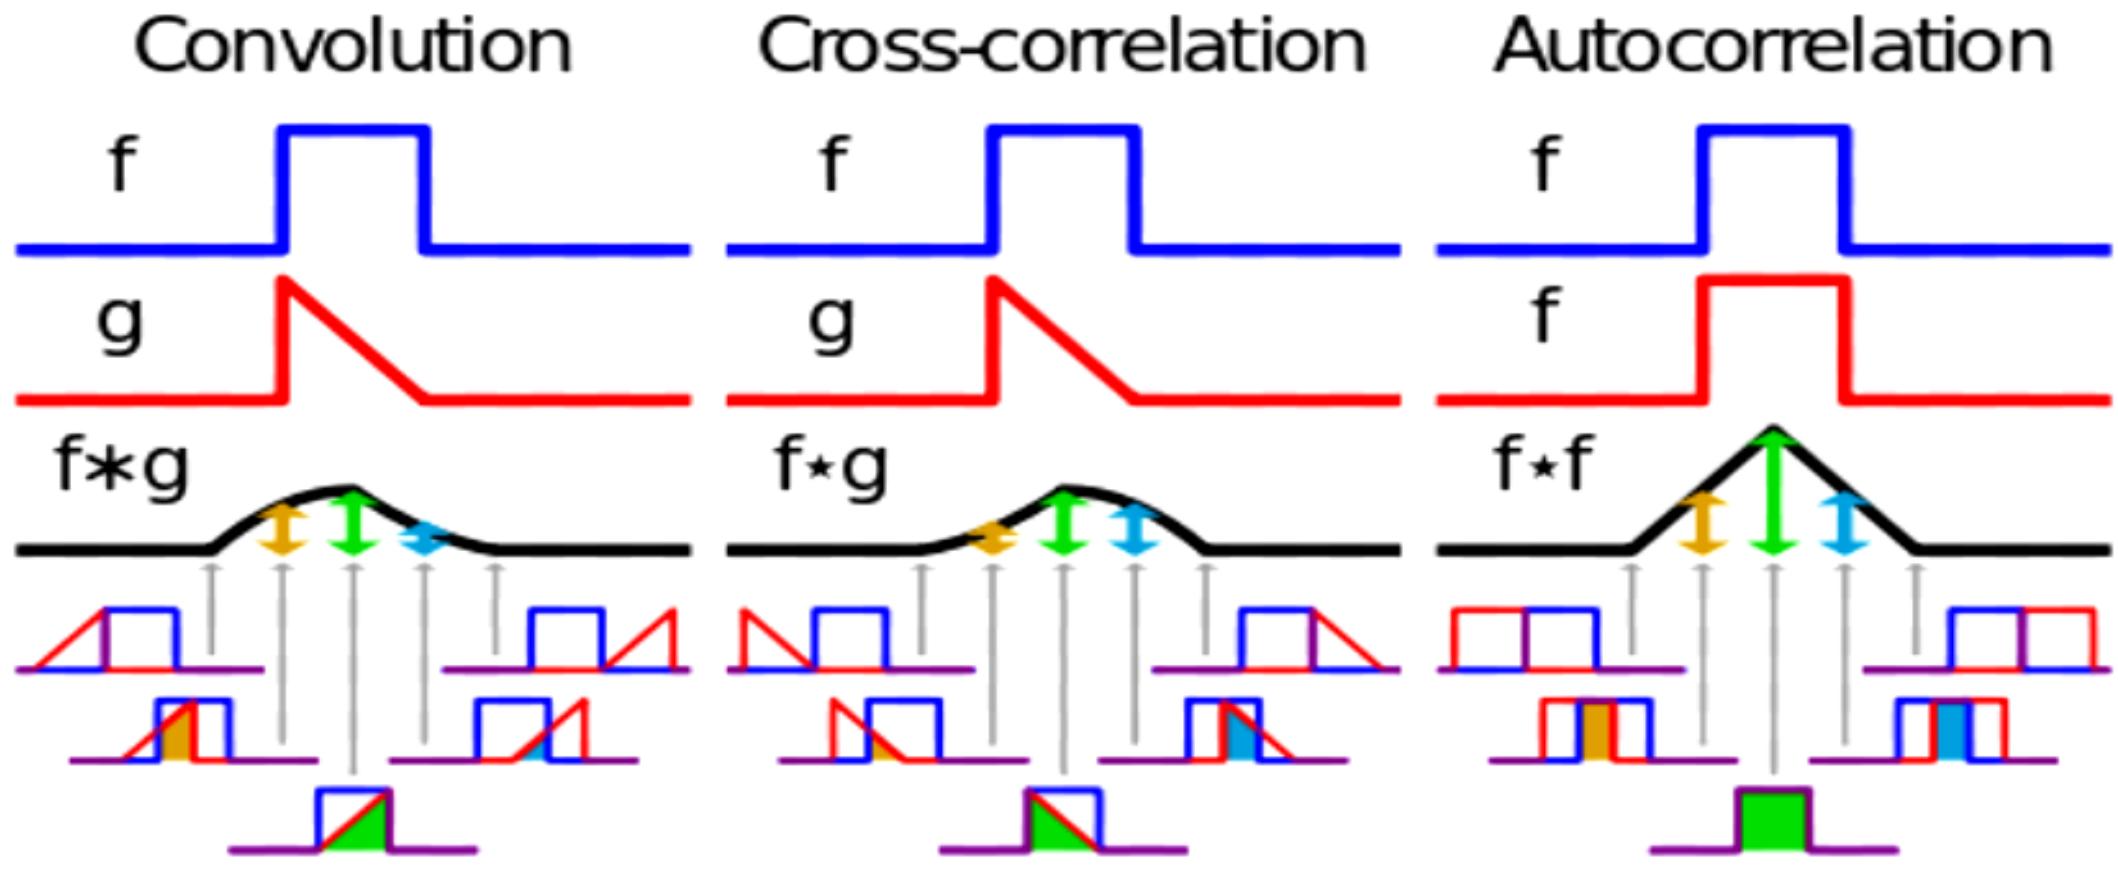
Spatial Filtering

- $w(\cdot)$ = convolution mask
- Convolution property
 - If convolution mask is symmetric about x and y axes, then convolved image = cross-correlated image !
 - See previous slide ...

Spatial Filtering

- **Auto-correlation**

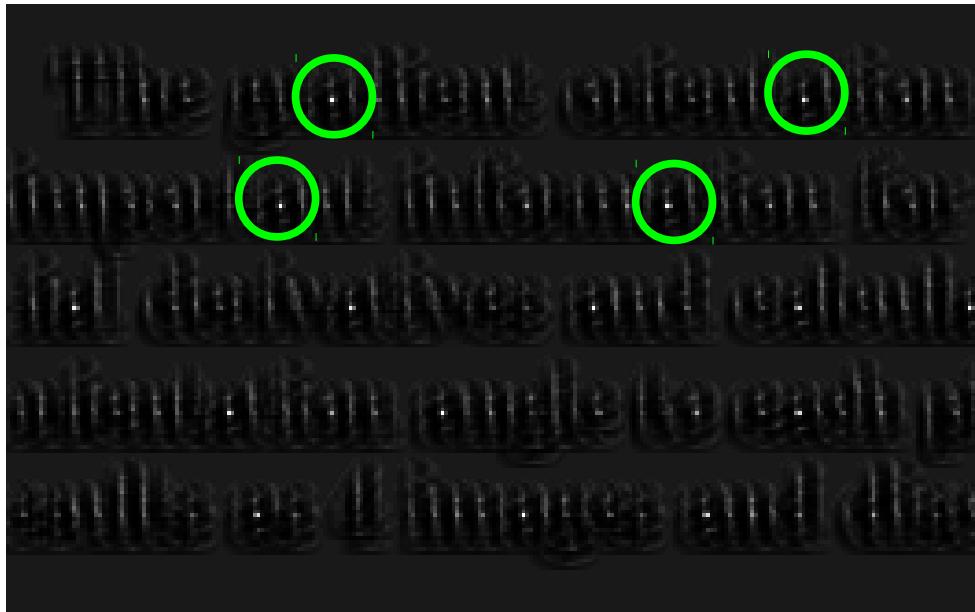
- Cross-correlation of a signal with itself
 - Similar to inner product of signal with itself
 - Reaches a maximum when shift = 0
 - Other shifts may also reach maximum (e.g., periodic signals)



Spatial Filtering

- Cross-correlation applied for object detection
 - Template matching
 - Template = a
 - Binary image $\{-1, 1\}$
 - Image = (below)
 - Cross correlation (down)

The gradient orientation is important information for spatial derivatives and calculating orientation angle to each pixel results as 4 images and disc



a	a	a
a	a	a
a	a	a
a	a	a

Input image



Template



Correlation

What is the problem here ?



Spatial Filtering

- **Normalized Cross-Correlation**

- Motivation: Want to measure similarity between 2 signals, without considering effects of :
 - (1) Intensity scaling
 - (2) Intensity shifts
- How to do this algorithmically ? $NCC(F, H) = ?$
 - (1) Subtract, from F , mean of $F \rightarrow F'$
 - Handles changes to intensity shifts
 - (2) Divide each value in F' by the root-mean-square of intensities in $F' \rightarrow F''$
 - Handles changes in intensity scale
 - (3) Do the same for $H \rightarrow H''$
 - (4) Perform cross-correlation between F'' and H''

normalization

Input image



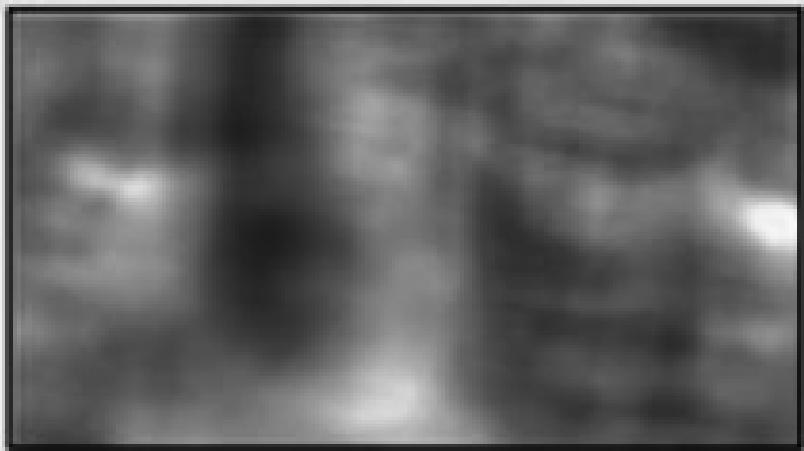
Template



Correlation



Normalized cross correlation



Spatial Filtering

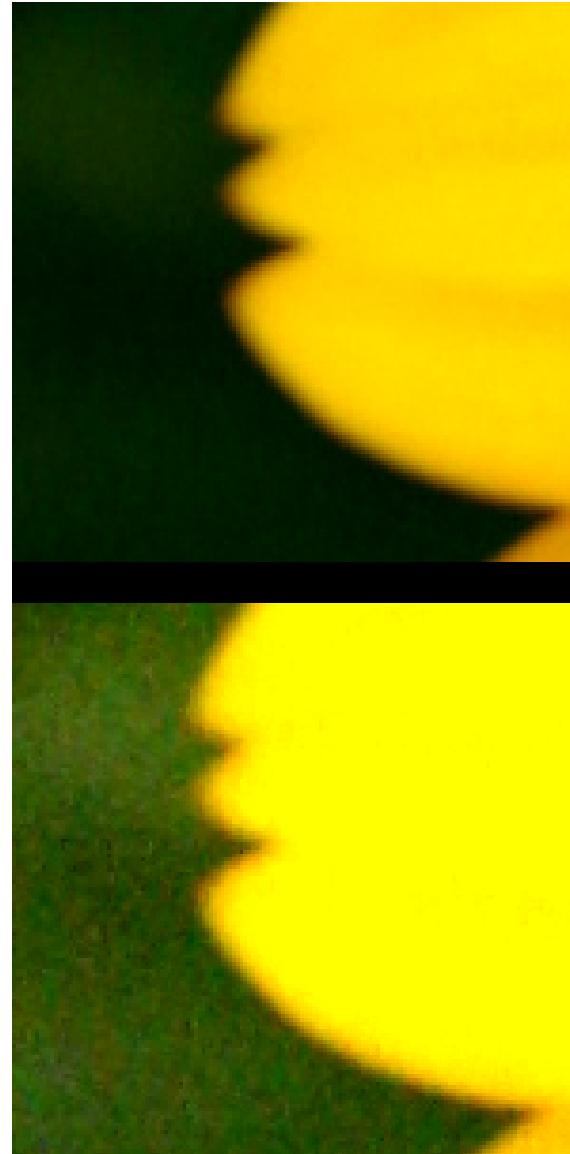
- Applications of spatial filtering
 - Blur / smooth image
 - Reduce image noise
 - Sharpen image
 - Edge detection
 - What is an edge ?
 - ...
 - How do we design weights / masks $w(i,j)$ for these tasks ?

Spatial Filtering

- Image noise
 - Corruption in the image in the form of **random fluctuations** in image intensities
 - Caused by:
 - Measurement errors
 - Fluctuations in signal itself

Spatial Filtering

- Types of image noise
 - Gaussian (additive) →
 - Mean zero
 - Salt and pepper / impulse
 - e.g., intensity spikes
 - e.g., data transmission error



Spatial Filtering

- Increasing levels of Gaussian noise
 - Increasing variance (or standard deviation)



Spatial Filtering

- Increasing levels of salt-and-pepper noise
 - Increasing the number of pixels that get assigned black or white intensities



Spatial Filtering

- **Mean filter (also called box filter)**

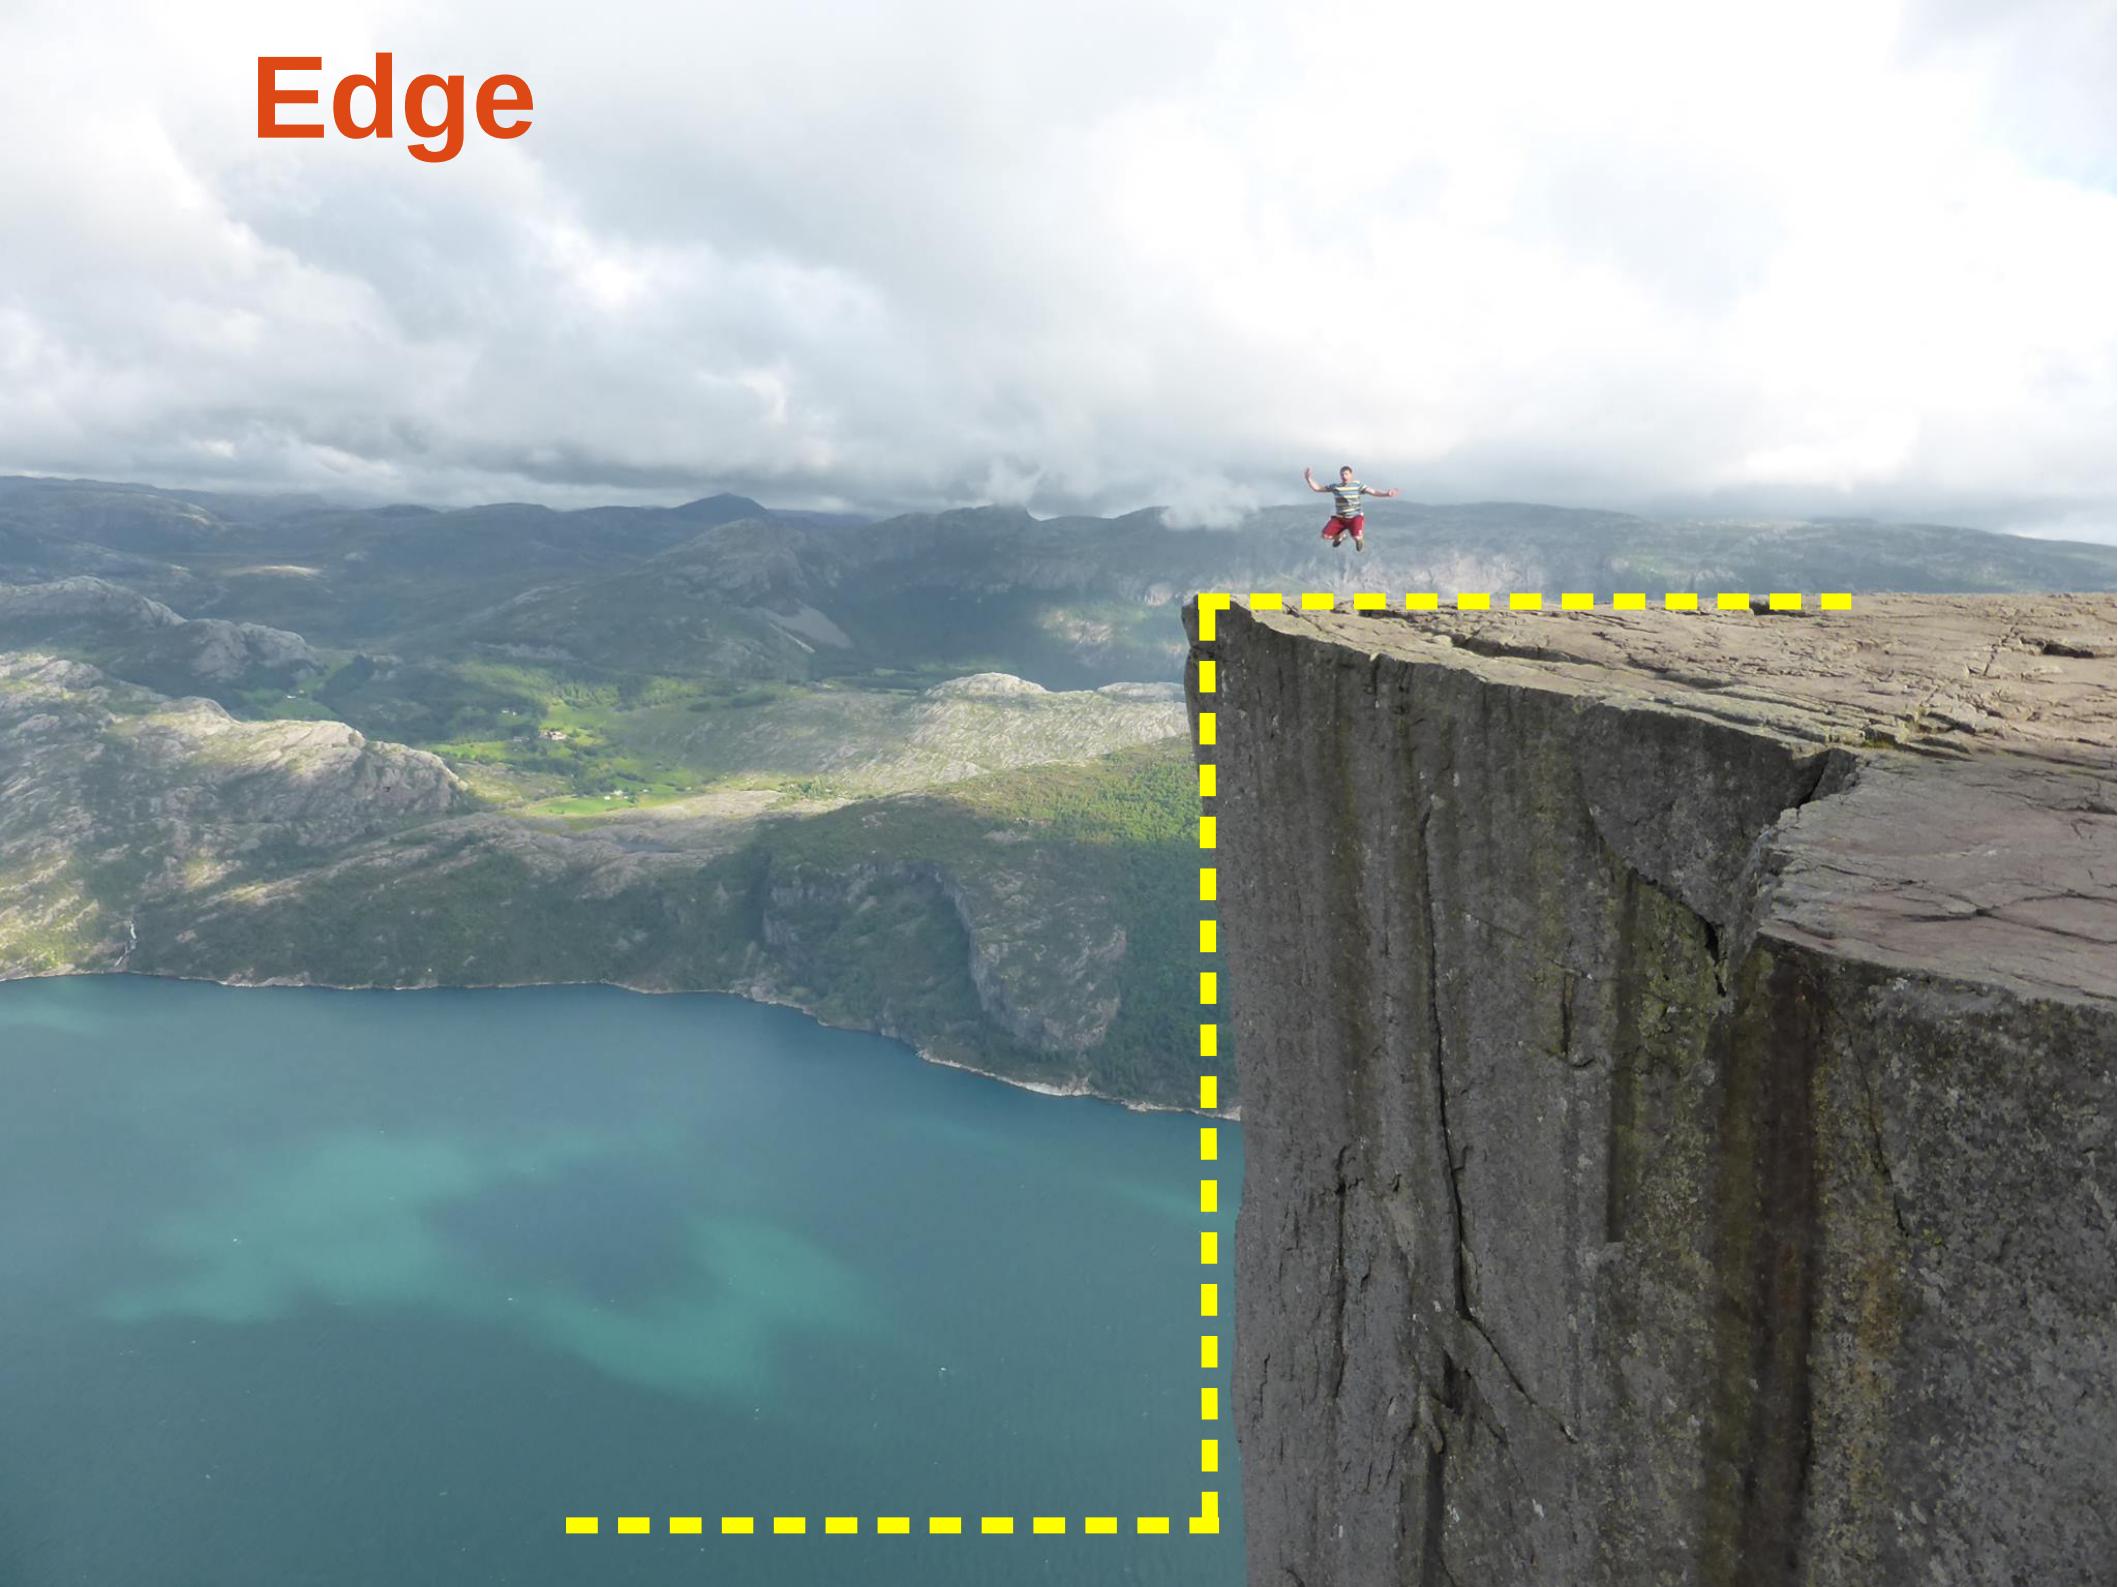
- Replace pixel value by **mean** of pixel values in window around pixel
- Mask for convolution
 - We want sum of weights = 1
 - Why ?
 - To preserve intensity range of the original image in the averaged image
 - How to handle boundary pixels ?
 - If mask cropped, then redefine scaling factor
 - How does the mean filter affect edges ?

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

$$\frac{1}{9} \times$$

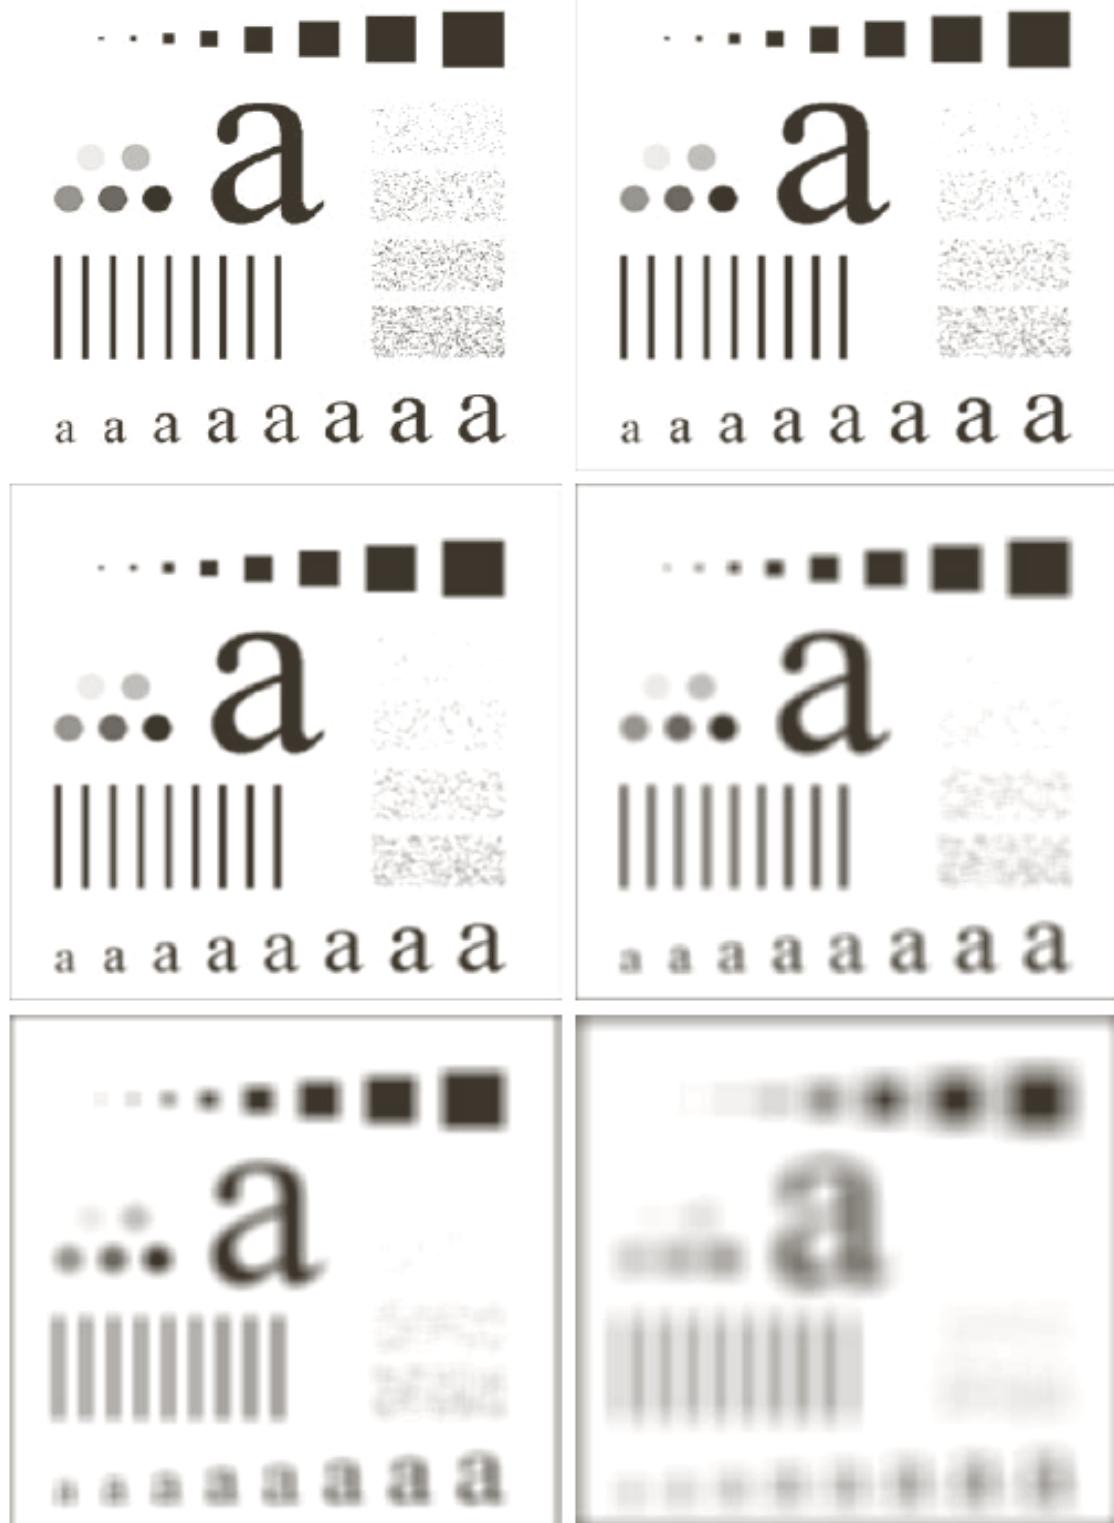
1	1	1
1	1	1
1	1	1

Edge



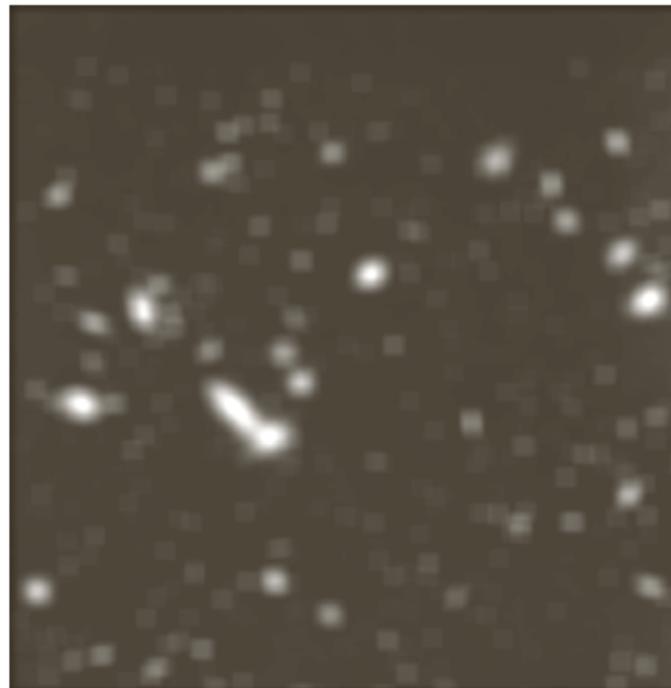
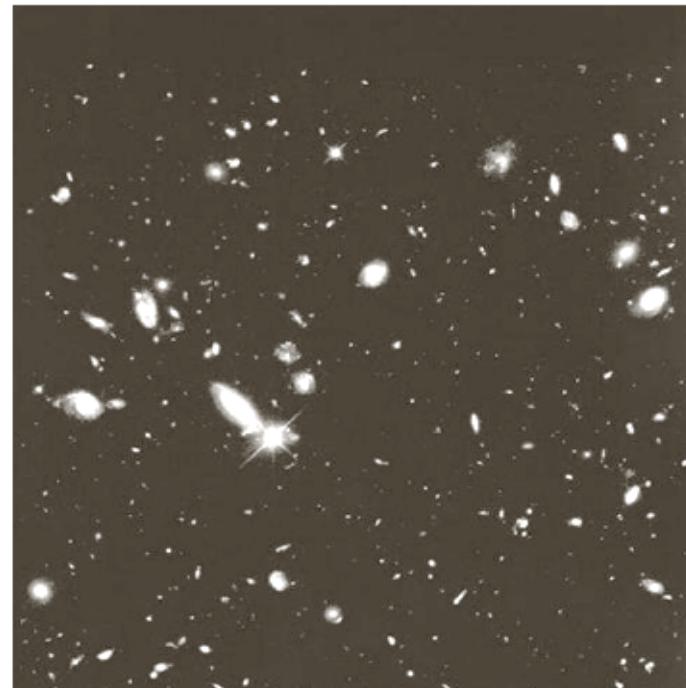
Spatial Filtering

- Mean filter
 - Different sizes
 - Original, box 3
box 5, box 9,
box 15, box 25
 - Circle diameter
(in image)
= 25 pixels



Spatial Filtering

- Mean filter
 - Left: image from Hubble space telescope
 - Middle: mean-filtered image
 - Right: thresholded image
 - Notice the **square-shaped blobs** (artifact)



Spatial Filtering

- How fast is convolution ?
 - If image size = $M \times N$
 - If mask size = $P \times Q$
 - Then, convolution takes $M N P Q$ operations
 - Can this be made faster ?
 - In some special cases ?
 - Yes ! How ?

Spatial Filtering

- How fast is convolution ?
 - Observe a trick
 - Let $U = 1D \text{ column image} = [u_1 \ u_2 \ u_3 \ u_4]^\top$
 - Let $V = 1D \text{ row image} = [v_1 \ v_2 \ v_3]$
 - Then, 2D convolution of U and V = ?
 - = (outer) product of U and V

$$U * V = U V = \begin{bmatrix} u_1 \\ u_2 \\ u_3 \\ u_4 \end{bmatrix} \begin{bmatrix} v_1 & v_2 & v_3 \end{bmatrix} = \begin{bmatrix} u_1 v_1 & u_1 v_2 & u_1 v_3 \\ u_2 v_1 & u_2 v_2 & u_2 v_3 \\ u_3 v_1 & u_3 v_2 & u_3 v_3 \\ u_4 v_1 & u_4 v_2 & u_4 v_3 \end{bmatrix}$$

Spatial Filtering

- How fast is convolution ?
 - Combine ideas from (1) trick (2) associativity of convolution
 - Let F be a 2D image
 - Let $U = 1D$ column image
 - Let $V = 1D$ row image
 - Then, $(U * V) * F = U * (V * F)$
 - $U * V$ is a 2D mask
 - $V * F$ = convolution of each row of F with V
 - Output $F1$ = same size (or slightly larger) than F
 - $U * F1$ = convolution of each column of $F1$ with U
 - Output $F2$ = same size (or slightly larger) than $F1$

Spatial Filtering

- How fast is convolution ?
 - Let F = image of size $M \times N$
 - Let H = convolution mask of size $P \times Q$
 - If $H = \text{outer product } U * V$
 - U is of size $P \times 1$
 - V is of size $1 \times Q$
 - Then, $H * F = (U * V) * F = U * (V * F)$
 - This can be computed faster than before:
 - Operations needed for $F1 = V * F = Q (M \ N)$
 - Operations needed for $U * F1 = P (M \ N)$
 - Total operations needed for $U * (V * F) = M \ N \ (P + Q)$

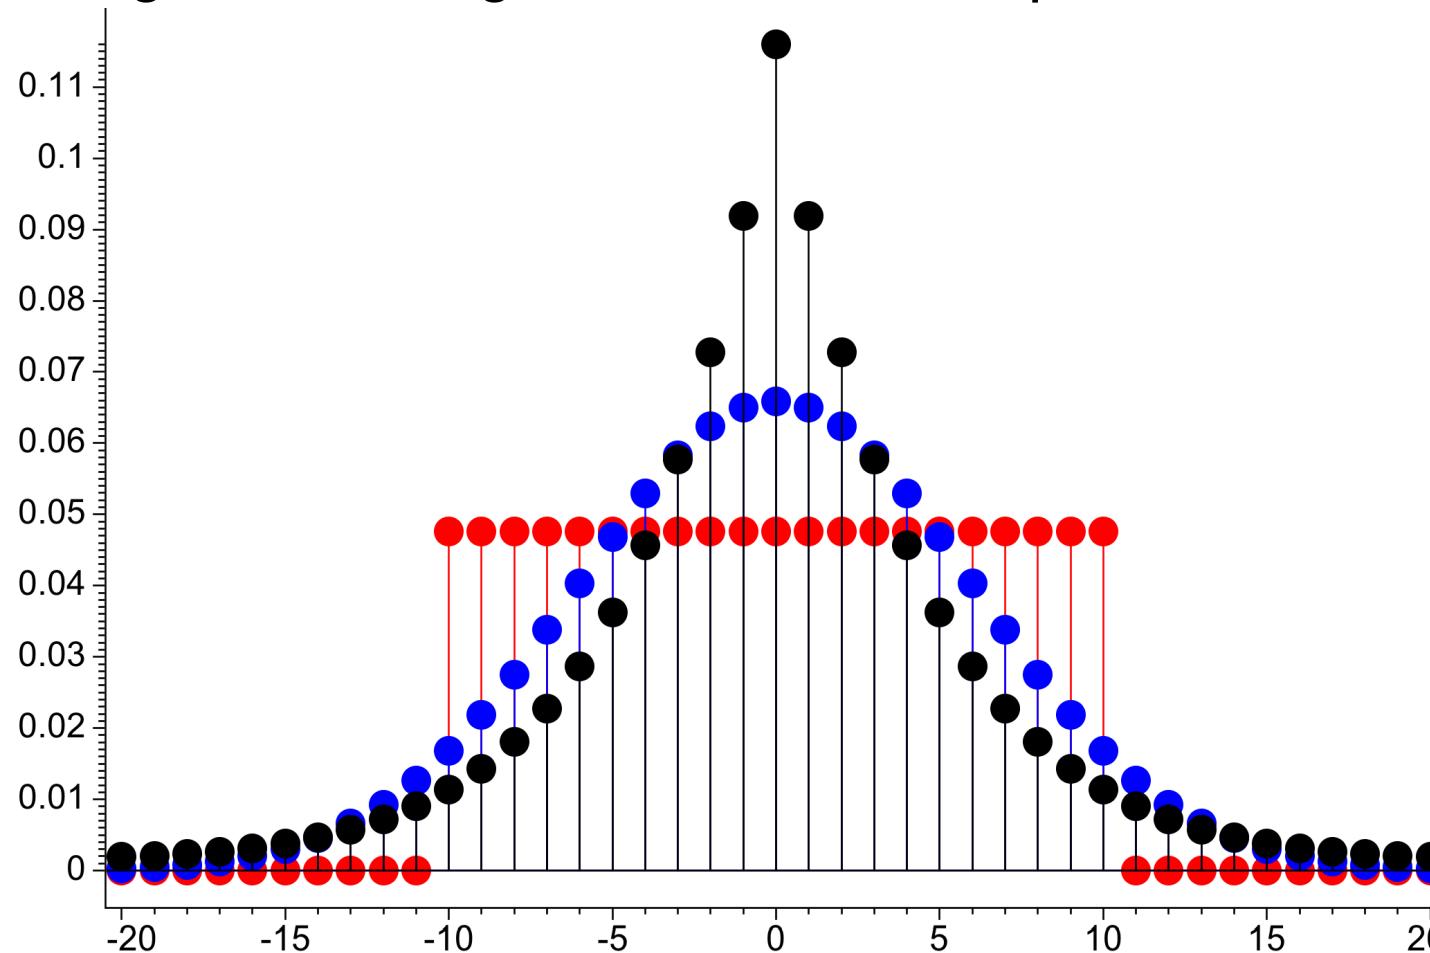
Spatial Filtering

- **Separable** filter = filter that can be written as an outer product
 - Separable filter leads to faster algorithm for convolution
 - Is the mean filter separable ?

$$\frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \frac{1}{3} [1 \ 1 \ 1] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

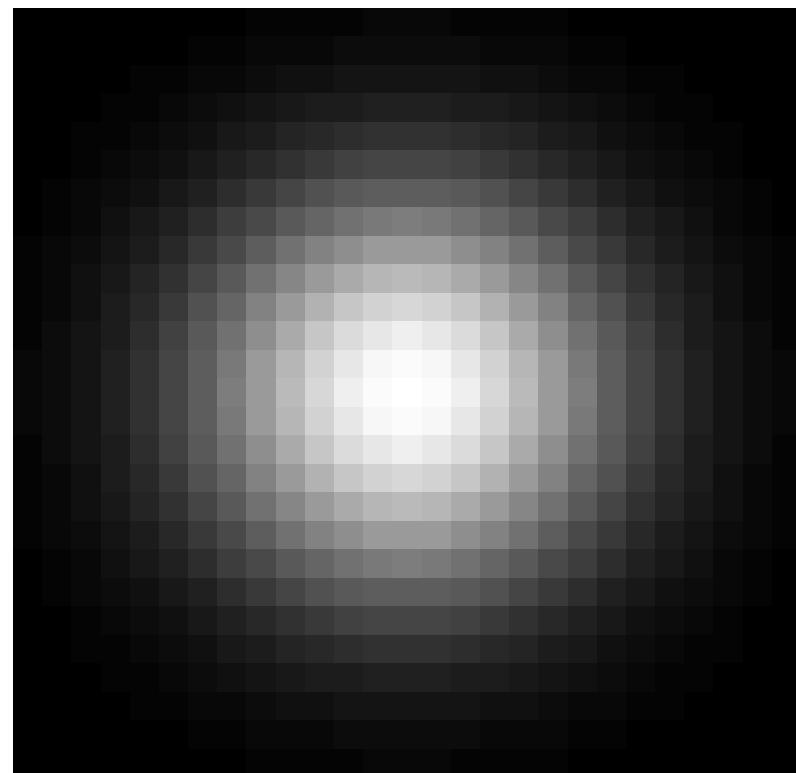
Spatial Filtering

- Weighted-mean filter
 - Mean filter = red
 - Gaussian-weighted filter = blue, black
 - Assign more weight to intensities at pixels closer to center pixel



Spatial Filtering

- Weighted-mean filter
 - 2D symmetric Gaussian mask (before rescaling)
 - Standard deviation = sigma = 2
 - $G(x,y) = \exp(-0.5(x^2 + y^2) / \sigma^2)$
 - Ignore density normalization factor
 - Is this filter separable ?

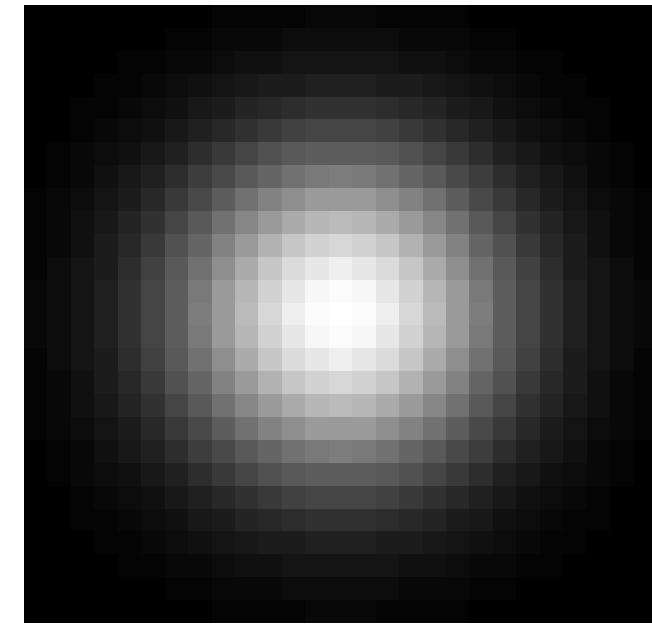


Spatial Filtering

- Weighted-mean filter
 - 2D symmetric Gaussian mask
(before rescaling)
 - Standard deviation = sigma = 2

$$\frac{-(x^2 + y^2)}{2\sigma^2}$$

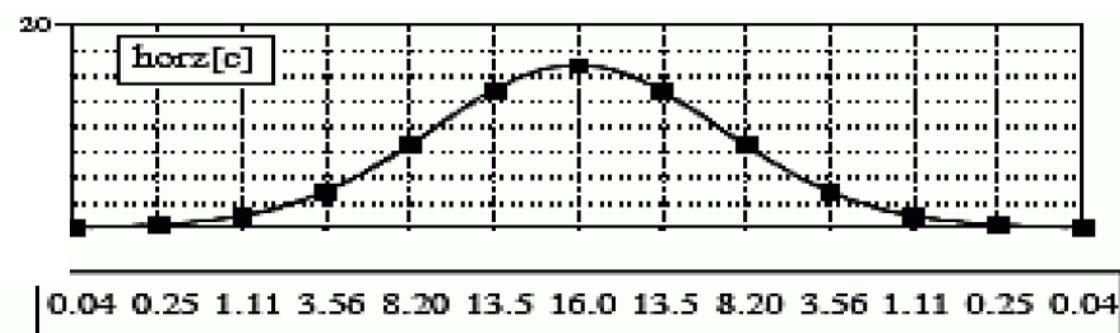
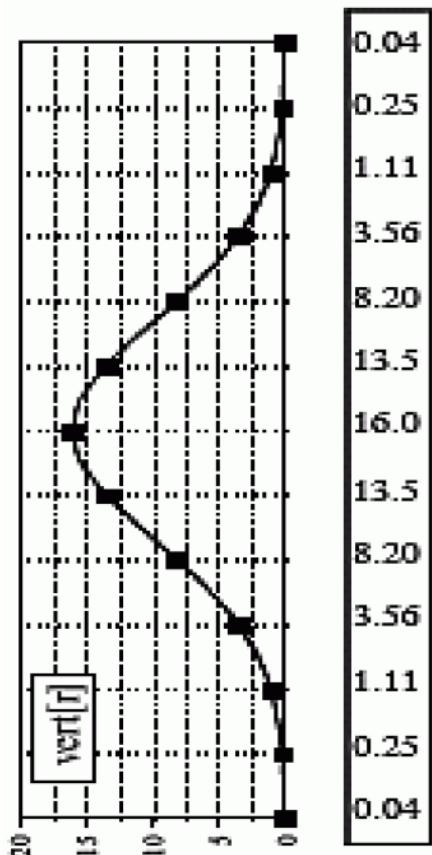
$$g(x, y) = e^{-\frac{(x^2 + y^2)}{2\sigma^2}}$$



0	0	0	0	1	2	2	2	1	0	0	0	0
0	0	1	3	6	9	11	9	6	3	1	0	0
0	1	4	11	20	30	34	30	20	11	4	1	0
0	3	11	26	50	73	82	73	50	26	11	3	0
1	6	20	50	93	136	154	136	93	50	20	6	1
2	9	30	73	136	198	225	198	136	73	30	9	2
2	11	34	82	154	225	255	225	154	82	34	11	2
2	9	30	73	136	198	225	198	136	73	30	9	2
1	6	20	50	93	136	154	136	93	50	20	6	1
0	3	11	26	50	73	82	73	50	26	11	3	0
0	1	4	11	20	30	34	30	20	11	4	1	0
0	0	1	3	6	9	11	9	6	3	1	0	0
0	0	0	0	1	2	2	2	1	0	0	0	0

Spatial Filtering

- Gaussian filter is separable
 - When covariance matrix is diagonal



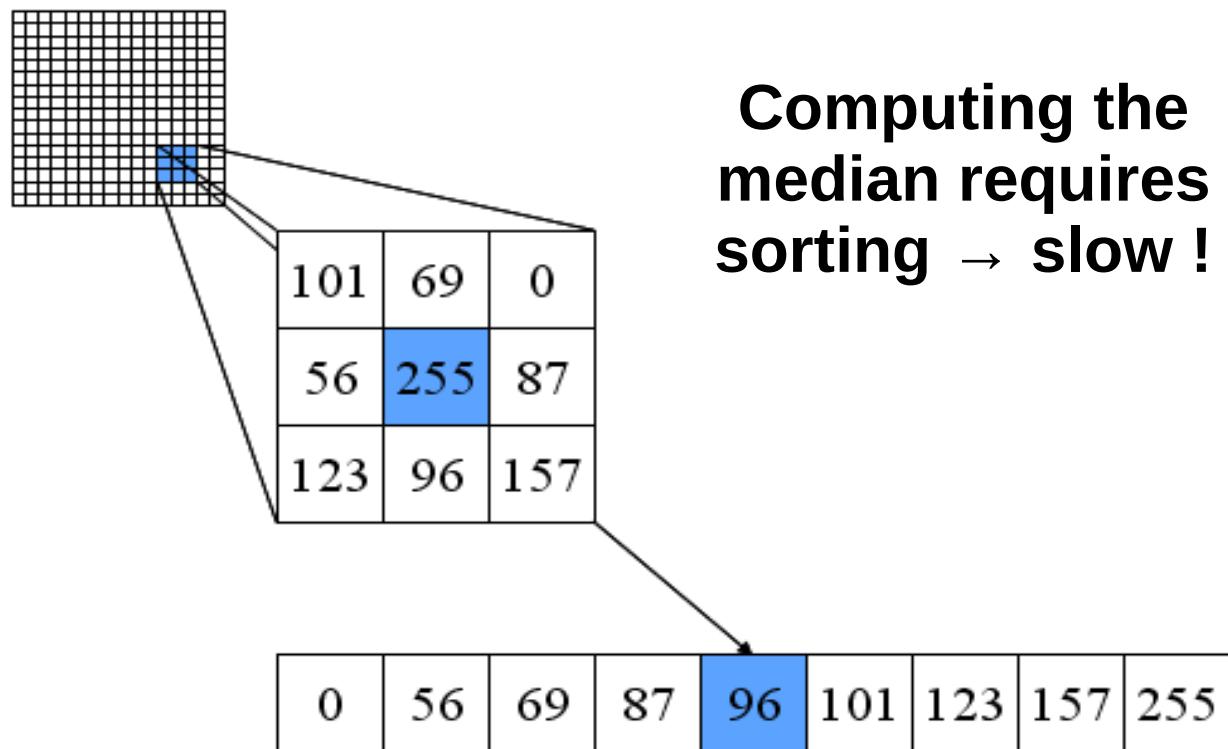
0	0	0	0	0	0	1	1	1	0	0	0	0	0
0	0	0	1	2	3	4	3	2	1	0	0	0	0
0	0	1	4	9	15	18	15	9	4	1	0	0	0
0	1	4	13	29	48	57	48	29	13	4	1	0	0
0	2	9	29	67	111	131	111	67	29	9	2	0	0
1	3	15	48	111	183	216	183	111	48	15	3	1	0
1	4	18	57	131	216	255	216	131	57	18	4	1	0
1	3	15	48	111	183	216	183	111	48	15	3	1	0
0	2	9	29	67	111	131	111	67	29	9	2	0	0
0	1	4	13	29	48	57	48	29	13	4	1	0	0
0	0	1	4	9	15	18	15	9	4	1	0	0	0
0	0	0	1	2	3	4	3	2	1	0	0	0	0
0	0	0	0	0	0	1	1	1	0	0	0	0	0

Spatial Filtering

- Weighted-mean filter versus mean filter
 - Trade offs
 - Assuming mask size is same, ...
 - Weighted-mean filter:
 - (1) blurs less and preserves more edges
 - (2) removes less noise in constant regions
 - (3) rotation-invariant filtering

Spatial Filtering

- Median filter
 - Replace pixel value by **median** of pixel values in window around pixel
 - For $P(X)$, median = m , s.t. $P (X \geq m) \geq 0.5$, $P (X \leq m) \geq 0.5$

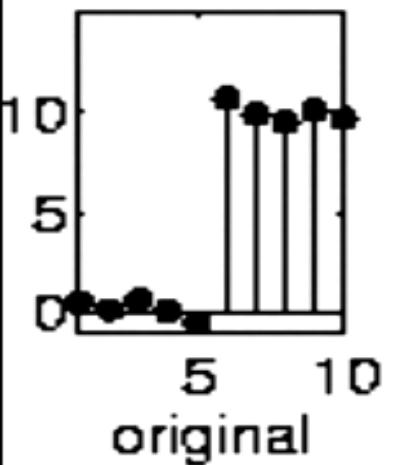


- Trick question
 - How can the median filter be described as a convolution ?

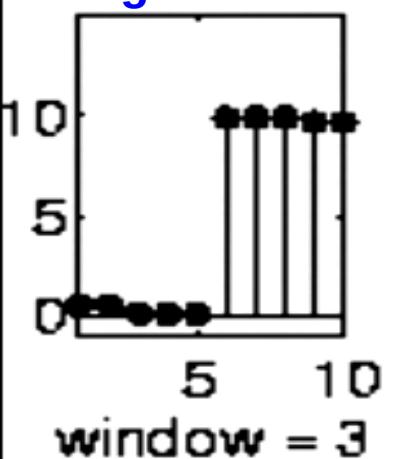
Spatial Filtering

- Median filter
 - Better for removing salt-and-pepper noise, as compared to mean filter
 - Why ?
 - Because (a small level of) salt-and-pepper noise **doesn't** affect the **median** of the original intensities
 - Because (any level of) salt-and-pepper noise **does** affect the **mean** of the original intensities
 - Median is **robust** to **outliers**. Mean isn't.
 - **Outlier** = an observation / outcome that is **extremely** different from others in population
 - How does median filter affect an edge within an image ?

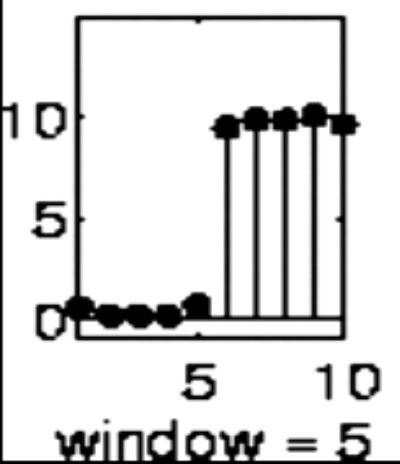
MEDIAN



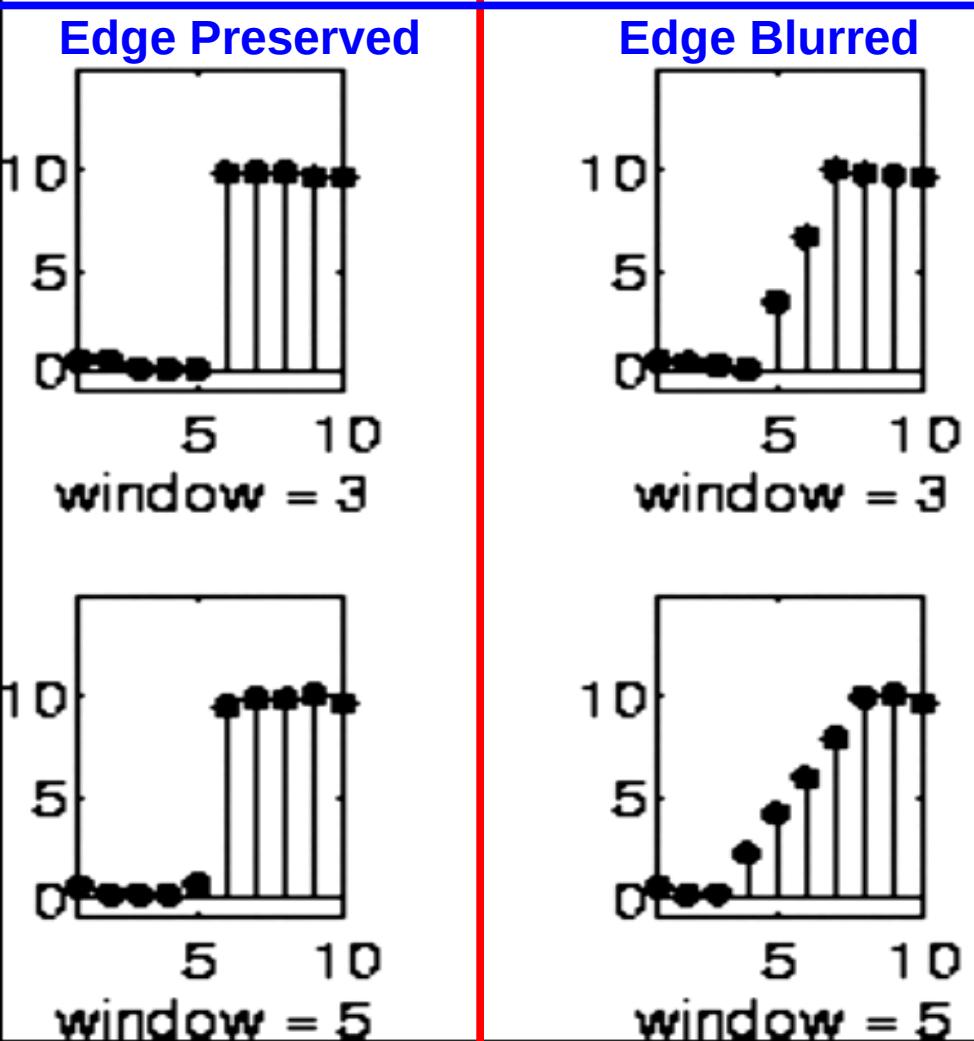
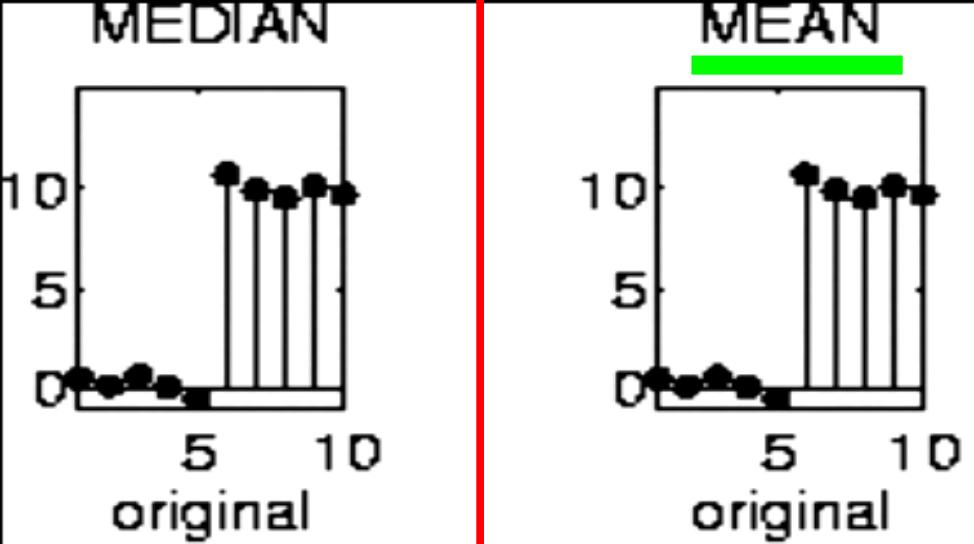
Edge Preserved

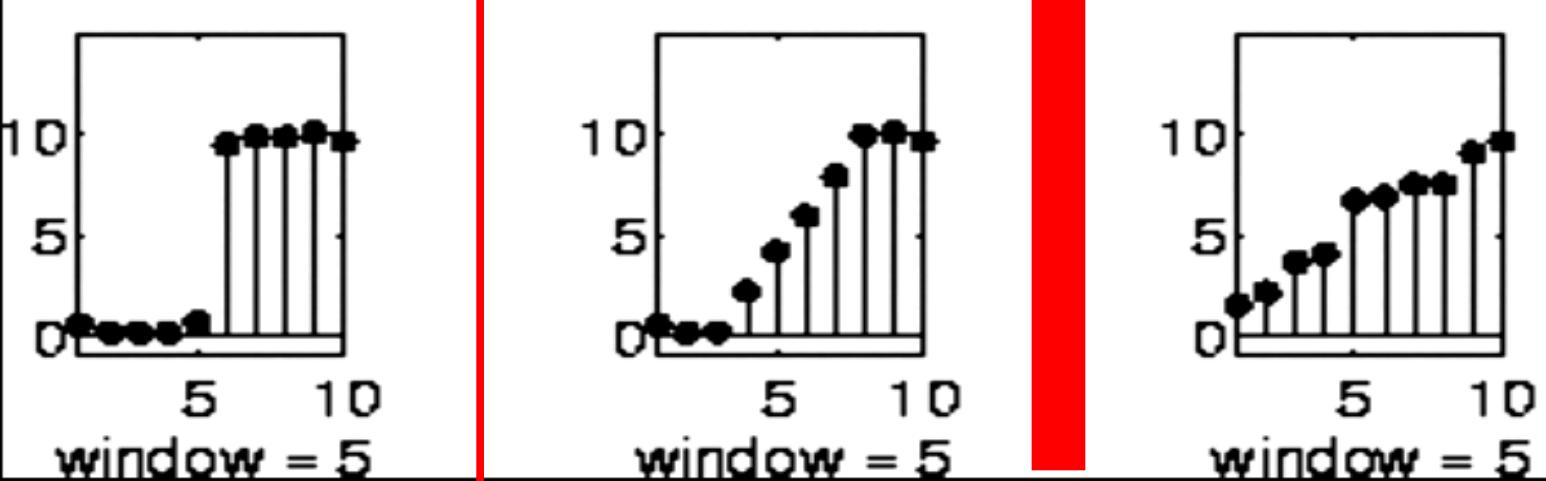
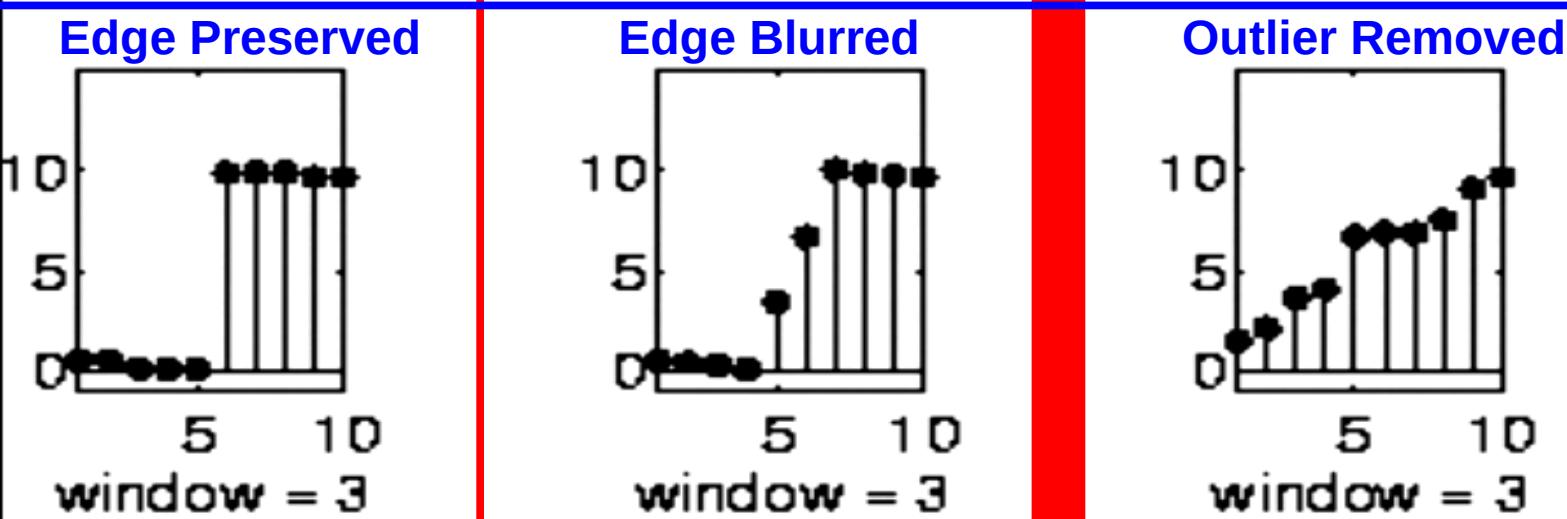
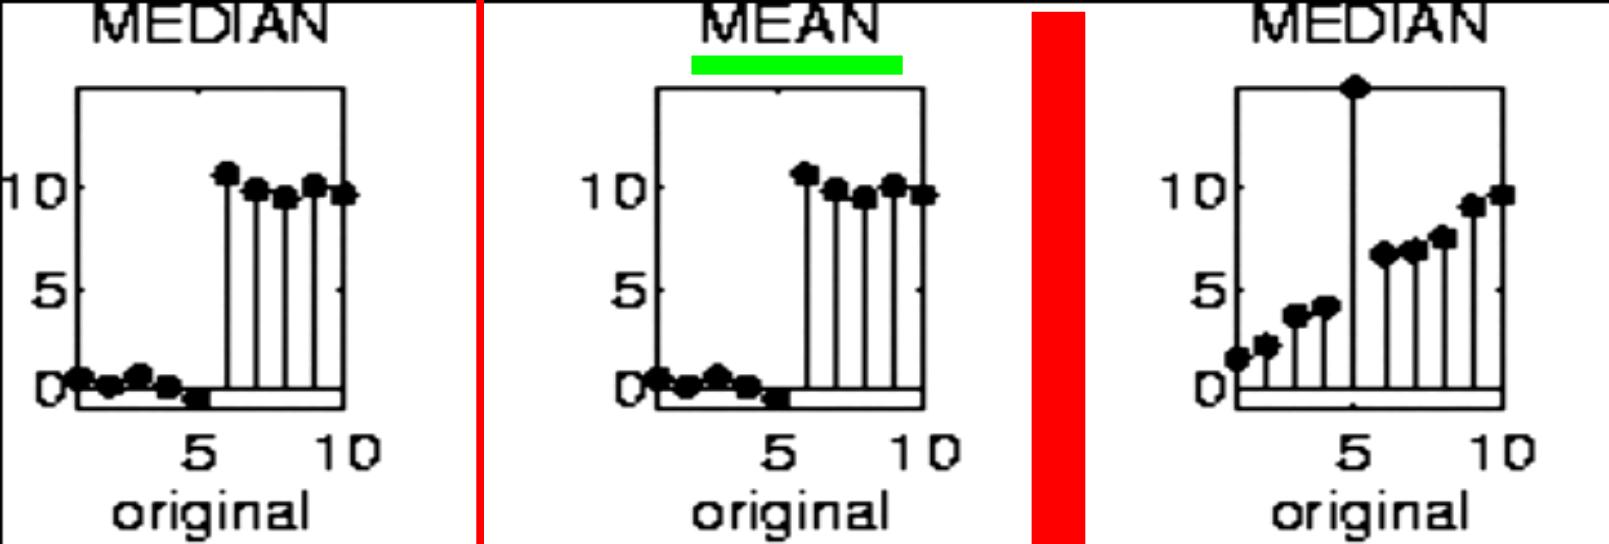


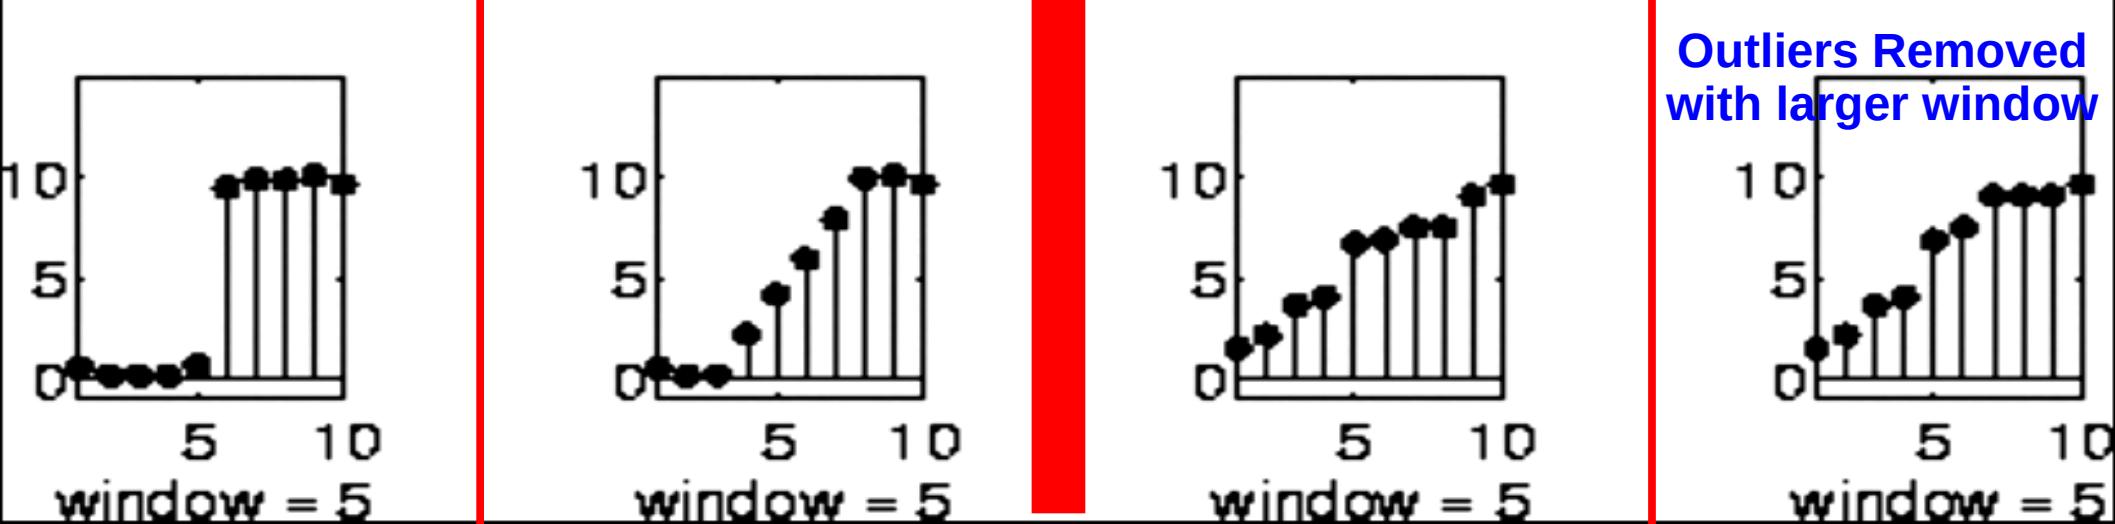
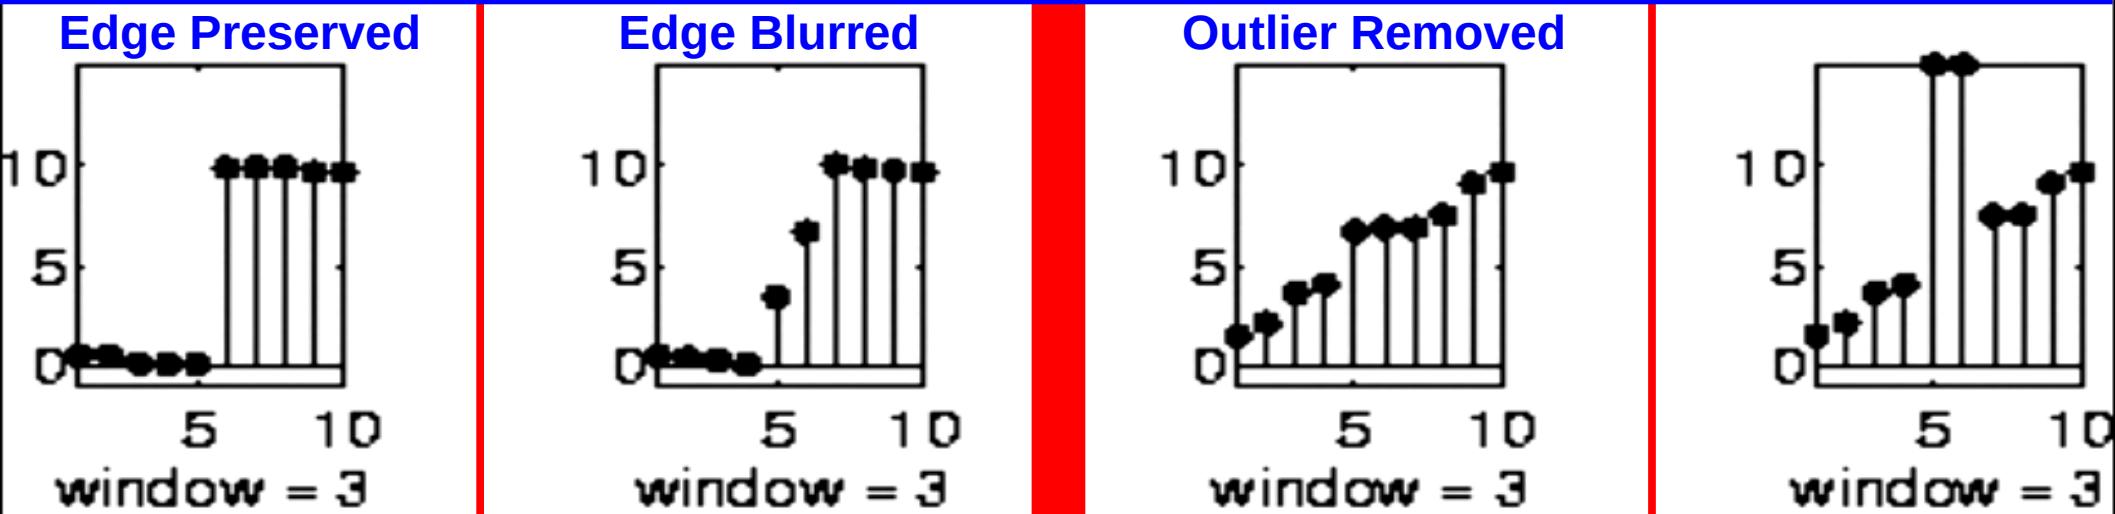
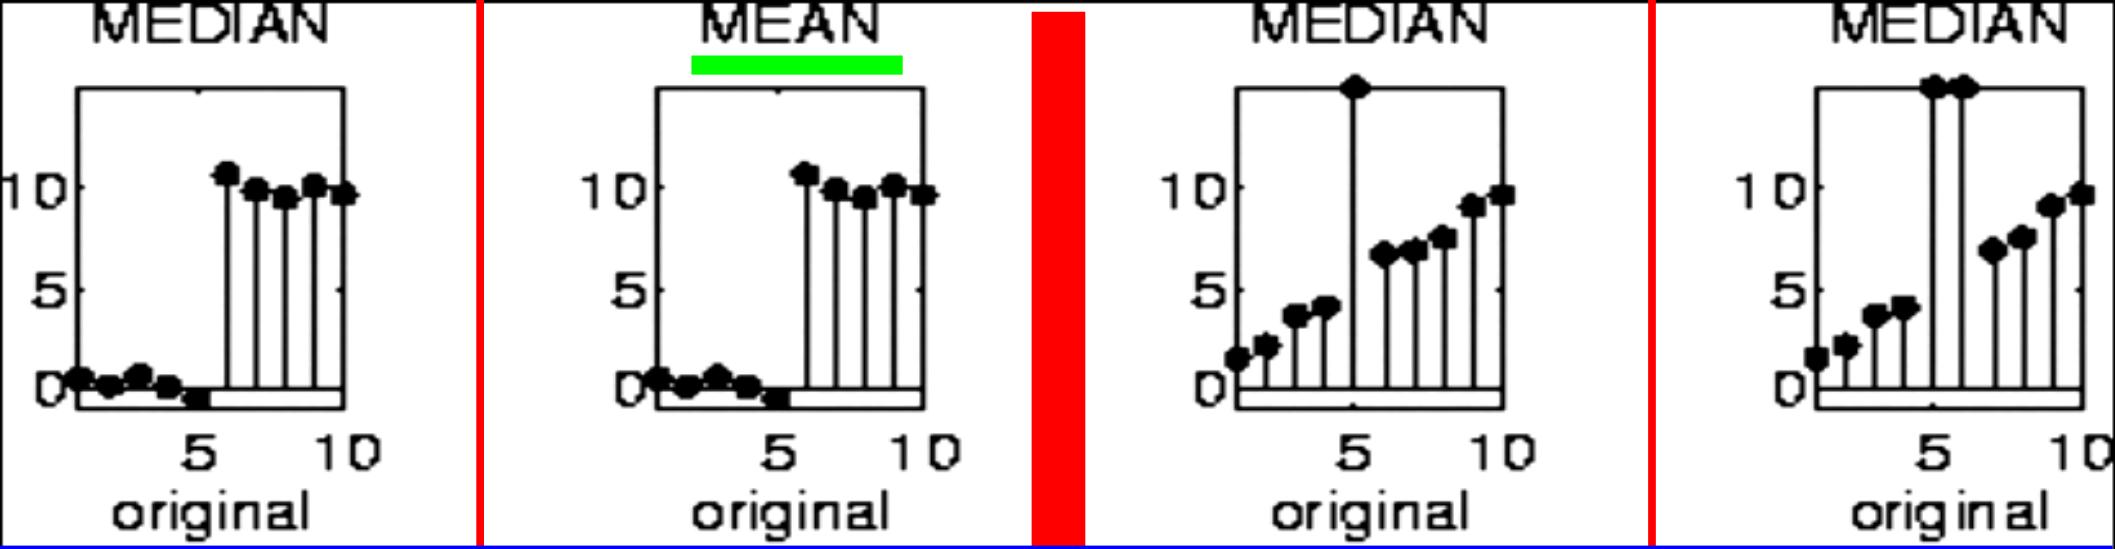
window = 3



window = 5







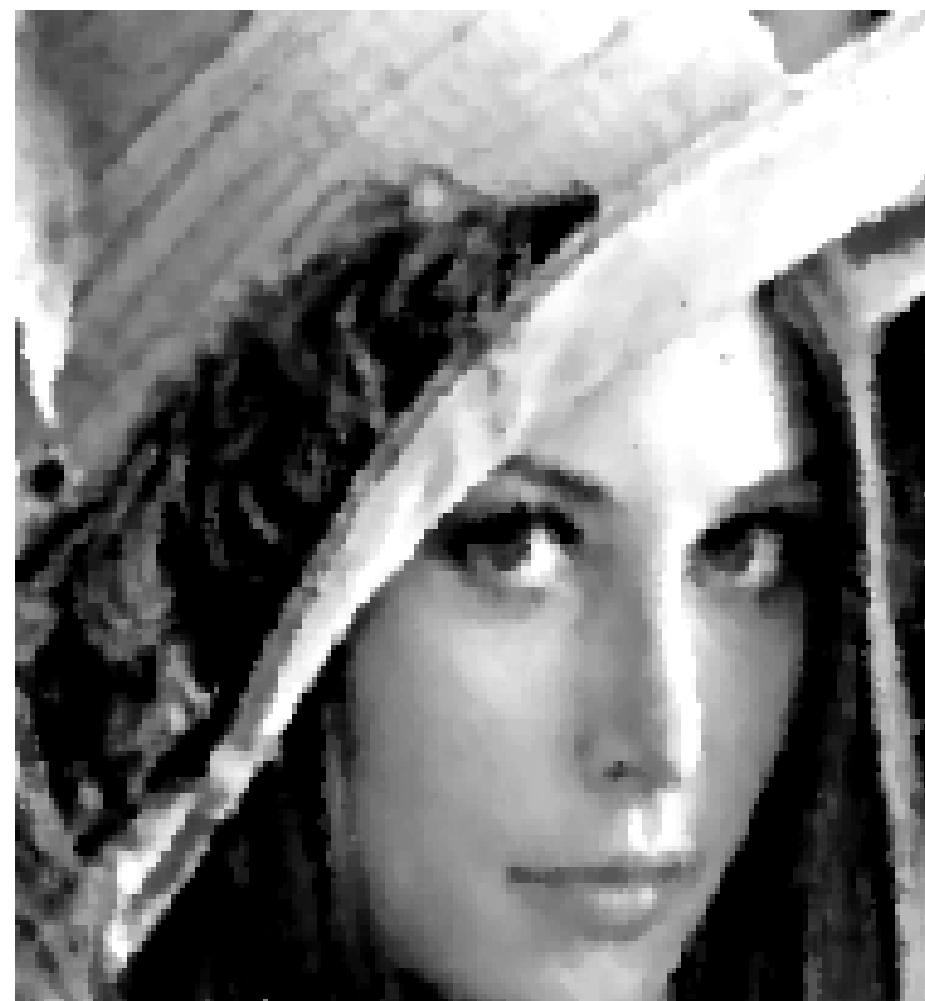
Spatial Filtering

- Median filter on salt-and-pepper noise



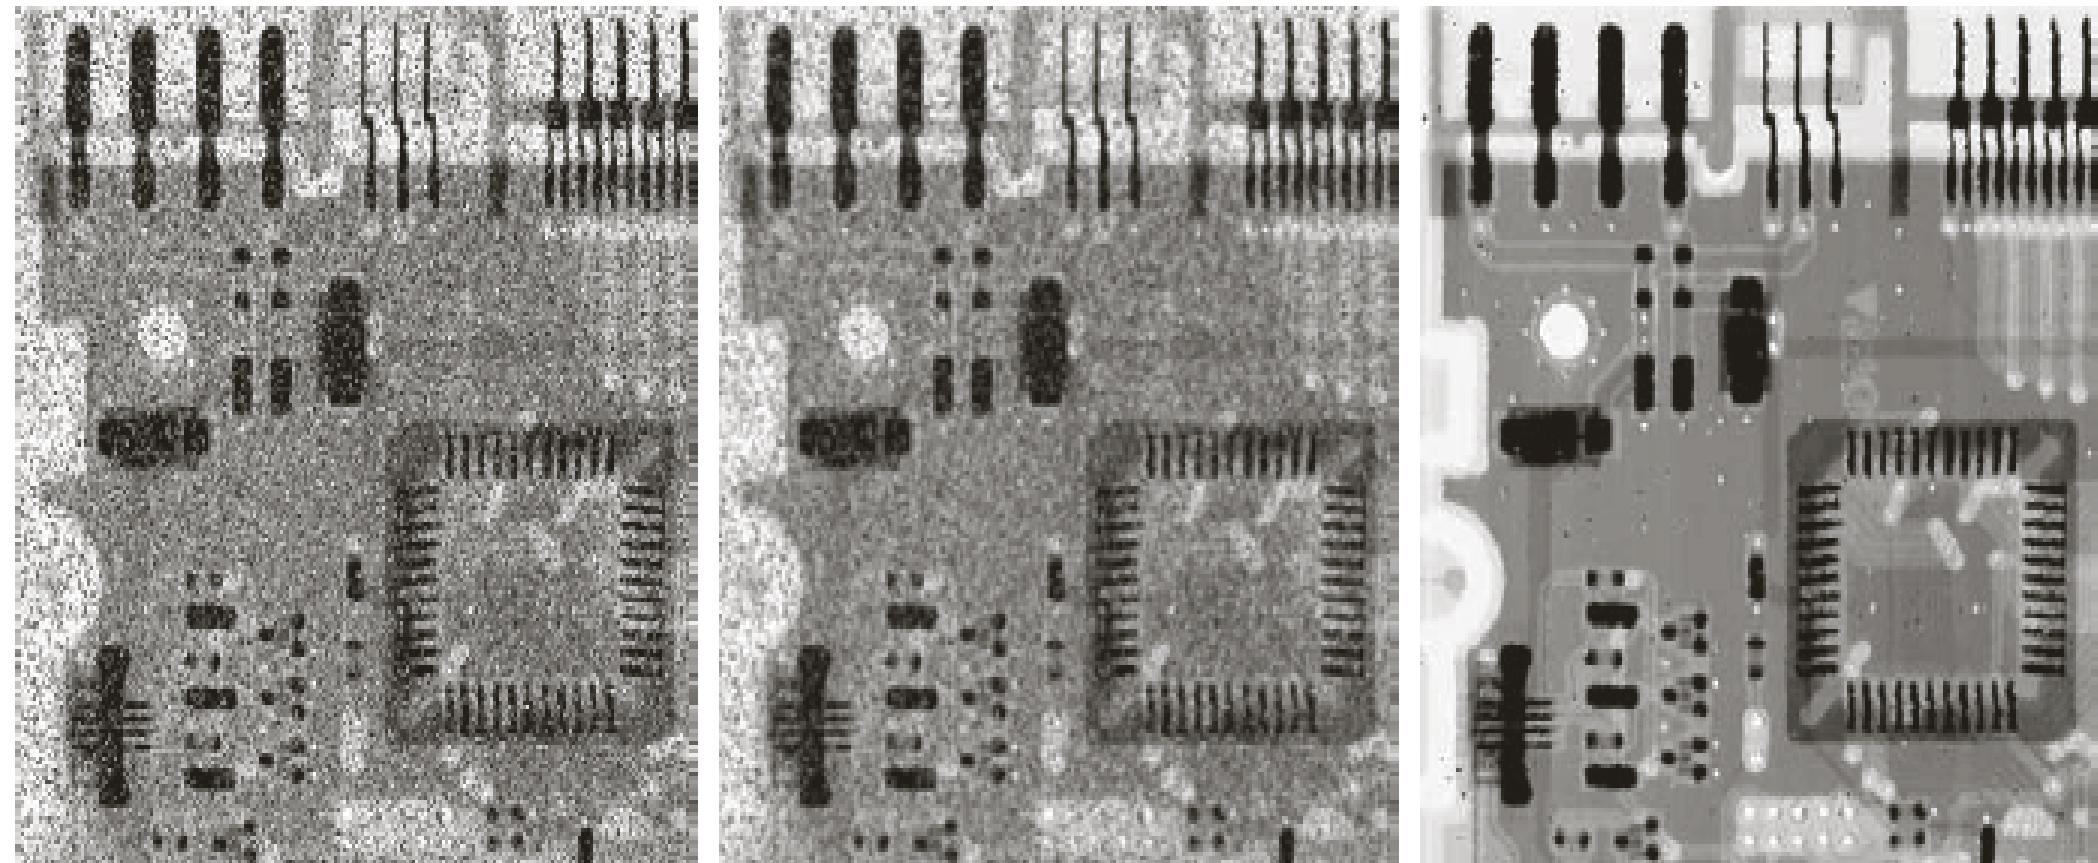
Spatial Filtering

- Median filter on salt-and-pepper noise
 - (Surprisingly) Good on large amounts of impulse noise



Spatial Filtering

- Median filter on salt-and-pepper noise



a | b | c

FIGURE 3.35 (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a 3×3 averaging mask. (c) Noise reduction with a 3×3 median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

Spatial Filtering

- Median filter on Gaussian noise
 - Can preserves edges better
 - But, can create artificial edges in smooth regions (e.g., forehead)

noisy



Gaussian filter

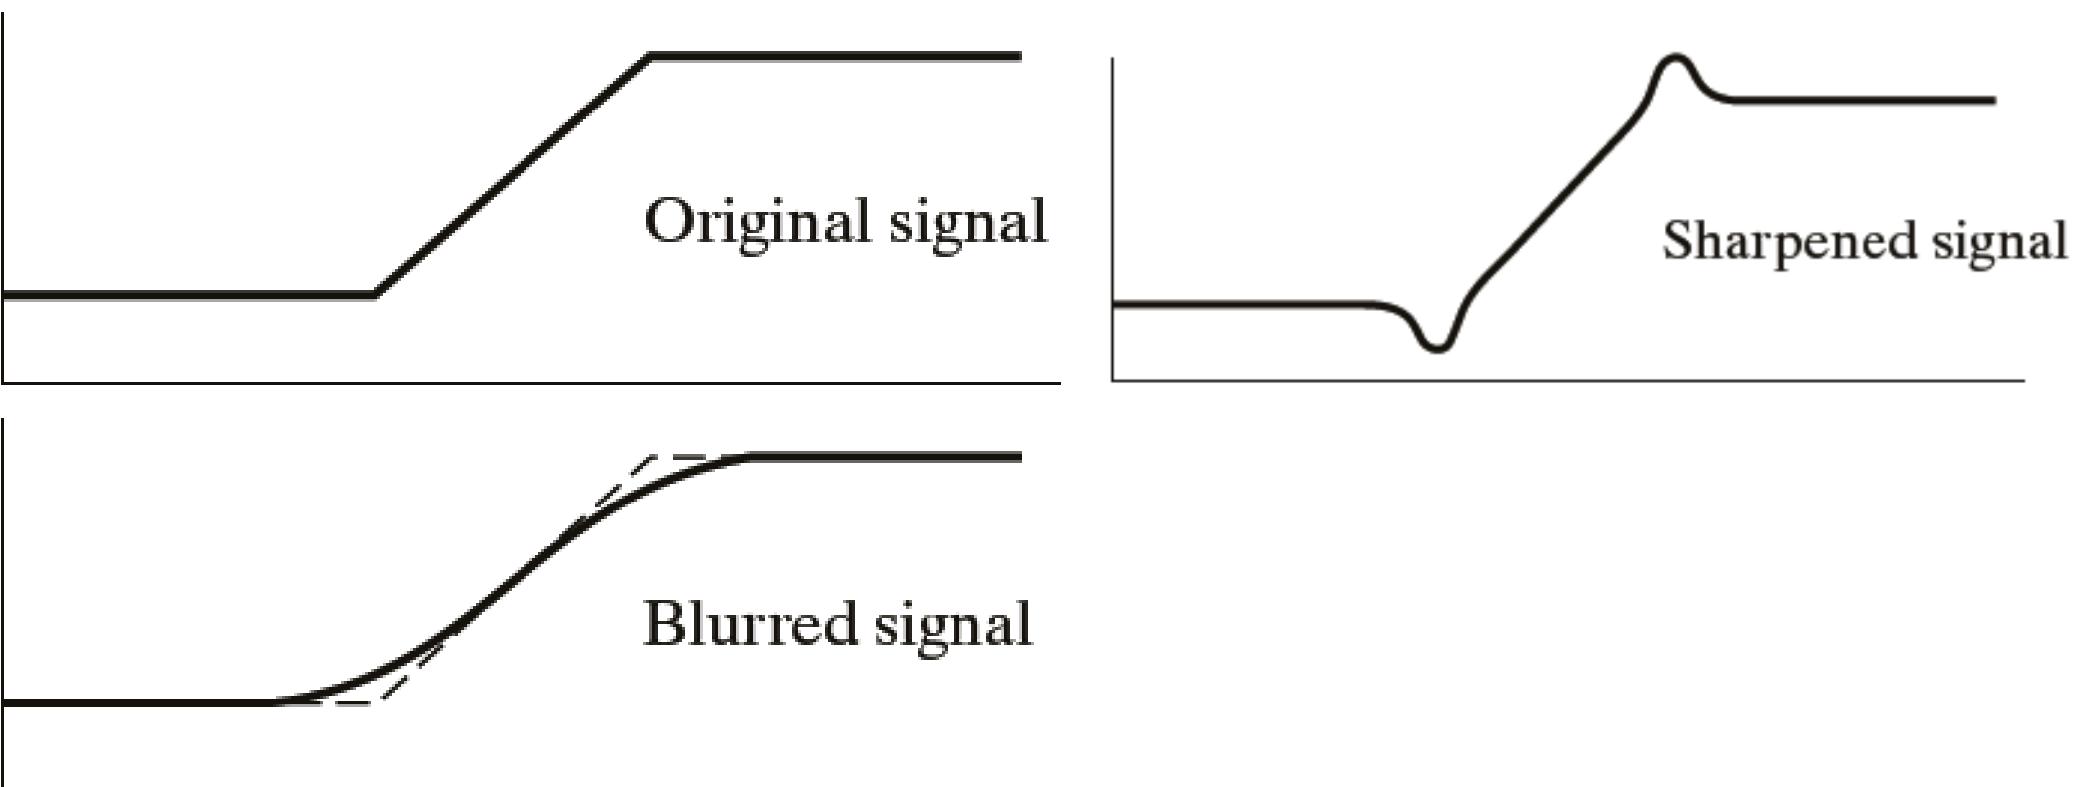


Median filter



Spatial Filtering

- Sharpening filter
 - “Opposite” of smoothing or blurring
 - Smoothing performs integration
 - Sharpening should do the opposite !



Spatial Filtering

- Sharpening filter examples



Spatial Filtering

- Image derivatives
 - 1st derivative across edge

- Gradient in 2D
 - For $f(x_1, x_2, \dots, x_n)$

$$\nabla f = \frac{\partial f}{\partial x_1} \mathbf{e}_1 + \dots + \frac{\partial f}{\partial x_n} \mathbf{e}_n$$

where $\mathbf{e}_1, \dots, \mathbf{e}_n$ are basis vectors

- 2nd derivative across edge
 - Laplacian in 2D
 - For $f(x_1, x_2, \dots, x_n)$

$$\Delta f = \sum_{i=1}^n \frac{\partial^2 f}{\partial x_i^2}$$

Edge



1st derivative

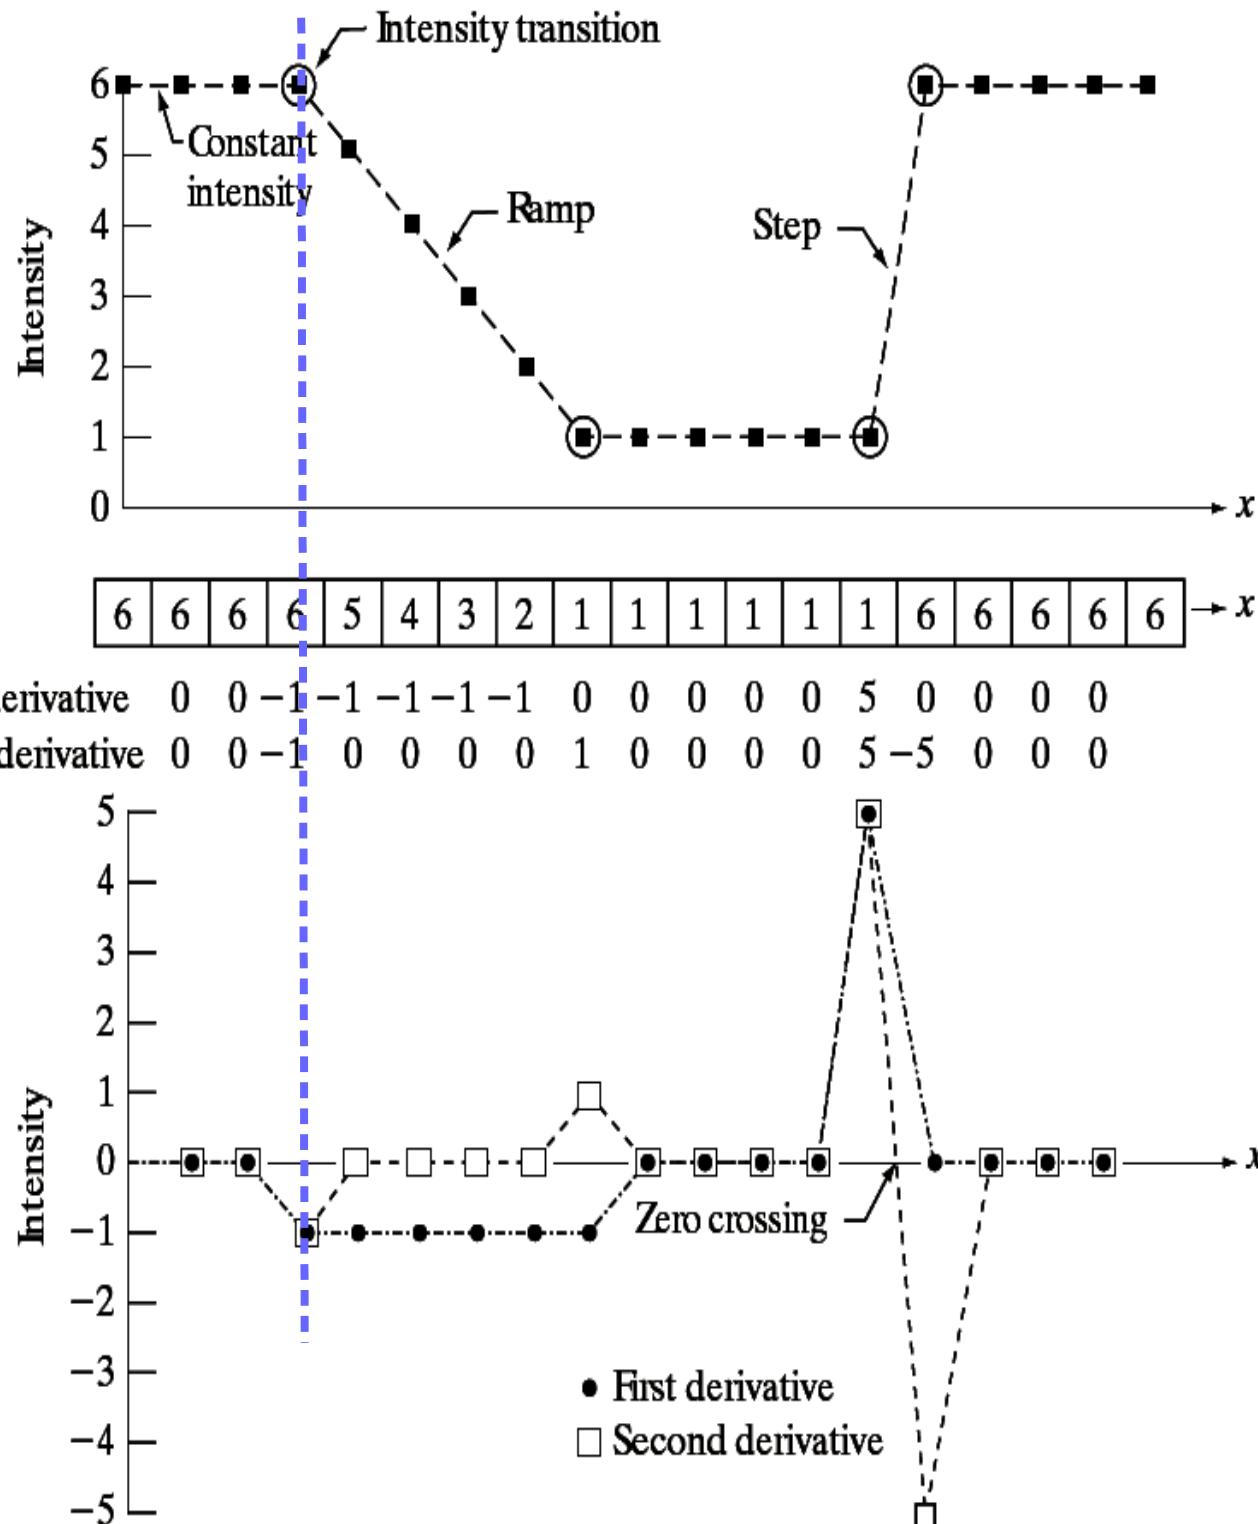


2nd derivative



Spatial Filteri

- Image derivatives
 - Very sensitive to noise
 - Example:
1-sided finite differences
 - “Zero crossing”

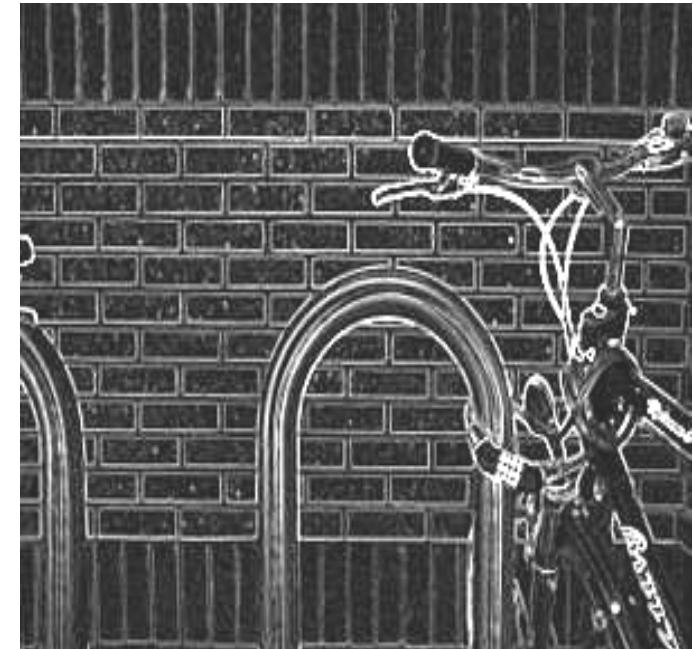
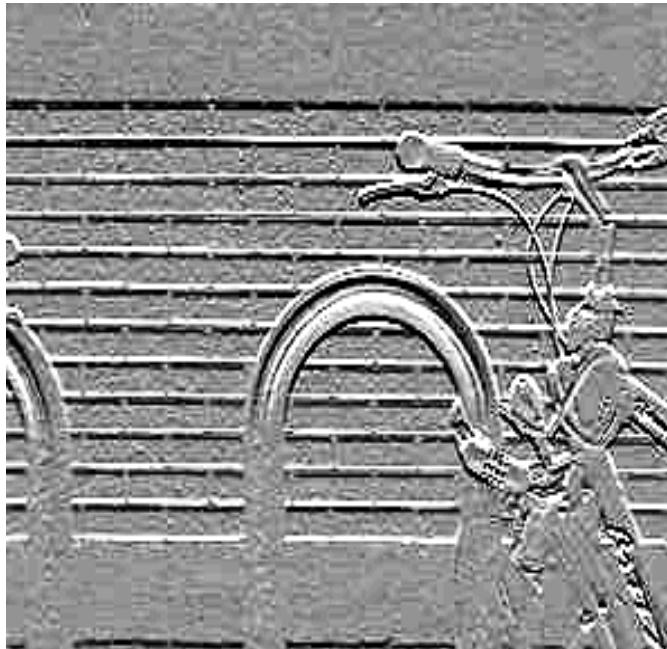
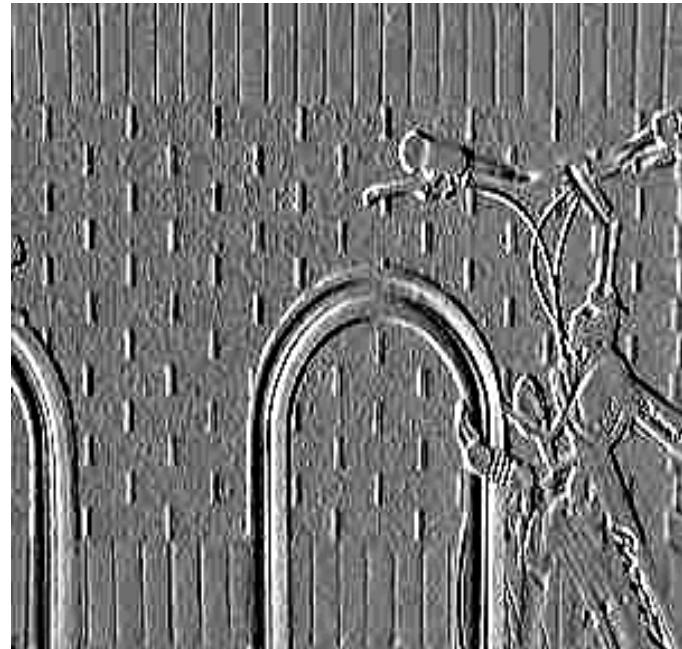


Spatial Filtering

- Image derivatives
 - Convolution mask for 1st derivative
 - Use standard 2-sided finite differences
 - Sobel operator
 - Derivative along X axis
 - Averages derivatives in neighborhood to reduce effect of noise
 - Derivative along Y axis
- $$\begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$
- $$\begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

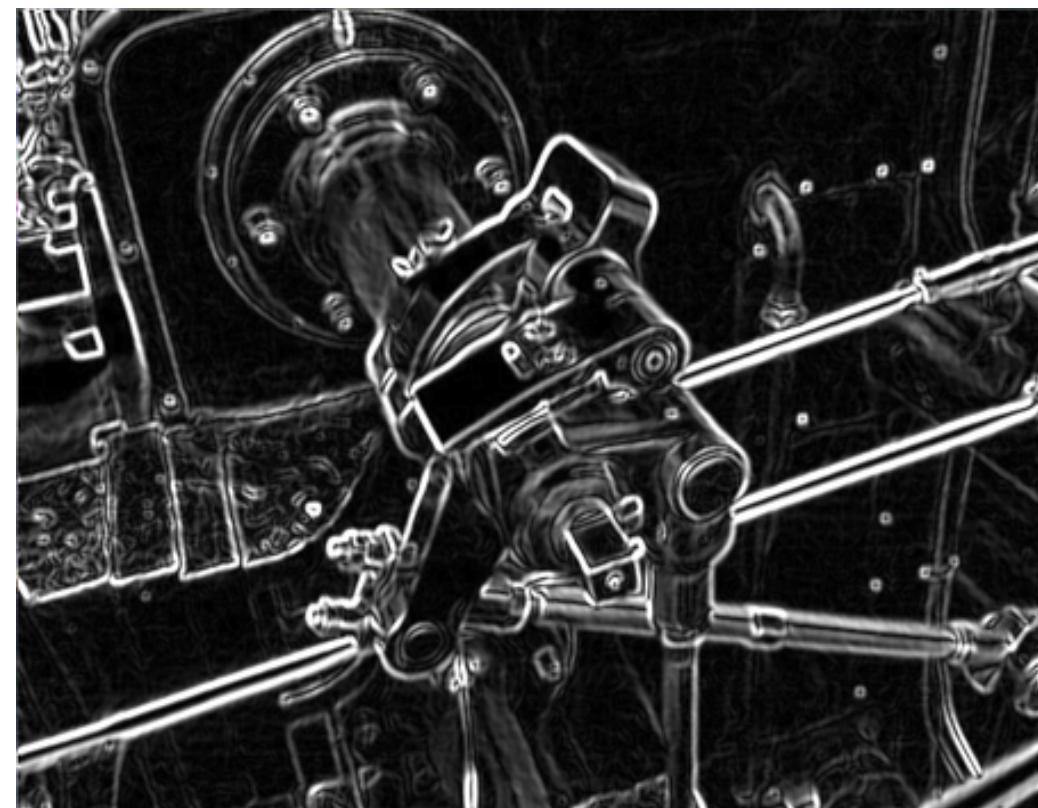
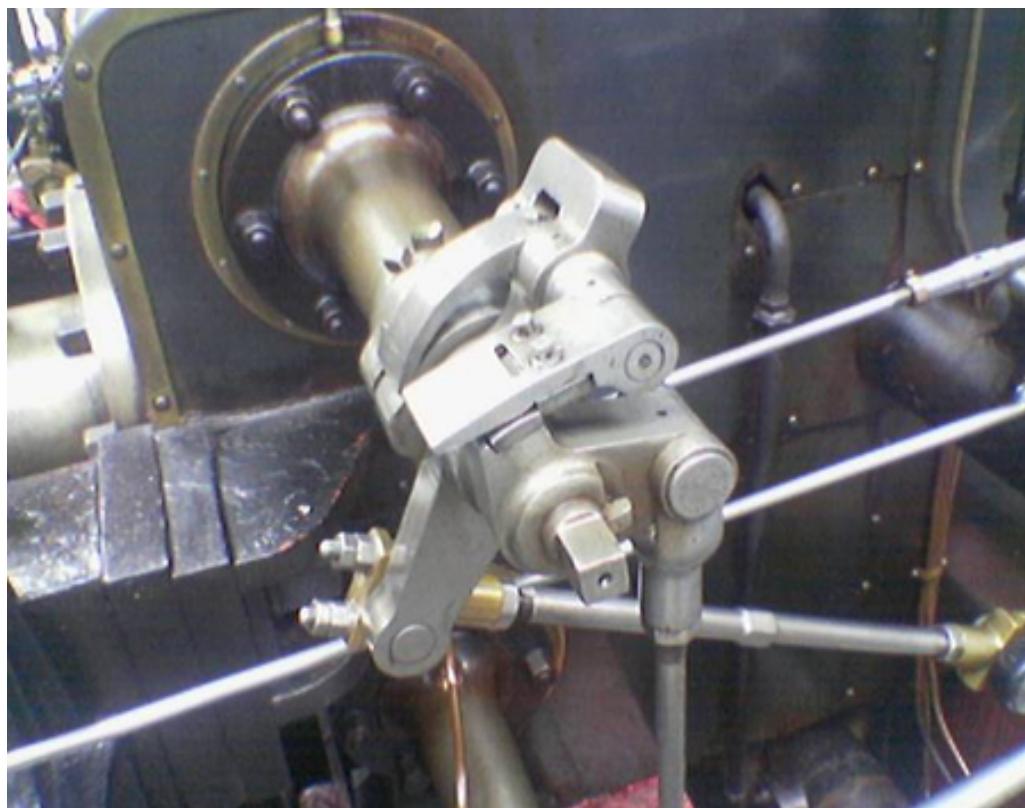
Spatial Filtering

- Image derivatives
 - Convolution mask for 1st derivative
 - Sobel operator
 - Averages derivatives in neighborhood to reduce effect of noise



Spatial Filtering

- Image derivatives
 - Convolution mask for 1st derivative
 - Sobel operator
 - How to apply it on a RGB color image ?



Spatial Filtering

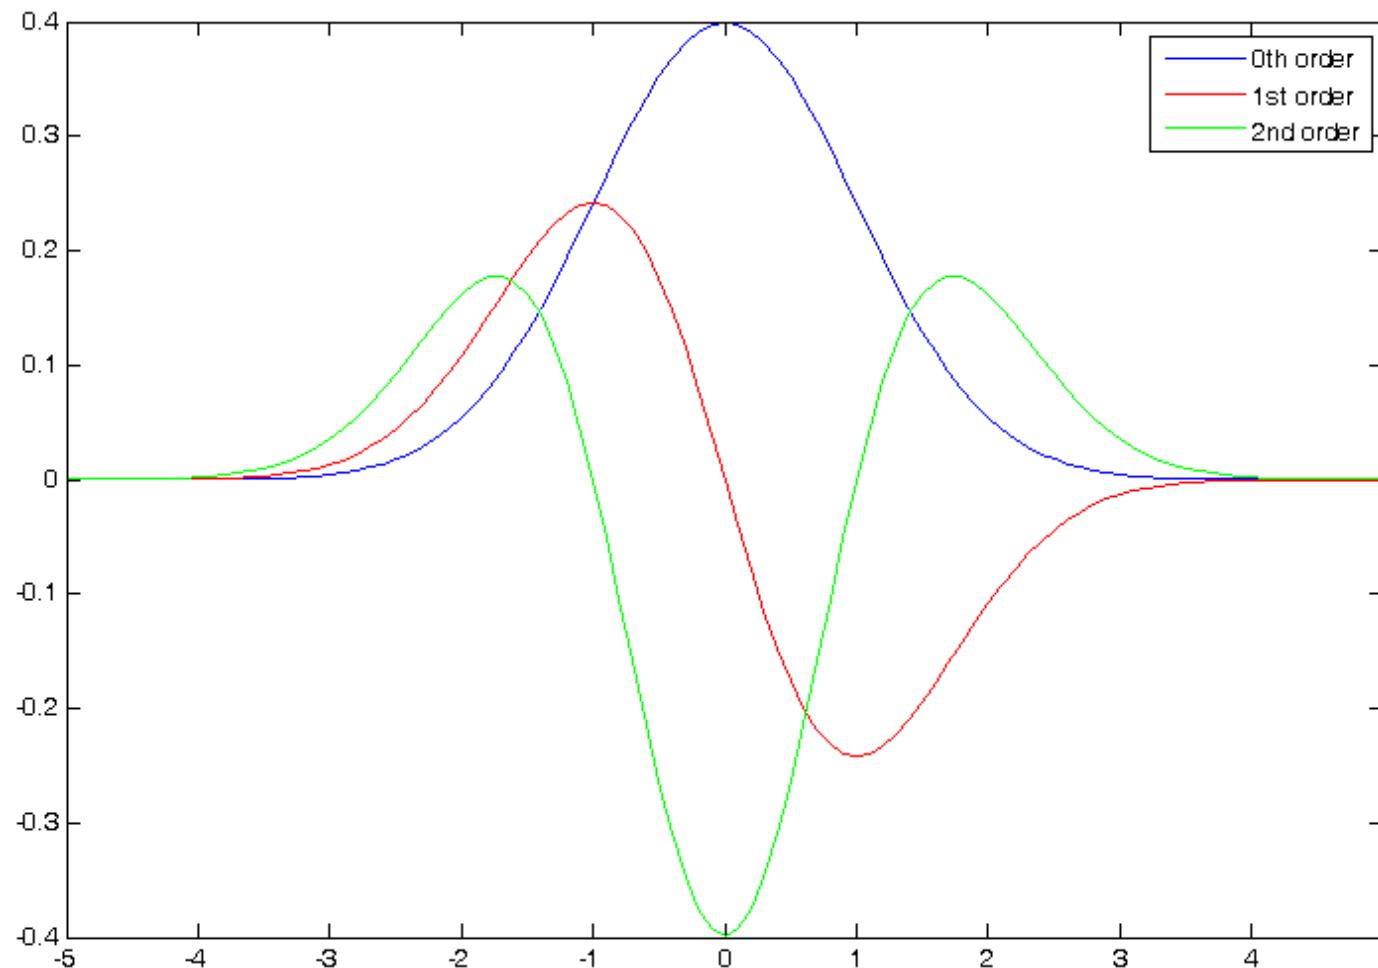
- Image derivatives
 - Convolution mask for second derivative
 - Use standard finite differences
 - In 1D : $[1 \ -2 \ 1]$
 - In 2D
 - Laplacian:
Add 2nd derivatives along X, Y directions
 - Add 2nd derivatives along X, Y, diagonal directions
 - In practice, this isn't preferred.
See next slide for what is done.

0	1	0
1	-4	1
0	1	0

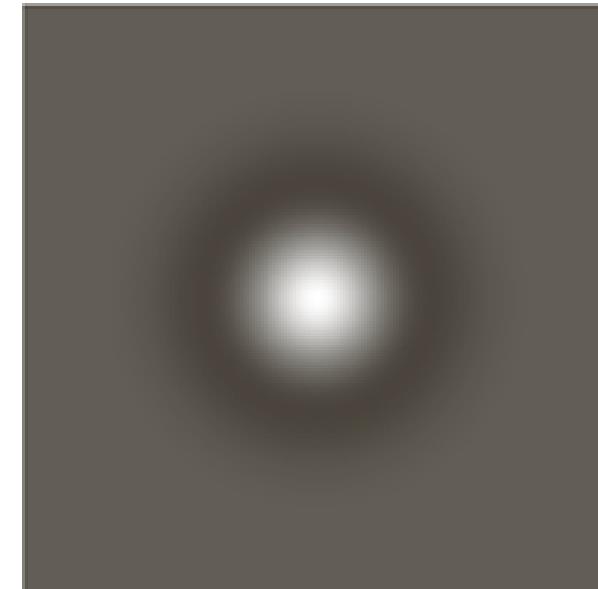
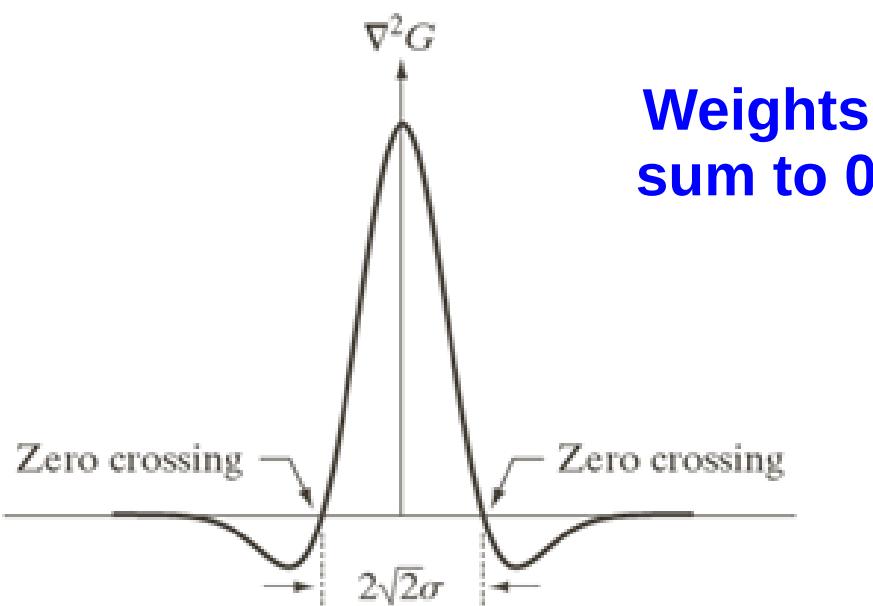
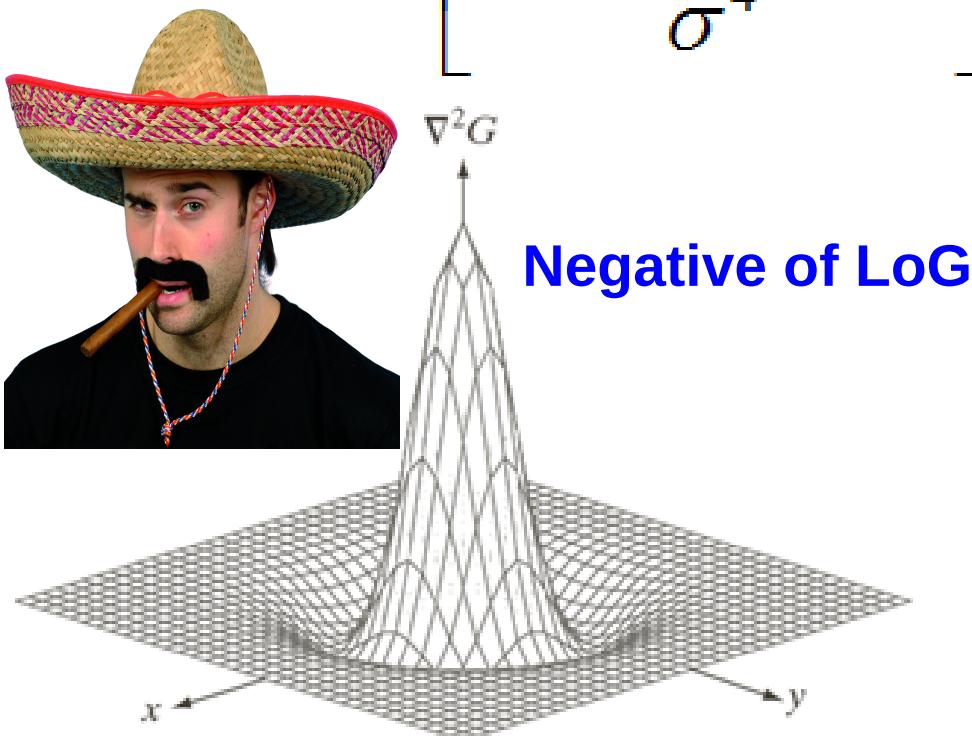
1	1	1
1	-8	1
1	1	1

Spatial Filtering

- Image derivatives
 - **Laplacian of Gaussian (LoG)**
 - (1) Apply a Gaussian filter to smooth image (reduce noise)
 - (2) Apply the Laplacian filter to the smoothed image



$$\nabla^2 G(x, y) = \left[\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right] e^{-\frac{x^2+y^2}{2\sigma^2}}$$



a	b
c	d

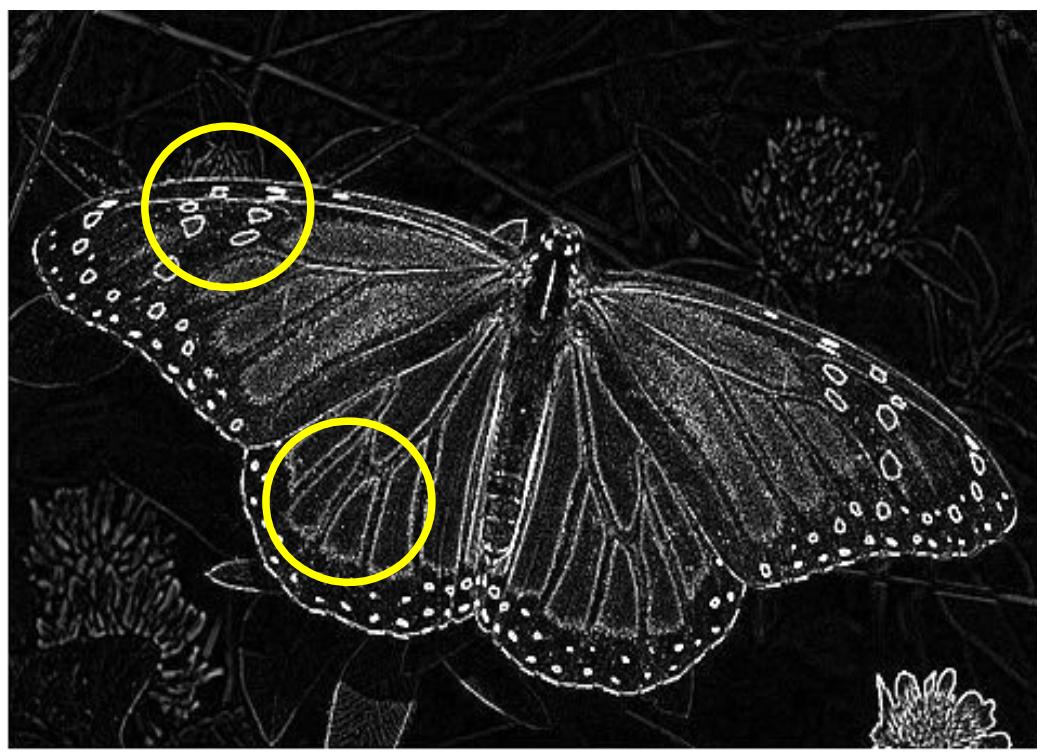
FIGURE 10.21

- (a) Three-dimensional plot of the *negative* of the LoG.
- (b) Negative of the LoG displayed as an image.
- (c) Cross section of (a) showing zero crossings.
- (d) 5×5 mask approximation to the shape in (a). The negative of this mask would be used in practice.

0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

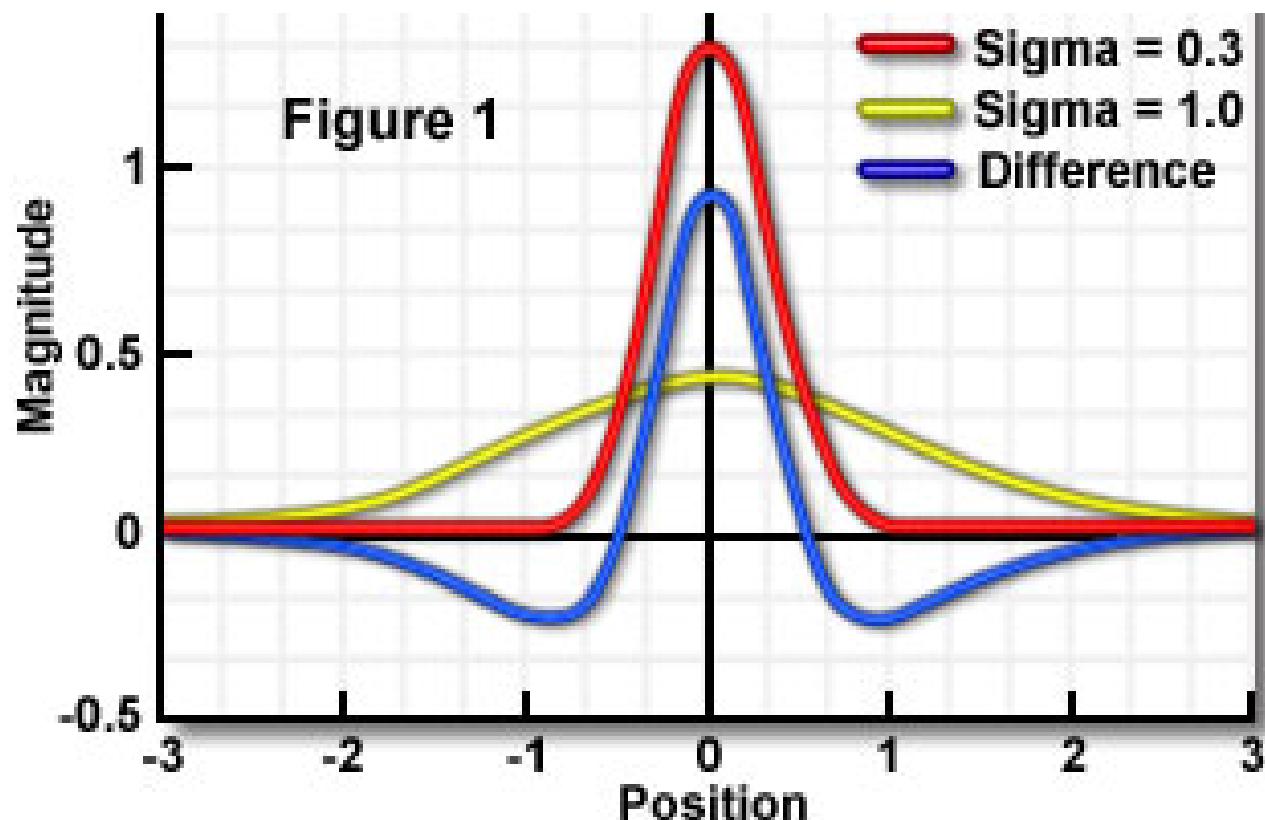
Spatial Filtering

- Image derivatives
 - Laplacian of Gaussian (LoG)
 - (1) “Blob” detector
 - Think: correlation with hat shape
 - (2) Gives strong positive and negative responses at edge
 - **Left:** 3x3 Gaussian + 3x3 Laplacian **versus** **Right:** 5x5 masks



Spatial Filtering

- Image derivatives
 - Difference of Gaussian (DoG)
 - Resembles (negative of) LoG when variances are carefully chosen

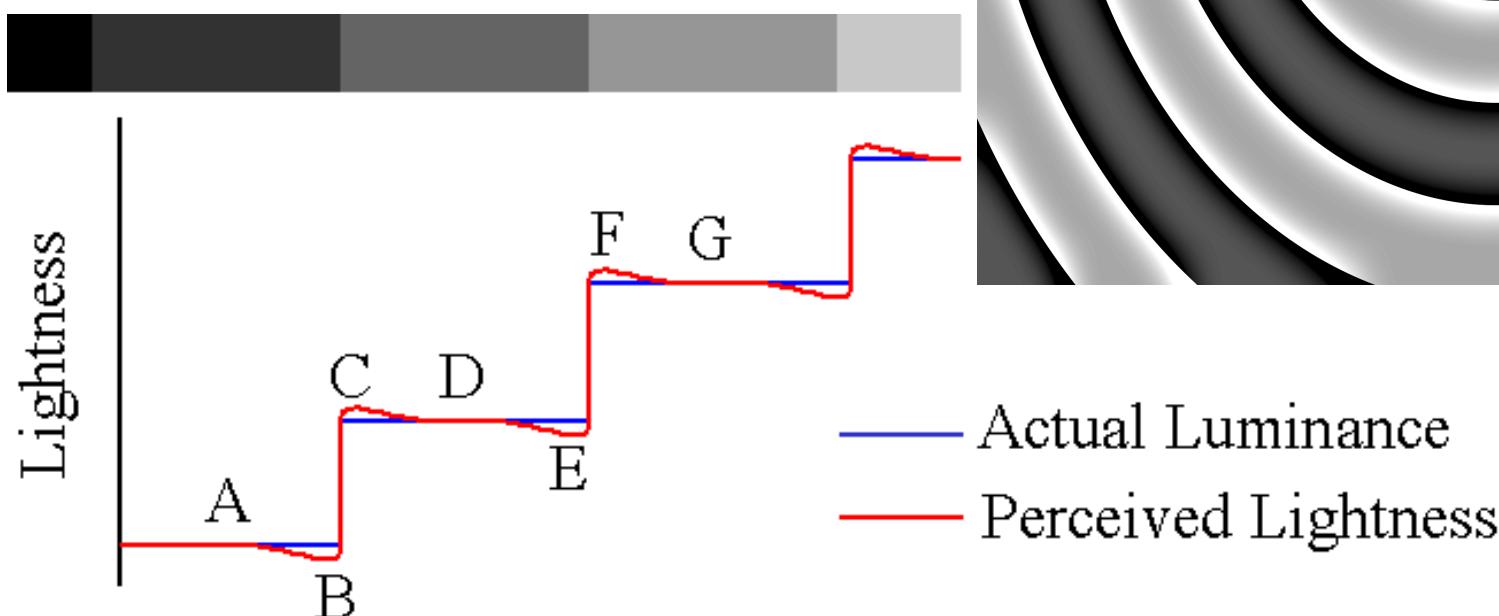


Spatial Filtering

- Sharpening
 - **Unsharp masking**
 - Unsharp \leftarrow blur
 - Mask \leftarrow negative
 - Common to apply to photos before printing
 - Because printed output is a blurred version of input
 - Sharpening helps compensate printer limitations

Spatial Filtering

- Sharpening
 - Unsharp masking
 - Relates to visual perception → Mach bands



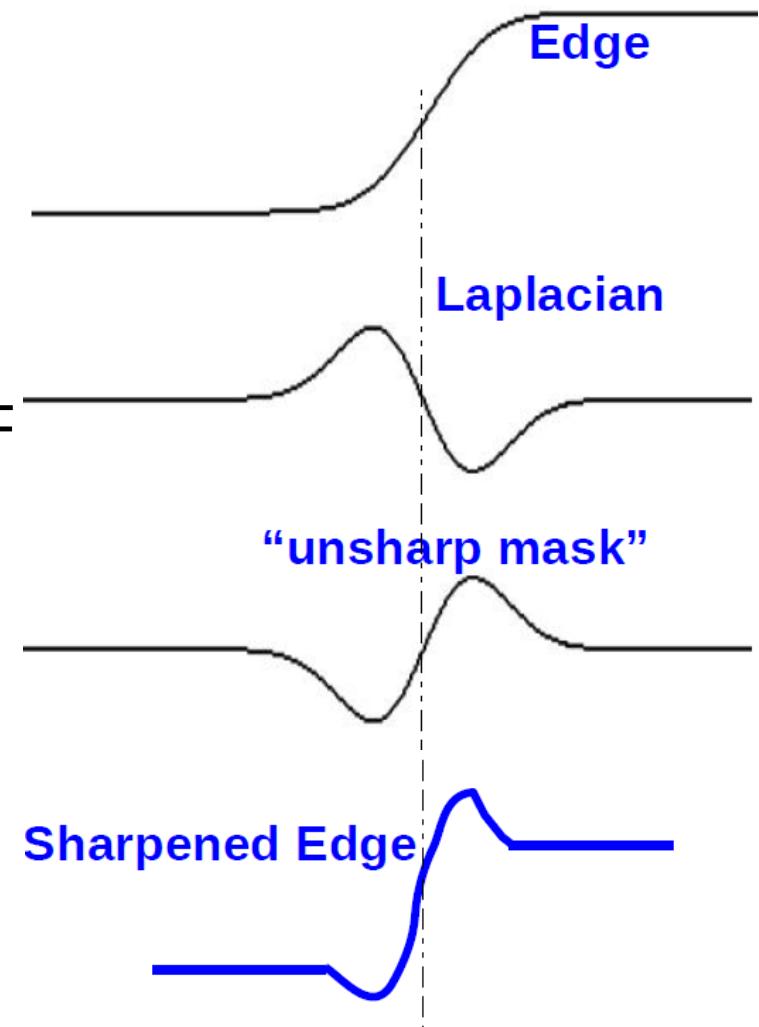
Spatial Filtering

- Sharpening
 - Unsharp masking



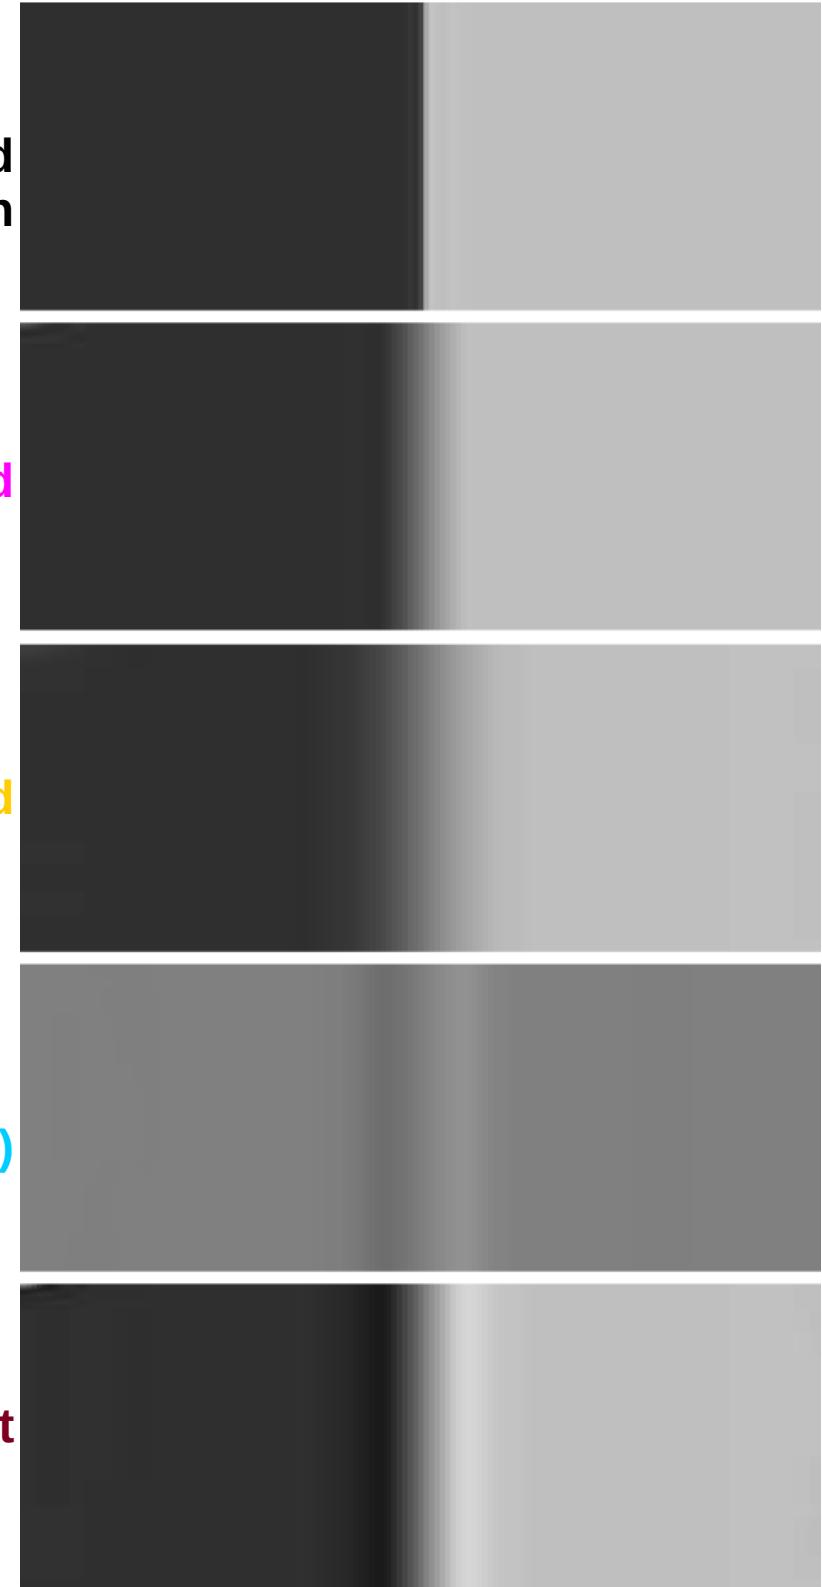
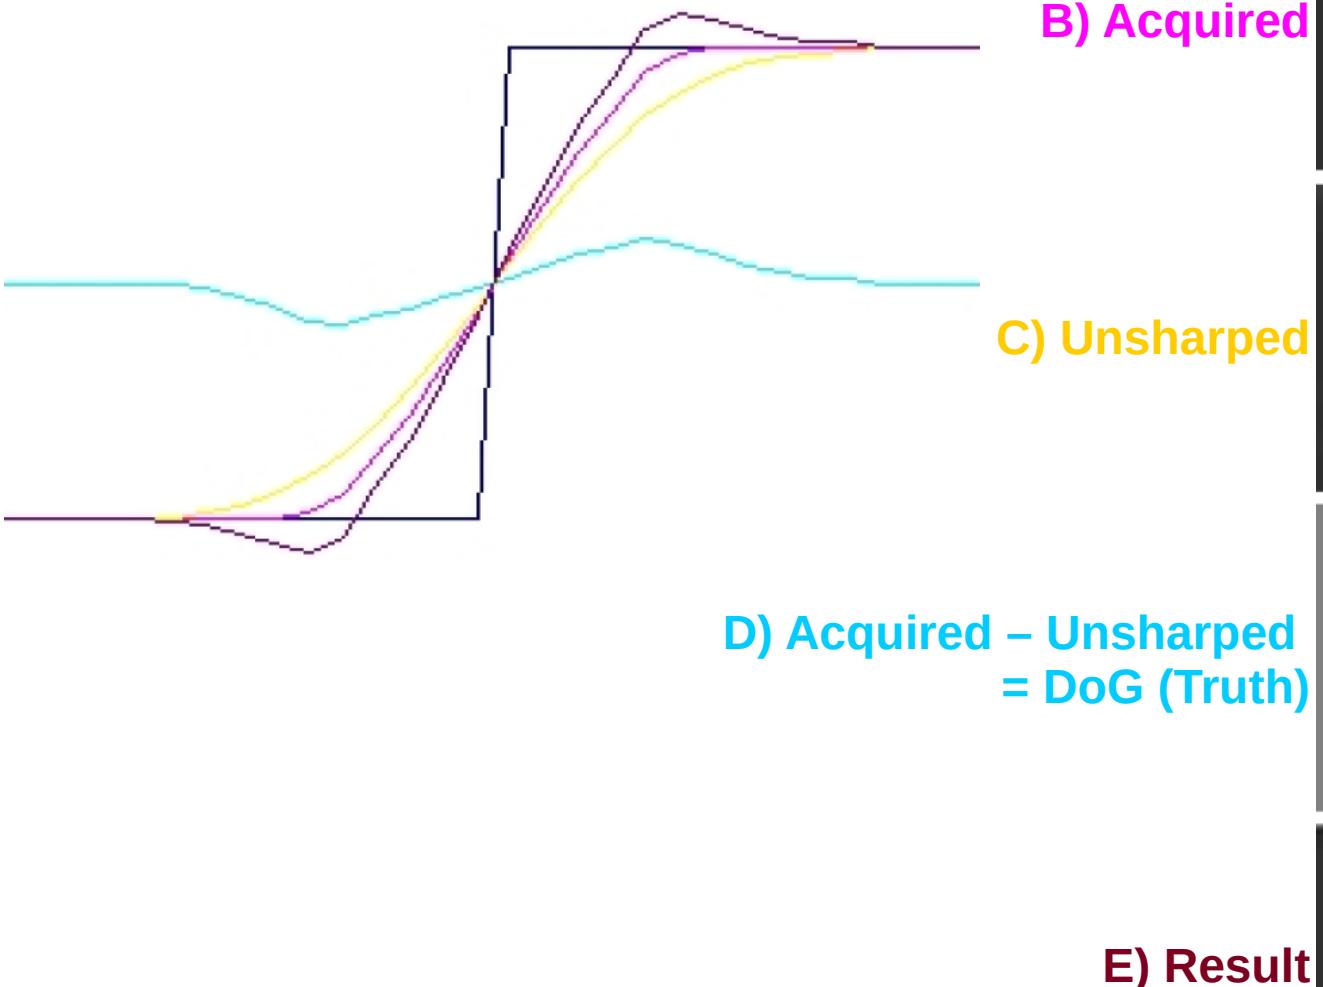
Spatial Filtering

- Sharpening
 - How is unsharp masking related to image derivatives ?
 - Input = Image F (already a bit blurred)
 - Further blur input image $\rightarrow G * F$,
 - G = Gaussian mask, $*$ = convolution
 - “Unsharp”
 - Negated-blurred image $= -G * F$
 - Subtract negated-blurred image from F $\rightarrow F - G * F$
 - Unsharp “mask”
 - ($\text{DoG} = -\text{LoG}$ of step edge)
 - Unsharp-mask filter $= F + s (F - G * F)$ $\sim F + s \text{DoG} * F$ $\sim F - s \text{LoG} * F$ $\sim F - \text{scaled-Laplacian-}F$



Spatial Filtering

- Sharpening
 - Unsharp masking



Spatial Filtering

- 1st derivatives versus 2nd derivatives
 - 1st derivatives used for:
 - Edge detection
 - 2nd derivatives used for:
 - Sharpening image (via LoG or DoG)
 - Blob detection

Spatial Filtering

- Applications in photography
 - Bokeh = aesthetic quality of the blur produced in the out-of-focus parts of an image, by a lens
 - In Japanese, “boke” = blur, haze

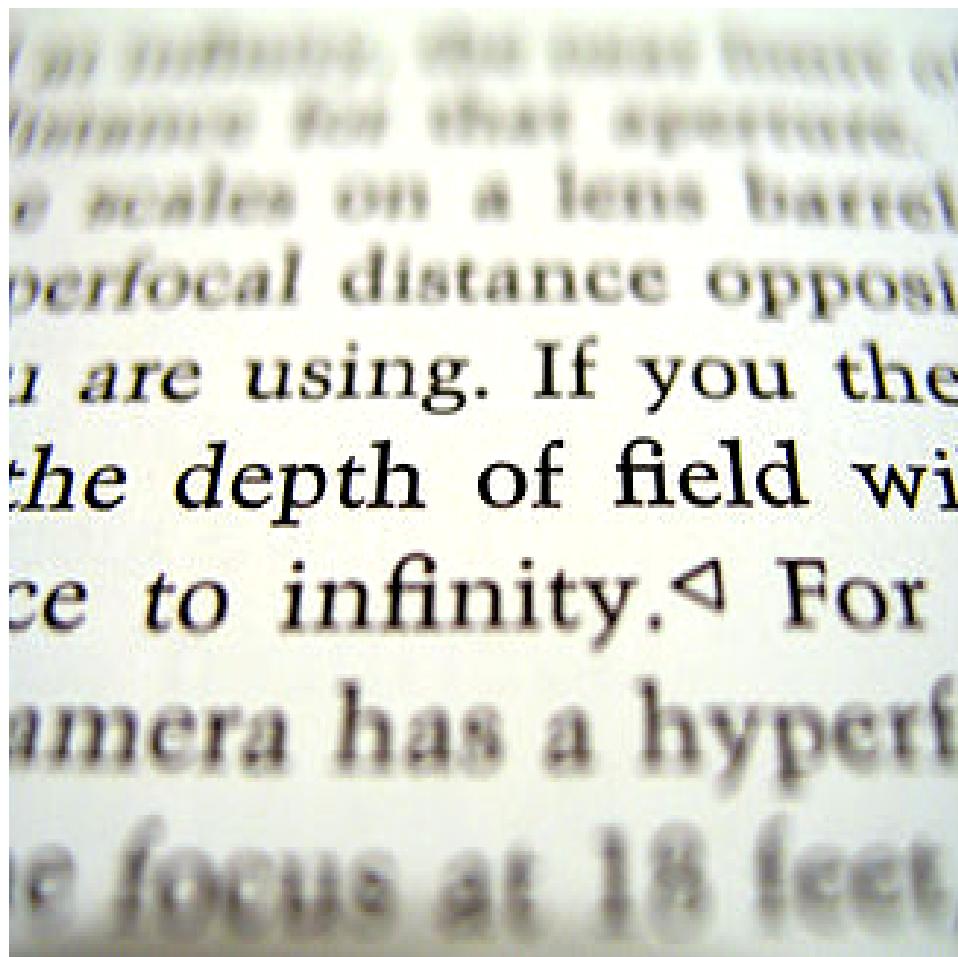


Spatial Filtering

- Applications in photography

- Bokeh

- Desire “**depth of field**” for artistic effects
 - Main subject should be in focus. All else should be out of focus.



and blurred background elements to
determine how often exposures
are made on a lens barrel.
perfocal distance opposite
you are using. If you then
the depth of field will
ce to infinity. For
amera has a hyperfocal
e focus at 18 feet,



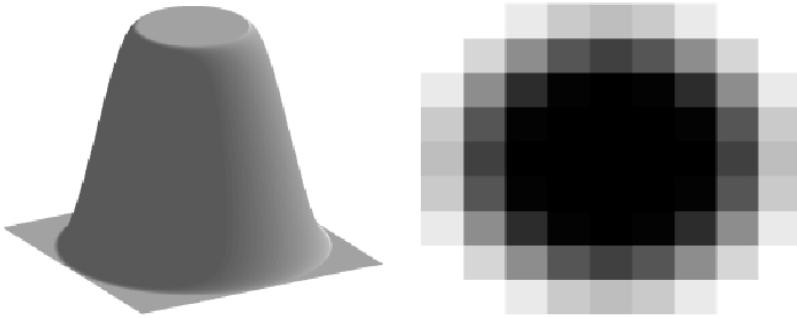
Spatial Filtering

- Applications in photography

- Bokeh effects in art



Spatial Filtering



- Applications in photography
 - Bokeh simulated in software in part of image
(1) original (2) bokeh background (3) Blurred bg



Spatial Filtering

- Boundary conditions
 - What happens when mask is placed to pixel close to image boundary ?
 - What are pixel intensities outside image boundary ?

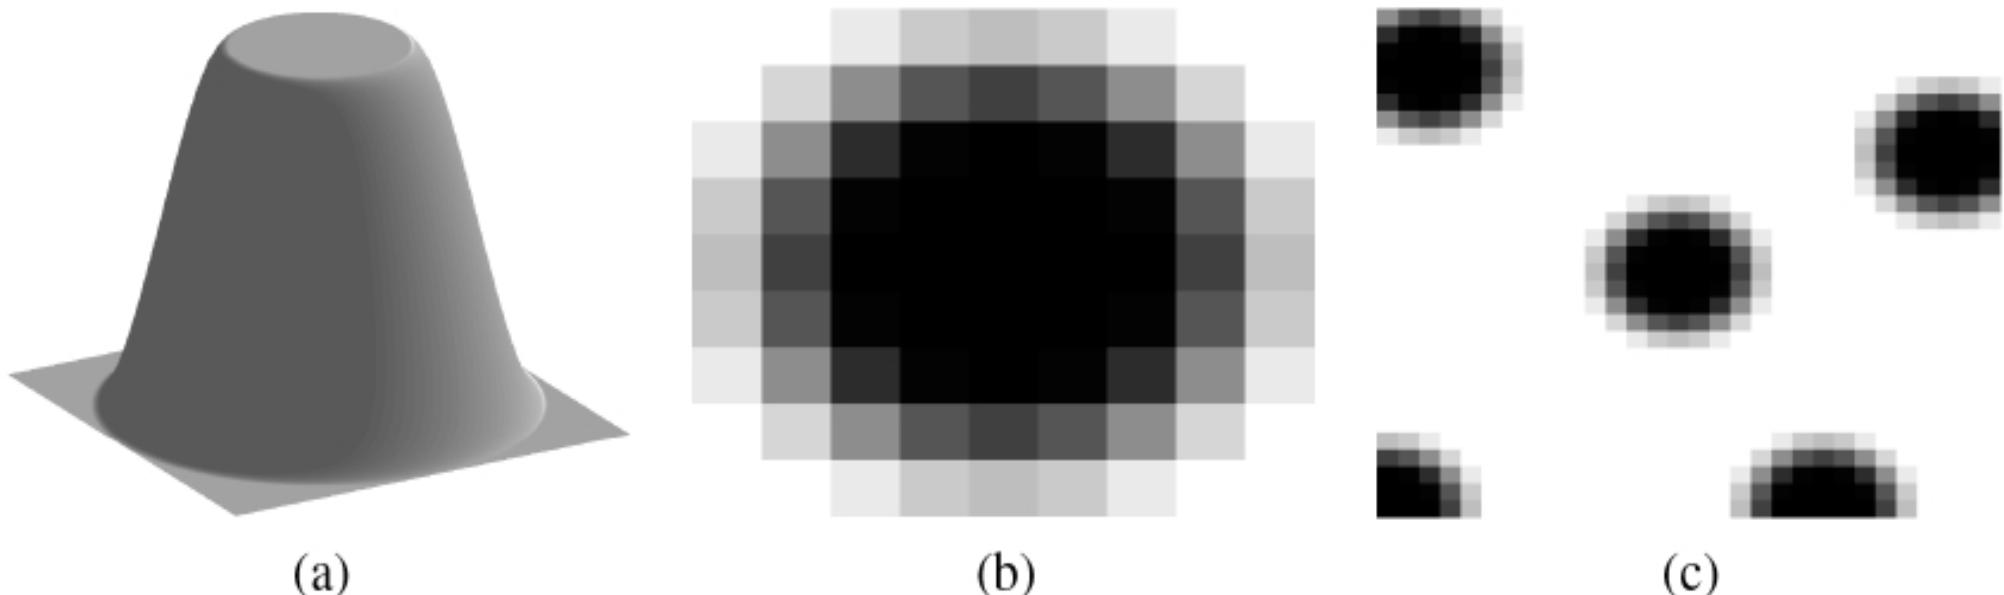


Figure 3.3. Neighborhood shapes. (a) Preserving rotational invariance via a neighborhood mask consisting of a flat central circular plateau with cubic splines on the sides. (b) The discrete sampling of the mask (black $\equiv 1$, white $\equiv 0$) for a 9×9 pixels neighborhood. (c) Anisotropic neighborhoods at boundaries.

Spatial Filtering

- Boundary conditions
 - What are pixel intensities outside image boundary ?
 - (1) Zero, or any fixed value
 - In the field of PDEs, Dirichlet / fixed boundary condition
 - (2) Use intensity of the closest pixel inside image
 - Normal derivative of function on boundary = 0
 - Derivative along a direction normal to the boundary
 - In the field of PDEs, Neumann boundary condition
 - (3) Use intensity of pixel determined by circular wrap around
 - Only makes sense in special cases when function is periodic
 - (4) Crop the mask (redefine the mask)
 - Redefine the neighborhood at each pixel