

# Digital Image Processing

## Image Smoothing for Noise Reduction: Bilateral Filtering

Suyash P. Awate

# Noise

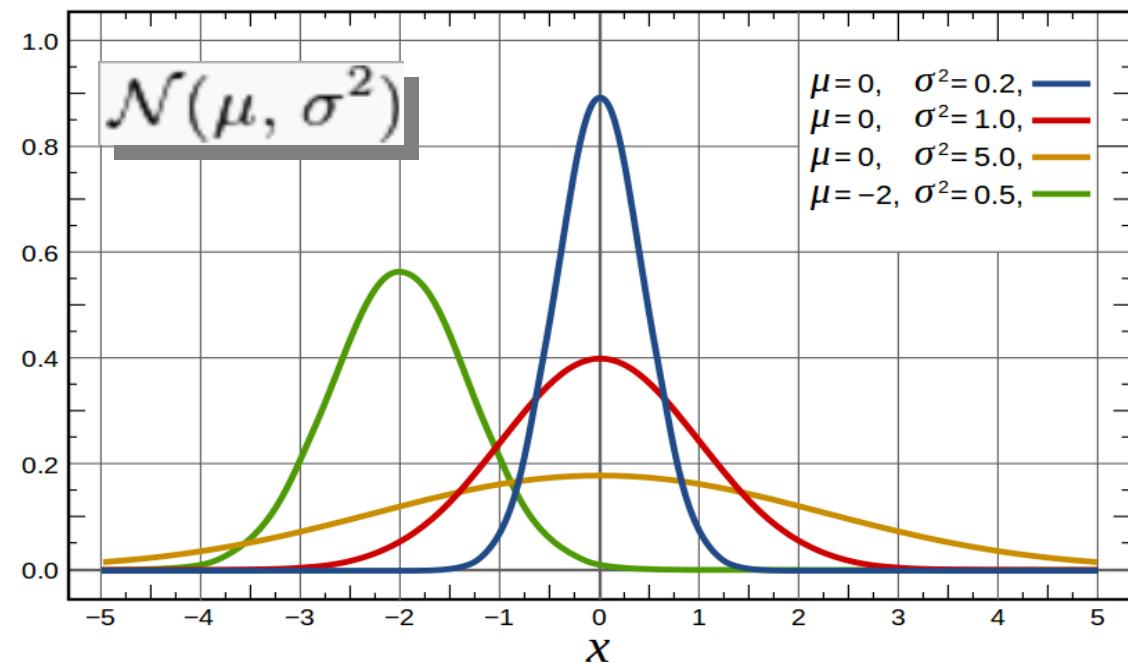
- A simple model
  - Additive independent-and-identically-distributed Gaussian noise
  - Corrupted image  $Y = \text{uncorrupted image } X + \text{noise } Z$

$$Z_i \sim \mathcal{N}(0, N)$$

$$Y_i = X_i + Z_i$$

- Normal PDF = Gaussian PDF

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



- In the real world, for some applications, noise can be difficult to model

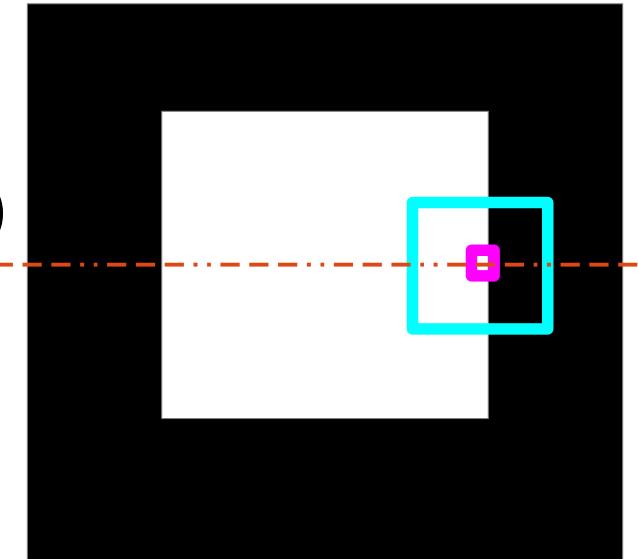
# Smoothing

- What is the problem with **weighted-mean** filters ?
  - Fail to preserve edges, while reducing noise
    - Oversmooth



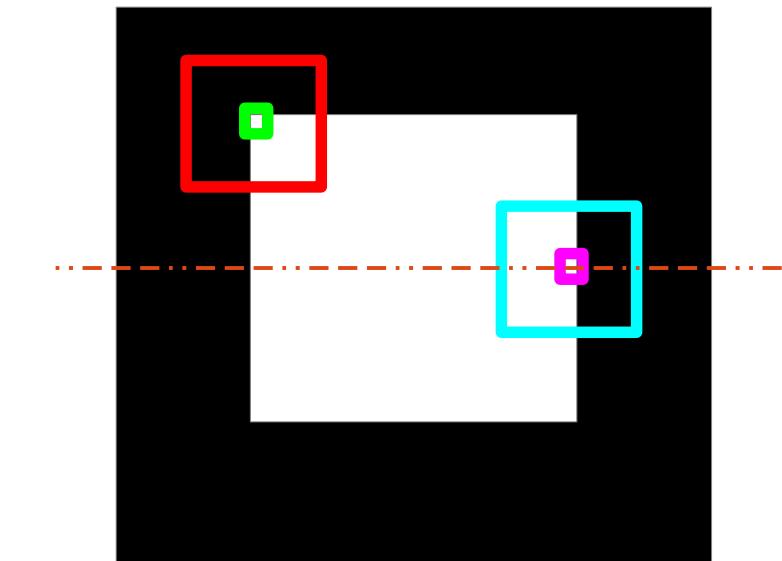
# Smoothing

- What is the problem with **median** filter ?
  - Preserves edges in 1D images
    - Consider the scan line (1D) crossing the “magenta” pixel (with blue neighborhood)
      - Color of pixel = white
      - In neighborhood,  
number of white pixels  
> number of black pixels
      - Median = white
  - Does it preserve all edges in 2D images ?



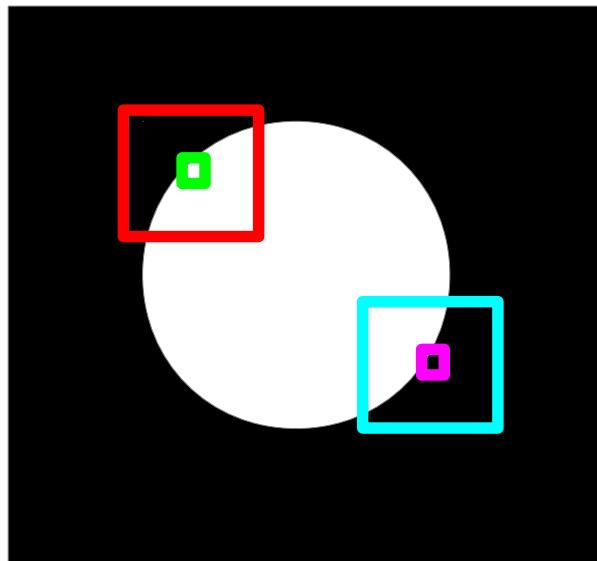
# Smoothing

- What is the problem with **median** filter ?
  - Does it preserve all edges in 2D images ?
    - Consider the “green” pixel (with red neighborhood)
      - Color of pixel = white
      - In neighborhood, number of white pixels < number of black pixels
      - Median = black !
      - Edge NOT preserved



# Smoothing

- What will happen with **repeated** application of:
  - Gaussian filter on a grayscale image ?
    - What will be the final image ?
  - Median filter on the square image in previous slide ?
    - What about the following image ?



# Edge-Preserving Smoothing

- We desire a filter that:
  - 1) Reduces noise
  - 2) Preserves object edges within image
- Price we are willing to pay
  - Nonlinearity of the filtering
    - Filtered value is a nonlinear function of pixel intensities
    - Filter can't be represented as convolution with original image
    - May increase computational load

# Edge-Preserving Smoothing

- General idea underlying any kind of smoothing
  - At pixel “p”, smoothed intensity = weighted average of all pixel intensities in image
  - Key questions
    - What are the weights ?
    - Are some weights zero ? = what neighborhood to choose ?
  - Examples
    - Mean filter
      - Weights  $\propto 1$  within neighborhood, 0 elsewhere
    - Weighted means filter
      - Weights  $\propto$  Gaussian within neighborhood
    - Median filter
      - Weight = 1 for only that neighboring pixel whose value is the median

# Edge-Preserving Smoothing

- Bilateral Filter
  - **2 key ideas used together**
  - At pixel location “p”,  
**weight** for another pixel “q” is based on :
  - (1) **Spatial distance** between “q” and “p”
    - Larger distance → lower weight
    - e.g., Gaussian over distance
      - Same as Gaussian smoothing filter
  - (2) **Dissimilarity between intensities** “ $I(q)$ ” and “ $I(p)$ ”
    - Larger dissimilarity → lower weight
    - e.g., Gaussian over intensity

# Edge-Preserving Smoothing

- Bilateral Filter
  - Updated pixel intensity =

$$I^{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} I(x_i) f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

where

$$W_p = \sum_{x_i \in \Omega} f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

$I^{\text{filtered}}$  is the filtered image;

$I$  is the original input image to be filtered;

$x$  are the coordinates of the current pixel to be filtered;

$\Omega$  is the window centered in  $x$ ;

$f_r$  is the range kernel for smoothing differences in intensities.

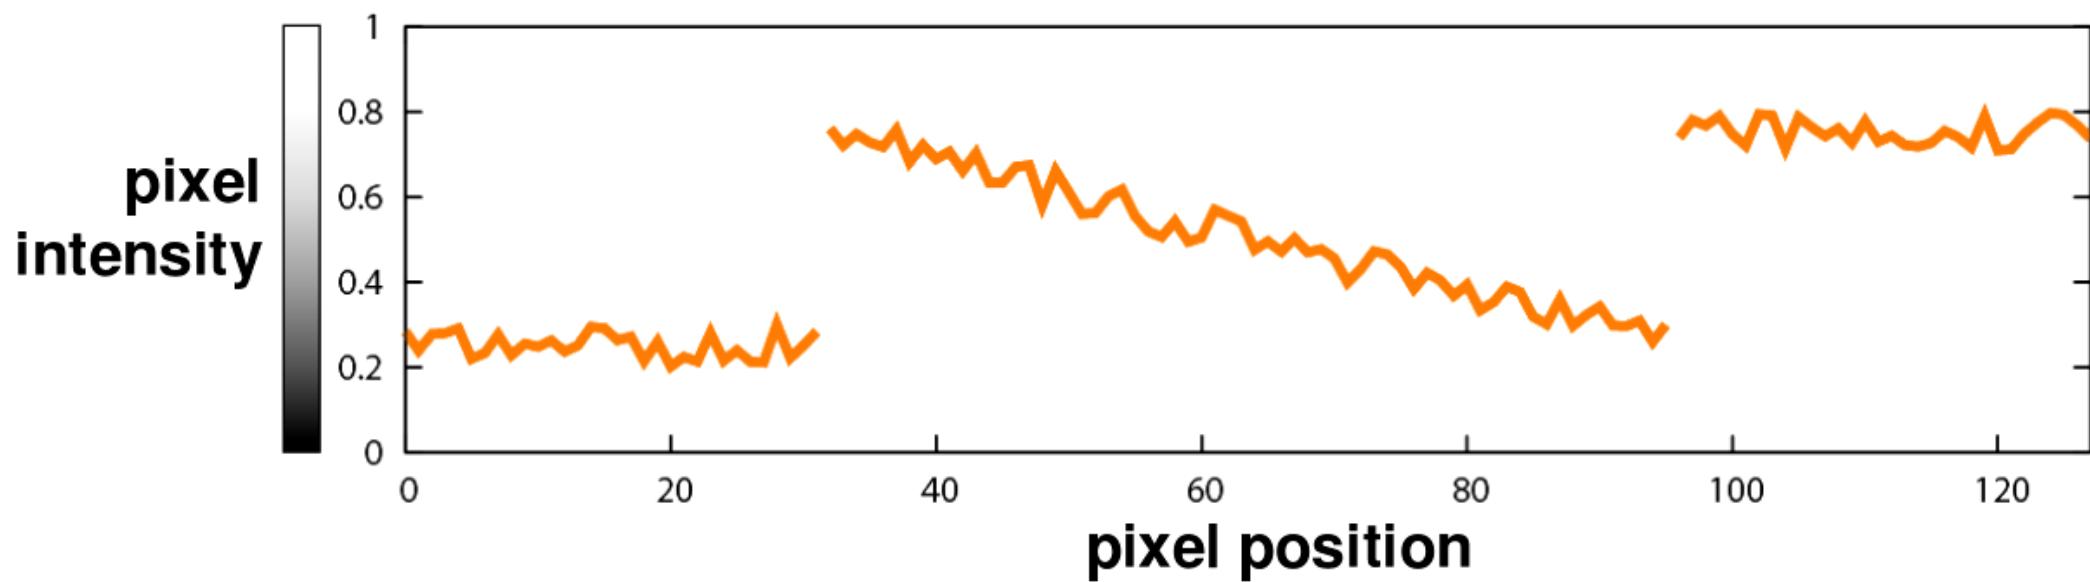
$g_s$  is the spatial kernel for smoothing differences in coordinates

# Edge-Preserving Smoothing

- Bilateral Filter
  - Functions  $f(\cdot)$  and  $g(\cdot)$  are typically Gaussian
  - Why can't this be written as convolution ?
    - Not a linear function of intensities
    - This defines a different weighting mask for every pixel

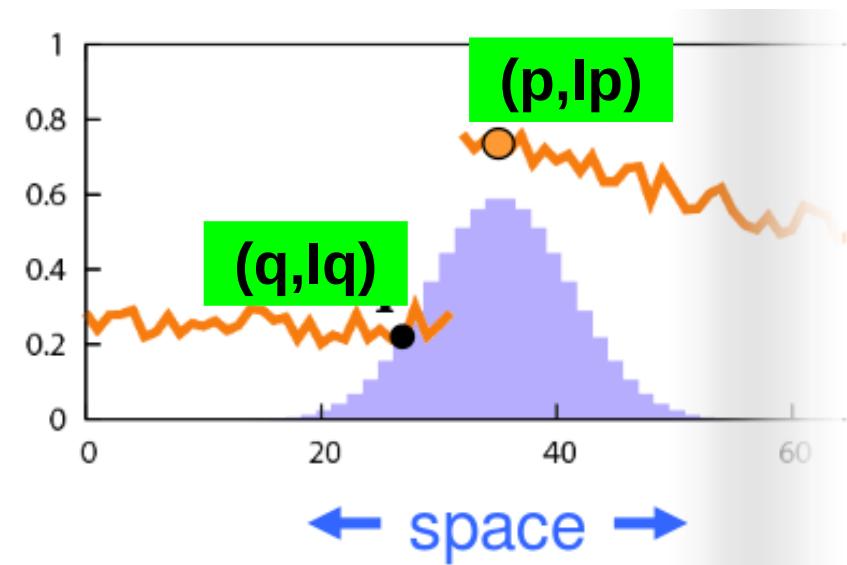
# Edge-Preserving Smoothing

- Bilateral Filter
  - What does Gaussian filter do to a noisy edge ?
    - 1D image



# Edge-Preserving Smoothing

- Bilateral Filter
  - Effect of **Gaussian** filter on weights at pixel “p”

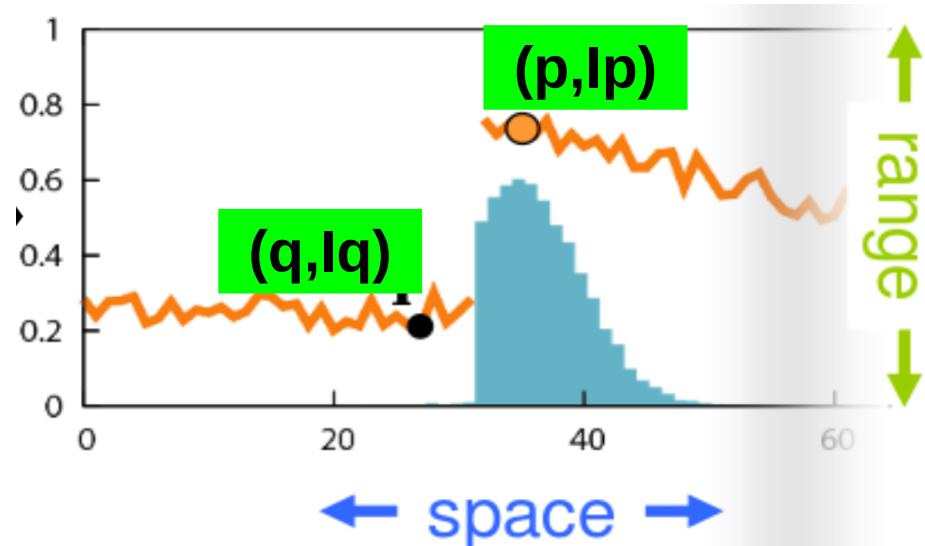
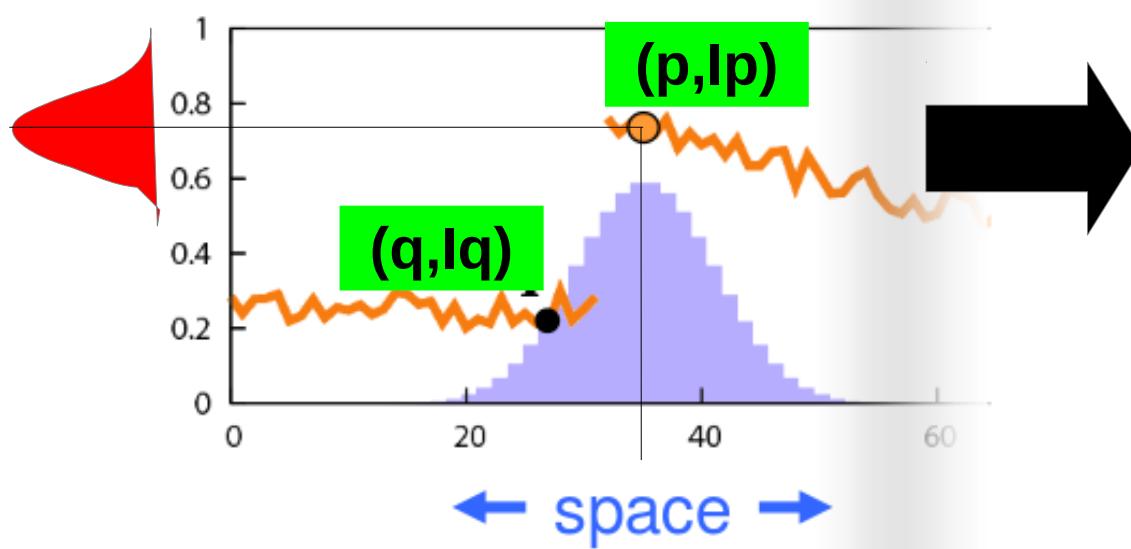
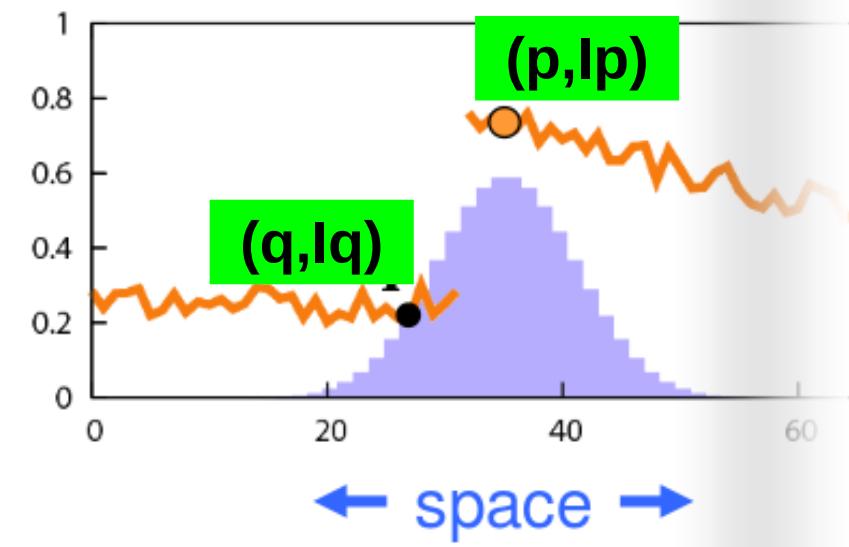


$$I^{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} I(x_i) f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

$$W_p = \sum_{x_i \in \Omega} f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

# Edge-Preserving Smoothing

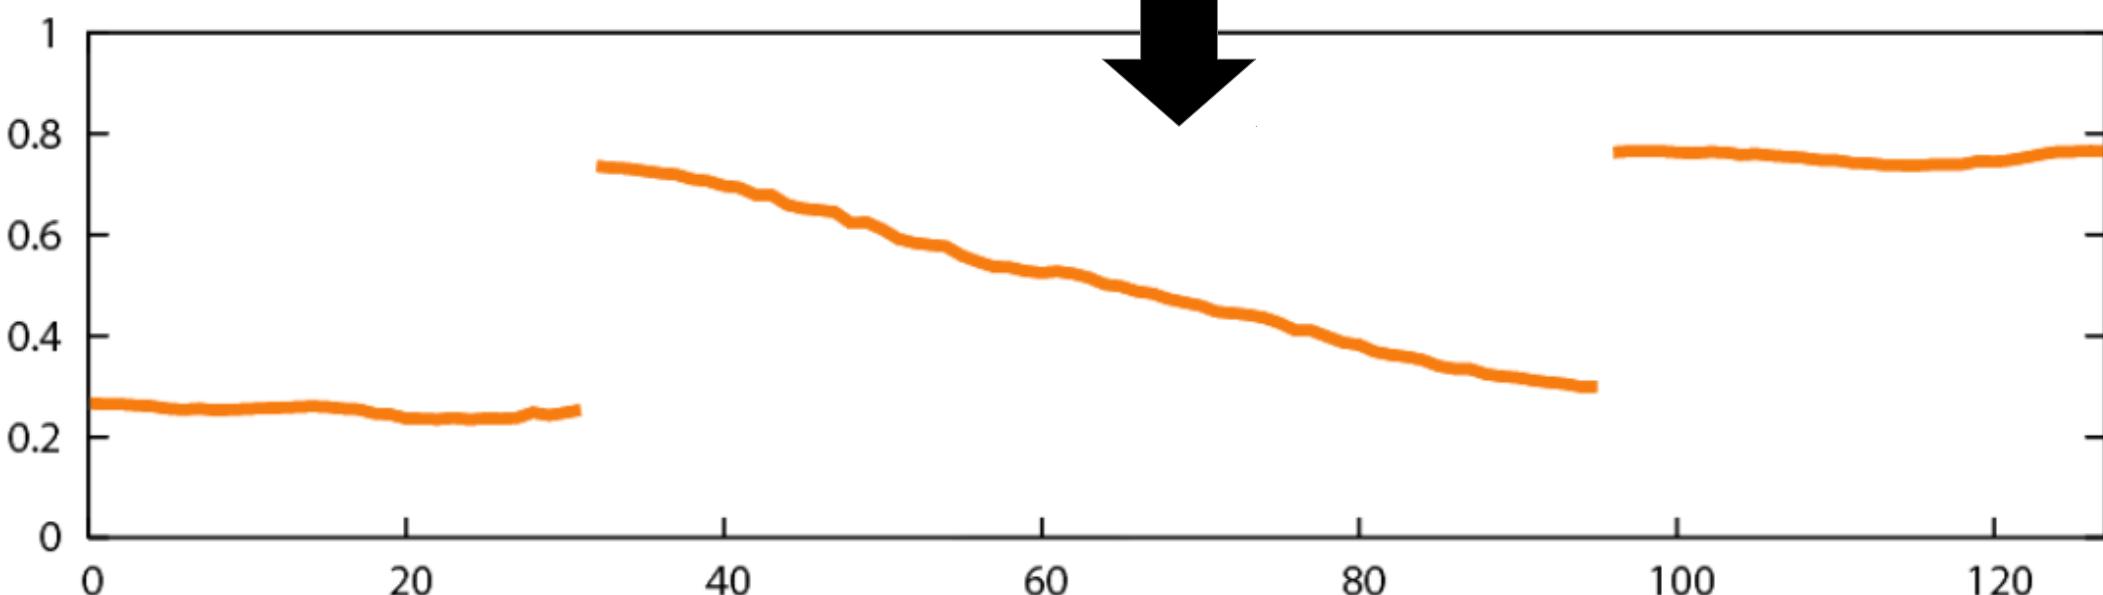
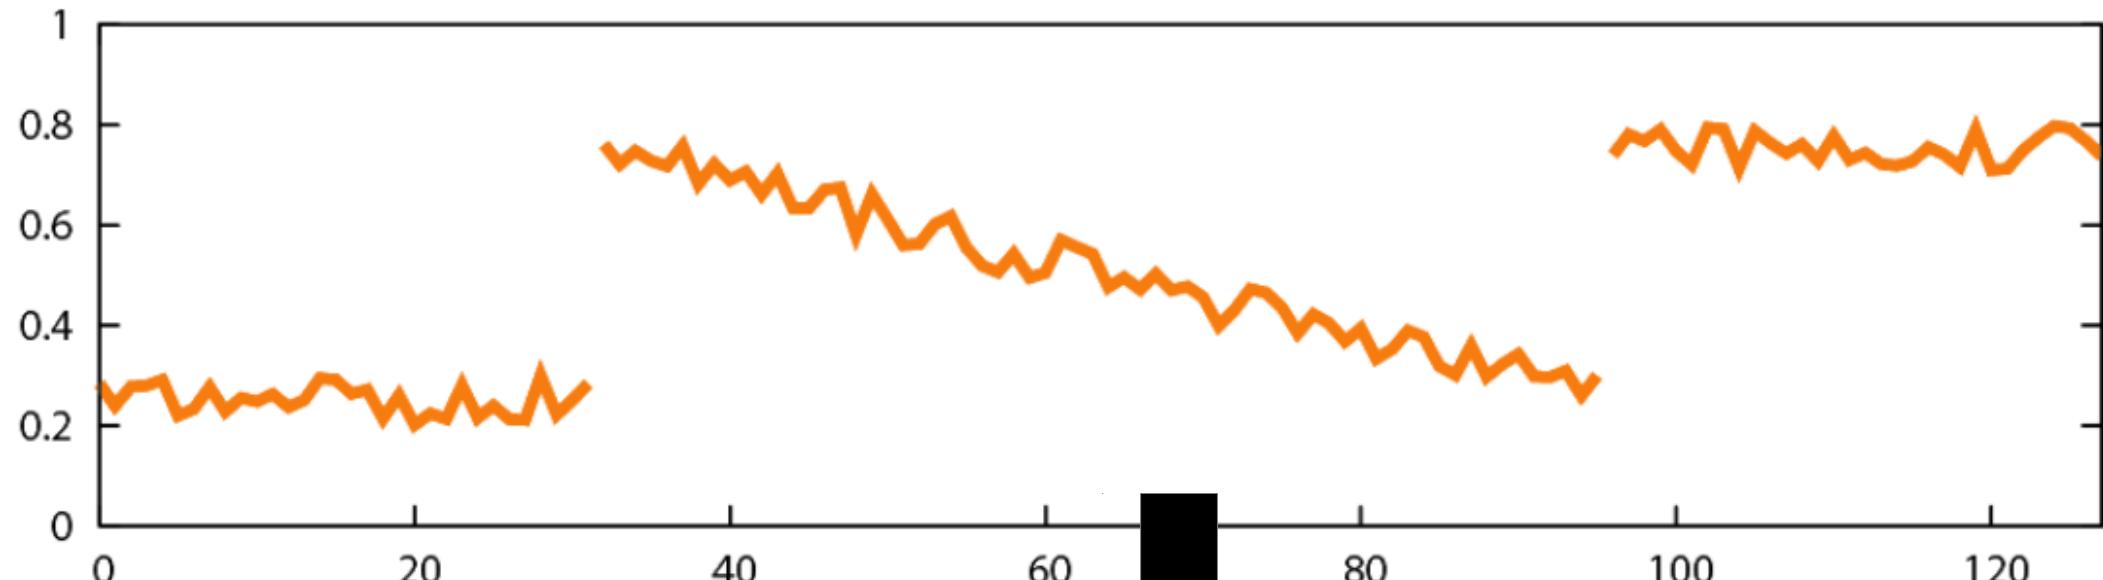
- Bilateral Filter
  - Effect of **Gaussian** filter on weights at pixel “p”
  - Effect of **bilateral** filter on weights at pixel “p”



$$I^{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} I(x_i) f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

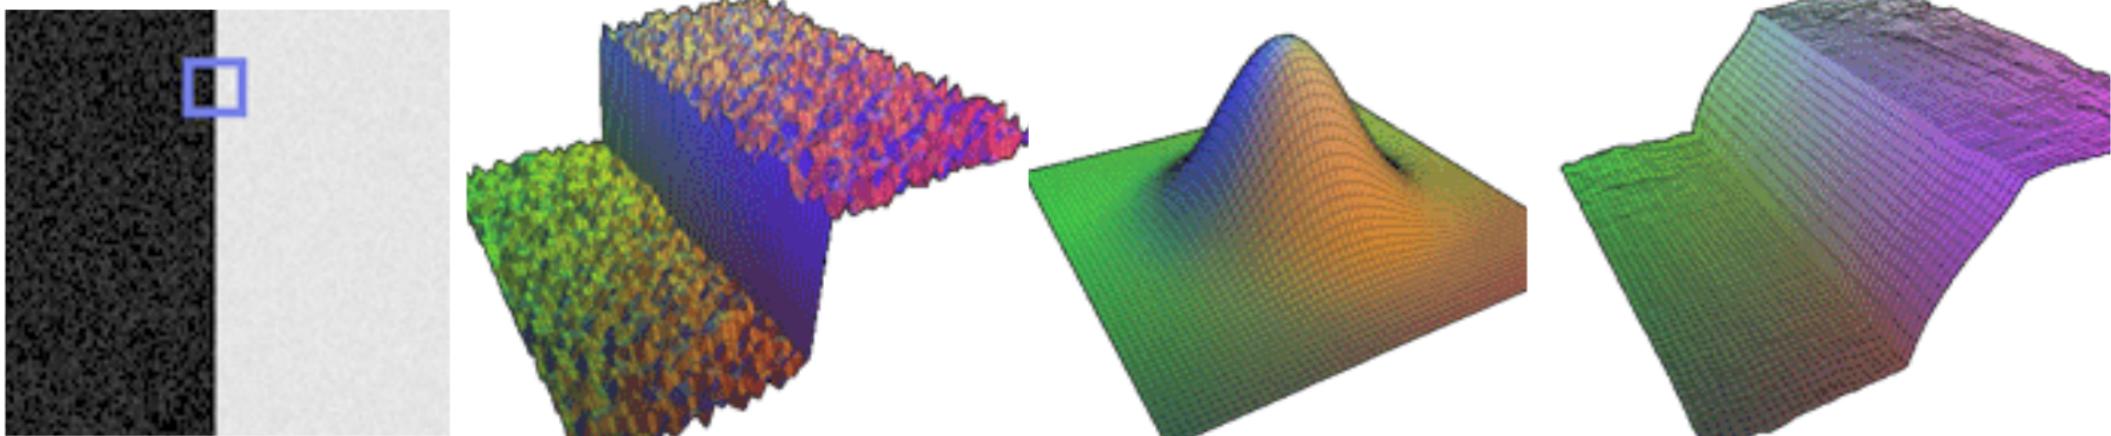
# Edge-Preserving Smoothing

- Bilateral Filter

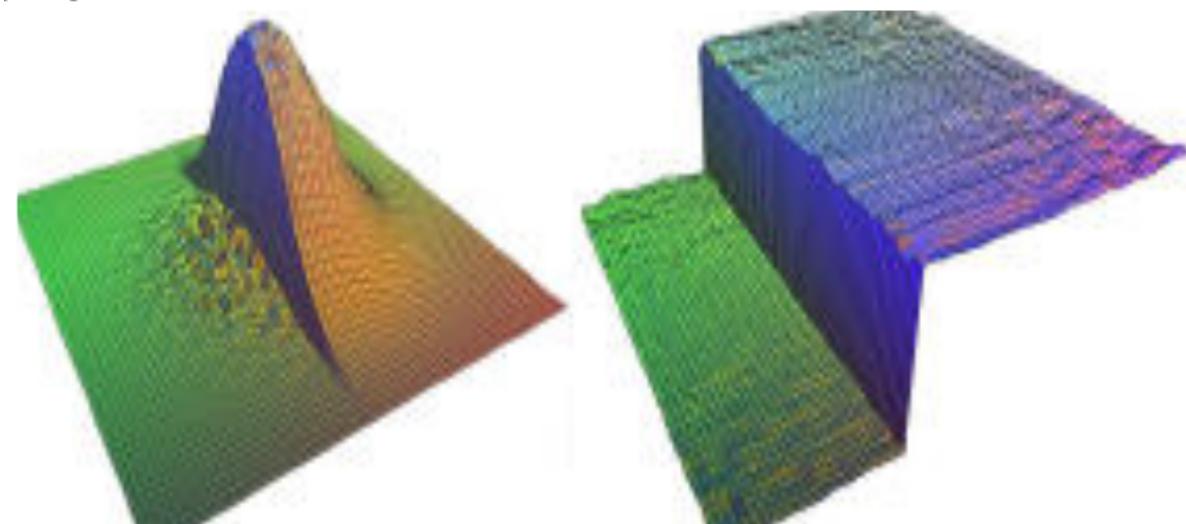


# Edge-Preserving Smoothing

- Bilateral Filter
  - What does Gaussian filter do to a noisy edge in 2D ?



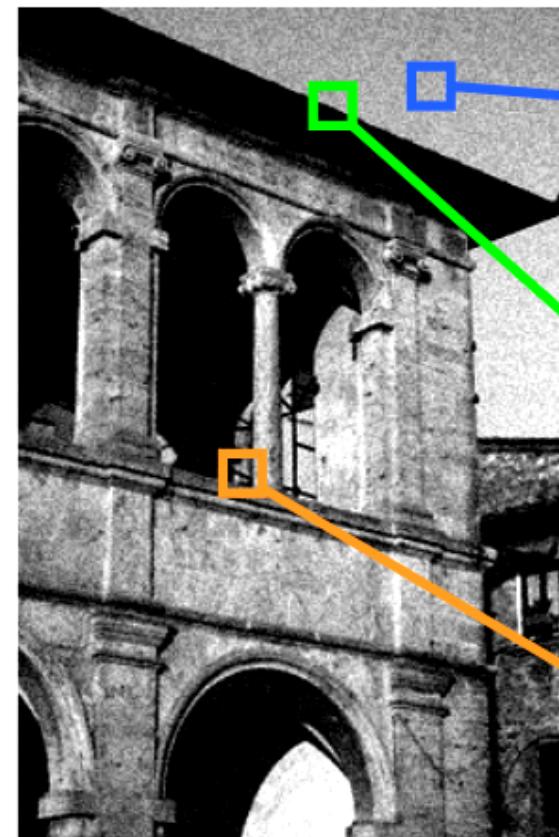
- What does bilateral filter do to a noisy edge in 2D ?
  - Mask for bright pixel =



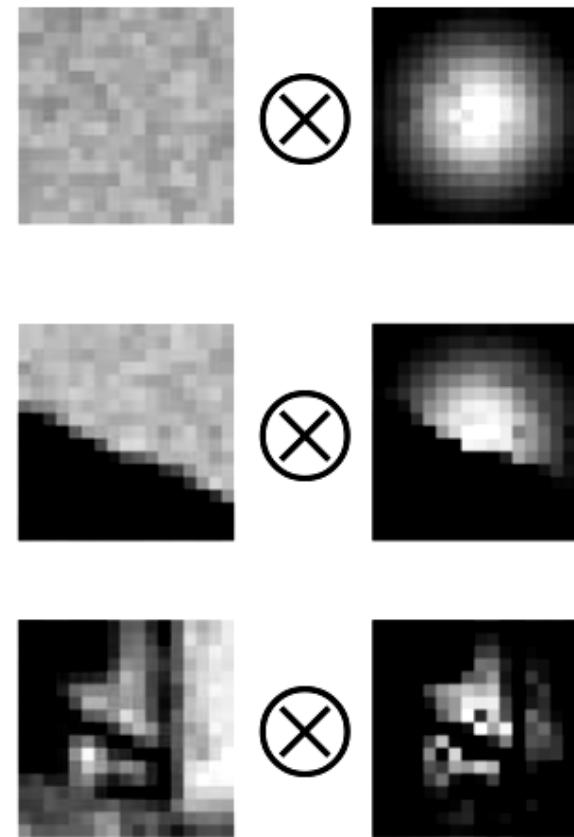
# Edge-Preserving Smoothing

- Bilateral Filter
  - On real images, weight masks can be quite complex

input



output



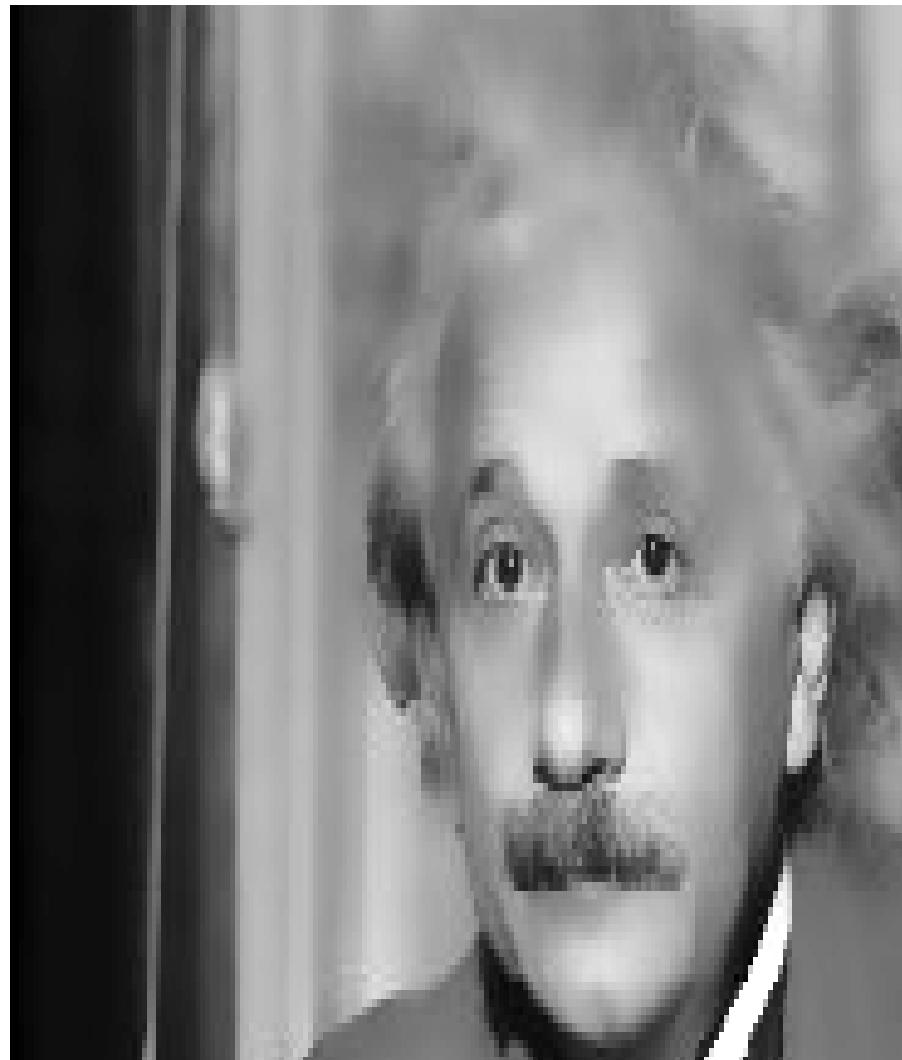
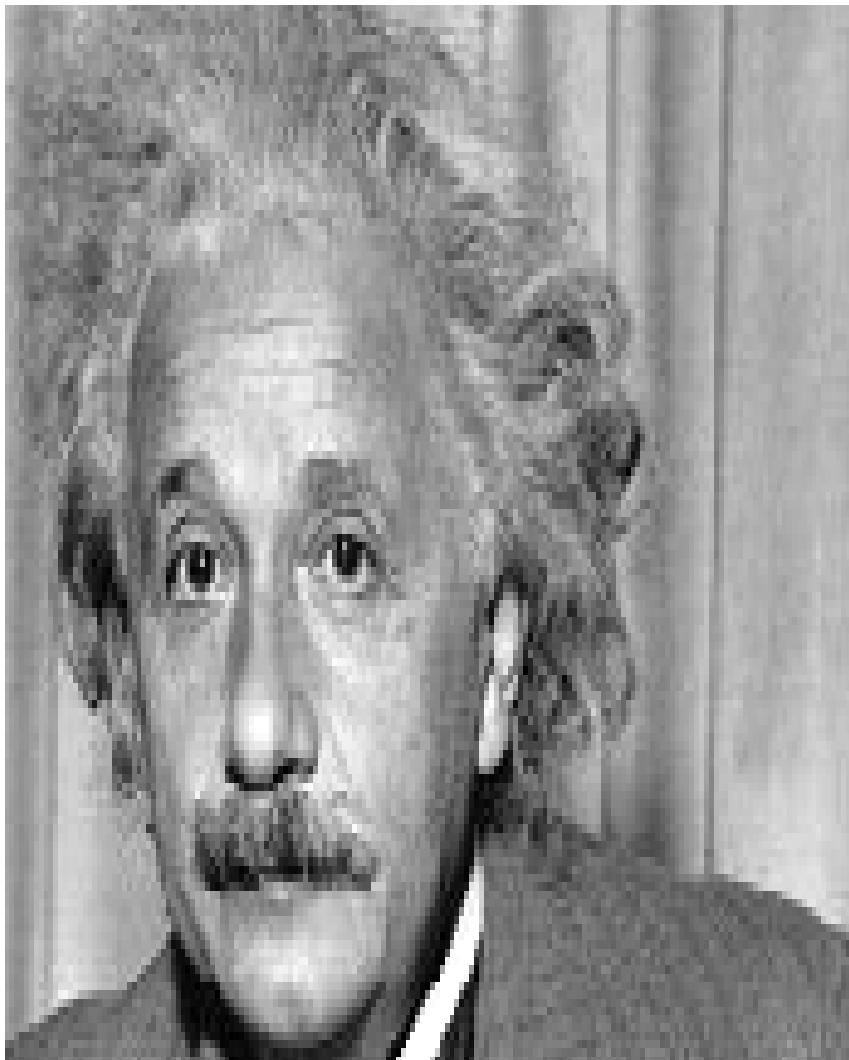
# Edge-Preserving Smoothing

- Bilateral Filter



# Edge-Preserving Smoothing

- Bilateral Filter



# Edge-Preserving Smoothing

- Bilateral Filter



# Edge-Preserving Smoothing

- Bilateral Filter

$$I^{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} I(x_i) f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$

- What happens when,  
for Gaussian  $f(\cdot)$  on **intensities**, variance  $\rightarrow \infty$  ?
- What happens when,  
for Gaussian  $f(\cdot)$  on intensities, variance  $\rightarrow 0$  ?
- What happens when,  
for Gaussian  $g(\cdot)$  on **space**, variance  $\rightarrow \infty$  ?
- What happens when,  
for Gaussian  $g(\cdot)$  on space, variance  $\rightarrow 0$  ?

# Edge-Preserving Smoothing

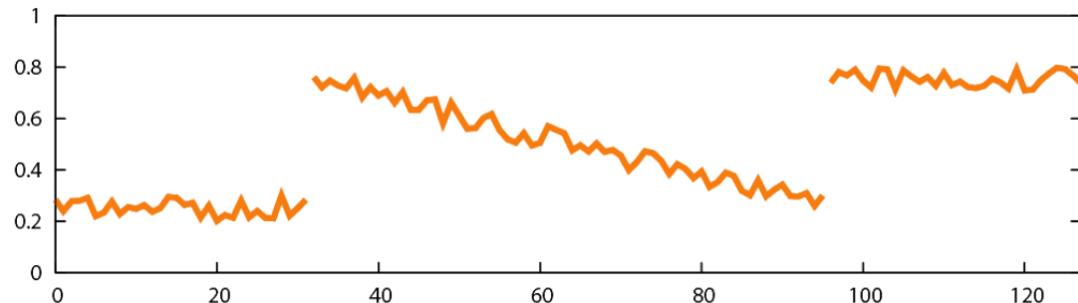
- Bilateral Filter
  - What happens when, for Gaussian  $f(\cdot)$  on **intensities**, variance  $\rightarrow \infty$ ?
    - Bilateral filter  $\rightarrow$  Gaussian smoothing
  - What happens when, for Gaussian  $f(\cdot)$  on intensities, variance  $\rightarrow 0$ ?
    - Filter doesn't modify image
  - What happens when, for Gaussian  $g(\cdot)$  on **space**, variance  $\rightarrow \infty$ ?
    - Filter averages over all pixels. Neighborhood = entire image.
  - What happens when, for Gaussian  $g(\cdot)$  on space, variance  $\rightarrow 0$ ?
    - Filter doesn't modify image

# Edge-Preserving Smoothing



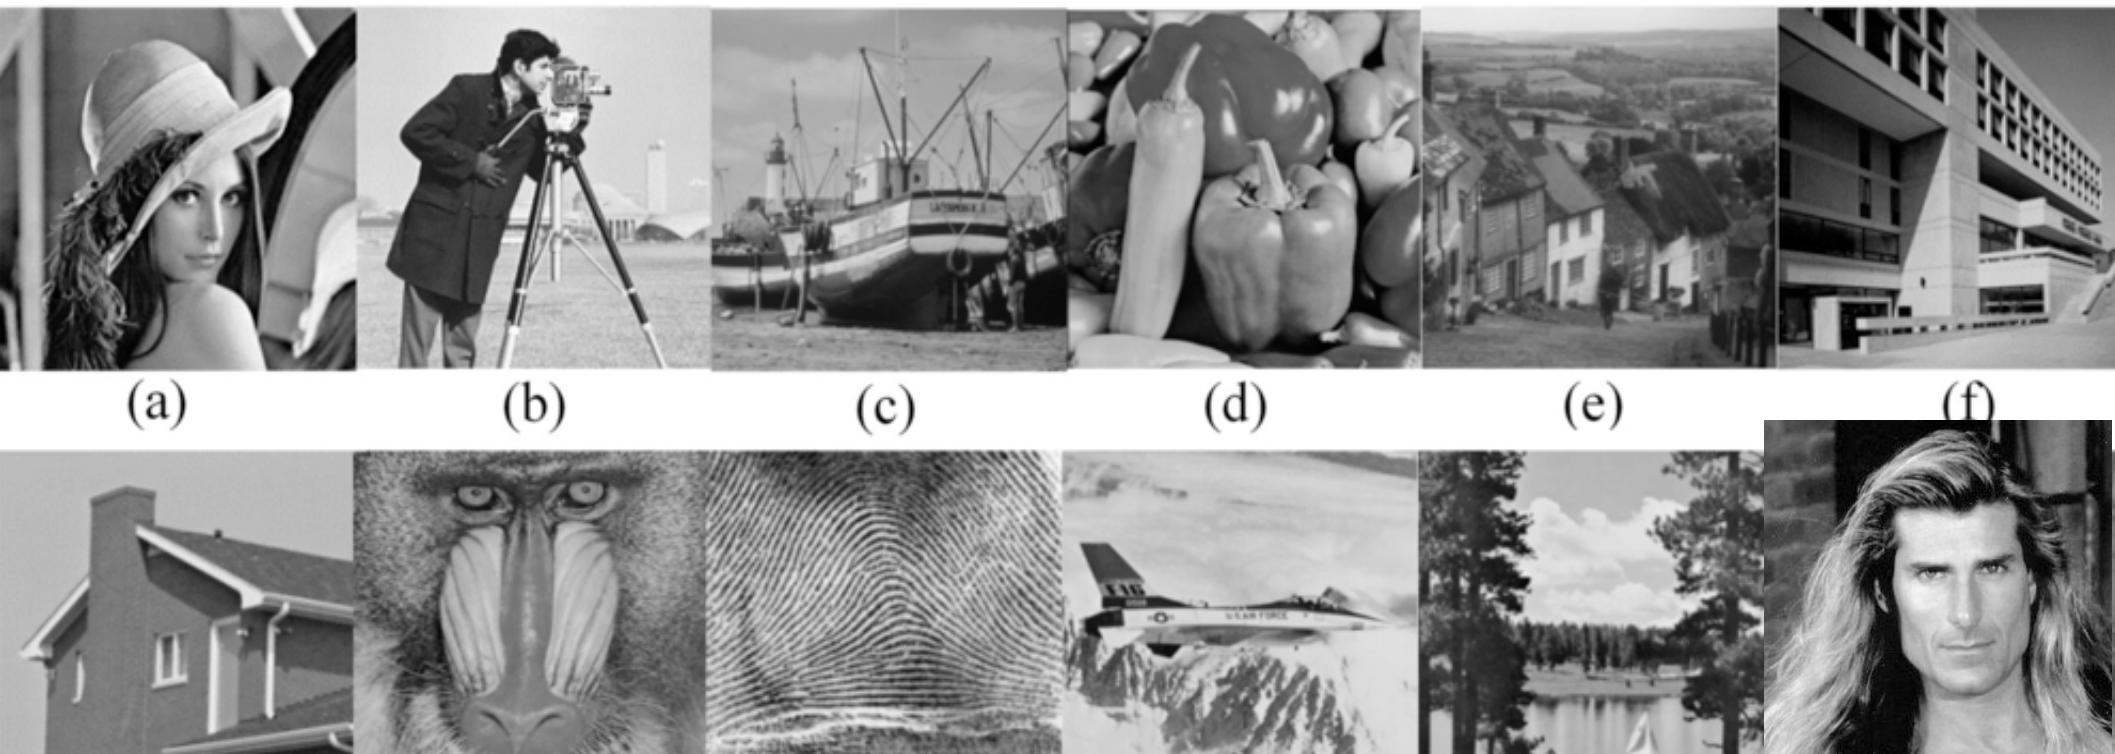
- Bilateral Filter
  - Parameter tuning ?
    - Standard deviation (SD) of Gaussian  $f(\cdot)$  on intensities
      - Edges with contrast  $\leq$  SD are blurred.  
We want these to correspond to noise !
      - Edges with contrast  $\gg$  SD are better preserved.  
We want these to correspond to signal !
    - SD of Gaussian  $g(\cdot)$  on space
      - Larger SD can allow more averaging  $\rightarrow$  more noise reduction
      - Larger SD can allow averaging intensities from different objects
        - May be undesirable

$$I^{\text{filtered}}(x) = \frac{1}{W_p} \sum_{x_i \in \Omega} I(x_i) f_r(\|I(x_i) - I(x)\|) g_s(\|x_i - x\|)$$



# Standard Test Images

- <http://sipi.usc.edu/database>
- [http://www.imageprocessingplace.com/root\\_files\\_V3/image\\_databases.htm](http://www.imageprocessingplace.com/root_files_V3/image_databases.htm)
- [https://en.wikipedia.org/wiki/Standard\\_test\\_image](https://en.wikipedia.org/wiki/Standard_test_image)
- And more ...



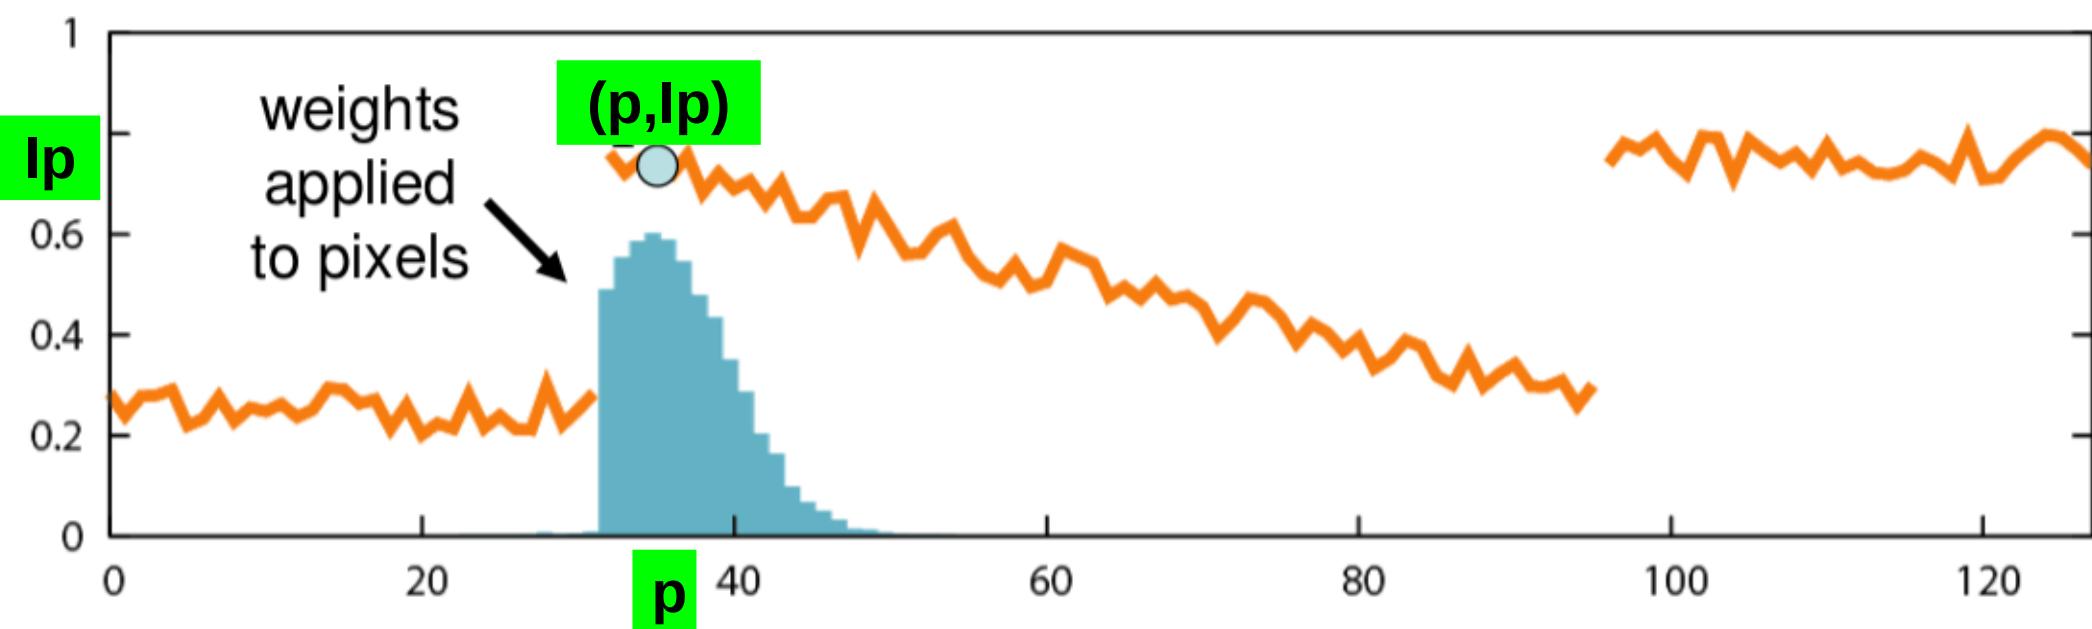
# Edge-Preserving Smoothing

- Bilateral Filter
  - Is nonlinear → NO convolution → no separability → slow
  - But can we still compute it fast ?
  - Yes !
  - How ?
    - Rewrite it using linear filtering in a different space (NOT image-coordinate space)
    - Fast and (sufficiently) accurate algorithm

[ [cs.brown.edu/courses/cs129/lectures/bf\\_course\\_Brown\\_Oct2012.pdf](http://cs.brown.edu/courses/cs129/lectures/bf_course_Brown_Oct2012.pdf) ]

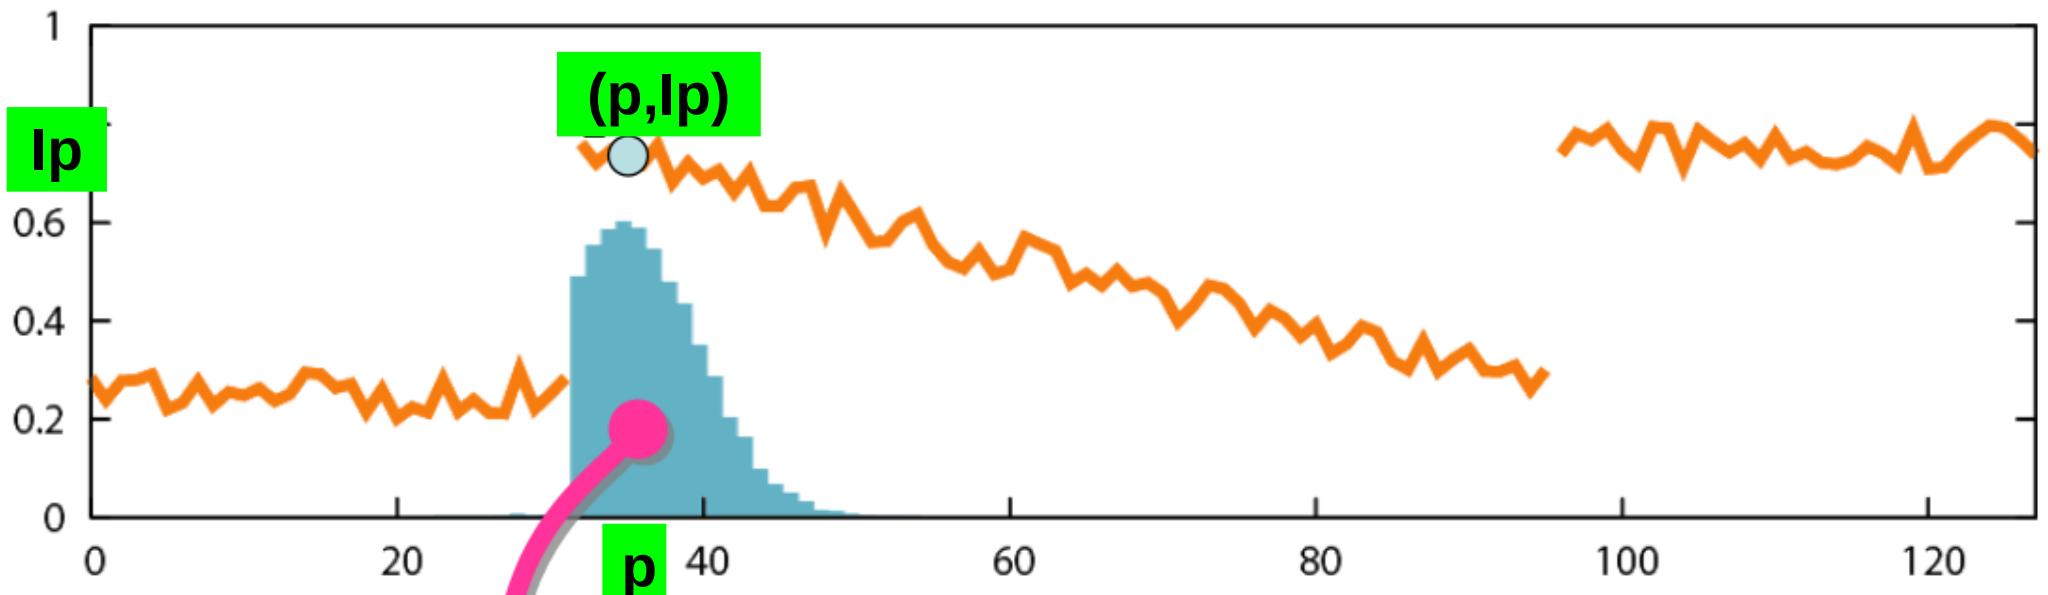
# Edge-Preserving Smoothing

- Bilateral Filter
  - Recap
  - Processing at pixel “p”



# Edge-Preserving Smoothing

- Bilateral Filter



$$I_p^{bf} = \frac{1}{W_p^{bf}} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

# Edge-Preserving Smoothing

- Bilateral Filter

- Two key equations for processing pixel “p”

$$I_p^{bf} = \frac{1}{W_p^{bf}} \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) I_q$$

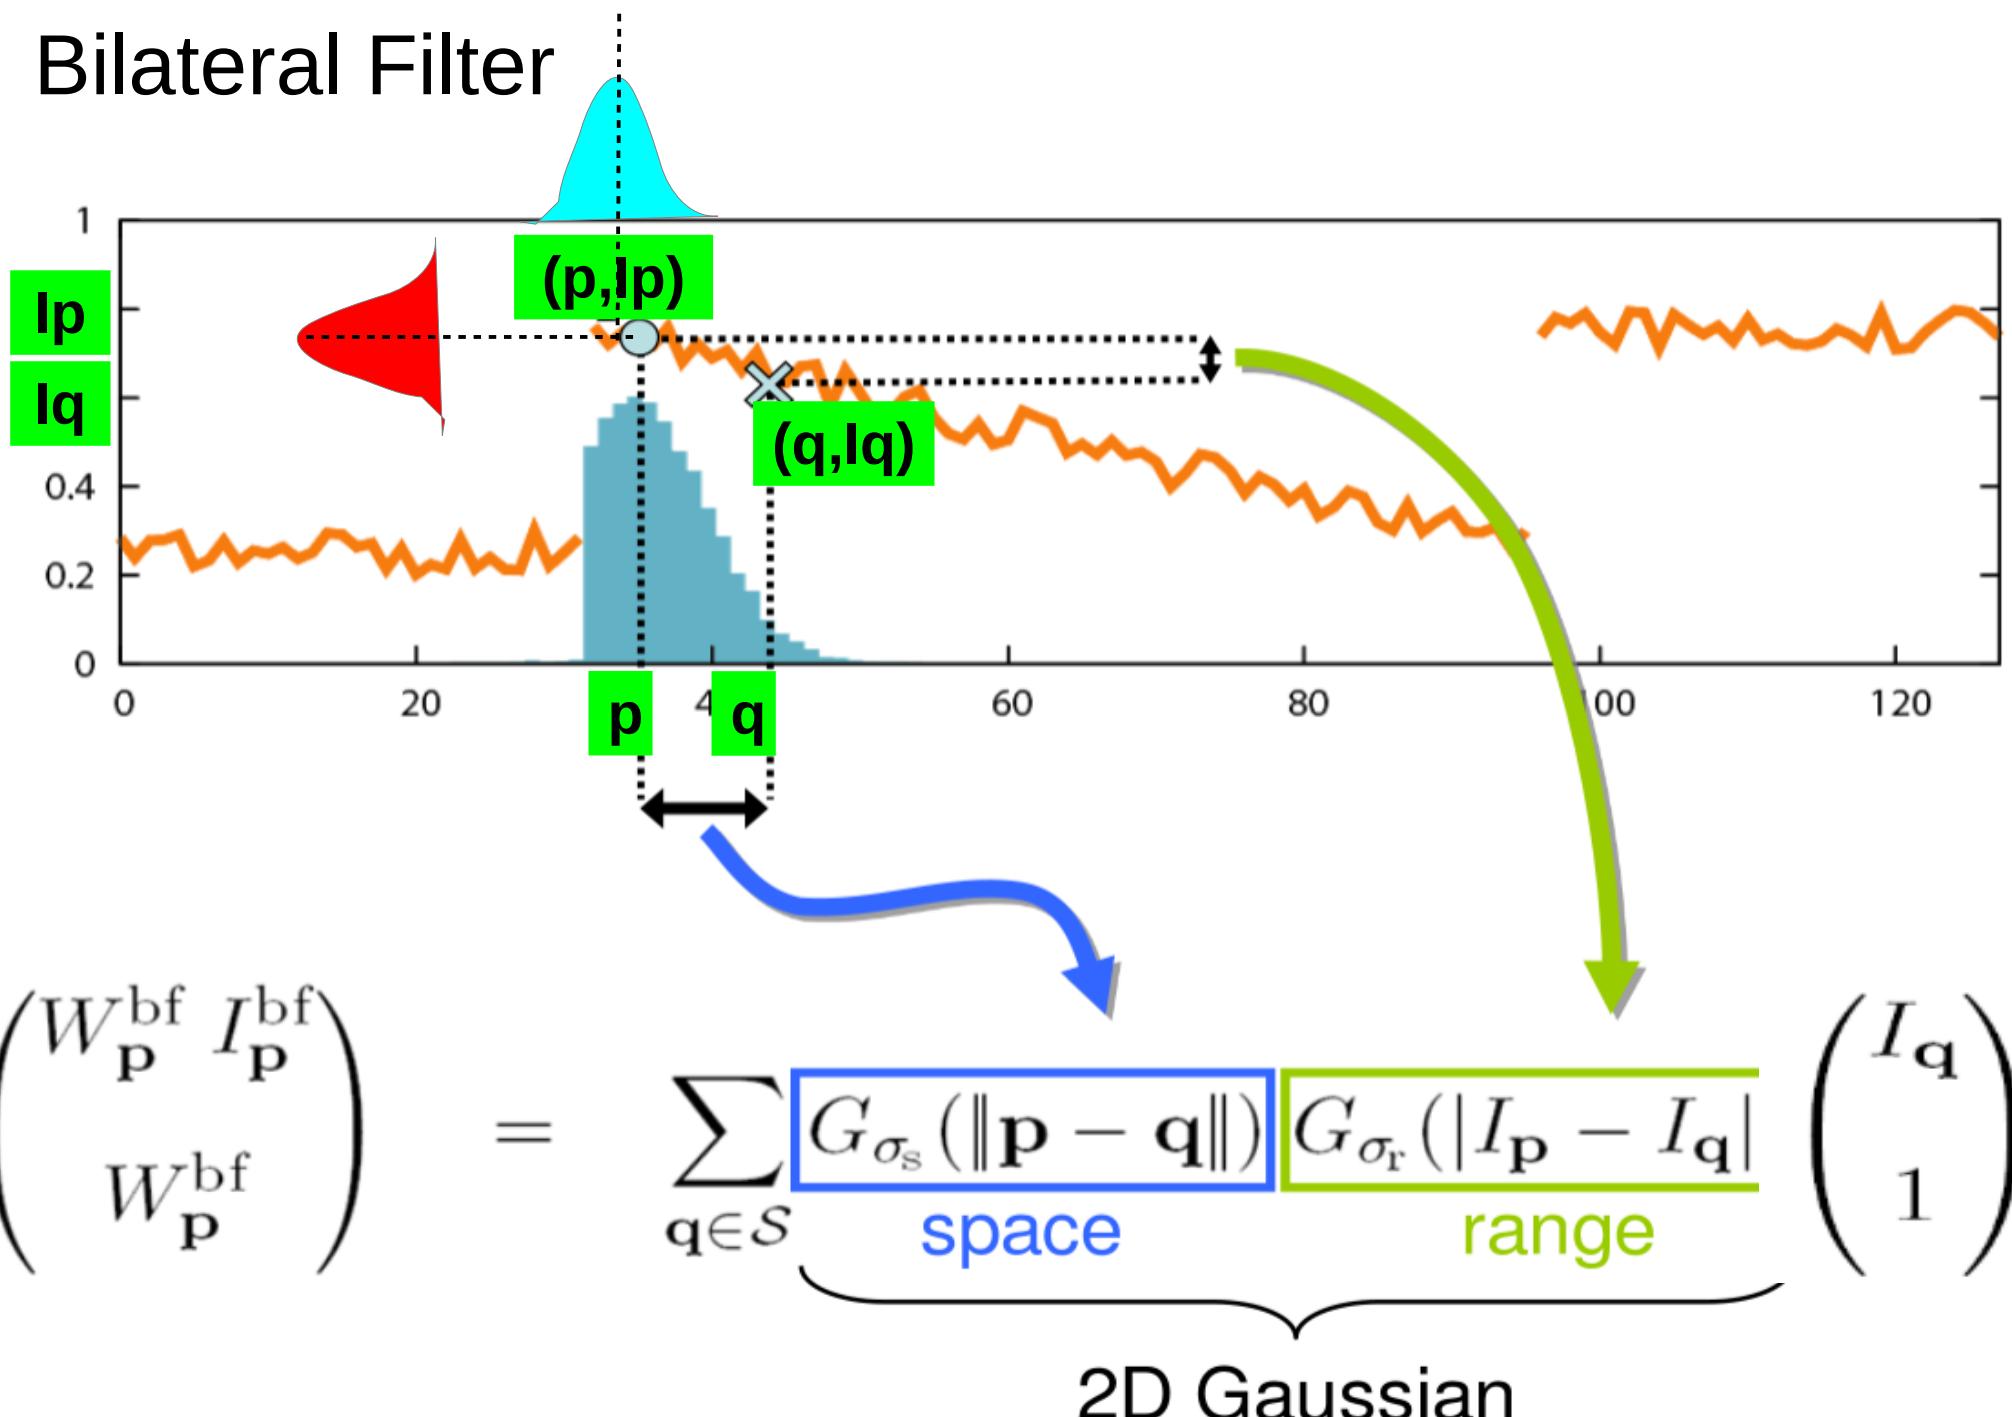
$$W_p^{bf} = \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|)$$

- In 1<sup>st</sup> equation, multiply both sides by normalization factor

$$\begin{pmatrix} W_p^{bf} & I_p^{bf} \\ W_p^{bf} & \end{pmatrix} = \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) \begin{pmatrix} I_q \\ 1 \end{pmatrix}$$

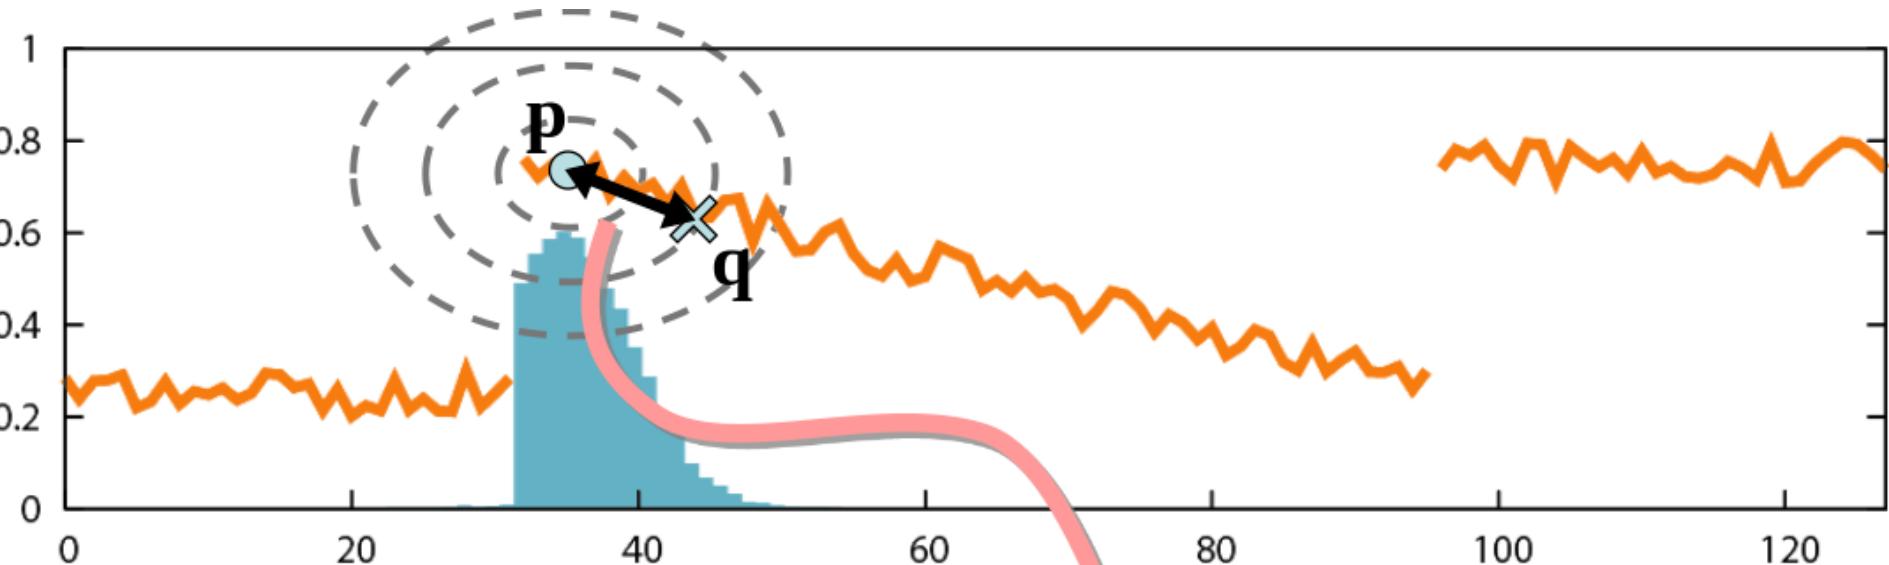
# Edge-Preserving Smoothing

- Bilateral Filter



# Edge-Preserving Smoothing

- Bilateral Filter



$$\begin{pmatrix} W_p^{\text{bf}} & I_p^{\text{bf}} \\ W_p^{\text{bf}} \end{pmatrix} = \sum_{q \in S} G_{\sigma_s}(\|p - q\|) G_{\sigma_r}(|I_p - I_q|) \begin{pmatrix} I_q \\ 1 \end{pmatrix}$$

space x range

# Edge-Preserving Smoothing

- Bilateral Filter : **Convolution in <location,intensity>**



sum all values

$$\begin{pmatrix} W_p^{bf} & I_p^{bf} \\ W_p^{bf} \end{pmatrix} = \boxed{\sum_{(\mathbf{q}, \zeta) \in \mathcal{S} \times \mathcal{R}}}$$

space-range Gaussian

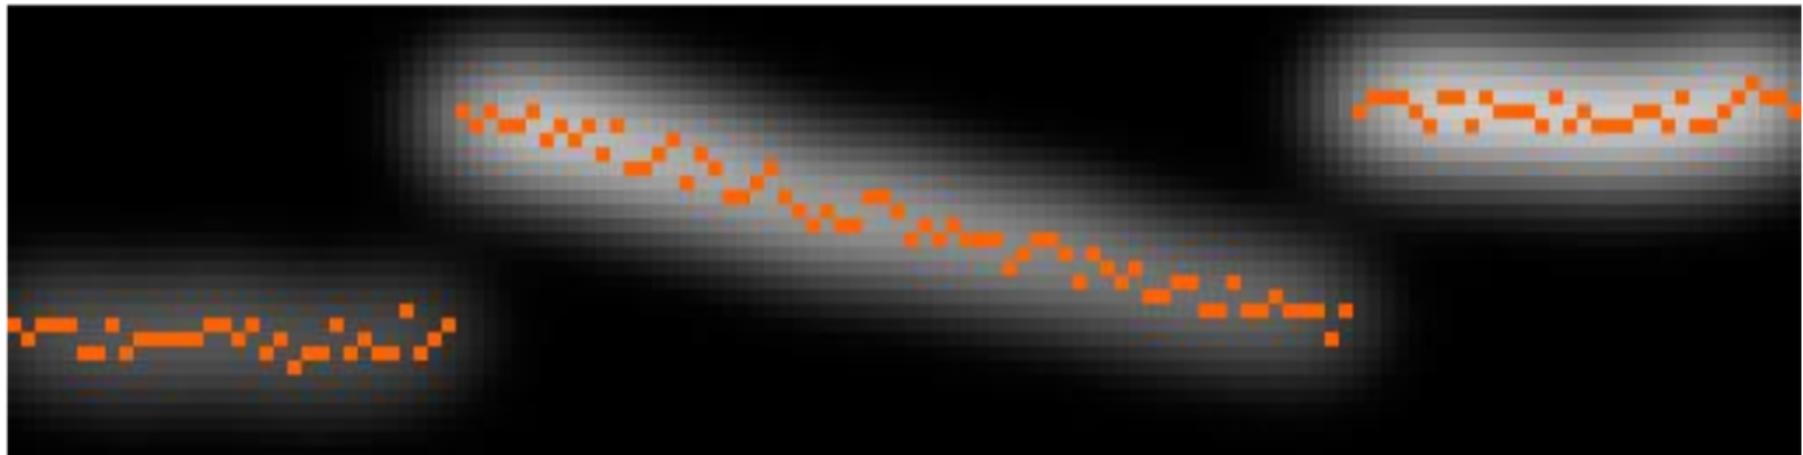
To compute **numerator**.  
Construct a **joint space**.  
1) At “location” ( $\mathbf{p}, I_p$ ),  
assign intensity  $I_p$   
2) Assign 0 elsewhere

$$\begin{pmatrix} I_q \\ 1 \end{pmatrix}$$

# Edge-Preserving Smoothing

- Bilateral Filter

Get numerator values by picking them from “locations” ( $q, I_q$ ) in convolved image



result of the convolution



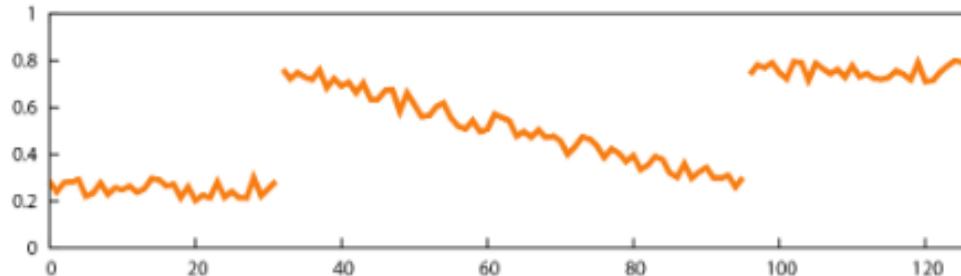
$$\begin{pmatrix} W_p^{bf} & I_p^{bf} \\ W_p^{bf} \end{pmatrix} = \sum_{(\mathbf{q}, \zeta) \in \mathcal{S} \times \mathcal{R}} \text{space-range Gaussian} \begin{pmatrix} I_{\mathbf{q}} \\ 1 \end{pmatrix}$$

# Edge-Preserving Smoothing

- Bilateral Filter
  - How do we compute values in the denominator ?

# Edge-Preserving Smoothing

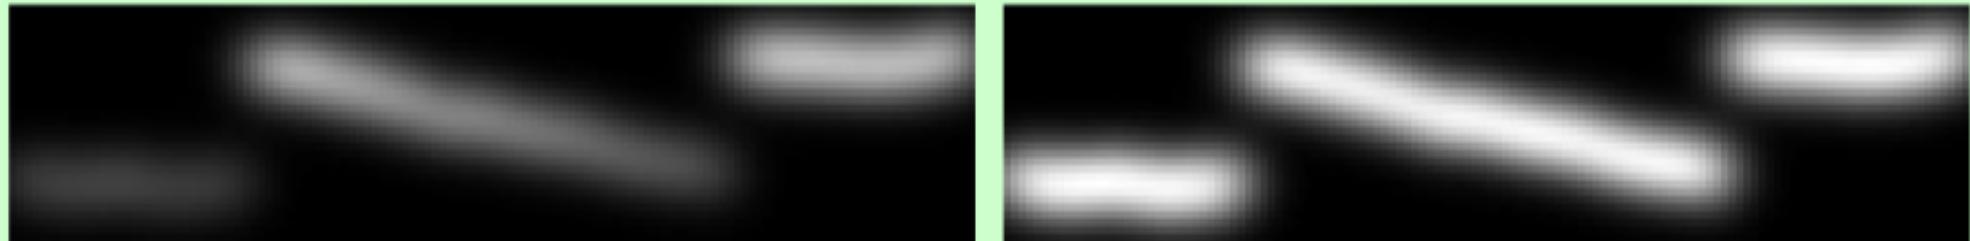
- Bilateral Filter: Get numerator, denominator values



higher dimensional functions

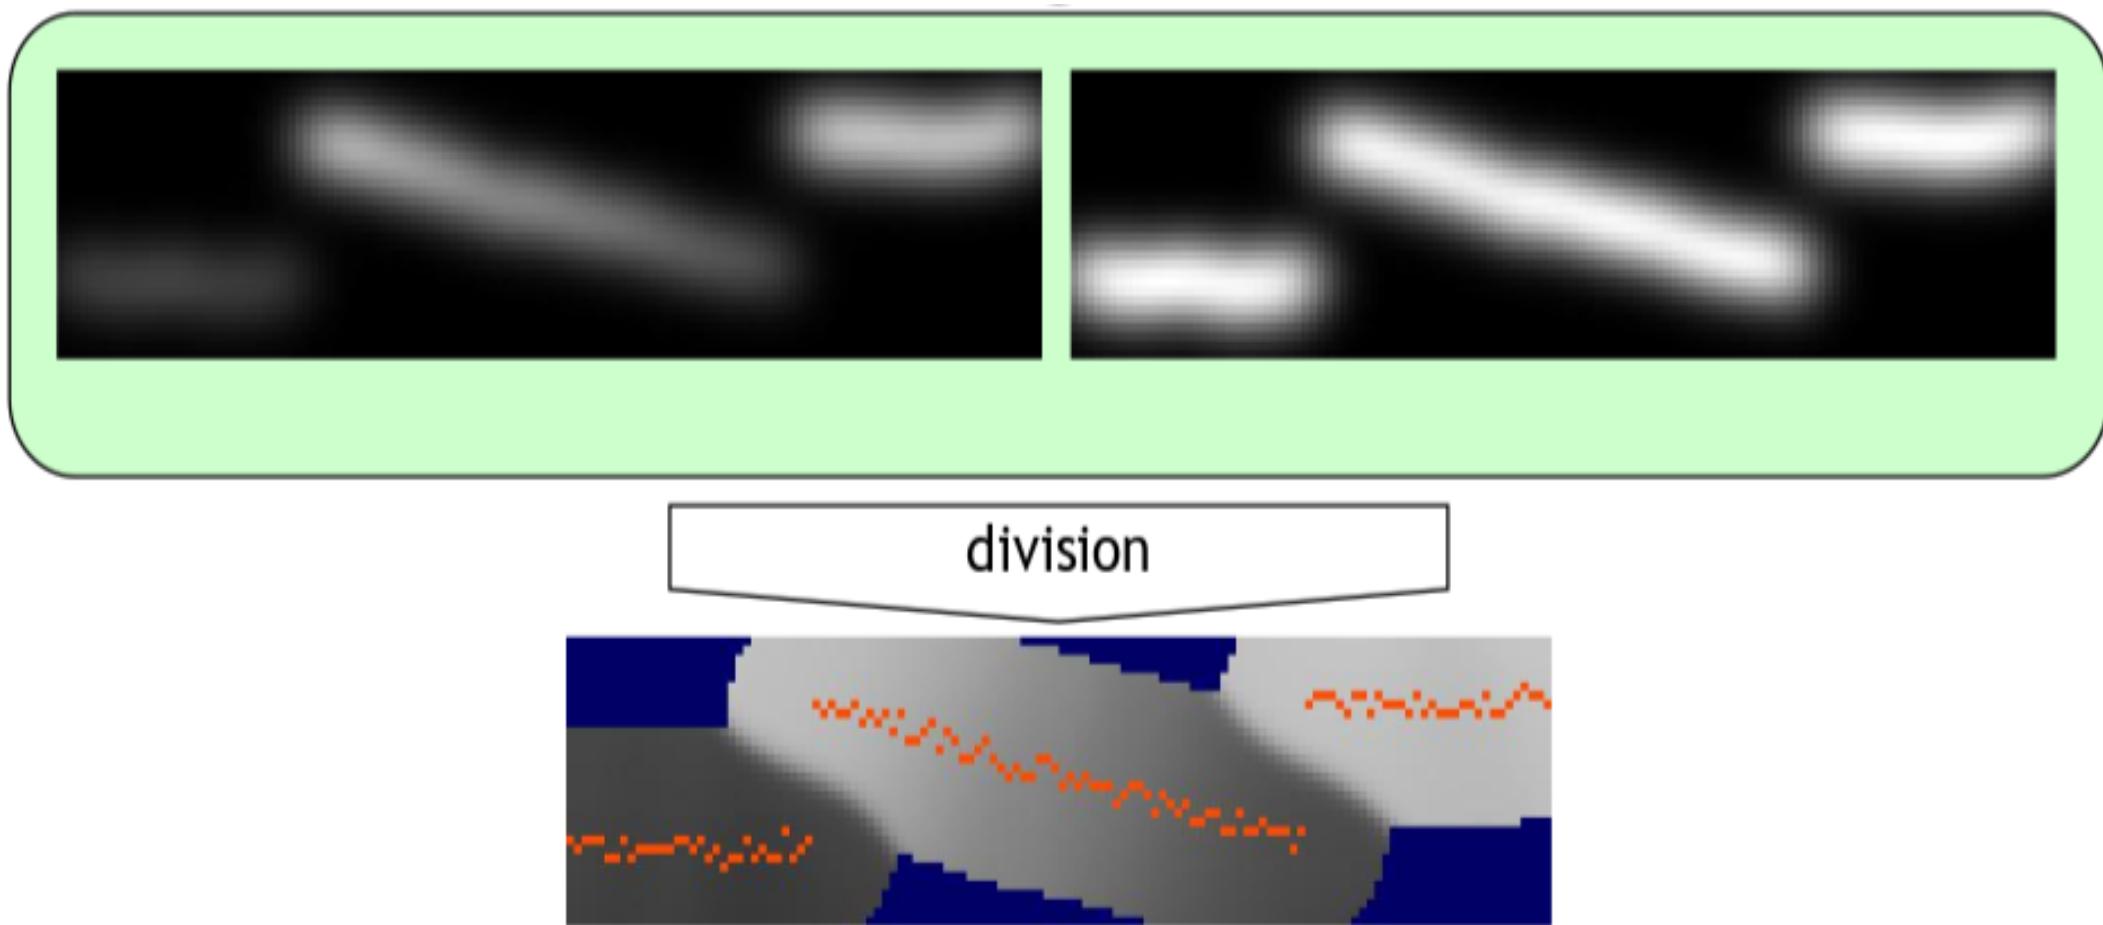


Gaussian convolution



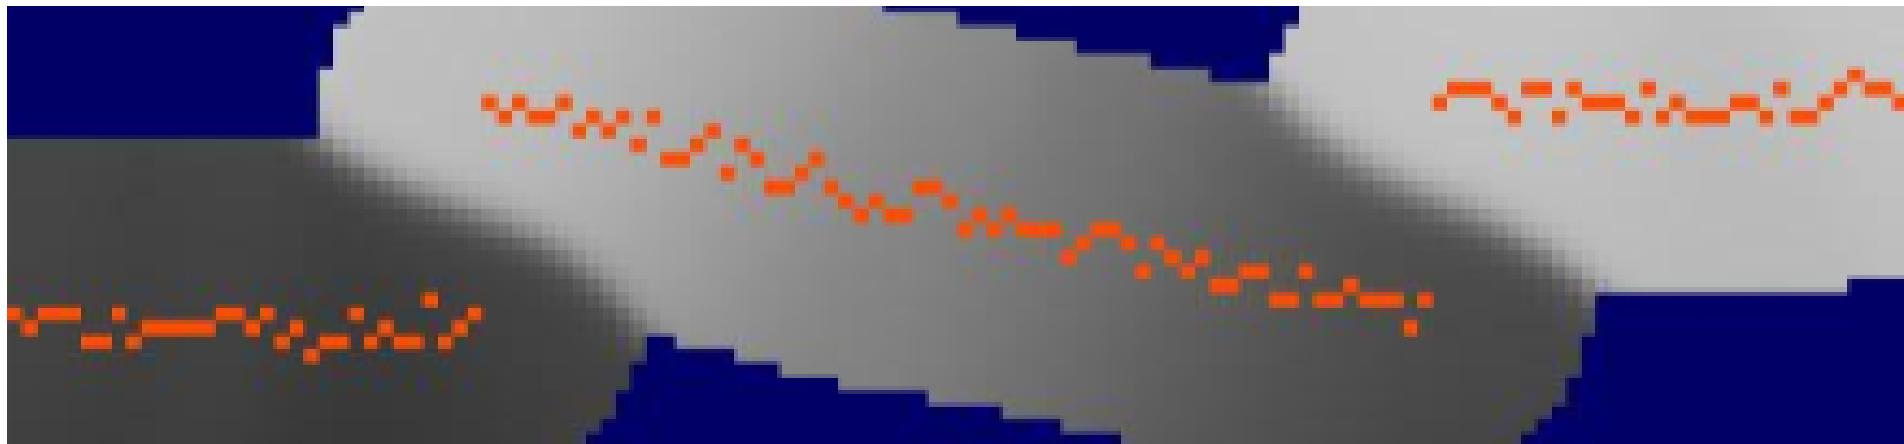
# Edge-Preserving Smoothing

- Bilateral Filter: Divide numerator by denominator
  - Consider values only at “locations” ( $p, l_p$ ) in numerator image and denominator image

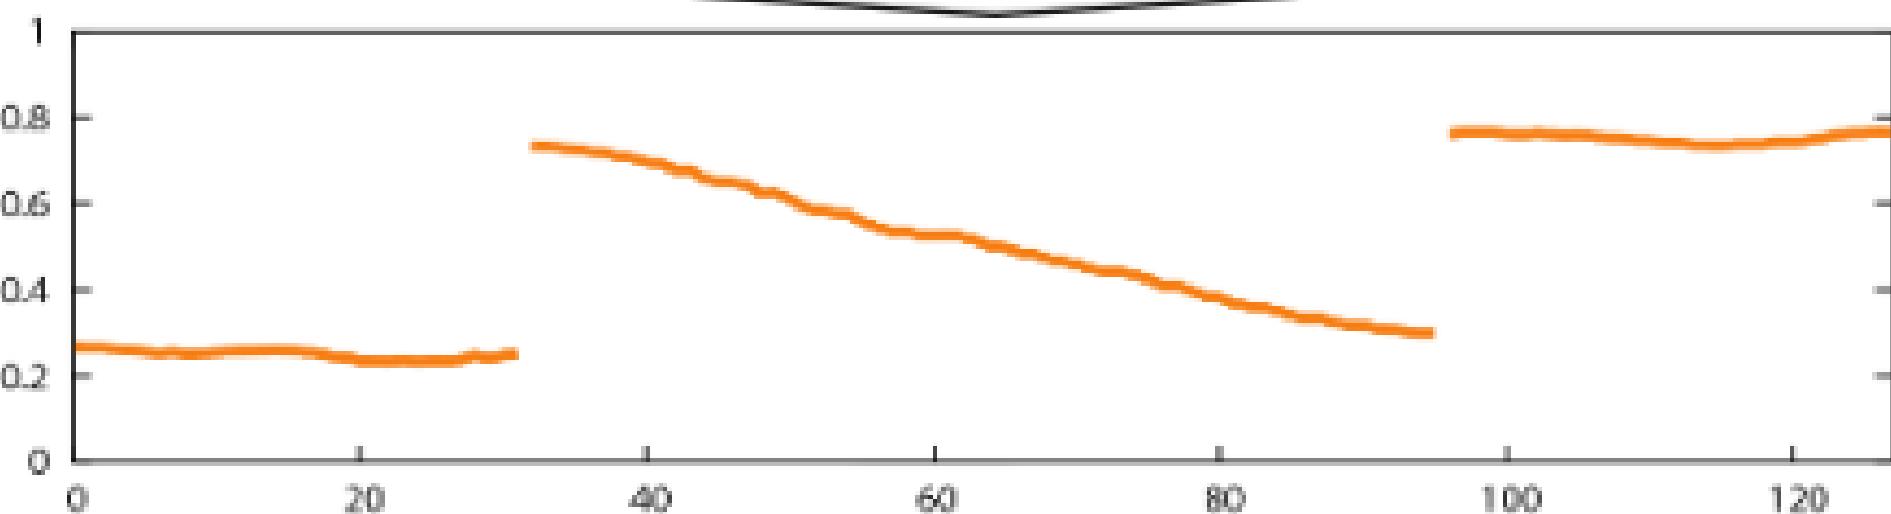


# Edge-Preserving Smoothing

- Bilateral Filter: Update image to get filtered result
  - At pixel “p”, filtered value is the division at “location”  $(p, l_p)$



slicing

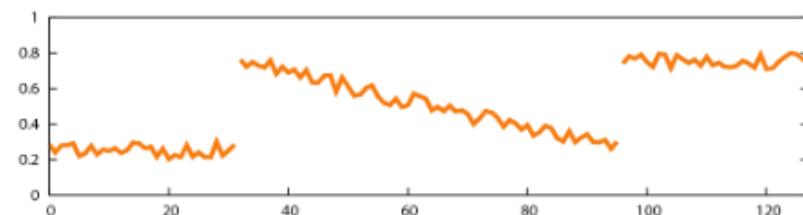


# Edge-Preserving Smoothing

- Bilateral Filter
  - What did we do ?
    - (1) Two convolutions in a higher-dimensional space  $\langle p, l_p \rangle$ 
      - Linear operation
      - Linear time complexity
    - (2) Pixel-by-pixel division between 2 images
      - Nonlinear operation
      - Simple, fast,  $O(\# \text{ pixels in image})$
  - Exact formulation (analytically)
  - Why is this faster than standard algorithm ?
    - Lets see next ...

## Recap:

- simple operations
- complex space



higher dimensional functions

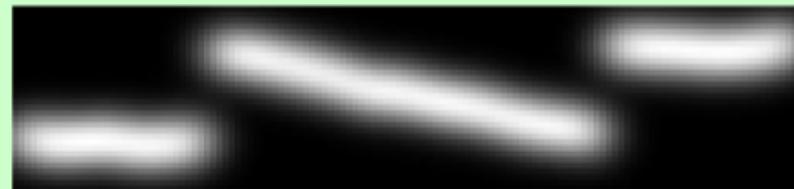
$w_i$



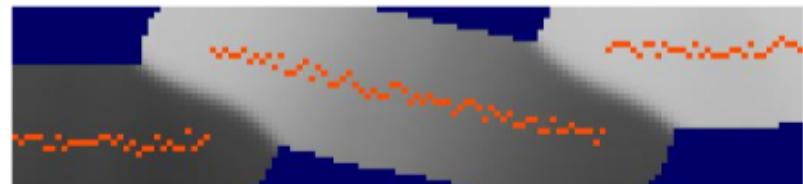
$W$



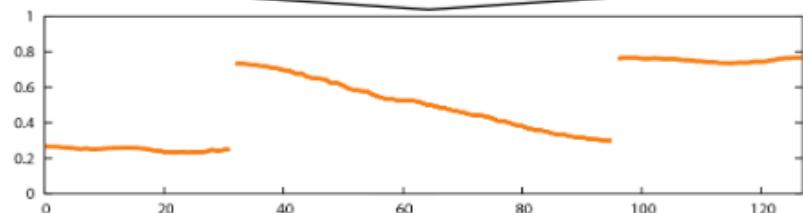
Gaussian convolution



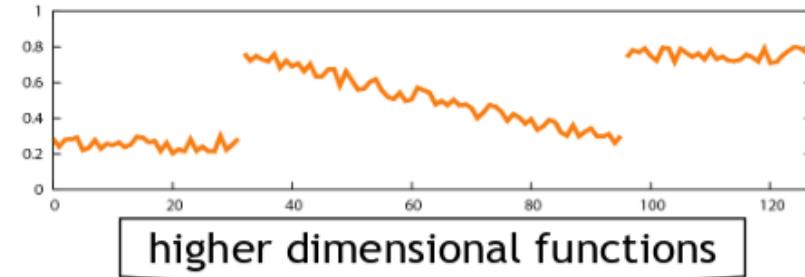
division



slicing

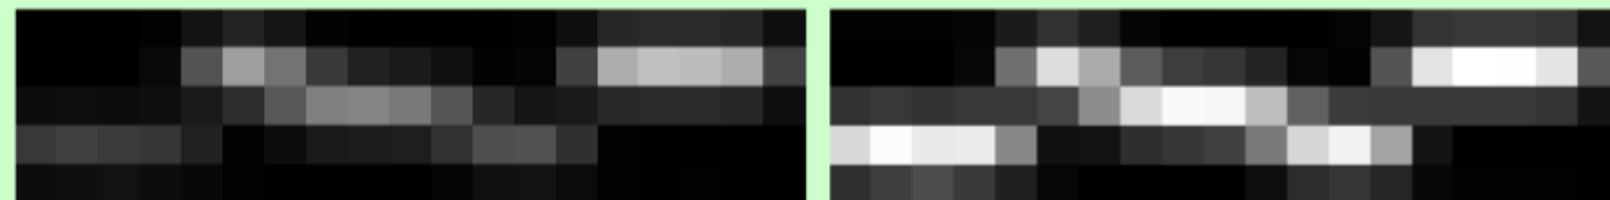


# Strategy: downsampled convolution



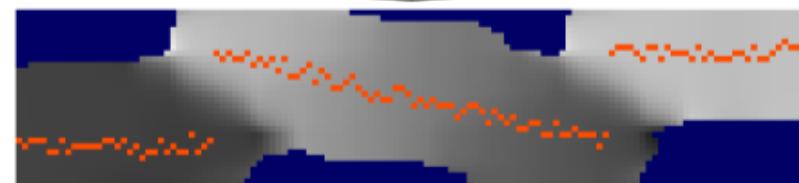
**DOWNSAMPLE** Averaging in square with one corner at pixel

Gaussian convolution

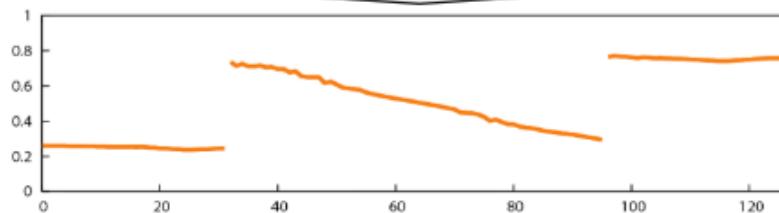


**UPSAMPLE**

division



slicing

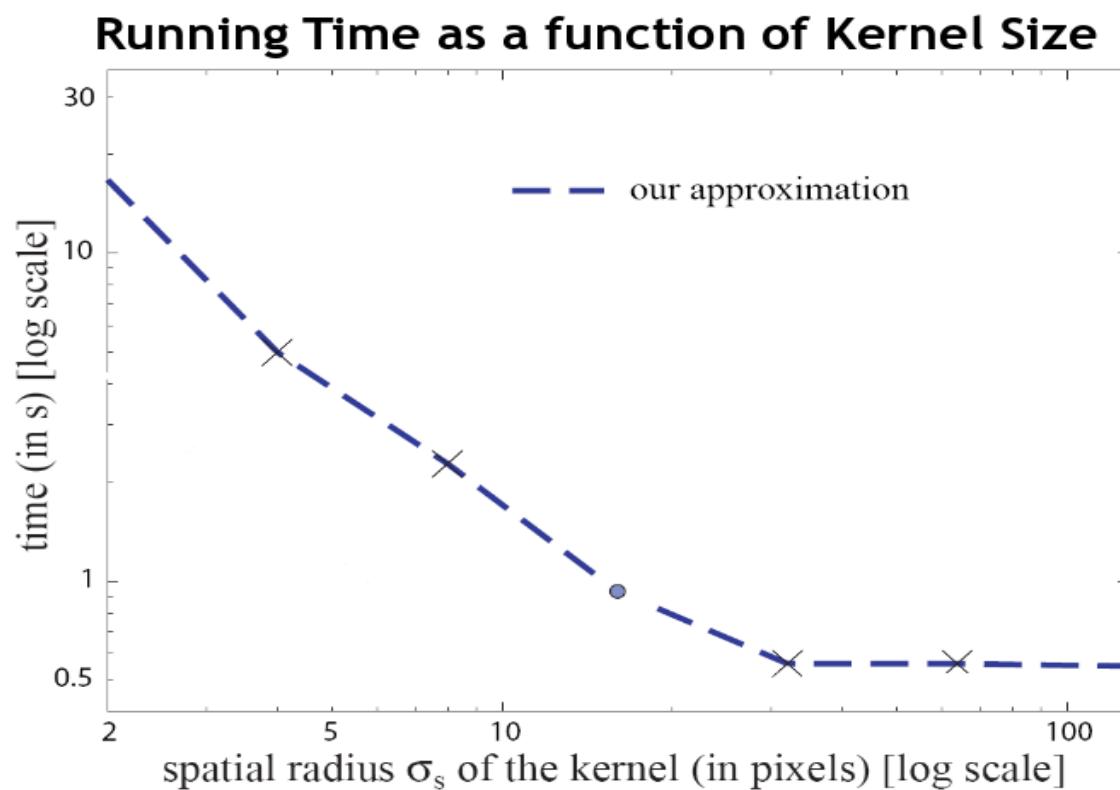


# Edge-Preserving Smoothing

- Bilateral Filter
  - Why did downsampling NOT change results much ?
    - Convolution blurs anyway
    - Downsampling must involve some “**smoothing**” / **averaging**
      - NOT selecting every n-th pixel !
    - But that averaging can be done in **MI operations** only
      - Not ( $MI * \text{box-filter-size}$ ) operations; No convolution !
  - Will this downsampling work for a general image ?
    - Yes
    - To understand why, we need concepts from:
      - Fourier analysis
      - Sampling theorem
      - Smoothing as low-pass filtering

# Edge-Preserving Smoothing

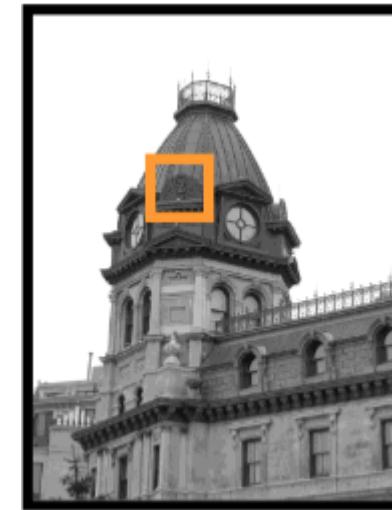
- Bilateral Filter
  - Effect of downsampling
    - Faster algorithm  $\leftarrow$  less data to process
    - Error  $\leftarrow$  approximation
  - Larger kernel  $\rightarrow$  Can afford more downsampling
    - So faster algorithm !
  - Typical trade off
    - Speed versus accuracy
  - Evaluate



# Edge-Preserving Smoothing

- Bilateral Filter

- Evaluate running time
- 1200 x 600 image with reasonable parameter settings
  - Original algorithm → 10 minutes
  - New algorithm → Few seconds
- Algorithmic complexity ( $M \times N$  image,  $P \times Q$  mask)
  - Original algorithm :  $M N P Q 3$  (find exp weights; multiply twice)
  - New algorithm :
    - Downsample image + mask:  $M N I \rightarrow M/a N/b I/c, P Q J \rightarrow P/a Q/b J/c$
    - Separable convolutions in joint space:  $2 (M/a N/b I/c) (P/a + Q/b + J/c)$ 
      - If  $a = b = c$ , then operations =  $(M N I) (P + Q + J) (2 / a^4)$
    - Upsampling:  $M N 4$  (bilinear interpolation at  $p-l_p$  locations only)
    - Division and update :  $M N$



1200 × 1600

# Edge-Preserving Smoothing

- Bilateral Filter

- Evaluate accuracy

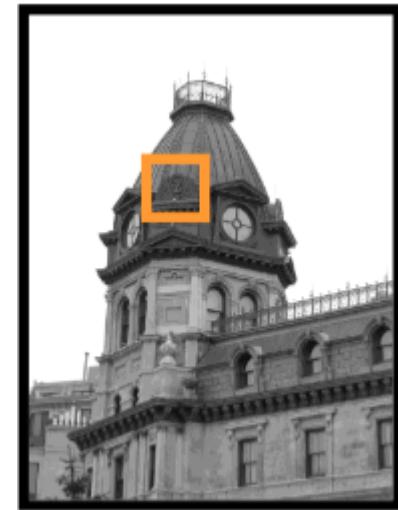
input



exact BF



our result



1200 × 1600

difference  
with exact  
computation  
(intensities in [0:1])



## NOVICE PROGRAMMER



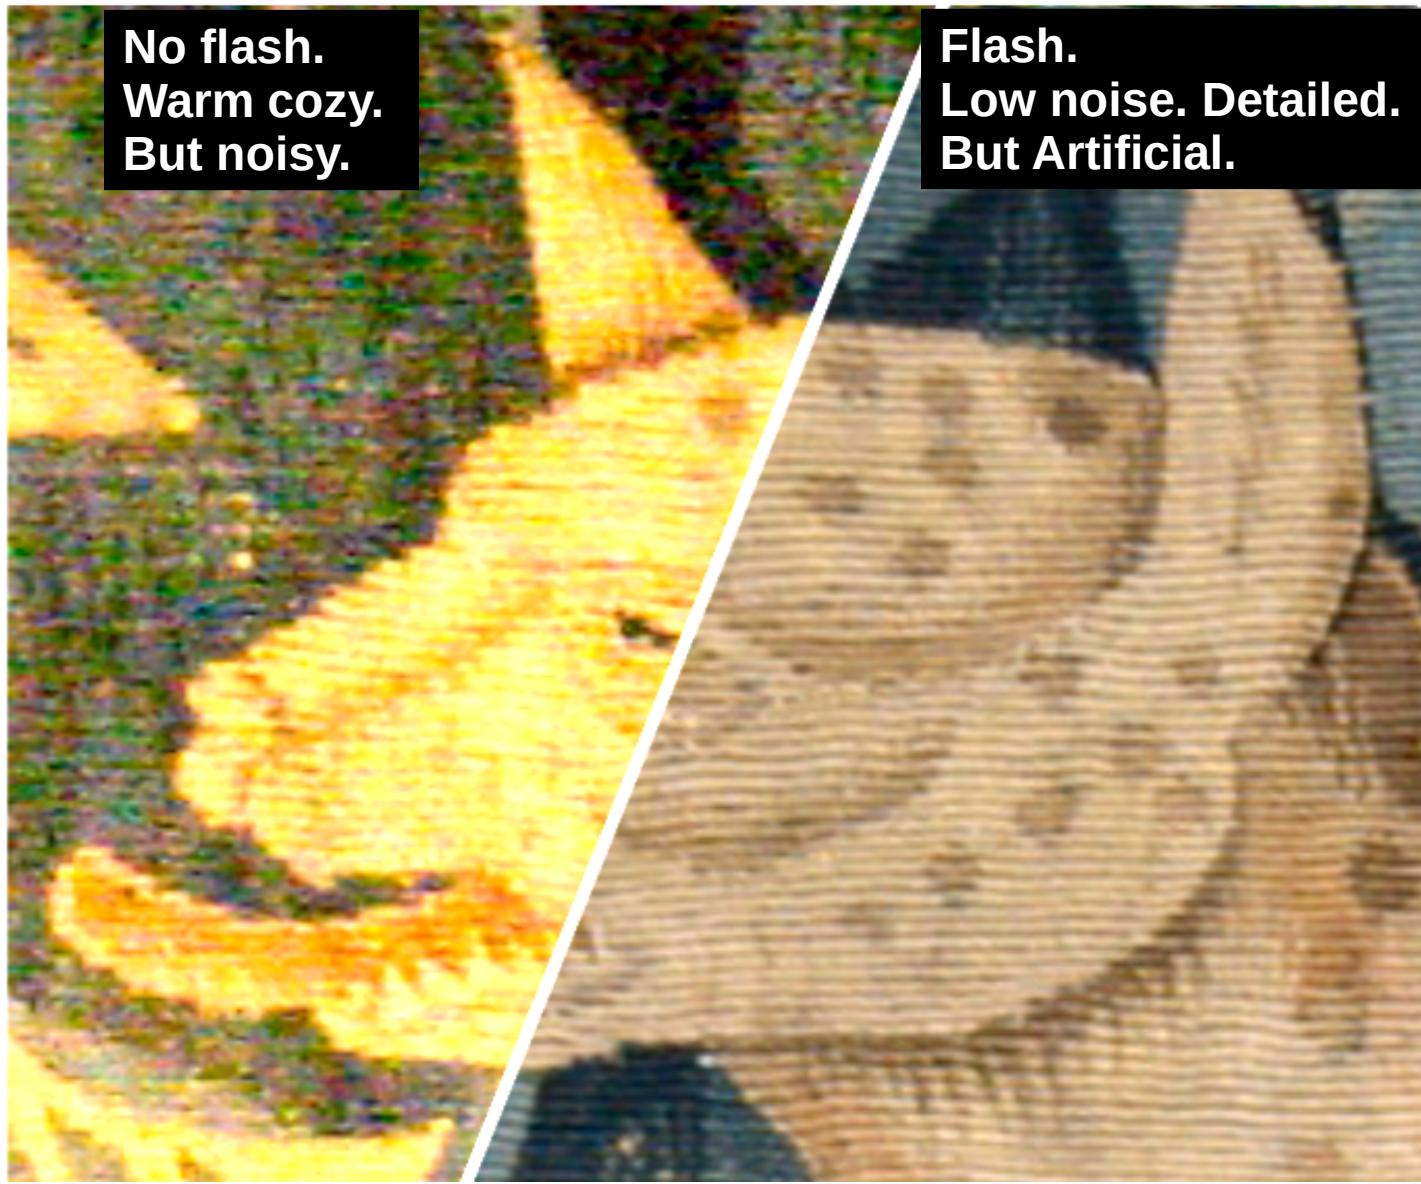
## EXPERIENCED PROGRAMMER



# Edge-Preserving Smoothing

- Cross / Joint Bilateral Filter
  - Application: Merging camera images with & without flash

- Images of a Belgian carpet



# Edge-Preserving Smoothing

- Cross / Joint Bilateral Filter

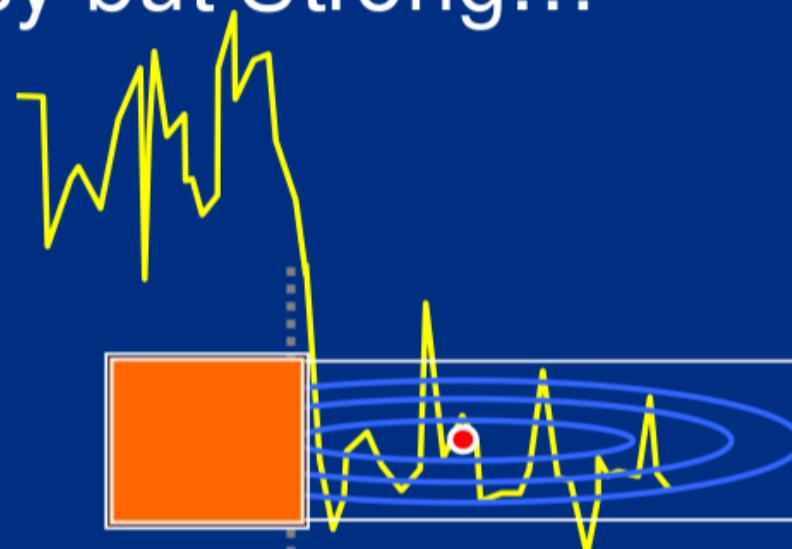
- 2 kinds of weights
    - Spatial Gaussian and Intensity Gaussian
    - A = image **without** flash
    - B = image **with** flash
  - Strategy
    - Smooth image A
    - 1) Weights based on spatial location → From A itself
    - 2) Weights based on intensity → From intensities in B



# Edge-Preserving Smoothing

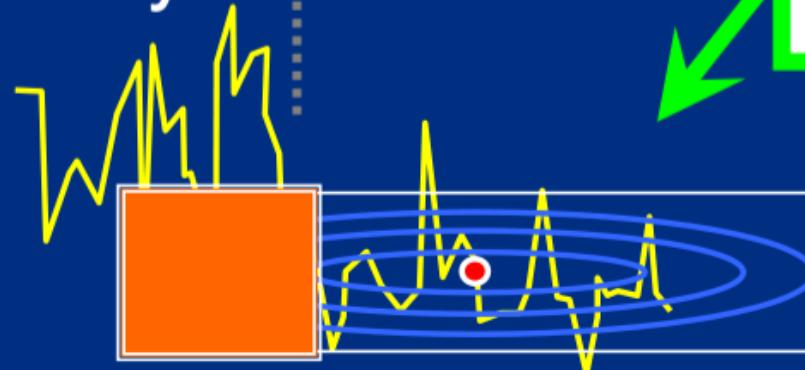
- Cross / Joint Bilateral Filter

Noisy but Strong...



Range filter preserves signal

Noisy and Weak...



Use stronger signal's range filter weights...

# Edge-Preserving Smoothing

- Cross / Joint Bilateral Filter
  - Enhances ability to find weak details in noisy image A (without flash)
    - Intensities in image B (with flash) inform which intensities to average in image A for noise reduction in image A

# Edge-

- Cross / Joint Bilateral Filter



(a) No-Flash

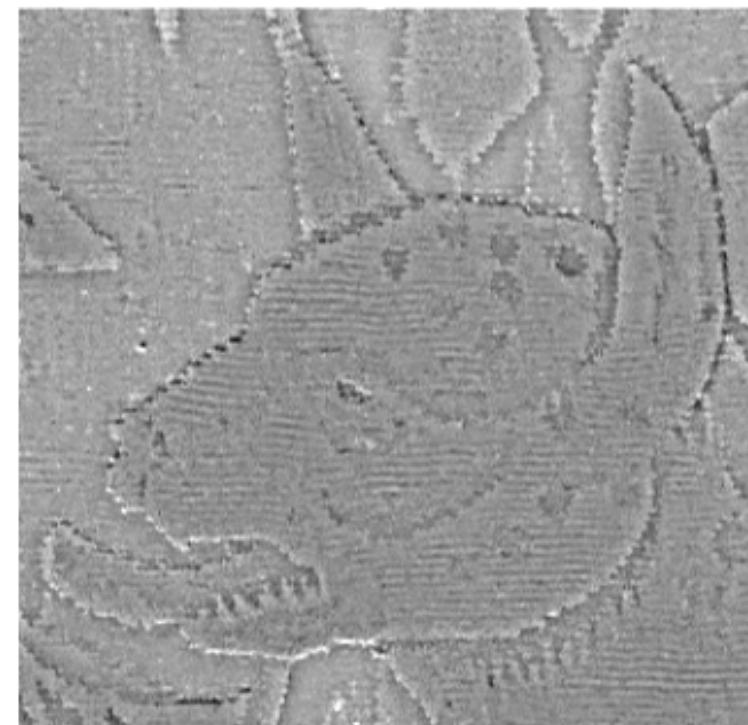
Flash



(b) Denoised via Bilateral Filter



(c) Denoised via Joint Bilateral Filter



(d) Difference

# Edge-Preserving Smoothing

- Bilateral Filter Limitations

- (1) Texture “softer” than intensity-kernel SD → removed



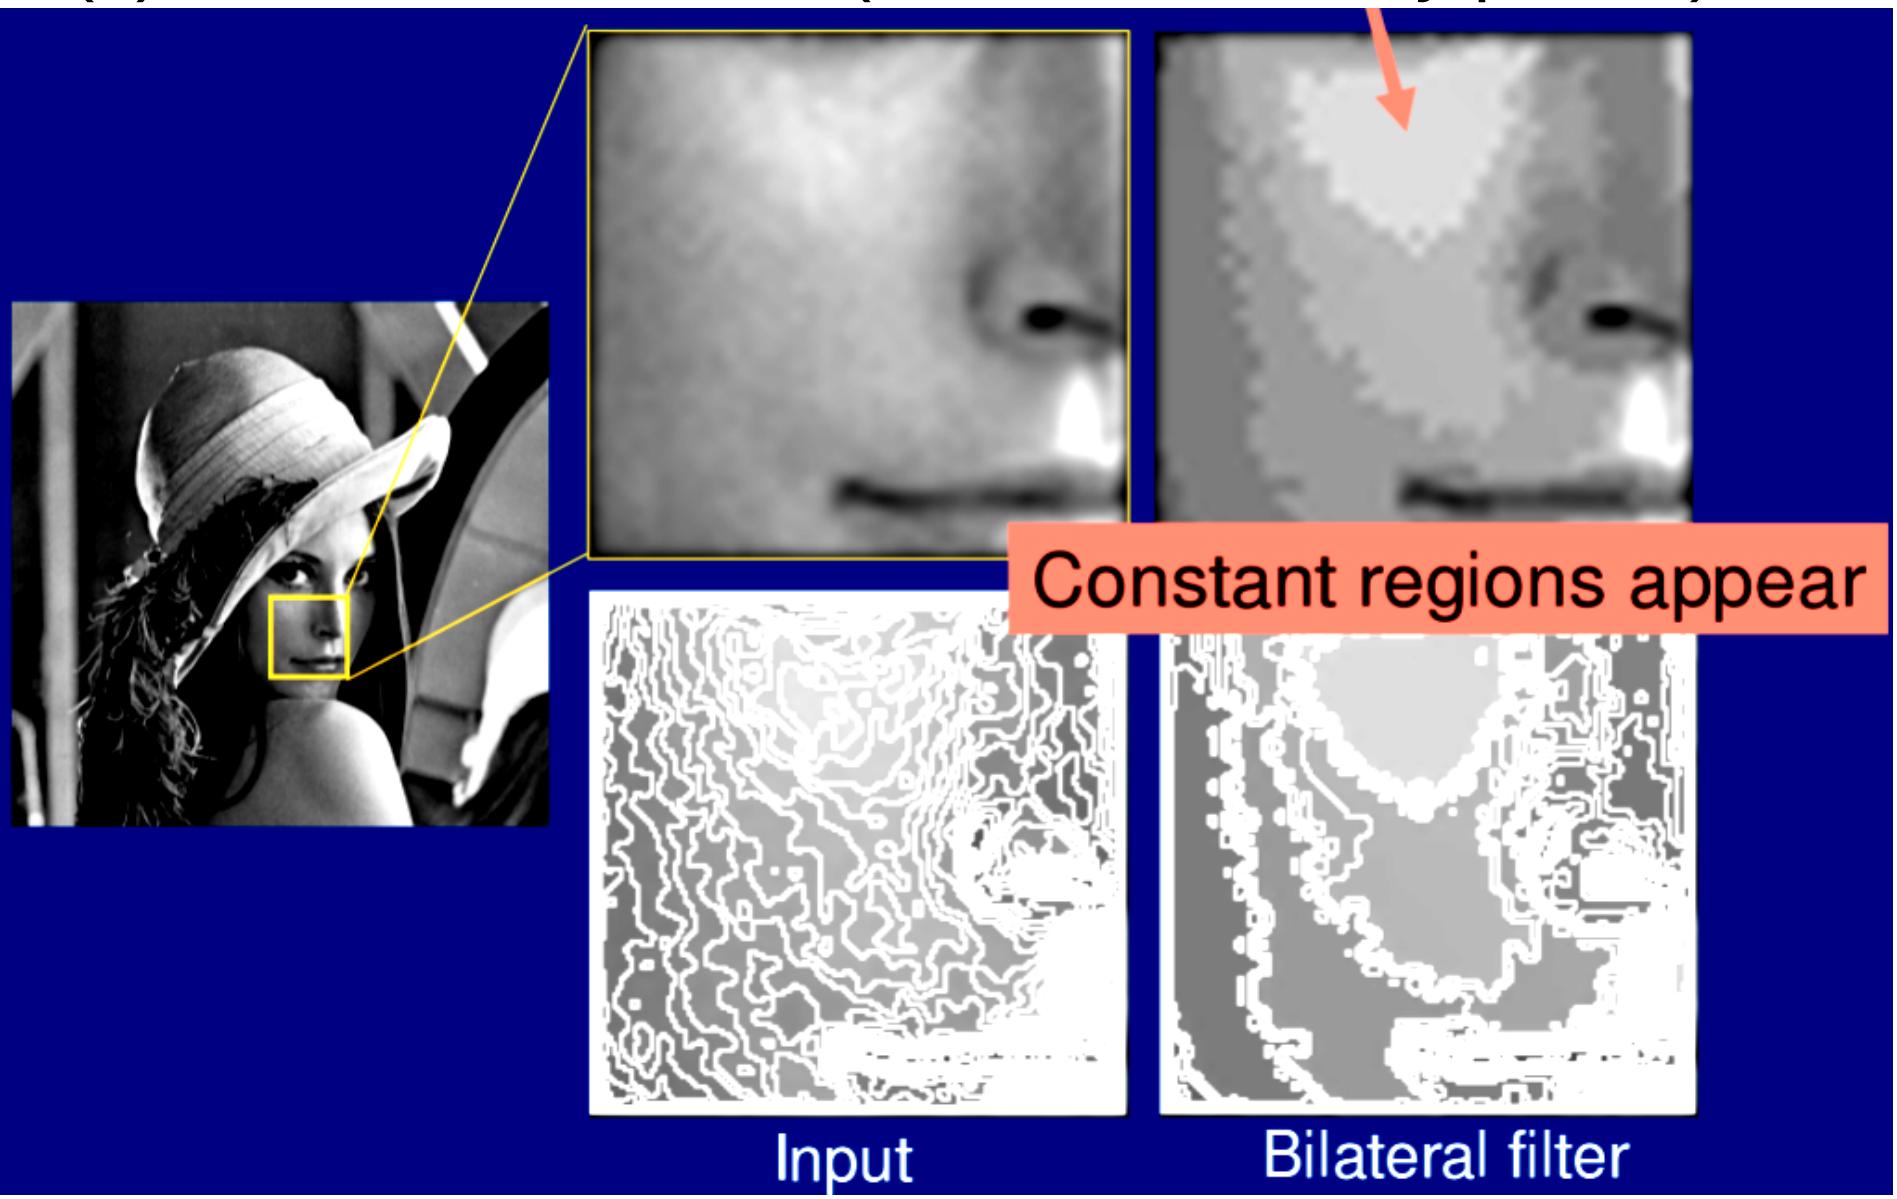
Input

Bilateral filter

# Edge-Preserving Smoothing

- Bilateral Filter Limitations

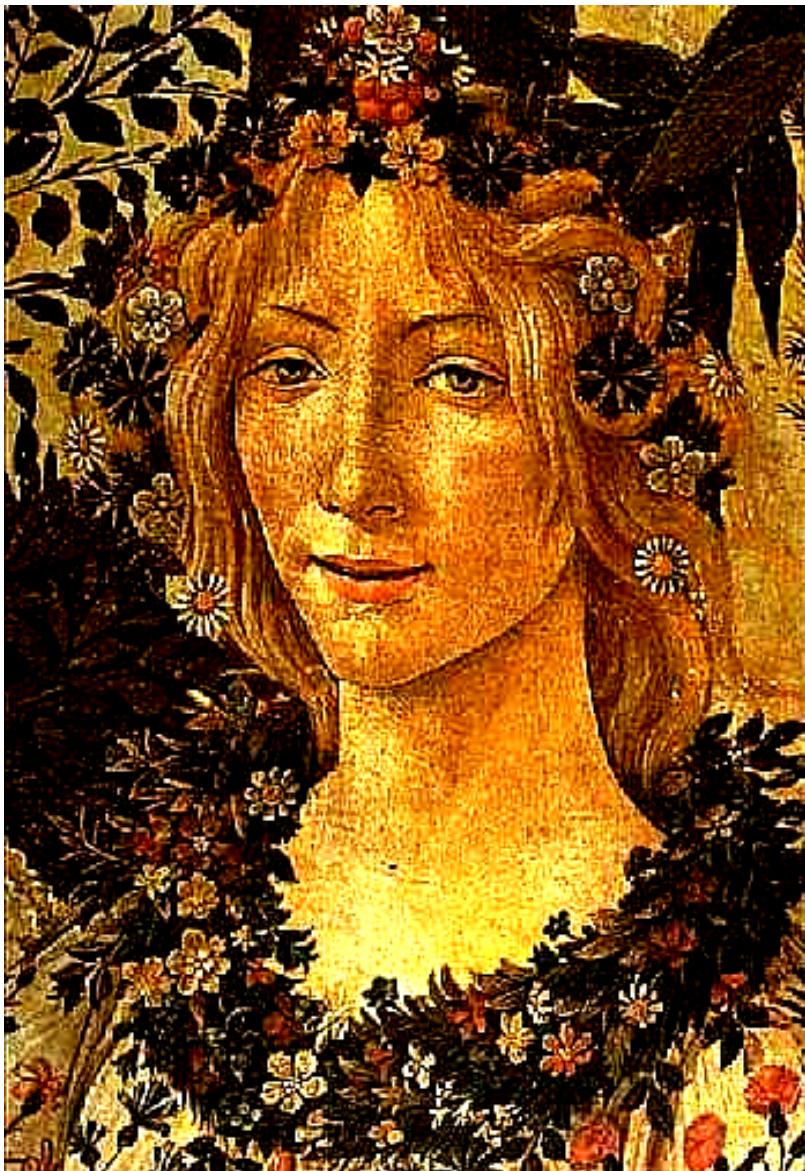
- (2) “Staircase” artifacts (constant-intensity pieces)



# Edge-Preserving Smoothing

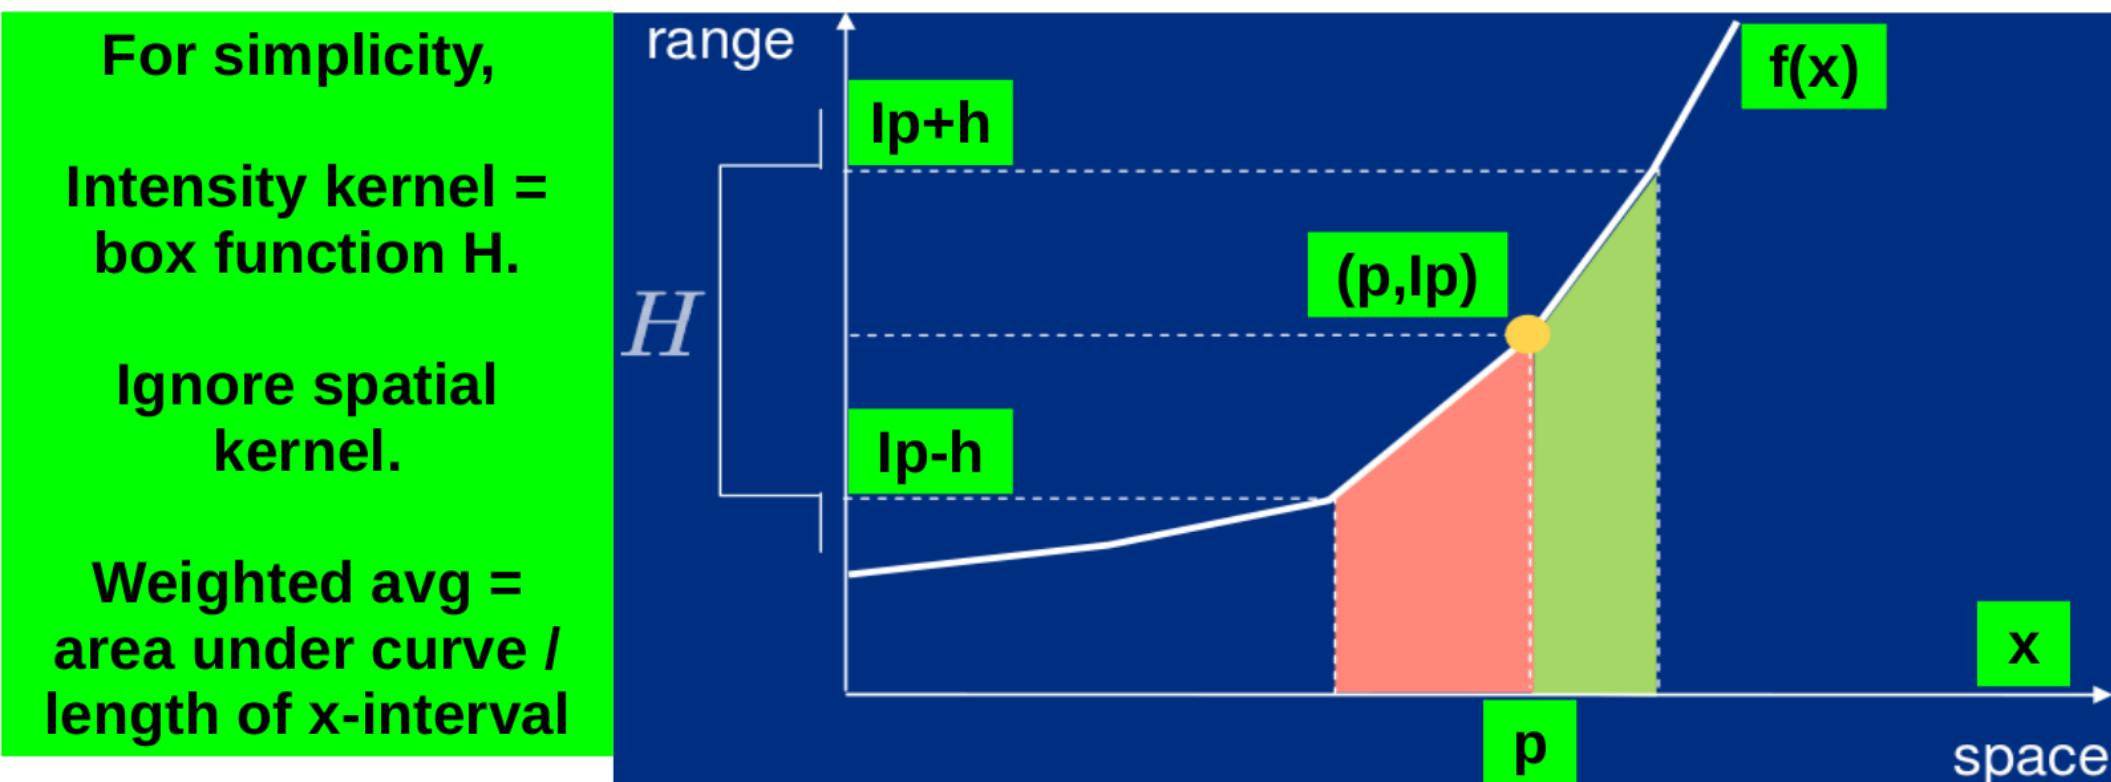
- Bilateral Filter Limitations

(2) “Staircase” artifacts



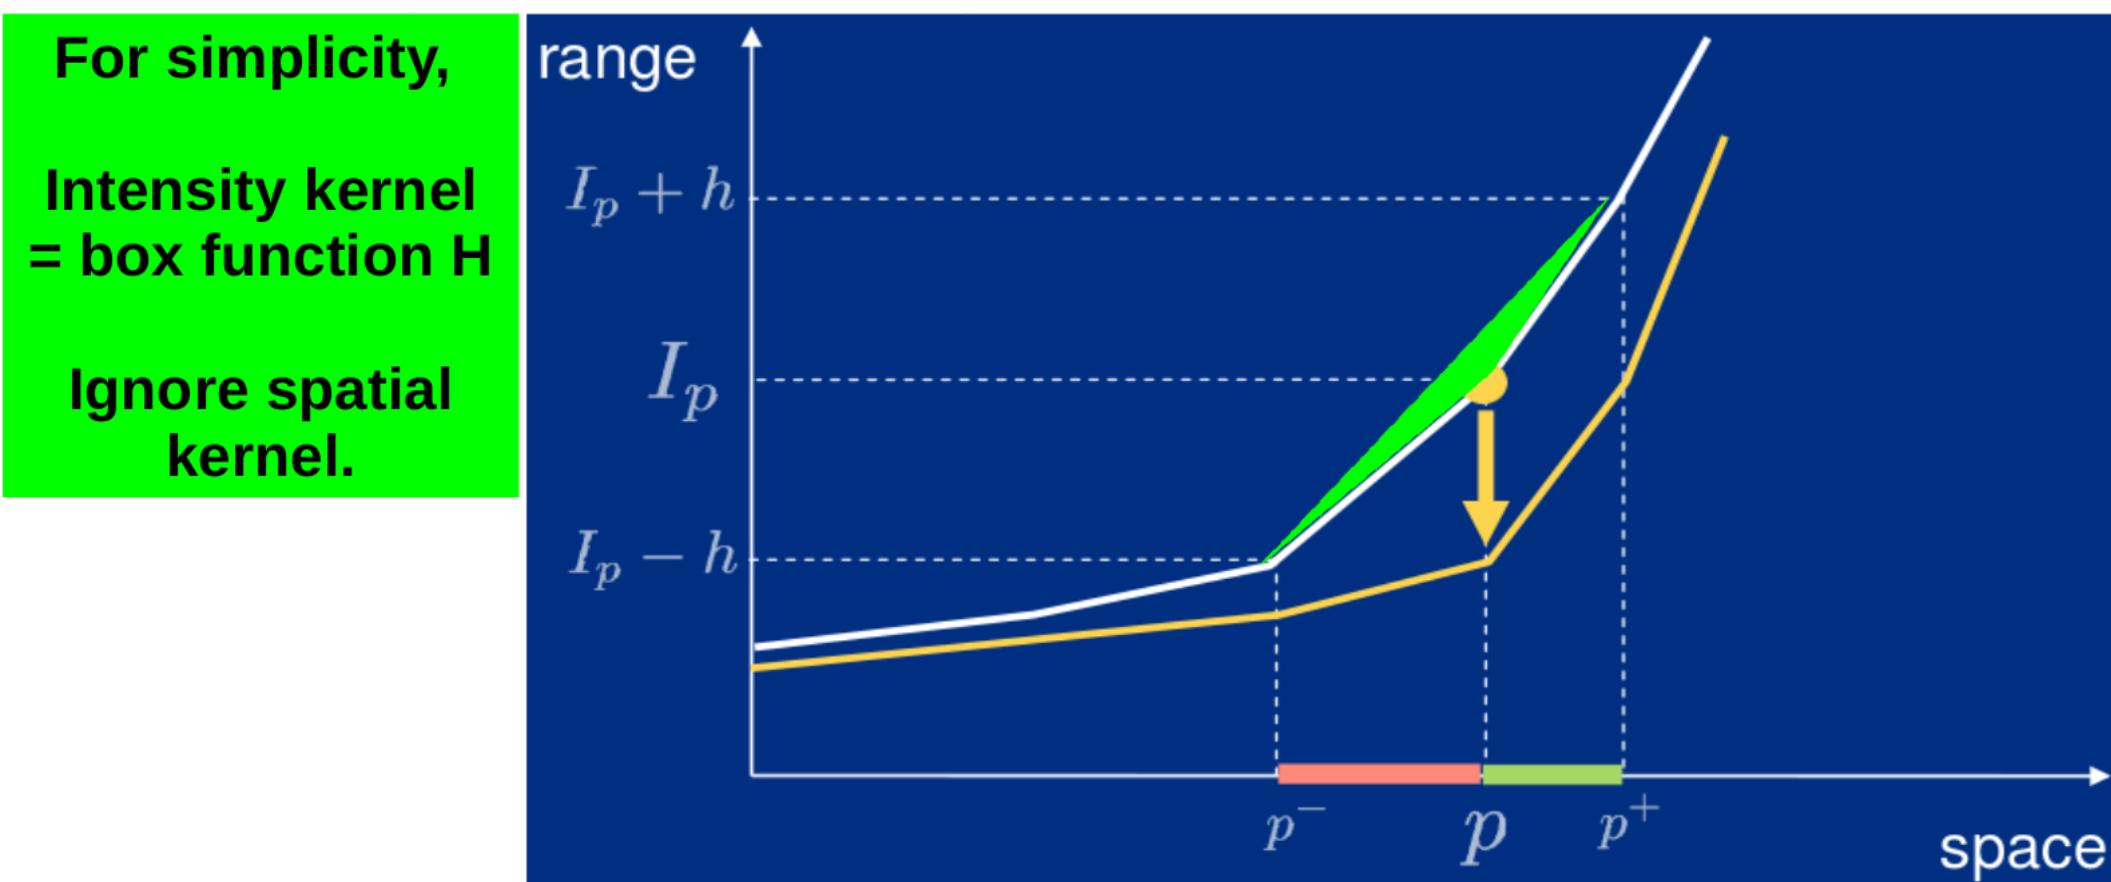
# Edge-Preserving Smoothing

- Bilateral Filter Limitations – Staircasing
  - Let intensity function  $f(x) =$  (locally) **monotonic + convex**
    - (1) Num. of pixels “q” s.t.  $I_q \in (I_p-h, I_p) >$  Num. of pixels “r” s.t.  $I_r \in (I_p, I_p+h)$
    - (2) All points “q” and “r” have equal weight



# Edge-Preserving Smoothing

- Bilateral Filter Limitations – Staircasing
  - (3) Weighted average
    - = (right-trapezoid area – green area) / length of x-interval
    - =  $I_p - \text{green area} / \text{length of x-interval} < I_p$



# Edge-Preserving Smoothing

- Bilateral Filter Limitations – Staircasing
  - Let intensity function  $f(x)$ = (locally) **monotonic + concave**
    - (1) Num. of pixels “q” s.t.  $I_q \in (I_p-h, I_p) <$  Num. of pixels “r” s.t.  $I_r \in (I_p, I_p+h)$
    - (2) All points “q” and “r” have equal weight
    - (3) **Weighted average at “p” will be  $> I_p$**
  - **Reverse smoothing (NOT “sharpening”) occurs !**



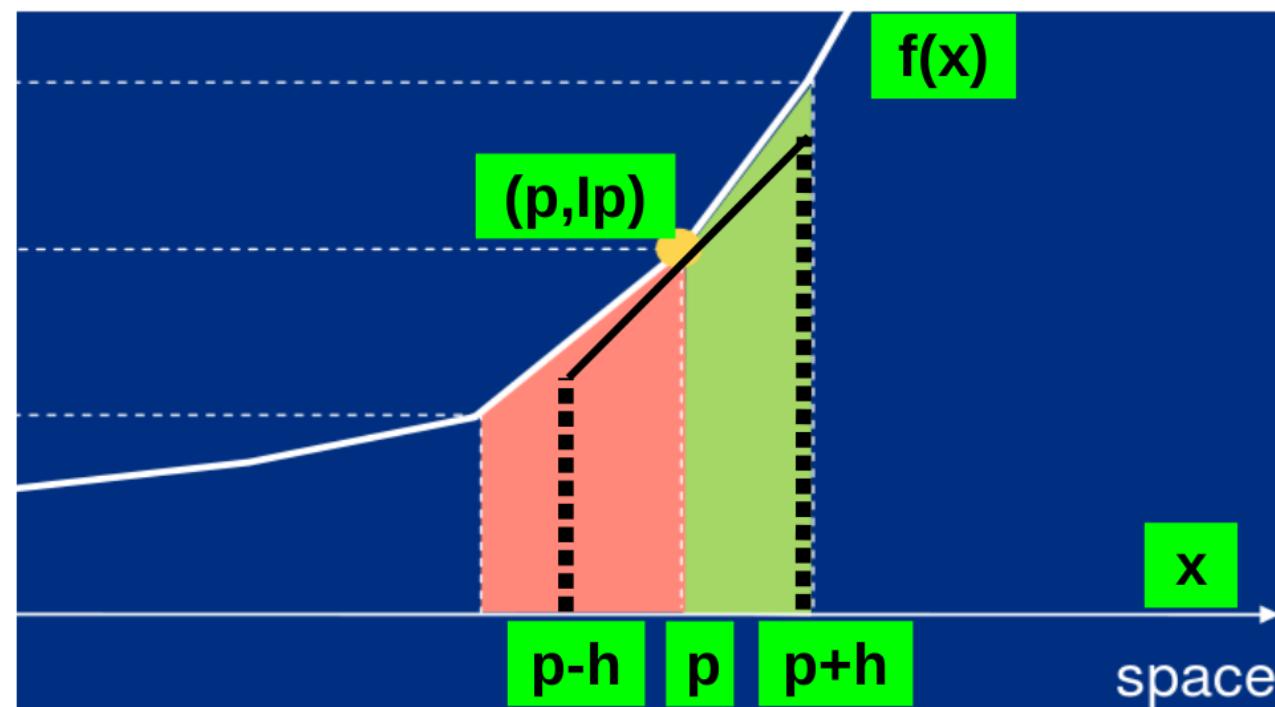
# Edge-Preserving Smoothing

- Gaussian convolution doesn't cause staircasing
  - Assume origin at  $(p, I_p)$ 
    - Box, Gaussian spatial-kernel  $k(x) \rightarrow$  even function over "x"
    - If intensity function = linear =  $\{f(x) = mx\} \rightarrow$  odd function over "x"
      - Weighted avg =  $I_p + \text{Integral}_{\{-h \rightarrow +h\}} mx k(x) dx = I_p$
    - Convex  $f(x) = mx + g(x)$  s.t.  $g(x) > 0 \rightarrow \text{Integral} > 0 \rightarrow w.\text{avg} > I_p$

Spatial kernel = box function H.

Ignore intensity kernel.

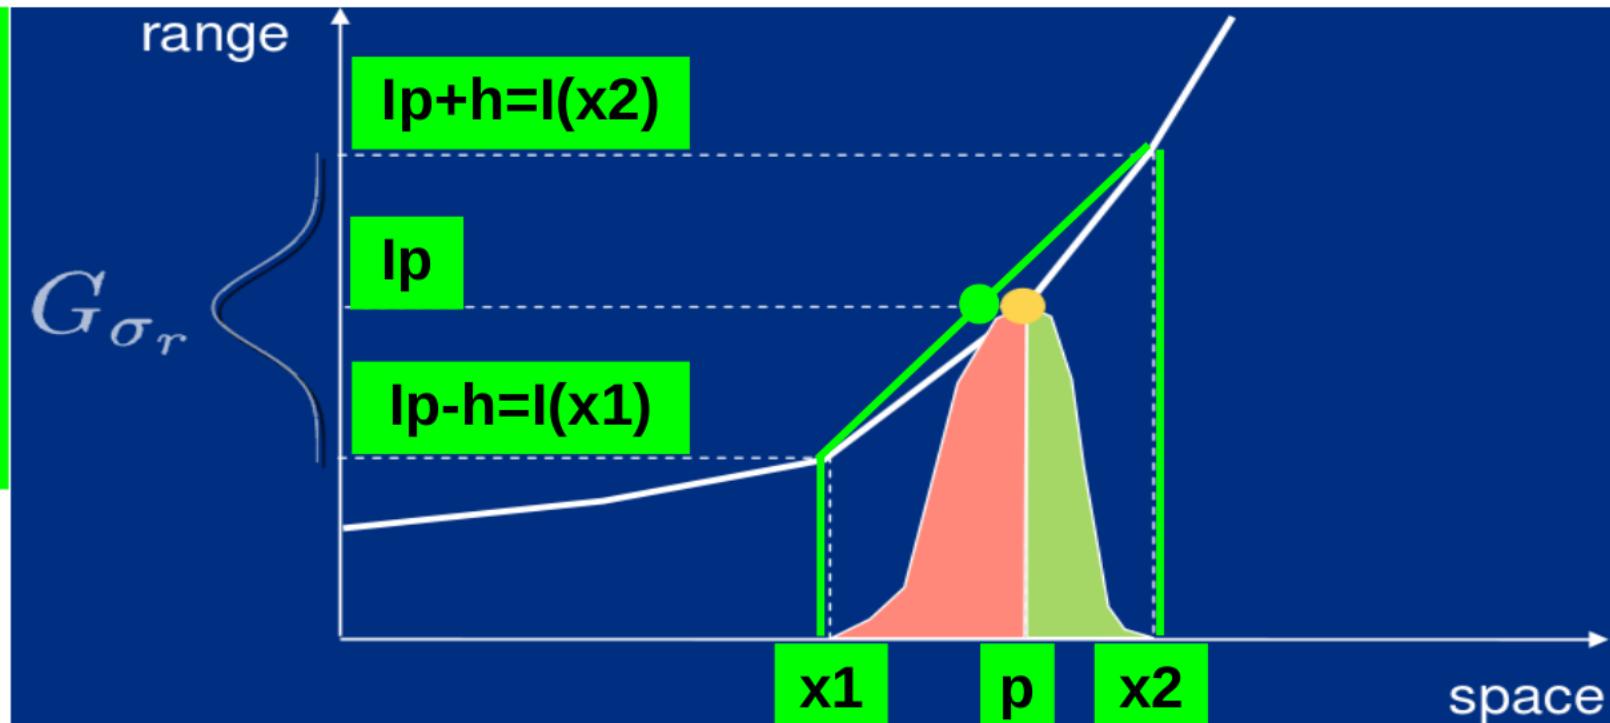
Weighted avg = area under curve / length of x-interval



# Edge-Preserving Smoothing

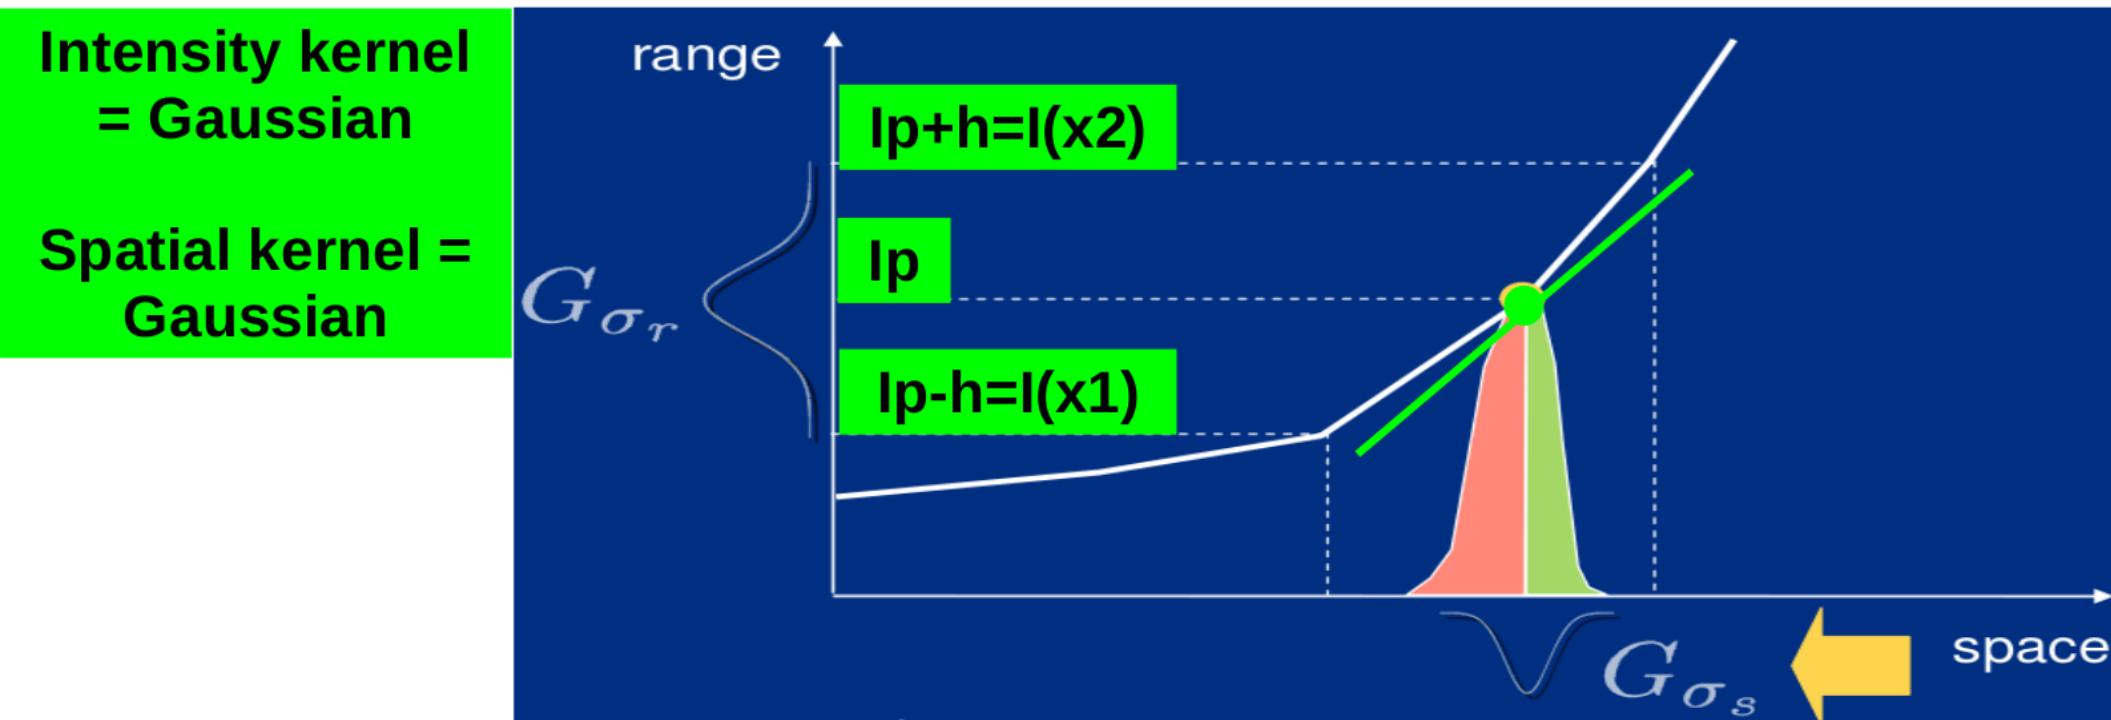
- Bilateral Filter Limitations – Staircasing
  - Assume **origin** at  $( (x_1 + x_2) / 2 , I_p )$ 
    - Gaussian intensity-kernel  $G(\cdot)$  → even function over “I” (& “x”)
    - If intensity function = linear =  $\{f(x) = mx\}$  → **odd** function over “x”
      - Weighted avg =  $I_p + \text{Integral}_{\{- (x_2 - x_1)/2 \rightarrow (x_2 - x_1)/2\}} mx G(mx) dx = I_p$
    - **Convex**  $f(x) = mx - g(x)$  s.t.  $g(x) > 0$  →  $\text{Integral} < 0 \rightarrow \text{W.Avg} < I_p$

For simplicity,  
Intensity kernel  
= Gaussian  
Ignore spatial  
kernel.



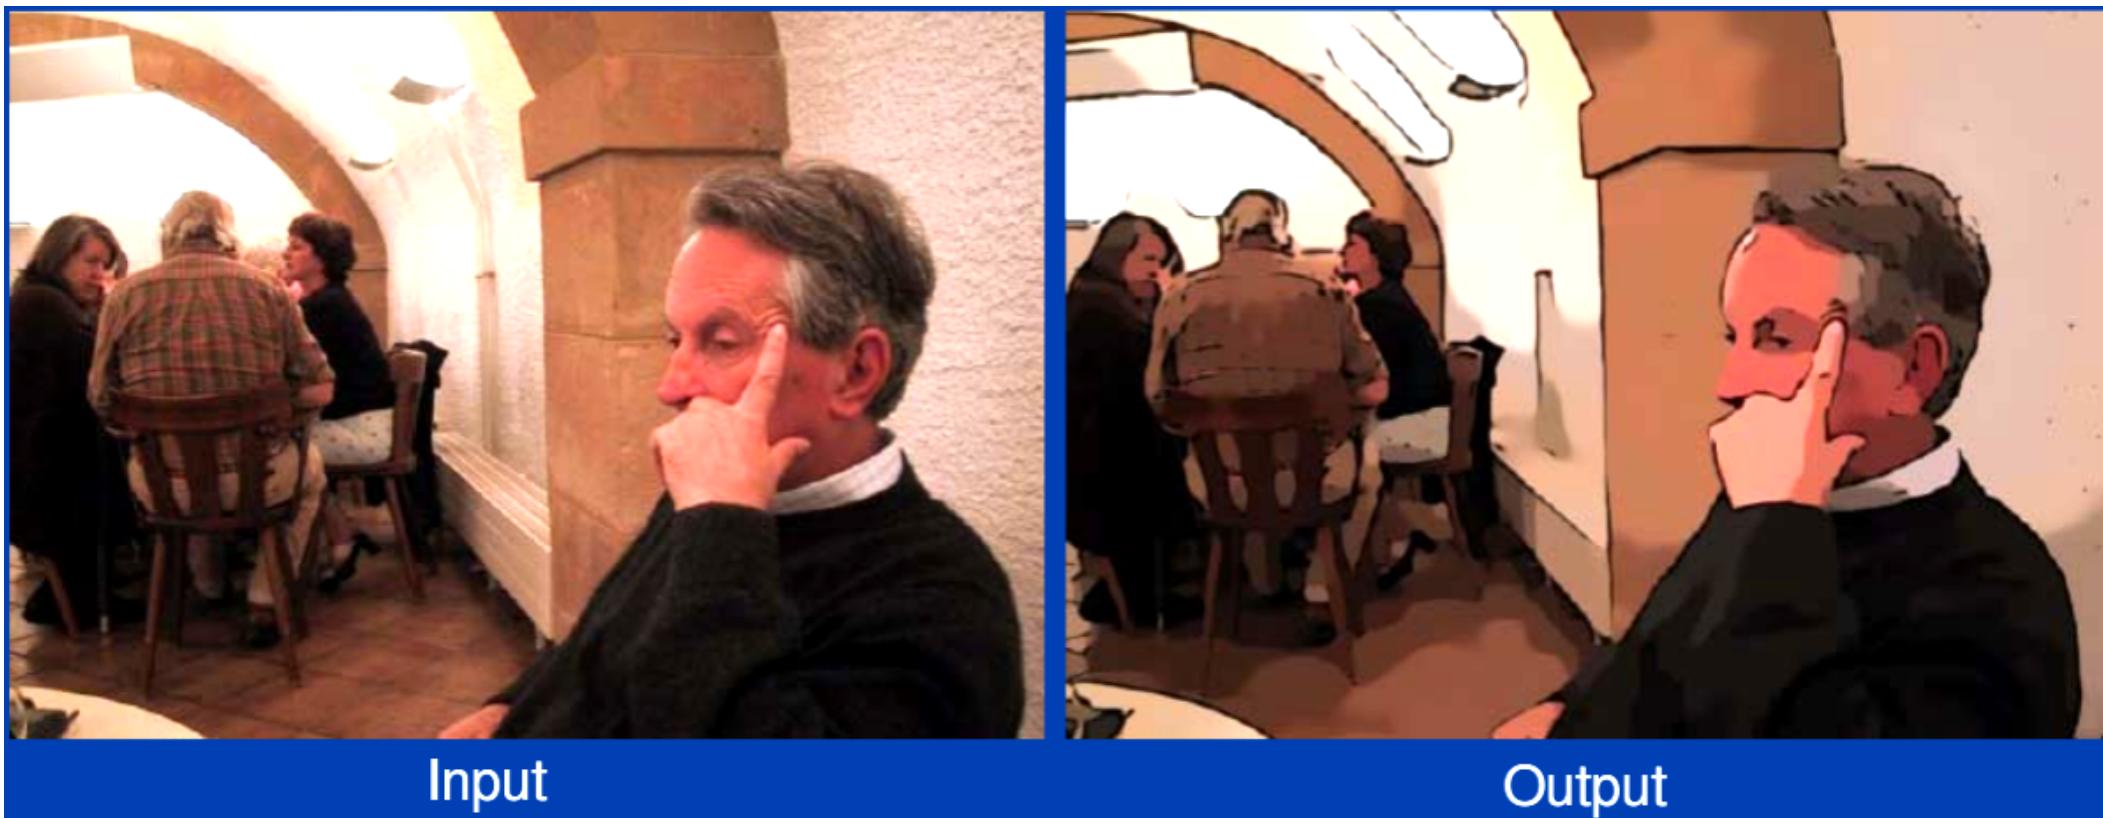
# Edge-Preserving Smoothing

- Bilateral Filter Limitations – Staircasing
  - Assume **origin at ( p , I<sub>p</sub> )**
    - Gaussian intensity-kernel + Gaussian spatial-kernel
    - If intensity function = linear → **odd** → Weighted avg = I<sub>p</sub>
    - Monotonic Convex f(x) → no guarantees (depends on I, weights)
      - In some cases, reverse smoothing can happen ← empirical evidence



# Edge-Preserving Smoothing

- Bilateral Filter Limitations
  - Can staircasing be useful ?
    - Cartoons !



# Edge-Preserving Smoothing

- Bilateral Filter Limitations

- Can this be useful ?      Cartoons !

[Winnemöller, Olsen, Gooch, 2006]

